

RISC-V Assembly Language Programming

(Draft v0.1-7-g4e7b603)

John Winans
jwinans@niu.edu

May 3, 2018

Copyright © 2018 John Winans

This document is made available under a Creative Commons Attribution 4.0 International License. See [Appendix A](#) for more information.

Download your own copy of this book from github here: <https://github.com/johnwinans/rvalp>.

This document may contain inaccuracies or errors. The author provides no guarantee regarding the accuracy of this document's contents. If you discover that this document contains errors, please notify the author.

Contents

Preface	vi
1 Number Systems	1
1.1 Integers	1
1.1.1 Converting Between Bases	3
1.1.1.1 From Binary to Decimal	3
1.1.1.2 From Binary to Hexadecimal	3
1.1.1.3 From Hexadecimal to Binary	3
1.1.1.4 From Decimal to Binary	4
1.1.1.5 From Decimal to Hex	4
1.1.2 Addition of Binary Numbers	4
1.1.3 Signed Numbers	5
1.1.3.1 Converting between Positive and Negative	5
1.1.4 Subtraction of Binary Numbers	6
1.1.5 Truncation and Overflow	7
1.1.6 Logical/Boolean Functions	7
1.1.6.1 NOT	7
1.1.6.2 AND	7
1.1.6.3 OR	8
1.1.6.4 XOR	8
1.2 IEEE-754 Floating Point Number Representation	8
1.2.1 Floating Point Number Accuracy	10
1.2.1.1 Powers Of Two	10
1.2.1.2 Clean Decimal Numbers	10
1.2.1.3 Accumulation of Error	11
1.2.2 Reducing Error Accumulation	12
2 The RISC-V GNU Toolchain	14
3 RV32 Machine Instructions	16

3.1	Introduction	16
3.2	Conventions and Terminology	16
3.2.1	XLEN	16
3.2.2	sx(val)	16
3.2.3	zx(val)	17
3.2.4	zr(val)	17
3.2.5	Sign Extended Left and Zero Extend Right	17
3.2.6	m8(addr)	18
3.2.7	m16(addr)	18
3.2.8	m32(addr)	19
3.2.9	m64(addr)	19
3.2.10	m128(addr)	19
3.2.11	.+offset	19
3.2.12	.-offset	19
3.2.13	pc	19
3.2.14	rd	20
3.2.15	rs1	20
3.2.16	rs2	20
3.2.17	imm	20
3.2.18	rsN[h:l]	20
3.3	Addressing Modes	20
3.4	Instruction Encoding Formats	20
3.4.1	U Type	21
3.4.2	J Type	22
3.4.3	R Type	23
3.4.4	I Type	23
3.4.5	S Type	23
3.4.6	B Type	23
3.4.7	CPU Registers	24
3.5	memory	24
3.6	RV32I Base Instruction Set	25
3.6.1	LUI rd, imm	25
3.6.2	AUIPC rd, imm	26
3.6.3	JAL rd, imm	26
3.6.4	JALR rd, rs1, imm	27
3.6.5	BEQ rs1, rs2, imm	28
3.6.6	BNE rs1, rs2, imm	29
3.6.7	BLT rs1, rs2, imm	29

3.6.8	BGE rs1, rs2, imm	30
3.6.9	BLTU rs1, rs2, imm	30
3.6.10	BGEU rs1, rs2, imm	30
3.6.11	LB rd, imm(rs1)	31
3.6.12	LH rd, imm(rs1)	31
3.6.13	LW rd, imm(rs1)	31
3.6.14	LBU rd, imm(rs1)	32
3.6.15	LHU rd, imm(rs1)	32
3.6.16	SB rs2, imm(rs1)	33
3.6.17	SH rs2, imm(rs1)	33
3.6.18	SW rs2, imm(rs1)	33
3.6.19	ADDI rd, rs1, imm	34
3.6.20	SLTI rd, rs1, imm	34
3.6.21	SLTIU rd, rs1, imm	35
3.6.22	XORI rd, rs1, imm	35
3.6.23	ORI rd, rs1, imm	36
3.6.24	ANDI rd, rs1, imm	36
3.6.25	SLLI rd, rs1, shamt	37
3.6.26	SRLI rd, rs1, shamt	37
3.6.27	SRAI rd, rs1, shamt	37
3.6.28	ADD rd, rs1, rs2	38
3.6.29	SUB rd, rs1, rs2	38
3.6.30	SLL rd, rs1, rs2	39
3.6.31	SLT rd, rs1, rs2	39
3.6.32	SLTU rd, rs1, rs2	40
3.6.33	XOR rd, rs1, rs2	40
3.6.34	SRL rd, rs1, rs2	41
3.6.35	SRA rd, rs1, rs2	41
3.6.36	OR rd, rs1, rs2	42
3.6.37	AND rd, rs1, rs2	42
3.6.38	FENCE predecessor, successor	43
3.6.39	FENCE.I	43
3.6.40	ECALL	44
3.6.41	EBREAK	44
3.6.42	CSR _{RR} W rd, csr, rs1	44
3.6.43	CSR _{RS} rd, csr, rs1	44
3.6.44	CSR _{RC} rd, csr, rs1	44
3.6.45	CSR _{RWI} rd, csr, imm	45

3.6.46 CSRRSI rd, csr, rs1	45
3.6.47 CSRRCI rd, csr, rs1	45
3.7 RV32M Standard Extension	45
3.8 RV32A Standard Extension	45
3.9 RV32F Standard Extension	46
3.10 RV32D Standard Extension	46
A Attribution 4.0 International	47
Bibliography	51
Index	52
Glossary	54

Preface

I set out to this book because I couldn't find it in a single volume elsewhere.

The closest thing to what I sought when deciding to collect my thoughts into this document would be select portions of *The RISC-V Instruction Set Manual, Volume I: User-Level ISA, Document Version 2.2*[\[1\]](#), *The RISC-V Reader*[\[2\]](#), and *Computer Organization and Design RISC-V Edition: The Hardware Software Interface*[\[3\]](#).

There *are* some terse guides around the Internet that are suitable for those that already know an assembly language. With all the (deserved) excitement brewing over system organization (and the need to compress the time out of university courses targeting assembly language programming [\[4\]](#)), it is no surprise that RISC-V texts for the beginning assembly programmer are not (yet) available.

When I got started in computing I learned how to count in binary in a high school electronics course using data sheets for integrated circuits such as the 74191[\[5\]](#) and 74154[\[6\]](#) prior to knowing that assembly language even existed.

I learned assembler from data sheets and texts (that are still sitting on my shelves) such as:

- The MCS-85 User's Manual[\[7\]](#)
- The EDTASM Manual[\[8\]](#)
- The MC68000 User's Manual[\[9\]](#)
- Assembler Language With ASSIST[\[10\]](#)
- IBM System/370 Principals of Operation[\[11\]](#)
- OS/VS-DOS/VSE-VM/370 Assembler Language[\[12\]](#)
- ... and several others

One way or another all of them discuss each CPU instruction in excruciating detail with both a logical and narrative description. For RISC-V this is also the case for the *RISC-V Reader*[\[2\]](#) and the *Computer Organization and Design RISC-V Edition*[\[3\]](#) books and is also present in this text (I consider that to be the minimal level of responsibility.)

Where I hope this text will differentiate itself from the existing RISC-V titles is in its attempt to address the needs of those learning assembly language for the first time. To this end I have primed this project with some of the material from old handouts I used when teaching assembly language programming in the late '80s.

Chapter 1

Number Systems

RISC-V systems represent information using binary values stored in little-endian order.¹

1.1 Integers

A binary integer is constructed with only 1s and 0s in the same manner as decimal numbers are constructed with values from 0 to 9.

Counting in binary is the same as in decimal. For example, when adding 1 to 9, the carry is added to the next place value. When subtracting 1 from 0, a borrow is required and so on.

Figure [Figure 1.1](#) shows an abridged table of the decimal, binary and hexadecimal values from 0 to 129.

One way to look at this table is on a per-row basis where each place value is represented by the base raised to the power of the place value position (shown in the column headings.) This is useful when converting arbitrary values between bases. For example to interpret the decimal value on the fourth row:

$$0 \times 10^2 + 0 \times 10^1 + 3 \times 10^0 = 3_{10}$$

And to interpret binary value on the same row by converting it to decimal:

$$0 \times 2^7 + 0 \times 2^6 + 0 \times 2^5 + 0 \times 2^4 + 0 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 1 \times 2^0 = 3_{10}$$

And the same for the hexadecimal value:

$$0 \times 16^1 + 3 \times 16^0 = 3_{10}$$

Another way to look at this table is on a per-column basis. When tasked with drawing such a table by hand, it might be useful to observe that, just as in decimal, the right-most column will cycle through all of the values represented in the chosen base then cycle back to zero and repeat. (For example, in binary this pattern is 0-1-0-1-0-1-0-...) The next column in each base will cycle in the same manner except each of the values is repeated as many times as is represented by the place value (in the case of decimal, 10^1 times, binary 2^1 times, hex 16^1 times. Again, the for binary numbers this pattern is 0-0-1-1-0-0-1-1-...) This continues for as many columns as are needed to represent the magnitude of the desired number.

¹See[\[13\]](#) for some history of the big/little-endian “controversy.”

Decimal			Binary								Hex	
10^2	10^1	10^0	2^7	2^6	2^5	2^4	2^3	2^2	2^1	2^0	16^1	16^0
100	10	1	128	64	32	16	8	4	2	1	16	1
0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	1	0	0	0	0	0	0	0	1	0	1
0	0	2	0	0	0	0	0	0	1	0	0	2
0	0	3	0	0	0	0	0	0	1	1	0	3
0	0	4	0	0	0	0	0	1	0	0	0	4
0	0	5	0	0	0	0	0	1	0	1	0	5
0	0	6	0	0	0	0	0	1	1	0	0	6
0	0	7	0	0	0	0	0	1	1	1	0	7
0	0	8	0	0	0	0	1	0	0	0	0	8
0	0	9	0	0	0	0	1	0	0	1	0	9
0	1	0	0	0	0	0	1	0	1	0	0	a
0	1	1	0	0	0	0	1	0	1	1	0	b
0	1	2	0	0	0	0	1	1	0	0	0	c
0	1	3	0	0	0	0	1	1	0	1	0	d
0	1	4	0	0	0	0	1	1	1	0	0	e
0	1	5	0	0	0	0	1	1	1	1	0	f
0	1	6	0	0	0	1	0	0	0	0	1	0
0	1	7	0	0	0	1	0	0	0	1	1	1
...			
1	2	5	0	1	1	1	1	1	0	1	7	d
1	2	6	0	1	1	1	1	1	1	0	7	e
1	2	7	0	1	1	1	1	1	1	1	7	f
1	2	8	1	0	0	0	0	0	0	0	8	0

Figure 1.1: Counting in decimal, binary and hexadecimal.

Another item worth noting is that any even binary number will always have a 0 LSB and odd numbers will always have a 1 LSB.

As is customary in decimal, leading zeroes are sometimes not shown for readability.

The relationship between binary and hex values is also worth taking note. Because $2^4 = 16$, there is a clean and simple grouping of 4 [bits](#) to 1 [bit](#). There is no such relationship between binary and decimal.

Writing and reading numbers in binary that are longer than 8 bits is cumbersome and prone to error. The simple conversion between binary and hex makes hex a convenient shorthand for expressing binary values in many situations.

For example, consider the following value expressed in binary, hexadecimal and decimal (spaced to show the relationship between binary and hex):

```

Binary value:    0010 0111 1011 1010 1100 1100 1111 0101
Hex Value:       2    7    B    A    C    C    F    5
Decimal Value:   666553589

```

Empirically we can see that grouping the bits into sets of four allows an easy conversion to hex and expressing it as such is $\frac{1}{4}$ as long as in binary while at the same time allowing for easy conversion back to binary.

The decimal value in this example does not easily convey a sense of the binary value.

1.1.1 Converting Between Bases

1.1.1.1 From Binary to Decimal

Alas, it is occasionally necessary to convert between decimal, binary and/or hex.

To convert from binary to decimal, put the decimal value of the place values ... 8 4 2 1 over the binary digits like this:

128	64	32	16	8	4	2	1
0	0	0	1	1	0	1	1

Now sum the place-values that are expressed in decimal for each bit with the value of 1: $16 + 8 + 2 + 1$. The integer binary value 00011011_2 represents the decimal value 27_{10} .

1.1.1.2 From Binary to Hexadecimal

Conversion from binary to hex involves grouping the bits into sets of four and then performing the same summing process as shown above. If there is not a multiple of four bits then extend the binary to the left with zeroes to make it so.

Grouping the bits into sets of four and summing:

Place:	8 4 2 1	8 4 2 1	8 4 2 1	8 4 2 1
Binary:	0 1 1 0	1 1 0 1	1 0 1 0	1 1 1 0
Decimal:	4+2 =6	8+4+ 1=13	8+ 2 =10	8+4+2 =14

After the summing, convert each decimal value to hex. The decimal values from 0–9 are the same values in hex. Because we don't have any more numerals to represent the values from 10–15, we use the first 6 letters (See the right-most column of [Figure 1.1.](#)) Fortunately there are only six hex mappings involving letters. Thus it is reasonable to memorize them.

Continuing this example:

Decimal:	6	13	10	14
Hex:	6	D	A	E

1.1.1.3 From Hexadecimal to Binary

Again, the four-bit mapping between binary and hex makes this task as straight forward as using a look-up table.

For each [hit](#) (Hex digIT), translate it to its unique four-bit pattern. Perform this task either by memorizing each of the 16 patterns or by converting each hit to decimal first and then converting each four-bit binary value to decimal using the place-value summing method discussed in [subsubsection 1.1.1.1](#).

For example:

Hex:	4	C	
Binary:	0 1 0 0	1 1 0 0	
Decimal:	128 64 32 16	8 4 2 1	
Sum:	64+	8+4	= 76

1.1.1.4 From Decimal to Binary

To convert arbitrary decimal numbers to binary, extend the list of binary place values until it exceeds the value of the decimal number being converted. Then make successive subtractions of each of the place values that would yield a non-negative result.

For example, to convert 1234_{10} to binary:

Place values: 2048-1024-512-256-128-64-32-16-8-4-2-1

0	2048	(too big)
1	1234 - 1024 = 210	
0	512	(too big)
0	256	(too big)
1	210 - 128 = 82	
1	82 - 64 = 18	
0	32	(too big)
1	18 - 16 = 2	
0	8	(too big)
0	4	(too big)
1	2 - 2 = 0	
0	1	(too big)

The answer using this notation is listed vertically in the left column with the **MSB** on the top and the **LSB** on the bottom line: 010011010010₂.

1.1.1.5 From Decimal to Hex

Conversion from decimal to hex can be done by using the place values for base-16 and the same math as from decimal to binary or by first converting the decimal value to binary and then from binary to hex by using the methods discussed above.

Because binary and hex are so closely related, performing a conversion by way of binary is quite straight forward.

1.1.2 Addition of Binary Numbers

The addition of binary numbers can be performed long-hand the same way decimal addition is taught in grade school. In fact binary addition is easier since it only involves adding 0 or 1.

The first thing to note that in any number base $0 + 0 = 0$, $0 + 1 = 1$, and $1 + 0 = 1$. Since there is no “two” in binary (just like there is no “ten” decimal) adding $1 + 1$ results in a zero with a carry as in: $1 + 1 = 10_2$ and in: $1 + 1 + 1 = 11_2$. Using these five sums, any two binary integers can be added.

For example:

```

    111111 1111 <== carries
    0110101111001111 <== addend
+   0000011101100011 <== addend
-----
    0111001100110010 <== sum

```

1.1.3 Signed Numbers

There are multiple methods used to represent signed binary integers. The method used by most modern computers is called “two’s complement.”

A two’s complement number is encoded in such a manner as to simplify the hardware used to add, subtract and compare integers.

A simple method of thinking about two’s complement numbers is to negate the place value of the [MSB](#). For example, the number one is represented the same as discussed before:

```

-128 64 32 16  8  4  2  1
   0  0  0  0  0  0  0  1

```

The [MSB](#) of any negative number in this format will always be 1. For example the value -1_{10} is:

```

-128 64 32 16  8  4  2  1
   1  1  1  1  1  1  1  1

```

... because: $-128 + 64 + 32 + 16 + 8 + 4 + 2 + 1 = -1$.

This format has the virtue of allowing the same addition logic discussed above to be used to calculate $-1 + 1 = 0$.

```

-128 64 32 16  8  4  2  1 <== place value
1  1  1  1  1  1  1  1  0 <== carries
   1  1  1  1  1  1  1  1 <== addend (-1)
+  0  0  0  0  0  0  0  1 <== addend (1)
-----
1  0  0  0  0  0  0  0  0 <== sum (0 with an overflow)

```

In order for this to work, the [overflow](#) carry out of the sum of the MSBs is ignored.

1.1.3.1 Converting between Positive and Negative

Changing the sign on two’s complement numbers can be described as inverting all of the bits (which is also known as the one’s complement) and then add one.

For example, inverting the number *four*:

```

-128 64 32 16  8  4  2  1
   0  0  0  0  0  1  0  0 <== 4

```

```

          1 1    <== carries
    1 1 1 1 1 0 1 1 <== one's complement of 4
+ 0 0 0 0 0 0 0 1 <== plus 1
-----
    1 1 1 1 1 1 0 0 <== -4

```

This can be verified by adding 5 to the result and observe that the sum is 1:

```

-128 64 32 16 8 4 2 1
    1 1 1 1 1      <== carries
    1 1 1 1 1 1 0 0 <== -4
+ 0 0 0 0 0 1 0 1 <== 5
-----
    1 0 0 0 0 0 0 1

```

Note that the changing of the sign using this method is symmetric in that it is identical when converting from negative to positive and when converting from positive to negative: flip the bits and add 1.

For example, changing the value -4 to 4 to illustrate the reverse of the conversion above:

```

-128 64 32 16 8 4 2 1
    1 1 1 1 1 1 0 0 <== -4

          1 1    <== carries
    0 0 0 0 0 0 1 1 <== one's complement of -4
+ 0 0 0 0 0 0 0 1 <== plus 1
-----
    0 0 0 0 0 1 0 0 <== 4

```

1.1.4 Subtraction of Binary Numbers

Subtraction of binary numbers is performed by first negating the subtrahend and then adding the two numbers. Due to the nature of two's complement numbers this will work for both signed and unsigned numbers.

To calculate $-4 - 8 = -12$

```

-128 64 32 16 8 4 2 1
    1 1 1 1 1 1 0 0 <== -4
- 0 0 0 0 0 1 0 0 <== 8
-----
    1 1 1 1 1 1 0 0 <== -8

          1 1 1    <== carries
    1 1 1 1 0 1 1 1 <== one's complement of -8
+ 0 0 0 0 0 0 0 1 <== plus 1
-----
    1 1 1 1 1 0 0 0 <== -8

```

➡ Fix Me:

This section needs more examples of subtracting signed and unsigned numbers and a discussion on how signedness is not relevant until the results are interpreted. For example adding $-4 + -8 = -12$ using two 8-bit numbers is the same as adding $252 + 248 = 500$ and truncating the result to 244.

```

  1  1  1  1          <== carries
  1  1  1  1  1  1  0  0 <== -4
+ 1  1  1  1  1  0  0  0 <== -8
-----
  1  1  1  1  1  0  1  0  0 <== -12

```

1.1.5 Truncation and Overflow

Discuss the details of truncation and overflow here.

➡ Fix Me:

This chapter should be made consistent in its use of truncation and overflow as occur with signed and unsigned addition and subtraction.

1.1.6 Logical/Boolean Functions

Unlike addition and subtraction, boolean functions apply on a per-bit basis. When applied to multi-bit values, each bit position is operated upon independently of the other bits.

➡ Fix Me:

This is unclear. Need to define bit positions and probably should add basic truth table diagrams.

1.1.6.1 NOT

The *NOT* operator applies to a single operand and represents the opposite of the input.

➡ Fix Me:

Need to define 1 as true and 0 as false somewhere.

If the input is 1 then the output is 0. If the input is 0 then the output is 1. In other words, the output value is *not* that of the input value.

➡ Fix Me:

Need to define unary, binary and ternary operators without confusing binary operators with binary numbers.

This text will use the operator used in the C language when discussing the *NOT* operator in symbolic form. Specifically the tilde: '~'.

```

~ 1 1 1 1 0 1 0 1 <== A
-----
  0 0 0 0 1 0 1 0 <== output

```

In a line of code the above might read like this: `output = ~A`

1.1.6.2 AND

The boolean *and* function has two or more inputs and the output is a single bit. The output is 1 if and only if all of the input values are 1. Otherwise it is 0.

This text will use the operator used in the C language when discussing the *AND* operator in symbolic form. Specifically the ampersand: '&'.

This function works like it does in spoken language. For example if A is 1 *AND* B is 1 then the output is 1 (true). Otherwise the output is 0 (false). For example:

```

  1 1 1 1 0 1 0 1 <== A
& 1 0 0 1 0 0 1 1 <== B
-----
  1 0 0 1 0 0 0 1 <== output

```

In a line of code the above might read like this: `output = A & B`

1.1.6.3 OR

The boolean *or* function has two or more inputs and the output is a single bit. The output is 1 if at least one of the input values are 1.

This text will use the operator used in the C language when discussing the *OR* operator in symbolic form. Specifically the pipe: '|'.

This function works like it does in spoken language. For example if A is 1 *OR* B is 1 then the output is 1 (true). Otherwise the output is 0 (false). For example:

```
  1 1 1 1 0 1 0 1  <== A
| 1 0 0 1 0 0 1 1  <== B
-----
  1 1 1 1 0 1 1 1  <== output
```

In a line of code the above might read like this: `output = A | B`

1.1.6.4 XOR

The boolean *exclusive or* function has two or more inputs and the output is a single bit. The output is 1 if only an odd number of inputs are 1. Otherwise the output will be 0.

This text will use the operator used in the C language when discussing the *XOR* operator in symbolic form. Specifically the carrot: '^'.

Note that when *XOR* is used with two inputs, the output is set to 1 (true) when the inputs have different values and 0 (false) when the inputs both have the same value.

For example:

```
  1 1 1 1 0 1 0 1  <== A
^ 1 0 0 1 0 0 1 1  <== B
-----
  0 1 1 0 0 1 1 0  <== output
```

In a line of code the above might read like this: `output = A ^ B`

1.2 IEEE-754 Floating Point Number Representation

This section provides an overview of the IEEE-754 32-bit binary floating point format.

- Recall that the place values for integer binary numbers are:

... 128 64 32 16 8 4 2 1

- We can extend this to the right in binary similar to the way we do for decimal numbers:

... 128 64 32 16 8 4 2 1 . 1/2 1/4 1/8 1/16 1/32 1/64 1/128 ...

The ‘.’ in a binary number is a binary point, not a decimal point.

- We use scientific notation as in 2.7×10^{-47} to express either small fractions or large numbers when we are not concerned every last digit needed to represent the entire, exact, value of a number.
- The format of a number in scientific notation is $\text{mantissa} \times \text{base}^{\text{exponent}}$
- In binary we have $\text{mantissa} \times 2^{\text{exponent}}$
- IEEE-754 format requires binary numbers to be *normalized* to $1.\text{significand} \times 2^{\text{exponent}}$ where the *significand* is the portion of the *mantissa* that is to the right of the binary-point.
 - The unnormalized binary value of -2.625 is 10.101
 - The normalized value of -2.625 is 1.0101×2^1
- We need not store the ‘1.’ because *all* normalized floating point numbers will start that way. Thus we can save memory when storing normalized values by adding 1 to the significand.

31	30	23	22	0
1	1 0 0 0 0 0 0 0	0 1 0 1 0		
sign	exponent	significand		

- $-((1 + \frac{1}{4} + \frac{1}{16}) \times 2^{128-127}) = -((1 + \frac{1}{4} + \frac{1}{16}) \times 2^1) = -(2 + \frac{1}{2} + \frac{1}{8}) = -(2 + .5 + .125) = -2.625$
- IEEE754 formats:

	IEEE754 32-bit	IEEE754 64-bit
sign	1 bit	1 bit
exponent	8 bits (excess-127)	11 bits (excess-1023)
mantissa	23 bits	52 bits
max exponent	127	1023
min exponent	-126	-1022

- When the exponent is all ones, the mantissa is all zeros, and the sign is zero, the number represents positive infinity.
- When the exponent is all ones, the mantissa is all zeros, and the sign is one, the number represents negative infinity.
- Note that the binary representation of an IEEE754 number in memory can be compared for magnitude with another one using the same logic as for comparing two’s complement signed integers because the magnitude of an IEEE number grows upward and downward in the same fashion as signed integers. This is why we use excess notation and locate the significand’s sign bit on the left of the exponent.
- Note that zero is a special case number. Recall that a normalized number has an implied 1-bit to the left of the significand. . . which means that there is no way to represent zero! Zero is represented by an exponent of all-zeros and a significand of all-zeros. This definition allows for a positive and a negative zero if we observe that the sign can be either 1 or 0.
- On the number-line, numbers between zero and the smallest fraction in either direction are in the *underflow* areas.
- On the number line, numbers greater than the mantissa of all-ones and the largest exponent allowed are in the *overflow* areas.
- Note that numbers have a higher resolution on the number-line when the exponent is smaller.

➡ Fix Me:

Need to add the standard lecture numberline diagram showing where the over/under-flow areas are and why.

1.2.1 Floating Point Number Accuracy

Due to the finite number of bits used to store the value of a floating point number, it is not possible to represent every one of the infinite values on the real number line. The following C programs illustrate this point.

1.2.1.1 Powers Of Two

Just like the integer numbers, the powers of two that have bits to represent them can be represented perfectly... as can their sums (provided that the significand requires no more than 23 bits.)

Listing 1.1: powersoftwo.c

Precise Powers of Two

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4
5 union floatbin
6 {
7     unsigned int    i;
8     float          f;
9 };
10 int main()
11 {
12     union floatbin  x;
13     union floatbin  y;
14     int             i;
15     x.f = 1.0;
16     while (x.f > 1.0/1024.0)
17     {
18         y.f = -x.f;
19         printf("%25.10f = %08x      %25.10f = %08x\n", x.f, x.i, y.f, y.i);
20         x.f = x.f/2.0;
21     }
22 }
```

Listing 1.2: powersoftwo.out

Output from powersoftwo.c

```

1 1.0000000000 = 3f800000          -1.0000000000 = bf800000
2 0.5000000000 = 3f000000          -0.5000000000 = bf000000
3 0.2500000000 = 3e800000          -0.2500000000 = be800000
4 0.1250000000 = 3e000000          -0.1250000000 = be000000
5 0.0625000000 = 3d800000          -0.0625000000 = bd800000
6 0.0312500000 = 3d000000          -0.0312500000 = bd000000
7 0.0156250000 = 3c800000          -0.0156250000 = bc800000
8 0.0078125000 = 3c000000          -0.0078125000 = bc000000
9 0.0039062500 = 3b800000          -0.0039062500 = bb800000
10 0.0019531250 = 3b000000          -0.0019531250 = bb000000
```

1.2.1.2 Clean Decimal Numbers

When dealing with decimal values, you will find that they don't map simply into binary floating point values.

Note how the decimal numbers are not accurately represented as they get larger. The decimal number on line 10 of Listing 1.4 can be perfectly represented in IEEE format. However, a problem arises in the 11th loop iteration. It is due to the fact that the binary number can not be represented accurately in IEEE format. Its least significant bits were truncated in a best-effort attempt at rounding the value off in order to fit the value into the bits provided. This is an example of *low order truncation*. Once this happens, the value of `x.f` is no longer as precise as it could be given more bits in which to save its value.

Listing 1.3: `cleandecimal.c`

Print Clean Decimal Numbers

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4
5 union floatbin
6 {
7     unsigned int    i;
8     float          f;
9 };
10 int main()
11 {
12     union floatbin  x, y;
13     int             i;
14
15     x.f = 10;
16     while (x.f <= 10000000000000.0)
17     {
18         y.f = -x.f;
19         printf("%25.10f = %08x    %25.10f = %08x\n", x.f, x.i, y.f, y.i);
20         x.f = x.f*10.0;
21     }
22 }
```

Listing 1.4: `cleandecimal.out`Output from `cleandecimal.c`

```

1      10.0000000000 = 41200000          -10.0000000000 = c1200000
2     100.0000000000 = 42c80000          -100.0000000000 = c2c80000
3    1000.0000000000 = 447a0000          -1000.0000000000 = c47a0000
4   10000.0000000000 = 461c4000          -10000.0000000000 = c61c4000
5  100000.0000000000 = 47c35000          -100000.0000000000 = c7c35000
6 1000000.0000000000 = 49742400          -1000000.0000000000 = c9742400
7 10000000.0000000000 = 4b189680          -10000000.0000000000 = cb189680
8 100000000.0000000000 = 4cbebc20          -100000000.0000000000 = ccbebc20
9 1000000000.0000000000 = 4e6e6b28          -1000000000.0000000000 = ce6e6b28
10 10000000000.0000000000 = 501502f9          -10000000000.0000000000 = d01502f9
11 99999997952.0000000000 = 51ba43b7          -99999997952.0000000000 = d1ba43b7
12 999999995904.0000000000 = 5368d4a5          -999999995904.0000000000 = d368d4a5
13 9999999827968.0000000000 = 551184e7          -9999999827968.0000000000 = d51184e7
```

1.2.1.3 Accumulation of Error

These rounding errors can be exaggerated when the number we multiply the `x.f` value by is, itself, something that can not be accurately represented in IEEE form.²

²Applications requiring accurate decimal values, such as financial accounting systems, can use a packed-decimal numeric format to avoid unexpected oddities caused by the use of binary numbers.

► Fix Me:

In a lecture one would show that one tenth is a repeating non-terminating binary number that gets truncated. This discussion should be reproduced here in text form.

For example, if we multiply our `x.f` value by $\frac{1}{10}$ each time, we can never be accurate and we start accumulating errors immediately.

Listing 1.5: `erroraccumulation.c`

Accumulation of Error

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4
5 union floatbin
6 {
7     unsigned int    i;
8     float          f;
9 };
10 int main()
11 {
12     union floatbin  x, y;
13     int             i;
14
15     x.f = .1;
16     while (x.f <= 2.0)
17     {
18         y.f = -x.f;
19         printf("%25.10f = %08x      %25.10f = %08x\n", x.f, x.i, y.f, y.i);
20         x.f += .1;
21     }
22 }
```

Listing 1.6: `erroraccumulation.out`Output from `erroraccumulation.c`

1	0.1000000015 = 3dcccccd	-0.1000000015 = bdcccccd
2	0.2000000030 = 3e4ccccd	-0.2000000030 = be4ccccd
3	0.3000000119 = 3e99999a	-0.3000000119 = be99999a
4	0.4000000060 = 3ecccccd	-0.4000000060 = becccccd
5	0.5000000000 = 3f000000	-0.5000000000 = bf000000
6	0.6000000238 = 3f19999a	-0.6000000238 = bf19999a
7	0.7000000477 = 3f333334	-0.7000000477 = bf333334
8	0.8000000715 = 3f4ccccce	-0.8000000715 = bf4ccccce
9	0.9000000954 = 3f666668	-0.9000000954 = bf666668
10	1.0000001192 = 3f800001	-1.0000001192 = bf800001
11	1.1000001431 = 3f8ccccce	-1.1000001431 = bf8ccccce
12	1.2000001669 = 3f99999b	-1.2000001669 = bf99999b
13	1.3000001907 = 3fa66668	-1.3000001907 = bfa66668
14	1.4000002146 = 3fb33335	-1.4000002146 = bfb33335
15	1.5000002384 = 3fc00002	-1.5000002384 = bfc00002
16	1.6000002623 = 3fcccccf	-1.6000002623 = bfcccccf
17	1.7000002861 = 3fd9999c	-1.7000002861 = bfd9999c
18	1.8000003099 = 3fe66669	-1.8000003099 = bfe66669
19	1.9000003338 = 3ff33336	-1.9000003338 = bff33336

1.2.2 Reducing Error Accumulation

In order to use floating point numbers in a program without causing excessive rounding problems an algorithm can be redesigned such that the accumulation is eliminated. This example is similar to the

previous one, but this time we recalculate the desired value from a known-accurate integer value. Some rounding errors remain present, but they can not accumulate.

Listing 1.7: errorcompensation.c

Accumulation of Error

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4
5 union floatbin
6 {
7     unsigned int    i;
8     float          f;
9 };
10 int main()
11 {
12     union floatbin  x, y;
13     int             i;
14
15     i = 1;
16     while (i <= 20)
17     {
18         x.f = i/10.0;
19         y.f = -x.f;
20         printf("%25.10f = %08x      %25.10f = %08x\n", x.f, x.i, y.f, y.i);
21         i++;
22     }
23     return(0);
24 }
```

Listing 1.8: errorcompensation.out

Output from erroraccumulation.c

1 0.1000000015 = 3dcccccd	-0.1000000015 = bdcccccd
2 0.2000000030 = 3e4ccccd	-0.2000000030 = be4ccccd
3 0.3000000119 = 3e99999a	-0.3000000119 = be99999a
4 0.4000000060 = 3ecccccd	-0.4000000060 = becccccd
5 0.5000000000 = 3f000000	-0.5000000000 = bf000000
6 0.6000000238 = 3f19999a	-0.6000000238 = bf19999a
7 0.6999999881 = 3f333333	-0.6999999881 = bf333333
8 0.8000000119 = 3f4ccccd	-0.8000000119 = bf4ccccd
9 0.8999999762 = 3f666666	-0.8999999762 = bf666666
10 1.0000000000 = 3f800000	-1.0000000000 = bf800000
11 1.1000000238 = 3f8ccccd	-1.1000000238 = bf8ccccd
12 1.2000000477 = 3f99999a	-1.2000000477 = bf99999a
13 1.2999999523 = 3fa66666	-1.2999999523 = bfa66666
14 1.3999999762 = 3fb33333	-1.3999999762 = bfb33333
15 1.5000000000 = 3fc00000	-1.5000000000 = bfc00000
16 1.6000000238 = 3fcccccd	-1.6000000238 = bfcccccd
17 1.7000000477 = 3fd9999a	-1.7000000477 = bfd9999a
18 1.7999999523 = 3fe66666	-1.7999999523 = bfe66666
19 1.8999999762 = 3ff33333	-1.8999999762 = bff33333
20 2.0000000000 = 40000000	-2.0000000000 = c0000000

Chapter 2

The RISC-V GNU Toolchain

This chapter discusses the GNU toolchain elements used to experiment with the material in this book.

The instructions and examples here were all implemented on Ubuntu 16.04 LTS.

Install custom code in a location that will not cause interference with other applications and allow for easy cleanup. These instructions install the toolchain in `/usr/local/riscv`. At any time you can remove the lot and start over by executing the following command:

```
rm -rf /usr/local/riscv/*
```

Tested on Ubuntu 16.04 LTS. 18.04 was just released... update accordingly.

These are the only commands that you should perform as root when installing the toolchain:

```
sudo apt-get install autoconf automake autotools-dev curl libmpc-dev \
libmpfr-dev libgmp-dev gawk build-essential bison flex texinfo gperf \
libtool patchutils bc zlib1g-dev libexpat-dev
sudo mkdir -p /usr/local/riscv/
chmod 777 /usr/local/riscv/
```

All other commands should be executed as a regular user. This will eliminate the possibility of clobbering system files that should not be touched when tinkering with the toolchain applicaitons.

To download, compile and “install” the toolchain:

```
# riscv toolchain:
#
# https://riscv.org/software-tools/risc-v-gnu-compiler-toolchain/

git clone --recursive https://github.com/riscv/riscv-gnu-toolchain
cd riscv-gnu-toolchain
./configure --prefix=/usr/local/riscv --with-arch=rv32im --with-abi=ilp32
make
make install
```

► Fix Me:

It would be good to find some Mac and Windows users to write and test proper variations on this section to address those systems. Pull requests, welcome!

Need to discuss augmenting the PATH environment variable.

Discuss the choice of ilp32 as well as what the other variations would do.

Discuss rv32im and note that the details are found in [chapter 3](#).

Discuss installing and using one of the RISC-V simulators here.

Describe the pre-processor, compiler, assembler and linker.

Source, object, and binary files

Assembly syntax (label: mnemonic op1, op2, op3 # comment).

text, data, bss, stack

Labels and scope.

Forward & backward references to throw-away labels.

The entry address of an application.

.s file contain assembler code. .S (or .sx) files contain assembler code that must be preprocessed. [[14](#), p. 29]

Pre-processing conditional assembly using #if.

Building with -mabi=ilp32 -march=rv32i -mno-fdiv -mno-div to match the config options on the toolchain.

Linker scripts.

Makefiles

objdump

nm

hexdump -C

Chapter 3

RV32 Machine Instructions

3.1 Introduction

► Fix Me:

Discuss what the IMAFD, G and other ISA extensions mean as well as the 32, 64 and 128-bit versions.

3.2 Conventions and Terminology

When discussing instructions, the following abbreviations/notations are used:

3.2.1 XLEN

XLEN represents the bit-length of an x register in the machine architecture. Possible values are 32, 64 and 128.

3.2.2 sx(val)

Sign extend *val* to the left.

This is used to convert a signed integer value expressed using some number of bits to a larger number of bits by adding more bits to the left. In doing so, the sign will be preserved. In this case *val* represents the least **MSBs** of the value. For more on binary numbers see [chapter 1](#).

[Figure 3.1](#) illustrates extending the negative sign bit of *val* to the left by replicating it. When *val* is negative, its **MSB** (bit 19 in this example) will be set to 1. Extending this value to the left will set all the new bits to the left of it to 1 as well.

[Figure 3.2](#) illustrates extending the positive sign bit of *val* to the left by replicating it. When *val* is positive, its **MSB** will be set to 0. Extending this value to the left will set all the new bits to the left of it to 0 as well.




Figure 3.1: Sign-extending a negative integer from 20 bits to 32 bits.




Figure 3.2: Sign-extending a positive integer from 20 bits to 32 bits.

3.2.3 `zx(val)`

Zero extend *val* to the left.

This is used to convert an unsigned integer value expressed using some number of bits to a larger number of bits by adding more bits to the left. In doing so, the new bits added will all be set to zero. As is the case with `sx(val)`, *val* represents the [LSBs](#) of the final value. [Figure 3.3](#) illustrates zero-extending a 20-bit *val* to the left to form a 32-bit fullword.

For more on binary numbers see [chapter 1](#).




Figure 3.3: Zero-extending an unsigned integer from 20 bits to 32 bits.

3.2.4 `zr(val)`

Zero extend *val* to the right.

Some times a binary value is encoded such that a set of bits represented by *val* are used to represent the [MSBs](#) of some longer (more bits) value. In this case it is necessary to append zeroes to the right to convert *val* to the longer value.

[Figure 3.4](#) illustrates converting a 20-bit *val* to a 32-bit fullword.

3.2.5 Sign Entended Left and Zero Extend Right

Some instructions such as the J-type (see [subsection 3.4.2](#)) include immediate operands that are extended in both directions.

[Figure 3.5](#) and [Figure 3.6](#) illustrates zero-extending a 20-bit negative number one bit to the right and




Figure 3.4: Zero-extending an integer to the right from 20 bits to 32 bits.

sign-extending it 11 bits to the left:




Figure 3.5: Sign-extending a positive 20-bit number 11 bits to the left and one bit to the right.




Figure 3.6: Sign-extending a negative 20-bit number 11 bits to the left and one bit to the right.

3.2.6 m8(addr)

The contents of an 8-bit value in memory at address *addr*.

Given the contents of the memory dump shown in [Figure 3.7](#), `m8(42)` refers to the memory location at address 42_{16} that currently contains the 8-bit value `fc16`.

The `mn(addr)` notation can be used to refer to memory that is being read or written depending on the context.

When memory is being written, the following notation is used to indicate that the least significant 8 bits of *source* will be written into memory at the address *addr*:

`m8(addr) ← source`

When memory is being read, the following notation is used to indicate that the 8 bit value at the address *addr* will be read and stored into *dest*:

`dest ← m8(addr)`

Note that *source* and *dest* are typically registers.

3.2.7 m16(addr)

The contents of an 16-bit little-endian value in memory at address *addr*.

```
00000030  2f 20 72 65 61 64 20 61 20 62 69 6e 61 72 79 20
00000040  66 69 fc 65 20 66 69 6c 6c 65 64 20 77 69 74 68
00000050  20 72 76 33 32 49 20 69 6e 73 74 72 75 63 74 69
00000060  6f 6e 73 20 61 6e 64 20 66 65 65 64 20 74 68 65
```

Figure 3.7: Sample memory contents.

Given the contents of the memory dump shown in [Figure 3.7](#), `m16(42)` refers to the memory location at address `4216` that currently contains `65fc16`. See also [subsection 3.2.6](#).

3.2.8 `m32(addr)`

The contents of an 32-bit little-endian value in memory at address *addr*.

Given the contents of the memory dump shown in [Figure 3.7](#), `m32(42)` refers to the memory location at address `4216` that currently contains `662065fc16`. See also [subsection 3.2.6](#).

3.2.9 `m64(addr)`

The contents of an 64-bit little-endian value in memory at address *addr*.

Given the contents of the memory dump shown in [Figure 3.7](#), `m64(42)` refers to the memory location at address `4216` that currently contains `656c6c69662065fc16`. See also [subsection 3.2.6](#).

3.2.10 `m128(addr)`

The contents of an 128-bit little-endian value in memory at address *addr*.

Given the contents of the memory dump shown in [Figure 3.7](#), `m128(42)` refers to the memory location at address `4216` that currently contains `7220687469772064656c6c69662065fc16`. See also [subsection 3.2.6](#).

3.2.11 `+.offset`

The address of the current instruction plus a numeric offset.

3.2.12 `-.offset`

The address of the current instruction minus a numeric offset.

3.2.13 `pc`

The current value of the program counter.

3.2.14 rd

An x-register used to store the result of instruction.

3.2.15 rs1

An x-register value used as a source operand for an instruction.

3.2.16 rs2

An x-register value used as a source operand for an instruction.

3.2.17 imm

An immediate numeric operand. The word *immediate* refers to the fact that the operand is stored within an instruction.

3.2.18 rsN[h:l]

The value of bits from h through l of x-register rsN. For example: rs1[15:0] refers to the contents of the 16 [LSBs](#) of rs1.

3.3 Addressing Modes

immediate, register, base-displacement, pc-relative

► Fix Me:

Write this section.

3.4 Instruction Encoding Formats

This document concerns itself with the following RISC-V instruction formats.

XXX Show and discuss a stack of formats explaining how the unnatural ordering of the *imm* fields reduces the number of possible locations that the hardware has to be prepared to *look* for various bits. For example, the opcode, rd, rs1, rs1, func3 and the sign bit (when used) are all always in the same position. Also note that imm[19:12] and imm[10:5] can only be found in one place. imm[4:0] can only be found in one of two places...

► Fix Me:

Should discuss types and sizes beyond the fundamentals. Will add if/when instruction details are added in the future.

The point to all this is that it is easier to build a machine if it does not have to accommodate many different ways to perform the same task. This simplification can also allow it operate faster.

[Figure 3.8](#) Shows the RISC-V instruction formats.

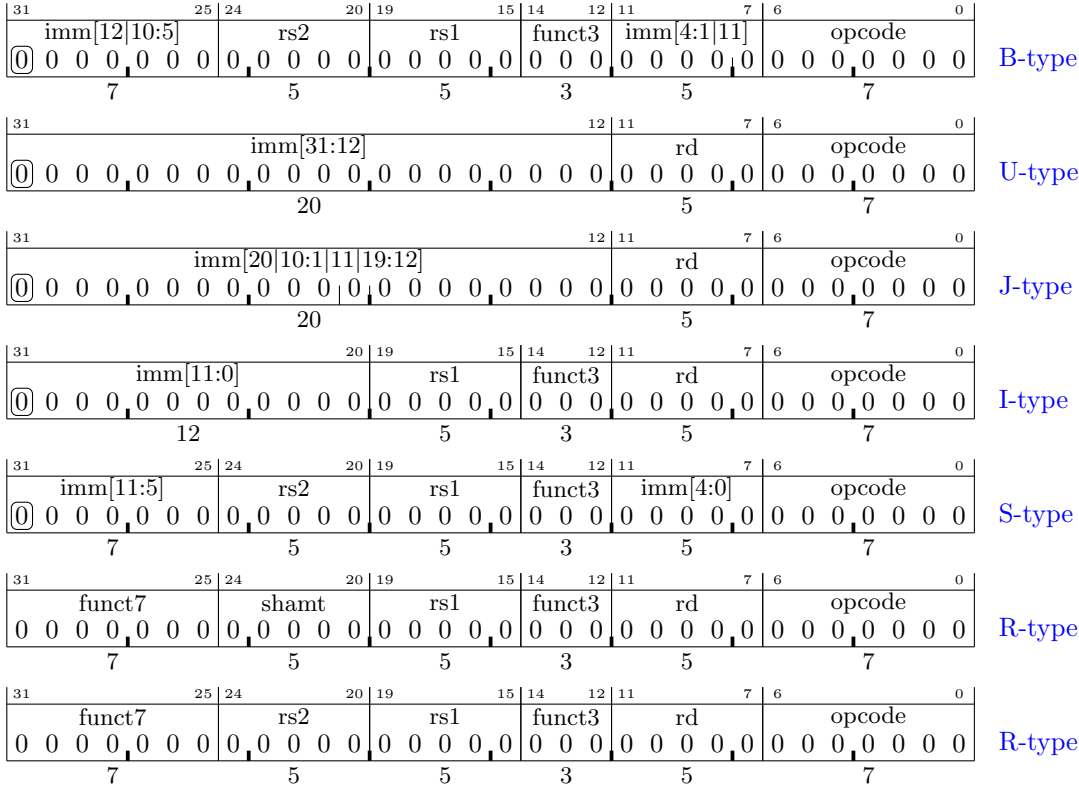



Figure 3.8: RISC-V instruction formats.

3.4.1 U Type


The U-Type format is used for instructions that use a 20-bit immediate operand and a destination register.



The **rd** field contains an **x** register number to be set to a value that depends on the instruction.

The **imm** field contains a 20-bit value that will be converted into **XLEN** bits by using the *imm* operand for bits 31:12 and then sign-extending it to the left¹ and zero-extending the LSBs as discussed in subsection 3.2.4.

If **XLEN**=32 then the **imm** value in this example will be converted as shown below.



Notice that the 20-bits of the **imm** field are mapped in the same order and in the same relative position

¹When **XLEN** is larger than 32.

If **XLEN**=64 then the imm value in this example will be converted to the same two's complement integer value by extending the sign to the left.

The J-type format is used for instructions that use a 20-bit immediate operand and a destination register. It is similar to the U-type. However, the immediate operand is constructed by arranging the *imm* bits in a different manner.

3.4.7 CPU Registers

The registers are names x0 through x31 and have aliases suited to their conventional use. The following table describes each register.

Note that the calling convention specifies that only some of the registers are to be saved by functions if they alter their contents. The idea being that accessing memory is time-consuming and that by classifying some registers as “temporary” (not saved by any function that alter its contents) it is possible to carefully implement a function with less need to store register values on the stack in order to use them to perform the operations of the function.

➡ Fix Me:
Need to add a section that discusses the calling conventions

The lack of grouping the temporary and saved registers is due to the fact that the C extension provides access to only the first 16 registers when executing instructions in the compressed format.

Reg	Alias	Description	Saved
x0	zero	Hard-wired zero	yes
x1	ra	Return address	
x2	sp	Stack pointer	
x3	gp	Global pointer	
x4	tp	Thread pointer	
x5	t0	Temporary/alternate link register	
x6	t1	Temporary	
x7	t2	Temporary	yes
x8	s0/fp	Saved register/frame pointer	
x9	s1	Saved register	
x10	a0	Function argument/return value	
x11	a1	Function argument/return value	
x12	a2	Function argument	
x13	a3	Function argument	
x14	a4	Function argument	
x15	a5	Function argument	
x16	a6	Function argument	
x17	a7	Function argument	
x18	s2	Saved register	yes
x19	s3	Saved register	yes
x20	s4	Saved register	yes
x21	s5	Saved register	yes
x22	s6	Saved register	yes
x23	s7	Saved register	yes
x24	s8	Saved register	yes
x25	s9	Saved register	yes
x26	s10	Saved register	yes
x27	s11	Saved register	yes
x28	t3	Temporary	
x29	t4	Temporary	
x30	t5	Temporary	
x31	t6	Temporary	

3.5 memory

Note that RISC-V is a little-endian machine.

RV32I
Instruction!LUI

If a RISC-V processor implements the C (compressed) extension then instructions may be aligned to 2-byte boundaries.[1, p. 68]

Data alignment is not necessary but unaligned data can be inefficient. Accessing unaligned data using any of the load or store instructions can also prevent a memory access from operating atomically. [1, p.19] See also ??.

3.6 RV32I Base Instruction Set

RV32I refers to the basic 32-bit integer instructions.

3.6.1 LUI rd, imm


Load Upper Immediate.

$$\text{rd} \leftarrow \text{zr}(\text{imm})$$

Copy the immediate value into bits 31:12 of the destination register and place zeroes into bits 11:0. When XLEN is 64 or 128, the immediate value is sign-extended to the left.


Instruction Format and Example:

LUI t0, 3



```
00010074: 000032b7 lui      x5, 0x3          // x5 = 0x3000
reg  0: 00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 00003000 f0f0f0f0 f0f0f0f0
reg  8: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
reg 16: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
reg 24: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
pc: 00010078
```

```
LUI t0, 0xffff
```



```
00010078: ffff2b7 lui      x5, 0xfffff // x5 = 0xfffff000
reg  0: 00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 ffff000 f0f0f0f0 f0f0f0f0
reg  8: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
reg 16: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
reg 24: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
      pc: 0001007c
```


3.6.2 AUIPC rd, imm

Instruction!AUIPC
Instruction!JAL


Add Upper Immediate to PC.

$$\text{rd} \leftarrow \text{pc} + \text{zr}(\text{imm})$$

Create a signed 32-bit value by zero-extending `imm[31:12]` to the right (see [subsection 3.2.4](#)) and add this value to the `pc` register, placing the result into `rd`.


When XLEN is 64 or 128, the immediate value is also sign-extended to the left prior to being added to the `pc` register.

AUIPC t0, 3



```
0001007c: 00003297  auipc      x5, 0x3           // x5 = 0x1307c = 0x1007c + 0x3000
reg  0: 00000000  f0f0f0f0  f0f0f0f0  f0f0f0f0-f0f0f0f0  0001307c  f0f0f0f0  f0f0f0f0
reg  8: f0f0f0f0  f0f0f0f0  f0f0f0f0  f0f0f0f0-f0f0f0f0  f0f0f0f0  f0f0f0f0  f0f0f0f0
reg 16: f0f0f0f0  f0f0f0f0  f0f0f0f0  f0f0f0f0-f0f0f0f0  f0f0f0f0  f0f0f0f0  f0f0f0f0
reg 24: f0f0f0f0  f0f0f0f0  f0f0f0f0  f0f0f0f0-f0f0f0f0  f0f0f0f0  f0f0f0f0  f0f0f0f0
pc: 00010080
```

AUIPC t0, 0x81000



```
00010080: 81000297  auipc    x5, 0x81000    // x5 = 0x81010080 = 0x10080 + 0x81000000
reg  0: 00000000  f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 81010080 f0f0f0f0 f0f0f0f0
reg  8: f0f0f0f0  f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
reg 16: f0f0f0f0  f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
reg 24: f0f0f0f0  f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
pc: 00010084
```

The AUIPC instruction supports two-instruction sequences to access arbitrary offsets from the PC for both control-flow transfers and data accesses. The combination of an AUIPC and the 12-bit immediate in a JALR can transfer control to any 32-bit PC-relative address, while an AUIPC plus the 12-bit immediate offset in regular load or store instructions can access any 32-bit PC-relative data address. [1, p. 14]

3.6.3 JAL rd, imm

Jump and link.


```
rd ← pc + 4
pc ← pc + sx(imm<<1)
```

This instruction saves the address of the next instruction that would otherwise execute (located at `pc+4`) into `rd` and then adds immediate value to the `pc` causing an unconditional branch to take place.

The standard software conventions for calling subroutines use `x1` as the return address (`rd` register in this case). [1, p. 16]

Encoding:

JAL `x7`, `+.16`



imm demultiplexed value = $0000000000000001000_2 \ll 1 = 16_{10}$

State of registers before execution:


`pc` = 0x11114444

State of registers after execution:

`pc` = 0x11114454 `x7` = 0x11114448

JAL provides a method to call a subroutine using a pc-relative address.

JAL `x7`, `-.16`



imm demultiplexed value = $1111111111111111000_2 \ll 1 = -16_{10}$

State of registers before execution:

`pc` = 0x11114444

State of registers after execution:

`pc` = 0x11114434 `x7` = 0x11114448

3.6.4 JALR `rd`, `rs1`, `imm`

Jump and link register.

$rd \leftarrow pc + 4$
 $pc \leftarrow (rs1 + sx(imm)) \& \sim 1$

This instruction saves the address of the next instruction that would otherwise execute (located at `pc+4`) into `rd` and then adds the immediate value to the `rs1` register and stores the sum into the `pc` register causing an unconditional branch to take place.


Note that the branch target address is calculated by sign-extending the `imm[11:0]` bits from the instruc-

tion, adding it to the `rs1` register and *then* the LSB of the sum is to zero and the result is stored into the `pc` register. The discarding of the LSB allows the branch to refer to any even address. Instruction!BEQ

The standard software conventions for calling subroutines use `x1` as the return address (`rd` register in this case). [1, p. 16]

Encoding:

JALR `x1, x7, 4`



Before:


`pc = 0x11114444`
`x7 = 0x44444444`

After

`pc = 0x5555888c`
`x1 = 0x11114448`

JALR provides a method to call a subroutine using a base-displacement address.

JALR `x1, x0, 5`



Note that the least significant bit in the result of `rs1+imm` is discarded/set to zero before the result is saved in the `pc`.

`pc = 0x11114444`

After

`pc = 0x00000004`
`x1 = 0x11114448`


3.6.5 BEQ `rs1, rs2, imm`

Branch if equal.

`pc ← (rs1 == rs2) ? pc+sx(imm[12:1]<<1) : pc+4`

Encoding:

BEQ `x3, x15, 2064`



Instruction!BNE
Instruction!BLT

$\text{imm}[12:1] = 010000001000_2 = 1032_{10}$

$\text{imm} = 2064_{10}$

$\text{funct3} = 000_2$

$\text{rs1} = \text{x3}$

$\text{rs2} = \text{x15}$


3.6.6 BNE rs1, rs2, imm

Branch if Not Equal.

$\text{pc} \leftarrow (\text{rs1} \neq \text{rs2}) ? \text{pc} + \text{sx}(\text{imm}[12:1] \ll 1) : \text{pc} + 4$

Encoding:

BNE x3, x15, 2064



$\text{imm}[12:1] = 010000001000_2 = 1032_{10}$

$\text{imm} = 2064_{10}$

$\text{funct3} = 001_2$

$\text{rs1} = \text{x3}$

$\text{rs2} = \text{x15}$


3.6.7 BLT rs1, rs2, imm

Branch if Less Than.

$\text{pc} \leftarrow (\text{rs1} < \text{rs2}) ? \text{pc} + \text{sx}(\text{imm}[12:1] \ll 1) : \text{pc} + 4$

Encoding:

BLT x3, x15, 2064



$\text{imm}[12:1] = 010000001000_2 = 1032_{10}$

$\text{imm} = 2064_{10}$

$\text{funct3} = 100_2$

$\text{rs1} = \text{x3}$

$\text{rs2} = \text{x15}$

3.6.8 BGE rs1, rs2, imm

Instruction!BGE
Instruction!BLTU
Instruction!BGEU

Branch if Greater or Equal.

$pc \leftarrow (rs1 \geq rs2) ? pc + sx(imm[12:1] \ll 1) : pc + 4$

Encoding:

BGE x3, x15, 2064

imm[12 10:5]					rs2			rs1			funct3		imm[4:1 11]		opcode						
0	0	0	0	0	0	0	1	1	1	1	0	0	0	1	1	1	0	0	0	1	1
7					5			5			3		5		7						

B-type

$imm[12:1] = 010000001000_2 = 1032_{10}$

$imm = 2064_{10}$

$funct3 = 101_2$

$rs1 = x3$

$rs2 = x15$

3.6.9 BLTU rs1, rs2, imm

Branch if Less Than Unsigned.

$pc \leftarrow (rs1 < rs2) ? pc + sx(imm[12:1] \ll 1) : pc + 4$

Encoding:

BLTU x3, x15, 2064

312524201915141211760																				
imm[12:10:5]							rs2			rs1			funct3		imm[4:1:11]		opcode			
0000000							01111			00011			1110		100001		11100011			
7							5			5			3		5		7			

B-type

$imm[12:1] = 010000001000_2 = 1032_{10}$

$imm = 2064_{10}$

$funct3 = 110_2$

$rs1 = x3$

$rs2 = x15$


3.6.10 BGEU rs1, rs2, imm

Branch if Greater or Equal Unsigned.

$pc \leftarrow (rs1 \geq rs2) ? pc + sx(imm[12:1] \ll 1) : pc + 4$

Encoding:

BGEU x3, x15, 2064



Instruction!LB
Instruction!LH
Instruction!LW

► Fix Me:

use symbols in branch examples

$\text{imm}[12:1] = 010000001000_2 = 1032_{10}$

$\text{imm} = 2064_{10}$

$\text{funct3} = 111_2$

$\text{rs1} = \text{x3}$

$\text{rs2} = \text{x15}$

3.6.11 LB rd, imm(rs1)

Load byte.


$\text{rd} \leftarrow \text{sx}(\text{m8}(\text{rs1} + \text{sx}(\text{imm})))$

$\text{pc} \leftarrow \text{pc} + 4$

Load an 8-bit value from memory at address $\text{rs1} + \text{imm}$, then sign-extend it to 32 bits before storing it in rd

Encoding:

LB x7, 4(x3)



3.6.12 LH rd, imm(rs1)

Load halfword.


$\text{rd} \leftarrow \text{sx}(\text{m16}(\text{rs1} + \text{sx}(\text{imm})))$

$\text{pc} \leftarrow \text{pc} + 4$

Load a 16-bit value from memory at address $\text{rs1} + \text{imm}$, then sign-extend it to 32 bits before storing it in rd

Encoding:

LH x7, 4(x3)



3.6.13 LW rd, imm(rs1)

Load word.


$rd \leftarrow sx(m32(rs1+sx(imm)))$
 $pc \leftarrow pc+4$

Instruction!LBU
 Instruction!LHU

Load a 32-bit value from memory at address $rs1+imm$, then store it in rd

Encoding:

LW x7, 4(x3)



3.6.14 LBU rd, imm(rs1)


Load byte unsigned.

$rd \leftarrow zx(m8(rs1+sx(imm)))$
 $pc \leftarrow pc+4$

Load an 8-bit value from memory at address $rs1+imm$, then zero-extend it to 32 bits before storing it in rd

Encoding:

LBU x7, 4(x3)



3.6.15 LHU rd, imm(rs1)


Load halfword unsigned.

$rd \leftarrow zx(m16(rs1+sx(imm)))$
 $pc \leftarrow pc+4$

Load an 16-bit value from memory at address $rs1+imm$, then zero-extend it to 32 bits before storing it in rd

Encoding:

LHU x7, 4(x3)



3.6.16 SB rs2, imm(rs1)

Instruction!SB
 Instruction!SH
 Instruction!SW


Store Byte.

```
m8(rs1+sx(imm)) ← rs2[7:0]
pc ← pc+4
```

Store the 8-bit value in **rs2**[7:0] into memory at address **rs1+imm**.

Encoding:

SB x3, 19(x15)

**3.6.17 SH rs2, imm(rs1)**


Store Halfword.

```
m16(rs1+sx(imm)) ← rs2[15:0]
pc ← pc+4
```

Store the 16-bit value in **rs2**[15:0] into memory at address **rs1+imm**.

Encoding:

SH x3, 19(x15)

**3.6.18 SW rs2, imm(rs1)**


Store Word

```
m16(rs1+sx(imm)) ← rs2[31:0]
pc ← pc+4
```

Store the 32-bit value in **rs1** into memory at address **rs2+imm**.

Encoding:

SW x3, 19(x15)



Show pos & neg imm examples.

Instruction!ADDI
Instruction!SLTI

3.6.19 ADDI rd, rs1, imm


Add Immediate

$rd \leftarrow rs1 + sx(imm)$

$pc \leftarrow pc + 4$

Encoding:

ADDI x1, x7, 4



Before:

$x7 = 0x11111111$

After:

$x1 = 0x11111115$

3.6.20 SLTI rd, rs1, imm

Set LessThan Immediate


$rd \leftarrow (rs1 < sx(imm)) ? 1 : 0$

$pc \leftarrow pc + 4$

If the sign-extended immediate value is less than the value in the **rs1** register then the value 1 is stored in the **rd** register. Otherwise the value 0 is stored in the **rd** register.

Encoding:

SLTI x1, x7, 4



Before:

$x7 = 0x11111111$

After:

$x1 = 0x00000000$

3.6.21 SLTIU rd, rs1, immInstruction!SLTIU
Instruction!XORI

Set LessThan Immediate Unsigned


$$rd \leftarrow (rs1 < sx(imm)) ? 1 : 0$$

$$pc \leftarrow pc+4$$

If the sign-extended immediate value is less than the value in the **rs1** register then the value 1 is stored in the **rd** register. Otherwise the value 0 is stored in the **rd** register. Both the immediate and **rs1** register values are treated as unsigned numbers for the purposes of the comparison.³

Encoding:

SLTIU x1, x7, 4



Before:

x7 = 0x81111111

After:

x1 = 0x00000001

3.6.22 XORI rd, rs1, imm

Exclusive Or Immediate


$$rd \leftarrow rs1 \wedge sx(imm)$$

$$pc \leftarrow pc+4$$

The logical XOR of the sign-extended immediate value and the value in the **rs1** register is stored in the **rd** register.

Encoding:

XORI x1, x7, 4



Before:

x7 = 0x81111111

After:

x1 = 0x81111115

³The immediate value is first sign-extended to XLEN bits then treated as an unsigned number. [1, p. 14]

3.6.23 ORI rd, rs1, immInstruction!ORI
Instruction!ANDI

Or Immediate

$$\text{rd} \leftarrow \text{rs1} \mid \text{sx}(\text{imm})$$

$$\text{pc} \leftarrow \text{pc} + 4$$

The logical OR of the sign-extended immediate value and the value in the **rs1** register is stored in the **rd** register.

Encoding:

ORI x1, x7, 4

31												20		19		15			14		12		11		7		6		0		
imm[11:0]												rs1					funct3			rd				opcode							
0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	0	0	0	0	0	1	0	0	1	0	0	1	1
12												5					3			5				7							

I-type

Before:

x7 = 0x81111111

After:

x1 = 0x81111115

3.6.24 ANDI rd, rs1, imm

And Immediate

$$\text{rd} \leftarrow \text{rs1} \& \text{sx}(\text{imm})$$

$$\text{pc} \leftarrow \text{pc} + 4$$

The logical AND of the sign-extended immediate value and the value in the **rs1** register is stored in the **rd** register.

Encoding:

ANDI x1, x7, 4

31												20				19		15			14		12		11		7			6		0			
imm[11:0]												rs1					funct3			rd					opcode										
0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	0	0	0	0	1	0	0	1	0	0	1	1			
12												5					3			5					7										

I-type

Before:

x7 = 0x81111111

After:

x1 = 0x81111115

3.6.25 SLLI rd, rs1, shamt

Instruction!SLLI
 Instruction!SRLI
 Instruction!SRAI

Shift Left Logical Immediate

$rd \leftarrow rs1 \ll shamt$
 $pc \leftarrow pc+4$

SLLI is a logical left shift operation (zeros are shifted into the lower bits). The value in rs1 shifted left shamt number of bits and the result placed into rd. [1, p. 14]

Encoding:

SLLI x7, x3, 2

312524201915141211760																					
funct7							shamt			rs1			funct3		rd			opcode			
0000000							00010			00011			001		00111			0100011			
7							5			5			3		5			7			

R-type

R-type

x3 = 0x81111111

After:

x7 = 0x04444444

3.6.26 SRLI rd, rs1, shamt

Shift Right Logical Immediate

$rd \leftarrow rs1 \gg shamt$
 $pc \leftarrow pc+4$

SRLI is a logical right shift operation (zeros are shifted into the higher bits). The value in rs1 shifted right shamt number of bits and the result placed into rd. [1, p. 14]

Encoding:

SRLI x7, x3, 2

312524201915141211760																						
funct7							shamt			rs1			funct3		rd			opcode				
0000000							00010			00011			101		00111			0010011				
7							5			5			3		5			7				R-type

R-type

x3 = 0x81111111

After:

x7 = 0x20444444

3.6.27 SRAI rd, rs1, shamt

Shift Right Arithmetic Immediate

$rd \leftarrow rs1 \gg shamt$
 $pc \leftarrow pc+4$

Instruction!ADD
 Instruction!SUB

SRAI is a logical right shift operation (zeros are shifted into the higher bits). The value in rs1 shifted right shamt number of bits and the result placed into rd. [1, p. 14]

Encoding:

SRAI x7, x3, 2

312524201915141211760							R-type
funct7shamtrs1funct3rdopcode							
010000000000000111001000111000100011							
755357							

R-type

x3 = 0x81111111

After:

x7 = 0xe0444444

3.6.28 ADD rd, rs1, rs2

Add

$rd \leftarrow rs1 + rs2$
 $pc \leftarrow pc+4$

ADD performs addition. Overflows are ignored and the low 32 bits of the result are written to rd. [1, p. 15]

Encoding:

ADD x7, x3, x31

312524201915141211760																							
funct7							rs2	rs1	funct3	rd	opcode												
0	0	0	0	0	0	0	1	1	1	1	0	0	0	1	1	0	1	1	0	0	1	1	
7							5		5		3	5			7					R-type			

R-type

x3 = 0x81111111 x31 = 0x22222222

After:

x7 = 0xa3333333

3.6.29 SUB rd, rs1, rs2

Subtract

$rd \leftarrow rs1 - rs2$
 $pc \leftarrow pc+4$


SUB performs subtraction. Underflows are ignored and the low 32 bits of the result are written to rd. [1,

p. 15]

Instruction!SLL
Instruction!SLT

Encoding:

SUB x7, x3, x31



x3 = 0x83333333 x31 = 0x01111111

After:

x7 = 0x82222222

3.6.30 SLL rd, rs1, rs2


Shift Left Logical

 $rd \leftarrow rs1 \ll rs2$ $pc \leftarrow pc+4$

SLL performs a logical left shift on the value in register rs1 by the shift amount held in the lower 5 bits of register rs2. [1, p. 15]

Encoding:

SLL x7, x3, x31



x3 = 0x83333333

x31 = 0x00000002

After:

x7 = 0x0ccccccc

3.6.31 SLT rd, rs1, rs2


Set Less Than

 $rd \leftarrow (rs1 < rs2) ? 1 : 0$ $pc \leftarrow pc+4$

SLT performs a signed compare, writing 1 to rd if rs1 < rs2, 0 otherwise. [1, p. 15]

Encoding:

SLT x7, x3, x31

Instruction!SLTU
Instruction!XOR

x3 = 0x83333333

x31 = 0x00000002

After:

x7 = 0x00000001

3.6.32 SLTU rd, rs1, rs2


Set Less Than Unsigned

 $rd \leftarrow (rs1 < rs2) ? 1 : 0$ $pc \leftarrow pc+4$

SLTU performs an unsigned compare, writing 1 to rd if rs1 < rs2, 0 otherwise. Note, SLTU rd, x0, rs2 sets rd to 1 if rs2 is not equal to zero, otherwise sets rd to zero (assembler pseudo-op SNEZ rd, rs). [1, p. 15]

Encoding:

SLTU x7, x3, x31



x3 = 0x83333333

x31 = 0x00000002

After:

x7 = 0x00000000

3.6.33 XOR rd, rs1, rs2

Exclusive Or

 $rd \leftarrow rs1 \oplus rs2$ $pc \leftarrow pc+4$

XOR performs a bit-wise exclusive or on rs1 and rs2. The result is stored on rd.

Encoding:

XOR x7, x3, x31

312524201915141211760																									
funct7							rs2				rs1				funct3			rd				opcode			
0000000							11111				00011				1000			00111				0110011			
7							5				5				3			5				7			

R-type

Instruction!SRL
Instruction!SRA

x3 = 0x83333333

x31 = 0x1888ffff

After:

x7 = 0x9bbbcccc

3.6.34 SRL rd, rs1, rs2

Shift Right Logical

 $rd \leftarrow rs1 \gg rs2$ $pc \leftarrow pc+4$

SRL performs a logical right shift on the value in register rs1 by the shift amount held in the lower 5 bits of register rs2. [1, p. 15]

Encoding:

SRL x7, x3, x31

31	25	24	20	19	15	14	12	11	7	6	0													
funct7							rs2			rs1			funct3		rd			opcode						
0	0	0	0	0	0	0	1	1	1	1	0	0	0	1	1	1	0	0	1	1	0	0	1	1
7							5			5			3		5			7						

R-type

x3 = 0x83333333

x31 = 0x00000010

After:

x7 = 0x00008333

3.6.35 SRA rd, rs1, rs2

Shift Right Arithmetic

 $rd \leftarrow rs1 \gg rs2$ $pc \leftarrow pc+4$

SRA performs an arithmetic right shift (the original sign bit is copied into the vacated upper bits) on the value in register rs1 by the shift amount held in the lower 5 bits of register rs2. [1, p. 14, 15]

Encoding:

SRA x7, x3, x31

Instruction!OR
Instruction!AND

```
x7 = 0xffff8333
```

Or

OR is a logical operation that performs a bit-wise OR on register rs1 and rs2 and then places the result in rd. [1, p. 14]

OR x7, x3, x31


R-type

x7 = 0x83333773

And

AND is a logical operation that performs a bit-wise AND on register rs1 and rs2 and then places the result in rd. [1, p. 14]

AND x7, x3, x31



Instruction!FENCE
Instruction!FENCE.I

x3 = 0x83333333
x31 = 0x00000fe2

After:

x7 = 0x00000322

3.6.38 FENCE predecessor, successor

The FENCE instruction is used to order device I/O and memory accesses as viewed by other RISC-V harts and external devices or co-processors. Any combination of device input (I), device output (O), memory reads (R), and memory writes (W) may be ordered with respect to any combination of the same. Informally, no other RISC-V hart or external device can observe any operation in the successor set following a FENCE before any operation in the predecessor set preceding the FENCE. The execution environment will define what I/O operations are possible, and in particular, which load and store instructions might be treated and ordered as device input and device output operations respectively rather than memory reads and writes. For example, memory-mapped I/O devices will typically be accessed with uncached loads and stores that are ordered using the I and O bits rather than the R and W bits. Instruction-set extensions might also describe new coprocessor I/O instructions that will also be ordered using the I and O bits in a FENCE. [1, p. 21]

Operation:

$pc \leftarrow pc+4$

Encoding:

0 0 0 1 1 1 1 0 0 0 1 1 1 1 1 1 1

3.6.39 FENCE.I

The FENCE.I instruction is used to synchronize the instruction and data streams. RISC-V does not guarantee that stores to instruction memory will be made visible to instruction fetches on the same RISC-V hart until a FENCE.I instruction is executed. A FENCE.I instruction only ensures that a subsequent instruction fetch on a RISC-V hart will see any previous data stores already visible to the same RISC-V hart. FENCE.I does not ensure that other RISC-V harts' instruction fetches will observe the local hart's stores in a multiprocessor system. To make a store to instruction memory visible to all RISC-V harts, the writing hart has to execute a data FENCE before requesting that all remote RISC-V harts execute a FENCE.I. [1, p. 21]

Operation:

$pc \leftarrow pc+4$

Encoding:

0 0 0 1 1 1 1 0 0 1 0 0 0 0 0 0 0

3.6.40 ECALL

Instruction!ECALL
Instruction!EBREAK
Instruction!CSRW
Instruction!CSRWS
Instruction!CSRRC

The ECALL instruction is used to make a request to the supporting execution environment, which is usually an operating system. The ABI for the system will define how parameters for the environment request are passed, but usually these will be in defined locations in the integer register file. [1, p. 24]

```
1 1 1 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
```

3.6.41 EBREAK

The EBREAK instruction is used by debuggers to cause control to be transferred back to a debugging environment. [1, p. 24]

```
1 1 1 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1
```

3.6.42 CSRRW rd, csr, rs1

The CSRRW (Atomic Read/Write CSR) instruction atomically swaps values in the CSRs and integer registers. CSRRW reads the old value of the CSR, zero-extends the value to XLEN bits, then writes it to integer register rd. The initial value in rs1 is written to the CSR. If rd=x0, then the instruction shall not read the CSR and shall not cause any of the side-effects that might occur on a CSR read. [1, p. 22]

```
1 1 1 0 0 1 1 0 0 0 1 1 0 0 1 0 1 1 1 0 0 0 0 0 0 0 0 0 0 0 1 0
```

3.6.43 CSRWS rd, csr, rs1

The CSRWS (Atomic Read and Set Bits in CSR) instruction reads the value of the CSR, zero-extends the value to XLEN bits, and writes it to integer register rd. The initial value in integer register rs1 is treated as a bit mask that specifies bit positions to be set in the CSR. Any bit that is high in rs1 will cause the corresponding bit to be set in the CSR, if that CSR bit is writable. Other bits in the CSR are unaffected (though CSRs might have side effects when written). [1, p. 22]

If rs1=x0, then the instruction will not write to the CSR at all, and so shall not cause any of the side effects that might otherwise occur on a CSR write, such as raising illegal instruction exceptions on accesses to read-only CSRs. Note that if rs1 specifies a register holding a zero value other than x0, the instruction will still attempt to write the unmodified value back to the CSR and will cause any attendant side effects. [1, p. 22]

```
1 1 1 0 0 1 1 0 0 0 1 1 0 1 0 0 1 1 1 0 0 0 0 0 0 0 0 0 0 0 1 0
```

3.6.44 CSRRC rd, csr, rs1

The CSRRC (Atomic Read and Clear Bits in CSR) instruction reads the value of the CSR, zero-extends the value to XLEN bits, and writes it to integer register rd. The initial value in integer register rs1 is treated as a bit mask that specifies bit positions to be cleared in the CSR. Any bit that is high in rs1 will cause the corresponding bit to be cleared in the CSR, if that CSR bit is writable. Other bits in the CSR are unaffected. [1, p. 22]

If $rs1=x0$, then the instruction will not write to the CSR at all, and so shall not cause any of the side effects that might otherwise occur on a CSR write, such as raising illegal instruction exceptions on accesses to read-only CSRs. Note that if $rs1$ specifies a register holding a zero value other than $x0$, the instruction will still attempt to write the unmodified value back to the CSR and will cause any attendant side effects. [1, p. 22]

Instruction!CSRRIWI
Instruction!CSRRIWI
Instruction!CSRRIWI
RV32M
RV32A

```
1 1 1 0 0 1 1 0 0 0 1 1 0 1 1 0 1 1 1 0 0 0 0 0 0 0 0 0 0 1 0
```

3.6.45 CSRRWI rd, csr, imm

This instruction is the same as CSRRW except a 5-bit unsigned (zero-extended) immediate value is used rather than the value from a register.

```
1 1 1 0 0 1 1 0 0 0 1 1 1 0 1 0 0 1 1 1 0 0 0 0 0 0 0 0 0 0 1 0
```

3.6.46 CSRRSI $rd, csr, rs1$

This instruction is the same as CSRRS except a 5-bit unsigned (zero-extended) immediate value is used rather than the value from a register.

If the $uimm[4:0]$ field is zero, then this instruction will not write to the CSR, and shall not cause any of the side effects that might otherwise occur on a CSR write. For CSRRWI, if $rd=x0$, then the instruction shall not read the CSR and shall not cause any of the side-effects that might occur on a CSR read. [1, p. 22]

3.6.47 CSRRCI $rd, csr, rs1$

This instruction is the same as CSRRC except a 5-bit unsigned (zero-extended) immediate value is used rather than the value from a register.

If the $uimm[4:0]$ field is zero, then this instruction will not write to the CSR, and shall not cause any of the side effects that might otherwise occur on a CSR write. For CSRRWI, if $rd=x0$, then the instruction shall not read the CSR and shall not cause any of the side-effects that might occur on a CSR read. [1, p. 22]

3.7 RV32M Standard Extension

32-bit integer multiply and divide instructions.

3.8 RV32A Standard Extension

32-bit atomic operations.

3.9 RV32F Standard Extension

RV32F
RV32D

32-bit IEEE floating point instructions.

3.10 RV32D Standard Extension

64-bit IEEE floating point instructions.

Appendix A

Attribution 4.0 International

Creative Commons Corporation ("Creative Commons") is not a law firm and does not provide legal services or legal advice. Distribution of Creative Commons public licenses does not create a lawyer-client or other relationship. Creative Commons makes its licenses and related information available on an "as-is" basis. Creative Commons gives no warranties regarding its licenses, any material licensed under their terms and conditions, or any related information. Creative Commons disclaims all liability for damages resulting from their use to the fullest extent possible.

Using Creative Commons Public Licenses

Creative Commons public licenses provide a standard set of terms and conditions that creators and other rights holders may use to share original works of authorship and other material subject to copyright and certain other rights specified in the public license below. The following considerations are for informational purposes only, are not exhaustive, and do not form part of our licenses.

Considerations for licensors: Our public licenses are intended for use by those authorized to give the public permission to use material in ways otherwise restricted by copyright and certain other rights. Our licenses are irrevocable. Licensors should read and understand the terms and conditions of the license they choose before applying it. Licensors should also secure all rights necessary before applying our licenses so that the public can reuse the material as expected. Licensors should clearly mark any material not subject to the license. This includes other CC-licensed material, or material used under an exception or limitation to copyright. More considerations for licensors: http://wiki.creativecommons.org/Considerations_for_licensors

Considerations for the public: By using one of our public licenses, a licensor grants the public permission to use the licensed material under specified terms and conditions. If the licensor's permission is not necessary for any reason—for example, because of any applicable exception or limitation to copyright—then that use is not regulated by the license. Our licenses grant only permissions under copyright and certain other rights that a licensor has authority to grant. Use of the licensed material may still be restricted for other reasons, including because others have copyright or other rights in the material. A licensor may make special requests, such as asking that all changes be marked or described. Although not required by our licenses, you are encouraged to respect those requests where reasonable. More considerations for the public: http://wiki.creativecommons.org/Considerations_for_licensees

Creative Commons Attribution 4.0 International Public License

By exercising the Licensed Rights (defined below), You accept and agree to be bound by the terms and conditions of this Creative Commons Attribution 4.0 International Public License ("Public License"). To the extent this Public License may be interpreted as a contract, You are granted the Licensed Rights in consideration of Your acceptance of these terms and conditions, and the Licensor grants You such rights in consideration of benefits the Licensor receives from making the Licensed Material available under these terms and conditions.

Section 1. Definitions

- a. Adapted Material means material subject to Copyright and Similar Rights that is derived from or based upon the Licensed Material and in which the Licensed Material is translated, altered, arranged, transformed, or otherwise modified in a manner requiring permission under the Copyright and Similar Rights held by the Licensor. For purposes of this Public License, where the Licensed Material is a musical work, performance, or sound recording, Adapted Material is always produced where the Licensed Material is synched in timed relation with a moving image.
- b. Adapter's License means the license You apply to Your Copyright and Similar Rights in Your contributions to Adapted Material in accordance with the terms and conditions of this Public License.

-
- c. Copyright and Similar Rights means copyright and/or similar rights closely related to copyright including, without limitation, performance, broadcast, sound recording, and Sui Generis Database Rights, without regard to how the rights are labeled or categorized. For purposes of this Public License, the rights specified in Section 2(b)(1)-(2) are not Copyright and Similar Rights.
 - d. Effective Technological Measures means those measures that, in the absence of proper authority, may not be circumvented under laws fulfilling obligations under Article 11 of the WIPO Copyright Treaty adopted on December 20, 1996, and/or similar international agreements.
 - e. Exceptions and Limitations means fair use, fair dealing, and/or any other exception or limitation to Copyright and Similar Rights that applies to Your use of the Licensed Material.
 - f. Licensed Material means the artistic or literary work, database, or other material to which the Licensor applied this Public License.
 - g. Licensed Rights means the rights granted to You subject to the terms and conditions of this Public License, which are limited to all Copyright and Similar Rights that apply to Your use of the Licensed Material and that the Licensor has authority to license.
 - h. Licensor means the individual(s) or entity(ies) granting rights under this Public License.
 - i. Share means to provide material to the public by any means or process that requires permission under the Licensed Rights, such as reproduction, public display, public performance, distribution, dissemination, communication, or importation, and to make material available to the public including in ways that members of the public may access the material from a place and at a time individually chosen by them.
 - j. Sui Generis Database Rights means rights other than copyright resulting from Directive 96/9/EC of the European Parliament and of the Council of 11 March 1996 on the legal protection of databases, as amended and/or succeeded, as well as other essentially equivalent rights anywhere in the world.
 - k. You means the individual or entity exercising the Licensed Rights under this Public License. Your has a corresponding meaning.

Section 2. Scope

- a. License grant.
 - 1. Subject to the terms and conditions of this Public License, the Licensor hereby grants You a worldwide, royalty-free, non-sublicensable, non-exclusive, irrevocable license to exercise the Licensed Rights in the Licensed Material to:
 - a. reproduce and Share the Licensed Material, in whole or in part; and
 - b. produce, reproduce, and Share Adapted Material.
 - 2. Exceptions and Limitations. For the avoidance of doubt, where Exceptions and Limitations apply to Your use, this Public License does not apply, and You do not need to comply with its terms and conditions.
 - 3. Term. The term of this Public License is specified in Section 6(a).
 - 4. Media and formats; technical modifications allowed. The Licensor authorizes You to exercise the Licensed Rights in all media and formats whether now known or hereafter created, and to make technical modifications necessary to do so. The Licensor waives and/or agrees not to assert any right or authority to forbid You from making technical modifications necessary to exercise the Licensed Rights, including technical modifications necessary to circumvent Effective Technological Measures. For purposes of this Public License, simply making modifications authorized by this Section 2(a) (4) never produces Adapted Material.
 - 5. Downstream recipients.
 - a. Offer from the Licensor – Licensed Material. Every recipient of the Licensed Material automatically receives an offer from the Licensor to exercise the Licensed Rights under the terms and conditions of this Public License.
 - b. No downstream restrictions. You may not offer or impose any additional or different terms or conditions on, or apply any Effective Technological Measures to, the Licensed Material if doing so restricts exercise of the Licensed Rights by any recipient of the Licensed Material.
 - 6. No endorsement. Nothing in this Public License constitutes or may be construed as permission to assert or imply that You are, or that Your use of the Licensed Material is, connected with, or sponsored, endorsed, or granted official status by, the Licensor or others designated to receive attribution as provided in Section 3(a)(1)(A)(i).
- b. Other rights.
 - 1. Moral rights, such as the right of integrity, are not licensed under this Public License, nor are publicity, privacy, and/or other similar personality rights; however, to the extent possible, the Licensor waives and/or agrees not to assert any such rights held by the Licensor to the limited extent necessary to allow You to exercise the Licensed Rights, but not otherwise.
 - 2. Patent and trademark rights are not licensed under this Public License.
 - 3. To the extent possible, the Licensor waives any right to collect royalties from You for the exercise of the Licensed Rights, whether directly or through a collecting society under any voluntary or waivable statutory or compulsory licensing scheme. In all other cases the Licensor expressly reserves any right to collect such royalties.

Section 3. License Conditions

Your exercise of the Licensed Rights is expressly made subject to the following conditions.

- a. Attribution.
 1. If You Share the Licensed Material (including in modified form), You must:
 - a. retain the following if it is supplied by the Licensor with the Licensed Material:
 - i. identification of the creator(s) of the Licensed Material and any others designated to receive attribution, in any reasonable manner requested by the Licensor (including by pseudonym if designated);
 - ii. a copyright notice;
 - iii. a notice that refers to this Public License;
 - iv. a notice that refers to the disclaimer of warranties;
 - v. a URI or hyperlink to the Licensed Material to the extent reasonably practicable;
 - b. indicate if You modified the Licensed Material and retain an indication of any previous modifications; and
 - c. indicate the Licensed Material is licensed under this Public License, and include the text of, or the URI or hyperlink to, this Public License.
 2. You may satisfy the conditions in Section 3(a)(1) in any reasonable manner based on the medium, means, and context in which You Share the Licensed Material. For example, it may be reasonable to satisfy the conditions by providing a URI or hyperlink to a resource that includes the required information.
 3. If requested by the Licensor, You must remove any of the information required by Section 3(a)(1)(A) to the extent reasonably practicable.
 4. If You Share Adapted Material You produce, the Adapter's License You apply must not prevent recipients of the Adapted Material from complying with this Public License.

Section 4. Sui Generis Database Rights

Where the Licensed Rights include Sui Generis Database Rights that apply to Your use of the Licensed Material:

- a. for the avoidance of doubt, Section 2(a)(1) grants You the right to extract, reuse, reproduce, and Share all or a substantial portion of the contents of the database;
- b. if You include all or a substantial portion of the database contents in a database in which You have Sui Generis Database Rights, then the database in which You have Sui Generis Database Rights (but not its individual contents) is Adapted Material; and
- c. You must comply with the conditions in Section 3(a) if You Share all or a substantial portion of the contents of the database.

For the avoidance of doubt, this Section 4 supplements and does not replace Your obligations under this Public License where the Licensed Rights include other Copyright and Similar Rights.

Section 5. Disclaimer of Warranties and Limitation of Liability

- a. UNLESS OTHERWISE SEPARATELY UNDERTAKEN BY THE LICENSOR, TO THE EXTENT POSSIBLE, THE LICENSOR OFFERS THE LICENSED MATERIAL AS-IS AND AS-AVAILABLE, AND MAKES NO REPRESENTATIONS OR WARRANTIES OF ANY KIND CONCERNING THE LICENSED MATERIAL, WHETHER EXPRESS, IMPLIED, STATUTORY, OR OTHER. THIS INCLUDES, WITHOUT LIMITATION, WARRANTIES OF TITLE, MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, NON-INFRINGEMENT, ABSENCE OF LATENT OR OTHER DEFECTS, ACCURACY, OR THE PRESENCE OR ABSENCE OF ERRORS, WHETHER OR NOT KNOWN OR DISCOVERABLE. WHERE DISCLAIMERS OF WARRANTIES ARE NOT ALLOWED IN FULL OR IN PART, THIS DISCLAIMER MAY NOT APPLY TO YOU.
- b. TO THE EXTENT POSSIBLE, IN NO EVENT WILL THE LICENSOR BE LIABLE TO YOU ON ANY LEGAL THEORY (INCLUDING, WITHOUT LIMITATION, NEGLIGENCE) OR OTHERWISE FOR ANY DIRECT, SPECIAL, INDIRECT, INCIDENTAL, CONSEQUENTIAL, PUNITIVE, EXEMPLARY, OR OTHER LOSSES, COSTS, EXPENSES, OR DAMAGES ARISING OUT OF THIS PUBLIC LICENSE OR USE OF THE LICENSED MATERIAL, EVEN IF THE LICENSOR HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH LOSSES, COSTS, EXPENSES, OR DAMAGES. WHERE A LIMITATION OF LIABILITY IS NOT ALLOWED IN FULL OR IN PART, THIS LIMITATION MAY NOT APPLY TO YOU.
- c. The disclaimer of warranties and limitation of liability provided above shall be interpreted in a manner that, to the extent possible, most closely approximates an absolute disclaimer and waiver of all liability.

Section 6. Term and Termination

- a. This Public License applies for the term of the Copyright and Similar Rights licensed here. However, if You fail to comply with this Public License, then Your rights under this Public License terminate automatically.
- b. Where Your right to use the Licensed Material has terminated under Section 6(a), it reinstates:
 - 1. automatically as of the date the violation is cured, provided it is cured within 30 days of Your discovery of the violation; or
 - 2. upon express reinstatement by the Licensor.

For the avoidance of doubt, this Section 6(b) does not affect any right the Licensor may have to seek remedies for Your violations of this Public License.

- c. For the avoidance of doubt, the Licensor may also offer the Licensed Material under separate terms or conditions or stop distributing the Licensed Material at any time; however, doing so will not terminate this Public License.
- d. Sections 1, 5, 6, 7, and 8 survive termination of this Public License.

Section 7. Other Terms and Conditions

- a. The Licensor shall not be bound by any additional or different terms or conditions communicated by You unless expressly agreed.
- b. Any arrangements, understandings, or agreements regarding the Licensed Material not stated herein are separate from and independent of the terms and conditions of this Public License.

Section 8. Interpretation

- a. For the avoidance of doubt, this Public License does not, and shall not be interpreted to, reduce, limit, restrict, or impose conditions on any use of the Licensed Material that could lawfully be made without permission under this Public License.
- b. To the extent possible, if any provision of this Public License is deemed unenforceable, it shall be automatically reformed to the minimum extent necessary to make it enforceable. If the provision cannot be reformed, it shall be severed from this Public License without affecting the enforceability of the remaining terms and conditions.
- c. No term or condition of this Public License will be waived and no failure to comply consented to unless expressly agreed to by the Licensor.
- d. Nothing in this Public License constitutes or may be interpreted as a limitation upon, or waiver of, any privileges and immunities that apply to the Licensor or You, including from the legal processes of any jurisdiction or authority.

Creative Commons is not a party to its public licenses. Notwithstanding, Creative Commons may elect to apply one of its public licenses to material it publishes and in those instances will be considered the Licensor. The text of the Creative Commons public licenses is dedicated to the public domain under the CC0 Public Domain Dedication. Except for the limited purpose of indicating that material is shared under a Creative Commons public license or as otherwise permitted by the Creative Commons policies published at <http://creativecommons.org/policies>, Creative Commons does not authorize the use of the trademark “Creative Commons” or any other trademark or logo of Creative Commons without its prior written consent including, without limitation, in connection with any unauthorized modifications to any of its public licenses or any other arrangements, understandings, or agreements concerning use of licensed material. For the avoidance of doubt, this paragraph does not form part of the public licenses.

Creative Commons may be contacted at <http://creativecommons.org>.

Bibliography

- [1] RISC-V Foundation, *The RISC-V Instruction Set Manual, Volume I: User-Level ISA, Document Version 2.2*, 5 2017. Editors Andrew Waterman and Krste Asanović. [vi](#), [25](#), [26](#), [27](#), [28](#), [35](#), [37](#), [38](#), [39](#), [40](#), [41](#), [42](#), [43](#), [44](#), [45](#)
- [2] D. Patterson and A. Waterman, *The RISC-V Reader: An Open Architecture Atlas*. Strawberry Canyon, 11 2017. ISBN: 978-0999249116. [vi](#)
- [3] D. Patterson and J. Hennessy, *Computer Organization and Design RISC-V Edition: The Hardware Software Interface*. Morgan Kaufmann, 4 2017. ISBN: 978-0128122754. [vi](#)
- [4] W. F. Decker, “A modern approach to teaching computer organization and assembly language programming,” *SIGCSE Bull.*, vol. 17, pp. 38–44, 12 1985. [vi](#)
- [5] Texas Instruments, *SN54190, SN54191, SN54LS190, SN54LS191, SN74190, SN74191, SN74LS190, SN74LS191 Synchronous Up/Down Counters With Down/Up Mode Control*, 3 1988. [vi](#)
- [6] Texas Instruments, *SN54154, SN74154 4-line to 16-line Decoders/Demultiplexers*, 12 1972. [vi](#)
- [7] Intel, *MCS-85 User’s Manual*, 9 1978. [vi](#)
- [8] Radio Shack, *TRS-80 Editor/Assembler Operation and Reference Manual*, 1978. [vi](#)
- [9] Motorola, *MC68000 16-bit Microprocessor User’s Manual*, 2nd ed., 1 1980. MC68000UM(AD2). [vi](#)
- [10] R. A. Overbeek and W. E. Singletary, *Assembler Language With ASSIST*. Science Research Associates, Inc., 2nd ed., 1983. [vi](#)
- [11] IBM, *IBM System/370 Principals of Operation*, 7th ed., 3 1980. [vi](#)
- [12] IBM, *OS/VS-DOS/VSE-VM/370 Assembler Language*, 6th ed., 3 1979. [vi](#)
- [13] D. Cohen, “[IEN 137, On Holy Wars and a Plea for Peace](#),” Apr. 1980. This note discusses the Big-Endian/Little-Endian byte/bit-order controversy, but did not settle it. A decade later, David V. James in “Multiplexed Buses: The Endian Wars Continue”, *IEEE Micro*, **10**(3), 9–21 (1990) continued the discussion. [1](#)
- [14] R. M. Stallman and the GCC Developer Community, *Using the GNU Compiler Collection (For GCC version 7.3.0)*. Free Software Foundation, 51 Franklin Street, Fifth Floor, Boston, MA 02110-1301 USA: GNU Press, 2017. [15](#)

Index

I

Instruction

ADD, 38
ADDI, 34
AND, 42
ANDI, 36
AUIPC, 26
BEQ, 28
BGE, 30
BGEU, 30
BLT, 29
BLTU, 30
BNE, 29
CSRRC, 44
CSRRCI, 45
CSRRS, 44
CSRRSI, 45
CSRRW, 44
CSRRWI, 45
EBREAK, 44
ECALL, 44
FENCE, 43
FENCE.I, 43
JAL, 26
JALR, 27
LB, 31
LBU, 32
LH, 31
LHU, 32
LUI, 25
LW, 31
OR, 42
ORI, 36
SB, 33
SH, 33
SLL, 39
SLLI, 37
SLT, 39
SLTI, 34
SLTIU, 35
SLTU, 40
SRA, 41
SRAI, 37
SRL, 41
SRLI, 37

SUB, 38

SW, 33

XOR, 40

XORI, 35

R

RV32, 16
RV32A, 45
RV32D, 46
RV32F, 46
RV32I, 25
RV32M, 45

Glossary

bit One binary digit. [2](#), [53](#)

byte A binary value represented by 8 bits. [53](#)

Doubleword A binary value represented by 64 bits. [53](#)

Fullword A binary value represented by 32 bits. [53](#)

Halfword A binary value represented by 16 bits. [53](#)

High order bits Some number of MSBs. [53](#)

hit One hex digit. [2](#), [3](#), [53](#)

ISA Instruction Set Architecture. [53](#)

LaTeX Is a mark up language specially suited for scientific documents. [53](#)

Low order bits Some number of LSBs. [53](#)

LSB Least Significant Bit. [4](#), [17](#), [20](#), [22](#), [53](#)

MSB Most Significant Bit. [4](#), [5](#), [16](#), [17](#), [22](#), [53](#)

overflow The situation where the result of an addition or subtraction operation is approaching positive or negative infinity and exceeds the number of bits allotted to contain the result. This is typically caused by high-order truncation. [5](#), [9](#), [53](#)

Quadword A binary value represented by 128 bits. [53](#)

RV32 Short for RISC-V 32. The number 32 refers to the XLEN. [25](#), [53](#)

RV64 Short for RISC-V 64. The number 64 refers to the XLEN. [53](#)

underflow The situation where the result of an addition or subtraction operation is approaching zero and exceeds the number of bits allotted to contain the result. This is typically caused by low-order truncation. [9](#), [53](#)

XLEN The number of bits a RISC-V x integer register (such as x0). For RV32 XLEN=32, RV64 XLEN=64 etc. [21](#), [22](#), [53](#)