

RISC-V Assembly Language Programming

(Draft v0.1-26-ge892461)

John Winans
jwinans@niu.edu

May 12, 2018

Copyright © 2018 John Winans

This document is made available under a Creative Commons Attribution 4.0 International License. See [Appendix C](#) for more information.

Download your own copy of this book from github here: <https://github.com/johnwinans/rvalp>.

This document may contain inaccuracies or errors. The author provides no guarantee regarding the accuracy of this document's contents. If you discover that this document contains errors, please notify the author.

ARM® is a registered trademark of ARM Limited in the EU and other countries.

► Fix Me:

*Need to say something
about trademarks for things
mentioned in this text*

Contents

Preface	vii
1 Introduction	1
1.1 The Digital Computer	1
1.1.1 Storage Systems	1
1.1.1.1 Volatile Storage	1
1.1.1.2 Non-Volatile Storage	2
1.1.2 CPU	2
1.1.2.1 Execution Unit	2
1.1.2.2 Arithmetic and Logic Unit	3
1.1.2.3 Registers	3
1.1.2.4 Harts	3
1.1.3 Peripherals	3
1.2 Instruction Set Architecture	3
1.2.1 RV Base Modules	4
1.2.2 Extension Modules	4
1.3 How the CPU Executes a Program	4
1.3.1 Instruction Fetch	4
1.3.2 Instruction Execute	5
2 Numbers and Storage Systems	6
2.1 Logical/Boolean Functions	6
2.1.1 NOT	6
2.1.2 AND	6
2.1.3 OR	7
2.1.4 XOR	7
2.2 Integers and Counting	8
2.2.1 Converting Between Bases	9
2.2.1.1 From Binary to Decimal	9
2.2.1.2 From Binary to Hexadecimal	9

2.2.1.3	From Hexadecimal to Binary	10
2.2.1.4	From Decimal to Binary	10
2.2.1.5	From Decimal to Hex	11
2.2.2	Addition of Binary Numbers	11
2.2.3	Signed Numbers	11
2.2.3.1	Converting between Positive and Negative	12
2.2.4	Subtraction of Binary Numbers	13
2.2.5	Truncation and Overflow	13
2.3	Main Memory Storage	13
2.3.1	Memory Dump	14
2.3.2	Big Endian Representation	14
2.3.3	Little Endian Representation	14
2.3.4	Character Strings and Arrays	14
2.3.5	Alignment	14
2.3.6	Instruction Alignment	14
3	The Elements of a Assembly Language Program	16
3.1	Assembly Language Statements	16
3.2	Memory Layout	16
3.3	A Sample Program Source Listing	16
3.4	Running a Program With rvddt	17
4	Using The RISC-V GNU Toolchain	20
5	RV32 Machine Instructions	22
5.1	Introduction	22
5.2	Conventions and Terminology	22
5.2.1	XLEN	22
5.2.2	sx(val)	22
5.2.3	zx(val)	23
5.2.4	zr(val)	23
5.2.5	Sign Extended Left and Zero Extend Right	23
5.2.6	m8(addr)	24
5.2.7	m16(addr)	24
5.2.8	m32(addr)	25
5.2.9	m64(addr)	25
5.2.10	m128(addr)	25
5.2.11	.+offset	25
5.2.12	.-offset	25

5.2.13	pc	25
5.2.14	rd	26
5.2.15	rs1	26
5.2.16	rs2	26
5.2.17	imm	26
5.2.18	rsN[h:l]	26
5.3	Addressing Modes	26
5.4	Instruction Encoding Formats	26
5.4.1	U Type	27
5.4.2	J Type	28
5.4.3	R Type	29
5.4.4	I Type	29
5.4.5	S Type	29
5.4.6	B Type	29
5.4.7	CPU Registers	30
5.5	memory	30
5.6	RV32I Base Instruction Set	31
5.6.1	LUI rd, imm	31
5.6.2	AUIPC rd, imm	32
5.6.3	JAL rd, imm	32
5.6.4	JALR rd, rs1, imm	33
5.6.5	BEQ rs1, rs2, imm	34
5.6.6	BNE rs1, rs2, imm	35
5.6.7	BLT rs1, rs2, imm	35
5.6.8	BGE rs1, rs2, imm	36
5.6.9	BLTU rs1, rs2, imm	36
5.6.10	BGEU rs1, rs2, imm	36
5.6.11	LB rd, imm(rs1)	37
5.6.12	LH rd, imm(rs1)	37
5.6.13	LW rd, imm(rs1)	37
5.6.14	LBU rd, imm(rs1)	38
5.6.15	LHU rd, imm(rs1)	38
5.6.16	SB rs2, imm(rs1)	39
5.6.17	SH rs2, imm(rs1)	39
5.6.18	SW rs2, imm(rs1)	39
5.6.19	ADDI rd, rs1, imm	40
5.6.20	SLTI rd, rs1, imm	40
5.6.21	SLTIU rd, rs1, imm	41

5.6.22	XORI rd, rs1, imm	41
5.6.23	ORI rd, rs1, imm	42
5.6.24	ANDI rd, rs1, imm	42
5.6.25	SLLI rd, rs1, shamt	43
5.6.26	SRLI rd, rs1, shamt	43
5.6.27	SRAI rd, rs1, shamt	43
5.6.28	ADD rd, rs1, rs2	44
5.6.29	SUB rd, rs1, rs2	44
5.6.30	SLL rd, rs1, rs2	45
5.6.31	SLT rd, rs1, rs2	45
5.6.32	SLTU rd, rs1, rs2	46
5.6.33	XOR rd, rs1, rs2	46
5.6.34	SRL rd, rs1, rs2	47
5.6.35	SRA rd, rs1, rs2	47
5.6.36	OR rd, rs1, rs2	48
5.6.37	AND rd, rs1, rs2	48
5.6.38	FENCE predecessor, successor	49
5.6.39	FENCE.I	49
5.6.40	ECALL	50
5.6.41	EBREAK	50
5.6.42	CSR _{RW} rd, csr, rs1	50
5.6.43	CSR _{RS} rd, csr, rs1	51
5.6.44	CSR _{RC} rd, csr, rs1	51
5.6.45	CSR _{RWI} rd, csr, imm	51
5.6.46	CSR _{RSI} rd, csr, rs1	52
5.6.47	CSR _{RCI} rd, csr, rs1	52
5.7	RV32M Standard Extension	52
5.7.1	MUL rd, rs1, rs2	52
5.7.2	MULH rd, rs1, rs2	53
5.7.3	MULHS rd, rs1, rs2	53
5.7.4	MULHU rd, rs1, rs2	53
5.7.5	DIV rd, rs1, rs2	53
5.7.6	DIVU rd, rs1, rs2	53
5.7.7	REM rd, rs1, rs2	54
5.7.8	REMU rd, rs1, rs2	54
5.8	RV32A Standard Extension	54
5.9	RV32F Standard Extension	54
5.10	RV32D Standard Extension	54

A	Installing a RISC-V Toolchain	55
A.1	The GNU Toolchain	55
A.2	rvddt	56
B	Floating Point Numbers	57
B.1	IEEE-754 Floating Point Number Representation	57
B.1.1	Floating Point Number Accuracy	58
B.1.1.1	Powers Of Two	58
B.1.1.2	Clean Decimal Numbers	59
B.1.1.3	Accumulation of Error	60
B.1.2	Reducing Error Accumulation	61
C	Attribution 4.0 International	63
	Bibliography	67
	Index	68
	Glossary	70

Preface

I set out to this book because I couldn't find it in a single volume elsewhere.

The closest thing to what I sought when deciding to collect my thoughts into this document would be select portions of *The RISC-V Instruction Set Manual, Volume I: User-Level ISA, Document Version 2.2*[\[1\]](#), *The RISC-V Reader*[\[2\]](#), and *Computer Organization and Design RISC-V Edition: The Hardware Software Interface*[\[3\]](#).

There *are* some terse guides around the Internet that are suitable for those that already know an assembly language. With all the (deserved) excitement brewing over system organization (and the need to compress the time out of university courses targeting assembly language programming [\[4\]](#)), it is no surprise that RISC-V texts for the beginning assembly programmer are not (yet) available.

When I got started in computing I learned how to count in binary in a high school electronics course using data sheets for integrated circuits such as the 74191[\[5\]](#) and 74154[\[6\]](#) prior to knowing that assembly language even existed.

I learned assembler from data sheets and texts (that are still sitting on my shelves) such as:

- The MCS-85 User's Manual[\[7\]](#)
- The EDTASM Manual[\[8\]](#)
- The MC68000 User's Manual[\[9\]](#)
- Assembler Language With ASSIST[\[10\]](#)
- IBM System/370 Principals of Operation[\[11\]](#)
- OS/VS-DOS/VSE-VM/370 Assembler Language[\[12\]](#)
- ... and several others

One way or another all of them discuss each CPU instruction in excruciating detail with both a logical and narrative description. For RISC-V this is also the case for the *RISC-V Reader*[\[2\]](#) and the *Computer Organization and Design RISC-V Edition*[\[3\]](#) books and is also present in this text (I consider that to be the minimal level of responsibility.)

Where I hope this text will differentiate itself from the existing RISC-V titles is in its attempt to address the needs of those learning assembly language for the first time. To this end I have primed this project with some of the material from old handouts I used when teaching assembly language programming in the late '80s.

Chapter 1

Introduction

At its core, a digital computer has at least one **Central Processing Unit** (CPU). A CPU executes a continuous stream of instructions called a **program**. These program instructions are expressed in what is called **machine language**. Each machine language instruction is a binary value. In order to provide a method to simplify the management of machine language programs a symbolic mapping is provided where a **mnemonic** can be used to specify each machine instruction and any of its parameters... rather than require that programs be expressed as a series of binary values. A set of mnemonics, parameters and rules for specifying their use for the purpose of programming a CPU is called an *Assembly Language*.

1.1 The Digital Computer

There are different types of computers. A *digital* computer is the type that most people think of when they hear the word *computer*. Other varieties of computers include *analog* and *quantum*.

A digital computer is one that processes data that are represented using numeric values (digits), most commonly expressed in binary (ones and zeros) form.

This text focuses on digital computing.

A typical digital computer is composed of storage systems (memory, disc drives, USB drives, etc.), a CPU (with one or more cores), input peripherals (a keyboard and mouse) and output peripherals (display, printer or speakers.)

1.1.1 Storage Systems

Computer storage systems are used to hold the data and instructions for the CPU.

Types of computer storage can be classified into two categories. Volatile and non-volatile.

1.1.1.1 Volatile Storage

Volatile storage is characterized by the fact that it will lose its contents (forget) any time that it is powered off.

One type of volatile storage is provided inside the CPU itself in small blocks called [registers](#). These CPU registers are used to hold individual data values that can be manipulated by the instructions that are executed by the CPU.

Another type of volatile storage is main memory. Main memory is connected to a computer's CPU and is used to hold the data and instructions that can not fit into the CPU registers.

Typically, a CPU's registers can hold tens of data values while the main memory can contain many billions of data values.

To keep track of the data values, each register is assigned a number and the main memory is broken up into small blocks called [bytes](#) that are also each assigned number called an [address](#) (an address is often referred to as a *location*).

A CPU can process data in a register at a speed that can be an order of magnitude faster than the rate that it can process (specifically, transfer data and instructions to and from) the main memory.

Register storage costs an order of magnitude more to manufacture than main memory. While it is desirable to have many registers the economics dictate that the vast majority of volatile computer storage be provided in its main memory. As a result, optimizing the copying of data between the registers and main memory is a desirable trait of good programs.

1.1.1.2 Non-Volatile Storage

Non-volatile storage is characterized by the fact that it will *NOT* lose its contents when it is powered off.

Common types of non-volatile storage are disc drives, flash cards and USB drives. Prices can vary widely depending on size and transfer speeds.

It is typical for a computer system's non-volatile storage to operate more slowly than its main memory.

This text is not particularly concerned with non-volatile storage.

1.1.2 CPU

The [CPU](#) is a collection of registers and circuitry designed manipulate the register data and to exchange data and instructions with the storage system. The instructions that it reads from the main memory tells the CPU to perform various mathematical and logical operations on the data in its registers and where to save the results of those operations.

►► Fix Me:
Add a block diagram of the CPU components described here.

1.1.2.1 Execution Unit

The part of a CPU that coordinates all aspects of the operations of each instruction is called the *execution unit*. It is what performs the transfers of instructions and data between the CPU and the main memory and tells the registers when they are supposed to either store or recall data being transferred. The execution unit also controls the ALU (Arithmetic and Logic Unit).

1.1.2.2 Arithmetic and Logic Unit

ALU
register
hart
ISA

When an instruction manipulates data by performing things like an *addition*, *subtraction*, *comparison* or other similar operations, the ALU is what will calculate the sum, difference, and so on.

1.1.2.3 Registers

In the RV32 CPU there are 31 general purpose registers that each contain 32 [bits](#) (where each bit is one [binary](#) digit value of one or zero) and a number of special-purpose registers. Each of the general purpose registers is given a name such as `x1`, `x2`, ... on up to `x31` (*general purpose* refers to the fact that the CPU itself does not prescribe any particular function to any these registers.) Two important special-purpose registers are `x0` and `pc`.

Register `x0` will always represent the value zero or logical *false* no matter what. If any instruction tries to change the value is `x0` value the operation will fail. The need for *zero* is so common that, other than the fact that it is hard-wired to zero, the `x0` register is made available as if it were otherwise a general purpose register.¹

The `pc` register is called the *program counter*. The CPU uses it to remember the memory address where its program instructions are located.

The number of bits in each register is defined by the [Instruction Set Architecture \(ISA\)](#).

1.1.2.4 Harts

Analogous to a *core* in other types of CPUs, a [hart](#) (hardware [thread](#)) in a RISC-V CPU refers to the collection of 32 registers, instruction execution unit and ALU.

When more than one hart is present in a CPU, a different stream of instructions can be executed on each hart all at the same time. Programs that are written to take advantage of this are called *multithreaded*.

This text will primarily focus on CPUs that have only one hart.

1.1.3 Peripherals

A peripheral is a device that is not a CPU or main memory. They are typically used to transfer information/data into and out of the main memory.

This text is not particularly concerned with the peripherals of a computer system other than in those sections where instructions are discussed whose purpose is to address the needs of a peripheral device. Such instructions are used to initiate, execute and/or synchronize data transfers.

1.2 Instruction Set Architecture

The catalog of rules that describes the details of the instructions and features that a given CPU provides is called its [Instruction Set Architecture \(ISA\)](#).

¹Having a special *zero* register allows the total set of instructions that the CPU can execute to be simplified. Thus reducing its complexity, power consumption and cost.

An ISA is typically expressed in terms of the specific meaning of each binary instruction that a CPU can recognize and how it will process each one.

The RISC-V ISA is defined as a set of modules. The purpose of dividing the ISA into modules is to allow an implementer to select which features to incorporate into a CPU design.

Any given RISC-V implementation must provide one of the *base* modules and zero or more of the *extension* modules.

RV32I
RV32M
RV32A
RV32F
RV32D
RV32Q
RV32C
RV32G
instruction cycle
instruction fetch

1.2.1 RV Base Modules

The base modules are RV32I (32-bit general purpose), RV32E (32-bit embedded), RV64I (64-bit general purpose) and RV128I (128-bit general purpose).

These base modules provide the minimal functional set of integer operations needed to execute a useful application. The differing bit-widths address the needs of different main-memory sizes.

This text primarily focuses on the RV32I base module and how to program it.

1.2.2 Extension Modules

RISC-V extension modules may be included by an implementer interested in optimizing a design for one or more purposes.

Available extension modules include M (integer math), A (atomic), F (32-bit floating point), D (64-bit floating point), Q (128-bit floating point), C (compressed size instructions) and others.

The extension name *G* is used to represent the combined set of IMAFD extensions as it is expected to be a common combination.

1.3 How the CPU Executes a Program

The process of executing a program is continuously repeating series of *instruction cycles* that are each comprised of an *instruction fetch* and an *instruction execute* phase.

The current status of a CPU is entirely embodied in the data values that are stored in its registers at any moment in time. Of particular interest to an executing a program is the *pc* register. The *pc* contains the memory address containing the instruction that the CPU will execute next.

For this to work, the instructions to be executed must have been previously stored in adjacent main memory locations and the address of the first instruction placed into the *pc* register.

1.3.1 Instruction Fetch

In order to *fetch* an instruction from the main memory the CPU must have a method to identify which instruction should be fetched and a method to fetch it.

Given that the main memory is broken up and that each of its bytes is assigned an address, the *pc* is used to hold the address of the location where the next instruction to execute is located.

Given an instruction address, the CPU can request that the main memory locate and return the value of the data stored there using what is called a *memory read* operation and then the CPU can treat that *fetch*ed value as an instruction and execute it.² instruction execute

Once an instruction has been fetched, it can be executed.

1.3.2 Instruction Execute

Typical instructions do things like add a number to the value currently stored in one of the registers or store the contents of a register into the main memory at some given address.

Also part of every instruction is a notion of what should be done next.

Most of the time an instruction will be complete by indicating that the CPU should proceed to fetch and execute the instruction at the next larger main memory address. In these cases the `pc` is incremented to point to the memory address after the current instruction.

Any parameters that an instruction requires must either be part of the instruction itself or read from (or stored into) one or more of the general purpose registers.

Some instructions can specify that the CPU proceed to execute an instruction at an address other than the one that follows itself. This class of instructions have names like *jump* and *branch* and are available in a variety of different styles.

The RISC-V ISA uses the word *jump* to refer to an *unconditional* change in the sequential processing of instructions and the word *branch* to refer to a *conditional* change.

For example, a (conditional) branch instruction might instruct the CPU to proceed to the instruction at the next main memory address if the value in register number 8 is currently less than the value in register number 24 *but otherwise* proceed to an instruction at a different address when it is not. This type of instruction can therefore result in having one of two different actions pending the resulting *condition* of the comparison.³

Once the instruction execution phase has completed, the next instruction cycle will be performed using the new `pc` register address.

²RV32I instructions are more than one byte in size, but this general description is suitable for now.

³This is the fundamental method used by a CPU to make decisions.

Chapter 2

Numbers and Storage Systems

This chapter discusses how data are represented and stored in a computer.

2.1 Logical/Boolean Functions

Unlike addition and subtraction, boolean functions apply on a per-bit basis. When applied to multi-bit values, each bit position is operated upon independently of the other bits.

► Fix Me:

This is unclear. Need to define bit positions and probably should add basic truth table diagrams.

2.1.1 NOT

The *NOT* operator applies to a single operand and represents the opposite of the input.

► Fix Me:

Need to define 1 as true and 0 as false somewhere.

If the input is 1 then the output is 0. If the input is 0 then the output is 1. In other words, the output value is *not* that of the input value.

► Fix Me:

Need to define unary, binary and ternary operators without confusing binary operators with binary numbers.

This text will use the operator used in the C language when discussing the *NOT* operator in symbolic form. Specifically the tilde: '~'.

```
~ 1 1 1 1 0 1 0 1  <== A
-----
 0 0 0 0 1 0 1 0  <== output
```

In a line of code the above might read like this: `output = ~A`

2.1.2 AND

The boolean *and* function has two or more inputs and the output is a single bit. The output is 1 if and only if all of the input values are 1. Otherwise it is 0.

This text will use the operator used in the C language when discussing the *AND* operator in symbolic form. Specifically the ampersand: '&'.

This function works like it does in spoken language. For example if A is 1 *AND* B is 1 then the output is 1 (true). Otherwise the output is 0 (false). For example:

```
  1 1 1 1 0 1 0 1  <== A
& 1 0 0 1 0 0 1 1  <== B
-----
  1 0 0 1 0 0 0 1  <== output
```

In a line of code the above might read like this: `output = A & B`

2.1.3 OR

The boolean *or* function has two or more inputs and the output is a single bit. The output is 1 if at least one of the input values are 1.

This text will use the operator used in the C language when discussing the *OR* operator in symbolic form. Specifically the pipe: '|'.

This function works like it does in spoken language. For example if A is 1 *OR* B is 1 then the output is 1 (true). Otherwise the output is 0 (false). For example:

```
  1 1 1 1 0 1 0 1  <== A
| 1 0 0 1 0 0 1 1  <== B
-----
  1 1 1 1 0 1 1 1  <== output
```

In a line of code the above might read like this: `output = A | B`

2.1.4 XOR

The boolean *exclusive or* function has two or more inputs and the output is a single bit. The output is 1 if only an odd number of inputs are 1. Otherwise the output will be 0.

This text will use the operator used in the C language when discussing the *XOR* operator in symbolic form. Specifically the carrot: '^'.

Note that when *XOR* is used with two inputs, the output is set to 1 (true) when the inputs have different values and 0 (false) when the inputs both have the same value.

For example:

```
  1 1 1 1 0 1 0 1  <== A
^ 1 0 0 1 0 0 1 1  <== B
-----
  0 1 1 0 0 1 1 0  <== output
```

In a line of code the above might read like this: `output = A ^ B`

2.2 Integers and Counting

A binary integer is constructed with only 1s and 0s in the same manner as decimal numbers are constructed with values from 0 to 9.

Counting in binary is the same as in decimal. For example, when adding 1 to 9, the carry is added to the next place value. When subtracting 1 from 0, a borrow is required and so on.

Figure [Figure 2.1](#) shows an abridged table of the decimal, binary and hexadecimal values from 0 to 129.

Decimal			Binary								Hex	
10^2	10^1	10^0	2^7	2^6	2^5	2^4	2^3	2^2	2^1	2^0	16^1	16^0
100	10	1	128	64	32	16	8	4	2	1	16	1
0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	1	0	0	0	0	0	0	0	1	0	1
0	0	2	0	0	0	0	0	0	1	0	0	2
0	0	3	0	0	0	0	0	0	1	1	0	3
0	0	4	0	0	0	0	0	1	0	0	0	4
0	0	5	0	0	0	0	0	1	0	1	0	5
0	0	6	0	0	0	0	0	1	1	0	0	6
0	0	7	0	0	0	0	0	1	1	1	0	7
0	0	8	0	0	0	0	1	0	0	0	0	8
0	0	9	0	0	0	0	1	0	0	1	0	9
0	1	0	0	0	0	0	1	0	1	0	0	a
0	1	1	0	0	0	0	1	0	1	1	0	b
0	1	2	0	0	0	0	1	1	0	0	0	c
0	1	3	0	0	0	0	1	1	0	1	0	d
0	1	4	0	0	0	0	1	1	1	0	0	e
0	1	5	0	0	0	0	1	1	1	1	0	f
0	1	6	0	0	0	1	0	0	0	0	1	0
0	1	7	0	0	0	1	0	0	0	1	1	1
...			
1	2	5	0	1	1	1	1	1	0	1	7	d
1	2	6	0	1	1	1	1	1	1	0	7	e
1	2	7	0	1	1	1	1	1	1	1	7	f
1	2	8	1	0	0	0	0	0	0	0	8	0

Figure 2.1: Counting in decimal, binary and hexadecimal.

One way to look at this table is on a per-row basis where each place value is represented by the base raised to the power of the place value position (shown in the column headings.) This is useful when converting arbitrary values between bases. For example to interpret the decimal value on the fourth row:

$$0 \times 10^2 + 0 \times 10^1 + 3 \times 10^0 = 3_{10}$$

And to interpret binary value on the same row by converting it to decimal:

$$0 \times 2^7 + 0 \times 2^6 + 0 \times 2^5 + 0 \times 2^4 + 0 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 1 \times 2^0 = 3_{10}$$

And the same for the hexadecimal value:

$$0 \times 16^1 + 3 \times 16^0 = 3_{10}$$

Another way to look at this table is on a per-column basis. When tasked with drawing such a table by hand, it might be useful to observe that, just as in decimal, the right-most column will cycle through

all of the values represented in the chosen base then cycle back to zero and repeat. (For example, in binary this pattern is 0-1-0-1-0-1-0-...) The next column in each base will cycle in the same manner except each of the values is repeated as many times as is represented by the place value (in the case of decimal, 10^1 times, binary 2^1 times, hex 16^1 times. Again, the for binary numbers this pattern is 0-0-1-1-0-0-1-1-...) This continues for as many columns as are needed to represent the magnitude of the desired number.

Another item worth noting is that any even binary number will always have a 0 LSB and odd numbers will always have a 1 LSB.

As is customary in decimal, leading zeros are sometimes not shown for readability.

The relationship between binary and hex values is also worth taking note. Because $2^4 = 16$, there is a clean and simple grouping of 4 bits to 1 hex. There is no such relationship between binary and decimal.

Writing and reading numbers in binary that are longer than 8 bits is cumbersome and prone to error. The simple conversion between binary and hex makes hex a convenient shorthand for expressing binary values in many situations.

For example, consider the following value expressed in binary, hexadecimal and decimal (spaced to show the relationship between binary and hex):

Binary value:	0010	0111	1011	1010	1100	1100	1111	0101
Hex Value:	2	7	B	A	C	C	F	5
Decimal Value:								666553589

Empirically we can see that grouping the bits into sets of four allows an easy conversion to hex and expressing it as such is $\frac{1}{4}$ as long as in binary while at the same time allowing for easy conversion back to binary.

The decimal value in this example does not easily convey a sense of the binary value.

2.2.1 Converting Between Bases

2.2.1.1 From Binary to Decimal

Alas, it is occasionally necessary to convert between decimal, binary and/or hex.

To convert from binary to decimal, put the decimal value of the place values ... 8 4 2 1 over the binary digits like this:

128	64	32	16	8	4	2	1
0	0	0	1	1	0	1	1

Now sum the place-values that are expressed in decimal for each bit with the value of 1: $16 + 8 + 2 + 1$. The integer binary value 00011011_2 represents the decimal value 27_{10} .

2.2.1.2 From Binary to Hexadecimal

Conversion from binary to hex involves grouping the bits into sets of four and then performing the same summing process as shown above. If there is not a multiple of four bits then extend the binary to the

left with zeros to make it so.

Grouping the bits into sets of four and summing:

Place:	8 4 2 1	8 4 2 1	8 4 2 1	8 4 2 1
Binary:	0 1 1 0	1 1 0 1	1 0 1 0	1 1 1 0
Decimal:	4+2 =6	8+4+ 1=13	8+ 2 =10	8+4+2 =14

After the summing, convert each decimal value to hex. The decimal values from 0–9 are the same values in hex. Because we don’t have any more numerals to represent the values from 10–15, we use the first 6 letters (See the right-most column of [Figure 2.1.](#)) Fortunately there are only six hex mappings involving letters. Thus it is reasonable to memorize them.

Continuing this example:

Decimal:	6	13	10	14
Hex:	6	D	A	E

2.2.1.3 From Hexadecimal to Binary

Again, the four-bit mapping between binary and hex makes this task as straight forward as using a look-up table.

For each [hit](#) (Hex digIT), translate it to its unique four-bit pattern. Perform this task either by memorizing each of the 16 patterns or by converting each hit to decimal first and then converting each four-bit binary value to decimal using the place-value summing method discussed in [subsection 2.2.1.1.](#)

For example:

Hex:	4	C
Binary:	0 1 0 0	1 1 0 0
Decimal:	128 64 32 16	8 4 2 1
Sum:	64+	8+4 = 76

2.2.1.4 From Decimal to Binary

To convert arbitrary decimal numbers to binary, extend the list of binary place values until it exceeds the value of the decimal number being converted. Then make successive subtractions of each of the place values that would yield a non-negative result.

For example, to convert 1234_{10} to binary:

Place values: 2048–1024–512–256–128–64–32–16–8–4–2–1

0	2048	(too big)
1	1234 – 1024 = 210	
0	512	(too big)
0	256	(too big)
1	210 – 128 = 82	

1	82 - 64	= 18	
0	32		(too big)
1	18 - 16	= 2	
0	8		(too big)
0	4		(too big)
1	2 - 2	= 0	
0	1		(too big)

The answer using this notation is listed vertically in the left column with the **MSB** on the top and the **LSB** on the bottom line: 010011010010₂.

2.2.1.5 From Decimal to Hex

Conversion from decimal to hex can be done by using the place values for base-16 and the same math as from decimal to binary or by first converting the decimal value to binary and then from binary to hex by using the methods discussed above.

Because binary and hex are so closely related, performing a conversion by way of binary is quite straight forward.

2.2.2 Addition of Binary Numbers

The addition of binary numbers can be performed long-hand the same way decimal addition is taught in grade school. In fact binary addition is easier since it only involves adding 0 or 1.

The first thing to note that in any number base $0 + 0 = 0$, $0 + 1 = 1$, and $1 + 0 = 1$. Since there is no “two” in binary (just like there is no “ten” decimal) adding $1 + 1$ results in a zero with a carry as in: $1 + 1 = 10_2$ and in: $1 + 1 + 1 = 11_2$. Using these five sums, any two binary integers can be added.

For example:

```

      111111 1111 <== carries
    0110101111001111 <== addend
+ 0000011101100011 <== addend
-----
    0111001100110010 <== sum

```

2.2.3 Signed Numbers

There are multiple methods used to represent signed binary integers. The method used by most modern computers is called “two’s complement.”

A two’s complement number is encoded in such a manner as to simplify the hardware used to add, subtract and compare integers.

A simple method of thinking about two’s complement numbers is to negate the place value of the **MSB**. For example, the number one is represented the same as discussed before:

```

-128 64 32 16  8  4  2  1

```

```
0 0 0 0 0 0 0 1
```

The **MSB** of any negative number in this format will always be 1. For example the value -1_{10} is:

```
-128 64 32 16 8 4 2 1
  1  1  1  1  1  1  1  1
```

... because: $-128 + 64 + 32 + 16 + 8 + 4 + 2 + 1 = -1$.

This format has the virtue of allowing the same addition logic discussed above to be used to calculate $-1 + 1 = 0$.

```
-128 64 32 16 8 4 2 1 <== place value
  1  1  1  1  1  1  1  0 <== carries
    1  1  1  1  1  1  1  1 <== addend (-1)
+  0  0  0  0  0  0  0  1 <== addend (1)
-----
  1  0  0  0  0  0  0  0 <== sum (0 with an overflow)
```

In order for this to work, the **overflow** carry out of the sum of the MSBs is ignored.

2.2.3.1 Converting between Positive and Negative

Changing the sign on two's complement numbers can be described as inverting all of the bits (which is also known as the one's complement) and then add one.

For example, inverting the number *four*:

```
-128 64 32 16 8 4 2 1
  0  0  0  0  0  1  0  0 <== 4

                1  1    <== carries
  1  1  1  1  1  0  1  1 <== one's complement of 4
+  0  0  0  0  0  0  0  1 <== plus 1
-----
  1  1  1  1  1  1  0  0 <== -4
```

This can be verified by adding 5 to the result and observe that the sum is 1:

```
-128 64 32 16 8 4 2 1
  1  1  1  1  1      <== carries
  1  1  1  1  1  1  0  0 <== -4
+  0  0  0  0  0  1  0  1 <== 5
-----
  1  0  0  0  0  0  0  1
```

Note that the changing of the sign using this method is symmetric in that it is identical when converting from negative to positive and when converting from positive to negative: flip the bits and add 1.

For example, changing the value -4 to 4 to illustrate the reverse of the conversion above:

```

-128 64 32 16 8 4 2 1
  1  1  1  1  1  1  0  0 <== -4

                1  1    <== carries
  0  0  0  0  0  0  1  1 <== one's complement of -4
+ 0  0  0  0  0  0  0  1 <== plus 1
-----
  0  0  0  0  0  1  0  0 <== 4

```

2.2.4 Subtraction of Binary Numbers

Subtraction of binary numbers is performed by first negating the subtrahend and then adding the two numbers. Due to the nature of two's complement numbers this will work for both signed and unsigned numbers.

To calculate $-4 - 8 = -12$

```

-128 64 32 16 8 4 2 1
  1  1  1  1  1  1  0  0 <== -4
- 0  0  0  0  1  0  0  0 <== 8

                1  1  1    <== carries
  1  1  1  1  0  1  1  1 <== one's complement of -8
+ 0  0  0  0  0  0  0  1 <== plus 1
-----
  1  1  1  1  1  0  0  0 <== -8

  1  1  1  1                <== carries
  1  1  1  1  1  1  0  0 <== -4
+ 1  1  1  1  1  0  0  0 <== -8
-----
  1  1  1  1  1  0  1  0 0 <== -12

```

Fix Me:

This section needs more examples of subtracting signed and unsigned numbers and a discussion on how signedness is not relevant until the results are interpreted. For example adding $-4 + -8 = -12$ using two 8-bit numbers is the same as adding $252 + 248 = 500$ and truncating the result to 244.

2.2.5 Truncation and Overflow

Discuss the details of truncation and overflow here.

Fix Me:

This chapter should be made consistent in its use of truncation and overflow as occur with signed and unsigned addition and subtraction.

2.3 Main Memory Storage

When transferring data between its registers registers and main memory a RISC-V system uses the little-endian byte order.¹

Fix Me:

Refactor this section and the memory discussion in RV32 reference chapter

Fix Me:

Discuss byte ordering, addressing and character strings.

¹ See[13] for some history of the big/little-endian “controversy.”

2.3.1 Memory Dump

Introduce the memory dump and how to read them here.

Discuss the pitfalls of assuming what a set of bytes is used for based on their contents!

2.3.2 Big Endian Representation

Using the memory dump contents in prior section, discuss how big endian values are stored.

2.3.3 Little Endian Representation

Using the memory dump contents in prior section, discuss how little endian values are stored.

2.3.4 Character Strings and Arrays

Define character strings and arrays.

Using the prior memory dump, discuss how and where things are stored and retrieved.

2.3.5 Alignment

Draw a diagram showing the overlapping data types when they are all aligned.

2.3.6 Instruction Alignment

Every possible instruction that an RV32I CPU can execute contains exactly 32 bits. Therefore each one must be stored in four bytes of the main memory.

To simplify the hardware, each instruction must be placed into four adjacent bytes whose numeric address sequence begins with a multiple four. For example, an instruction might be located in bytes 4, 5, 6 and 7 (but not in 5, 6, 7 and 8 nor in 9, 3, 1, and 0...).

This sort of addressing requirement is common and is referred to as [alignment](#). An aligned instruction begins at a memory address that is a multiple of four. An *unaligned* instruction would be one beginning at any other address and is *illegal*.

An attempt to fetch an instruction from an unaligned address will result in an error referred to as an alignment [exception](#). This and other exceptions cause the CPU to stop executing the current instruction and start executing a different set of instructions that are prepared to handle the problem. Often an exception is handled by completely stopping the program in a way that is commonly referred to as a system or application *crash*.

Given a properly aligned instruction address, the CPU can request that the main memory locate and deliver the values of the four bytes in the address sequence to the CPU using what is called a memory read operation. Some systems can deliver four (or more) bytes at the same time while others might only

► Fix Me:

Rewrite this section for data rather than instructions and then note here that instructions must be naturally aligned. For RV32 that is on a 4-byte boundary

be capable of delivering one or two bytes at a time. These differences in hardware typically impact the cost and performance of a system.²

²The design and implementation choices that determine how any given system operates are part of what is called a system's *organization* and is beyond the scope of this text. See [3] for more information on computer organization.

Chapter 3

The Elements of a Assembly Language Program

3.1 Assembly Language Statements

Introduce the assembly language grammar. Statement = 1 line of text containing an instruction or directive.

Instruction = label, mnemonic, operands, comment.

Directive = Used to control the operation of the assembler.

3.2 Memory Layout

Is this a good place to introduce the text, data, bss, heap and stack regions?

Or does that belong in a new section/chapter that discusses addressing modes?

3.3 A Sample Program Source Listing

A simple program that illustrates how this text presents program source code is seen in [Listing 3.1](#). This program will place a zero in each of the 4 registers named x28, x29, x30 and x31.

Listing 3.1: `zero4regs.S`
Setting four registers to zero.

```
1  .text                # put this into the text section
2  .align 2             # align to 2^2
3  .globl _start
4  _start:
5      addi x28, x0, 0    # set register x28 to zero
6      addi x29, x0, 0    # set register x29 to zero
7      addi x30, x0, 0    # set register x30 to zero
8      addi x31, x0, 0    # set register x31 to zero
```


This program listing illustrates a number of things:

rvddt

- Listings are identified by the name of the file within which they are stored. This listing is from a file named: `zero4regs.S`.
- The assembly language programs discussed in this text will be saved in files that end with: `.S` (Alternately you can use `.sx` on systems that don't understand the difference between upper and lowercase letters.¹)
- A description of the listing's purpose appears under the name of the file. The description of [Listing 3.1](#) is *Setting four registers to zero*.
- The lines of the listing are numbered on the left margin for easy reference.
- An assembly program consists of lines of plain text.
- The RISC-V ISA does not provide an operation that will simply set a register to a numeric value. To accomplish our goal this program will add zero to zero and place the sum in each of the four registers.
- The lines that start with a dot `'.'` (on lines 1, 2 and 3) are called *assembler directives* as they tell the assembler itself how we want it to translate the following *assembly language instructions* into *machine language instructions*.
- Line 4 shows a *label* named `_start`. The colon at the end is the indicator to the assembler that causes it to recognize the preceding characters as a label.
- Lines 5-8 are the four assembly language instructions that make up the program. Each instruction in this program consists of four *fields*. (Different instructions can have a different number of fields.) The fields on line 5 are:

`addi` The instruction mnemonic. It indicates the operation that the CPU will perform.

`x28` The *destination* register that will receive the sum when the *addi* instruction is finished. The names of the 32 registers are expressed as `x0 – x31`.

`x0` One of the addends of the sum operation. (The `x0` register will always contain the value zero. It can never be changed.)

`0` The second addend is the number zero.

`# set ...` Any text anywhere in a RISC-V assembly language program that starts with the pound-sign is ignored by the assembler. They are used to place a *comment* in the program to help the reader better understand the motive of the programmer.

3.4 Running a Program With rvddt

To illustrate what a CPU does when it executes instructions this text will use the [rvddt](#) simulator to display shows sequence of events and the binary values involved. This simulator supports the RV32I ISA and has a configurable amount of memory.²

[Listing 3.2](#) shows the operation of the four *addi* instructions from [Listing 3.1](#) when it is executed in trace-mode.

¹The author of this text prefers to avoid using such systems.

²The *rvddt* simulator was written to generate the listings for this text. It is similar to the fancier *spike* simulator. Given the simplicity of the RV32I ISA, *rvddt* is less than 1700 lines of C++ and was written in one (long) afternoon.

Listing 3.2: zero4regs.out

Running a program with the rvddt simulator

```

1 [winans@w510 src]$ ./rvddt -f ../t1/load4regs.bin
2 Loading '../t1/load4regs.bin' to 0x0
3 ddt> t4
4   x0: 00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
5   x8: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
6  x16: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
7  x24: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
8   pc: 00000000
9 00000000: 00000e13 addi    x28, x0, 0    # x28 = 0x00000000 = 0x00000000 + 0x00000000
10  x0: 00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
11  x8: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
12  x16: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
13  x24: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0
14   pc: 00000004
15 00000004: 00000e93 addi    x29, x0, 0    # x29 = 0x00000000 = 0x00000000 + 0x00000000
16  x0: 00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
17  x8: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
18  x16: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
19  x24: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-00000000 00000000 f0f0f0f0 f0f0f0f0
20   pc: 00000008
21 00000008: 00000f13 addi    x30, x0, 0    # x30 = 0x00000000 = 0x00000000 + 0x00000000
22  x0: 00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
23  x8: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
24  x16: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
25  x24: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-00000000 00000000 00000000 f0f0f0f0
26   pc: 0000000c
27 0000000c: 00000f93 addi    x31, x0, 0    # x31 = 0x00000000 = 0x00000000 + 0x00000000
28 ddt> r
29  x0: 00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
30  x8: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
31  x16: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
32  x24: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-00000000 00000000 00000000 00000000
33   pc: 00000010
34 ddt> x
35 [winans@w510 src]$

```

- ℓ 1 This listing includes the command-line that shows how the simulator was executed to load a file containing the machine instructions (aka machine code) from the assembler.
- ℓ 2 A message from the simulator indicating that it loaded the machine code into simulated memory at address 0.
- ℓ 3 This line shows the prompt from the debugger and the command `t4` that the user entered to request that the simulator trace the execution of four instructions.
- ℓ 4-8 Prior to executing the first instruction, the state of the CPU registers is displayed.
- ℓ 4 The values in registers 0, 1, 2, 3, 4, 5, 6 and 7 are printed from left to right in [big endian](#), [hexadecimal](#) form. The dash ‘-’ character in the middle of the line is a reference to make it easier to visually navigate across the line without being forced to count the values from the far left when seeking the value of, say, x5.
- ℓ 5-7 The values of registers 8–31 are printed.
- ℓ 8 The *program counter* (pc) register is printed. It contains the address of the instruction that the CPU will execute. After each instruction, the pc will either advance four bytes ahead or be set to another value by a branch instruction as discussed above.
- ℓ 9 A four-byte instruction is fetched from memory at the address in the pc register, is decoded and printed. From left to right the fields shown on this line are:

00000000 The memory address from which the instruction was fetched. This address is displayed in [big endian](#), [hexadecimal](#) form.

00000e13 The machine code of the instruction displayed in [big endian](#), [hexadecimal](#) form.

addi The mnemonic for the machine instruction.

x28 The `rd` field of the `addi` instruction.

x0 The `rs1` field of the `addi` instruction that holds one of the two addends of the operation.

0 The `imm` field of the `addi` instruction that holds the second of the two addends of the operation.

... A simulator-generated comment that explains what the instruction is doing. For this instruction it indicates that `x28` will have the value zero stored into it as a result of performing the addition: $0 + 0$.

ℓ 10-14 These lines are printed as the prelude while tracing the second instruction. Lines 7 and 13 show that `x28` has changed from `f0f0f0f0` to `00000000` as a result of executing the first instruction and lines 8 and 14 show that the `pc` has advanced from zero (the location of the first instruction) to four, where the second instruction will be fetched. None of the rest of the registers have changed values.

ℓ 15 The second instruction decoded executed and described. This time register `x29` will be assigned a value.

ℓ 16-27 The third and fourth instructions are traced.

ℓ 28 Tracing has completed. The simulator prints its prompt and the user enters the ‘r’ command to see the register state after the fourth instruction has completed executing.

ℓ 29-33 Following the fourth instruction it can be observed that registers `x28`, `x29`, `x30` and `x31` have been set to zero and that the `pc` has advanced from zero to four, then eight, then 12 (the hex value for 12 is `c`) and then to 16 (which, in hex, is `10`).

ℓ 34 The simulator exit command ‘x’ is entered by the user and the terminal displays the shell prompt.

Chapter 4

Using The RISC-V GNU Toolchain

This chapter discusses using the GNU toolchain elements to experiment with the material in this book.

See [Appendix A](#) if you do not already have the GNU crosscompiler toolchain available on your system.

Discuss the choice of ilp32 as well as what the other variations would do.

Discuss rv32im and note that the details are found in [chapter 5](#).

Discuss installing and using one of the RISC-V simulators here.

Describe the pre-processor, compiler, assembler and linker.

Source, object, and binary files

Assembly syntax (label: mnemonic op1, op2, op3 # comment).

text, data, bss, stack

Labels and scope.

Forward & backward references to throw-away labels.

The entry address of an application.

.s file contain assembler code. .S (or .sx) files contain assembler code that must be preprocessed. [[14](#), p. 29]

Pre-processing conditional assembly using #if.

Building with `-mabi=ilp32 -march=rv32i -mno-fdiv -mno-div` to match the config options on the toolchain.

Linker scripts.

Makefiles

objdump

nm

hexdump -C

Chapter 5

RV32 Machine Instructions

5.1 Introduction

► Fix Me:

Discuss what the IMAFD, G and other ISA extensions mean as well as the 32, 64 and 128-bit versions.

5.2 Conventions and Terminology

When discussing instructions, the following abbreviations/notations are used:

5.2.1 XLEN

XLEN represents the bit-length of an `x` register in the machine architecture. Possible values are 32, 64 and 128.

5.2.2 `sx(val)`

Sign extend `val` to the left.

This is used to convert a signed integer value expressed using some number of bits to a larger number of bits by adding more bits to the left. In doing so, the sign will be preserved. In this case `val` represents the least **MSBs** of the value. For more on binary numbers see [Appendix B](#).

[Figure 5.1](#) illustrates extending the negative sign bit of `val` to the left by replicating it. When `val` is negative, its **MSB** (bit 19 in this example) will be set to 1. Extending this value to the left will set all the new bits to the left of it to 1 as well.

[Figure 5.2](#) illustrates extending the positive sign bit of `val` to the left by replicating it. When `val` is positive, its **MSB** will be set to 0. Extending this value to the left will set all the new bits to the left of it to 0 as well.




Figure 5.1: Sign-extending a negative integer from 20 bits to 32 bits.




Figure 5.2: Sign-extending a positive integer from 20 bits to 32 bits.

5.2.3 `zx(val)`

Zero extend *val* to the left.

This is used to convert an unsigned integer value expressed using some number of bits to a larger number of bits by adding more bits to the left. In doing so, the new bits added will all be set to zero. As is the case with `sx(val)`, *val* represents the **LSBs** of the final value. [Figure 5.3](#) illustrates zero-extending a 20-bit *val* to the left to form a 32-bit fullword.

For more on binary numbers see [Appendix B](#).




Figure 5.3: Zero-extending an unsigned integer from 20 bits to 32 bits.

5.2.4 `zr(val)`

Zero extend *val* to the right.

Some times a binary value is encoded such that a set of bits represented by *val* are used to represent the **MSBs** of some longer (more bits) value. In this case it is necessary to append zeros to the right to convert *val* to the longer value.

[Figure 5.4](#) illustrates converting a 20-bit *val* to a 32-bit fullword.

5.2.5 Sign Extended Left and Zero Extend Right

Some instructions such as the J-type (see [subsection 5.4.2](#)) include immediate operands that are extended in both directions.

[Figure 5.5](#) and [Figure 5.6](#) illustrates zero-extending a 20-bit negative number one bit to the right and




Figure 5.4: Zero-extending an integer to the right from 20 bits to 32 bits.

sign-extending it 11 bits to the left:




Figure 5.5: Sign-extending a positive 20-bit number 11 bits to the left and one bit to the right.




Figure 5.6: Sign-extending a negative 20-bit number 11 bits to the left and one bit to the right.

5.2.6 m8(addr)

The contents of an 8-bit value in memory at address *addr*.

Given the contents of the memory dump shown in [Figure 5.7](#), `m8(42)` refers to the memory location at address 42_{16} that currently contains the 8-bit value `fc16`.

The `mn(addr)` notation can be used to refer to memory that is being read or written depending on the context.

When memory is being written, the following notation is used to indicate that the least significant 8 bits of *source* will be written into memory at the address *addr*:

`m8(addr) ← source`

When memory is being read, the following notation is used to indicate that the 8 bit value at the address *addr* will be read and stored into *dest*:

`dest ← m8(addr)`

Note that *source* and *dest* are typically registers.

5.2.7 m16(addr)

The contents of an 16-bit little-endian value in memory at address *addr*.


```
00000030  2f 20 72 65 61 64 20 61 20 62 69 6e 61 72 79 20
00000040  66 69 fc 65 20 66 69 6c 6c 65 64 20 77 69 74 68
00000050  20 72 76 33 32 49 20 69 6e 73 74 72 75 63 74 69
00000060  6f 6e 73 20 61 6e 64 20 66 65 65 64 20 74 68 65
```

Figure 5.7: Sample memory contents.

Given the contents of the memory dump shown in [Figure 5.7](#), `m16(42)` refers to the memory location at address `4216` that currently contains `65fc16`. See also [subsection 5.2.6](#).

5.2.8 `m32(addr)`

The contents of an 32-bit little-endian value in memory at address *addr*.

Given the contents of the memory dump shown in [Figure 5.7](#), `m32(42)` refers to the memory location at address `4216` that currently contains `662065fc16`. See also [subsection 5.2.6](#).

5.2.9 `m64(addr)`

The contents of an 64-bit little-endian value in memory at address *addr*.

Given the contents of the memory dump shown in [Figure 5.7](#), `m64(42)` refers to the memory location at address `4216` that currently contains `656c6c69662065fc16`. See also [subsection 5.2.6](#).

5.2.10 `m128(addr)`

The contents of an 128-bit little-endian value in memory at address *addr*.

Given the contents of the memory dump shown in [Figure 5.7](#), `m128(42)` refers to the memory location at address `4216` that currently contains `7220687469772064656c6c69662065fc16`. See also [subsection 5.2.6](#).

5.2.11 `.+offset`

The address of the current instruction plus a numeric offset.

5.2.12 `.-offset`

The address of the current instruction minus a numeric offset.

5.2.13 `pc`

The current value of the program counter.

5.2.14 rd

An x-register used to store the result of instruction.

5.2.15 rs1

An x-register value used as a source operand for an instruction.

5.2.16 rs2

An x-register value used as a source operand for an instruction.

5.2.17 imm

An immediate numeric operand. The word *immediate* refers to the fact that the operand is stored within an instruction.

5.2.18 rsN[h:l]

The value of bits from h through l of x-register rsN. For example: rs1[15:0] refers to the contents of the 16 [LSBs](#) of rs1.

5.3 Addressing Modes

immediate, register, base-displacement, pc-relative

► Fix Me:
Write this section.

5.4 Instruction Encoding Formats

This document concerns itself with the following RISC-V instruction formats.

XXX Show and discuss a stack of formats explaining how the unnatural ordering of the *imm* fields reduces the number of possible locations that the hardware has to be prepared to *look* for various bits. For example, the opcode, rd, rs1, rs1, func3 and the sign bit (when used) are all always in the same position. Also note that imm[19:12] and imm[10:5] can only be found in one place. imm[4:0] can only be found in one of two places...

► Fix Me:
Should discuss types and sizes beyond the fundamentals. Will add if/when instruction details are added in the future.

The point to all this is that it is easier to build a machine if it does not have to accommodate many different ways to perform the same task. This simplification can also allow it operate faster.

[Figure 5.8](#) Shows the RISC-V instruction formats.





Figure 5.8: RISC-V instruction formats.

5.4.1 U Type


The U-Type format is used for instructions that use a 20-bit immediate operand and a destination register.



The **rd** field contains an **x** register number to be set to a value that depends on the instruction.

The **imm** field contains a 20-bit value that will be converted into **XLEN** bits by using the *imm* operand for bits 31:12 and then sign-extending it to the left¹ and zero-extending the LSBs as discussed in subsection 5.2.4.

If **XLEN**=32 then the **imm** value in this example will be converted as shown below.




Notice that the 20-bits of the **imm** field are mapped in the same order and in the same relative position

¹When **XLEN** is larger than 32.



If **XLEN**=64 then the imm value in this example will be converted to the same two's complement integer value by extending the sign to the left.

The J-type format is used for instructions that use a 20-bit immediate operand and a destination register. It is similar to the U-type. However, the immediate operand is constructed by arranging the *imm* bits in a different manner.




The example above shows that the bit positions in the *imm* field description. We see that the 20 *imm* bits are re-ordered according to: [20|10:1|11|19:12]. This means that the **MSB** of the *imm* field is to be placed into bit 20 of the immediate integer value ultimately used by the instruction when it is converted into **XLEN** bits. The next bit to the right in the *imm* field is to be placed into bit 10 of the immediate value and so on.

If **XLEN**=32 then the *imm* value in this example will be converted as shown below.



```
~/rvalp/book/./rv32/chapter.tex
v0.1-26-ge892461 2018-05-12 17:31:17 -0500
```



5.4.7 CPU Registers

The registers are names x0 through x31 and have aliases suited to their conventional use. The following table describes each register.

Note that the calling convention specifies that only some of the registers are to be saved by functions if they alter their contents. The idea being that accessing memory is time-consuming and that by classifying some registers as “temporary” (not saved by any function that alter its contents) it is possible to carefully implement a function with less need to store register values on the stack in order to use them to perform the operations of the function.

➡ Fix Me:
Need to add a section that discusses the calling conventions

The lack of grouping the temporary and saved registers is due to the fact that the C extension provides access to only the first 16 registers when executing instructions in the compressed format.

Reg	Alias	Description	Saved
x0	zero	Hard-wired zero	yes
x1	ra	Return address	
x2	sp	Stack pointer	
x3	gp	Global pointer	
x4	tp	Thread pointer	
x5	t0	Temporary/alternate link register	
x6	t1	Temporary	
x7	t2	Temporary	yes
x8	s0/fp	Saved register/frame pointer	
x9	s1	Saved register	
x10	a0	Function argument/return value	
x11	a1	Function argument/return value	
x12	a2	Function argument	
x13	a3	Function argument	
x14	a4	Function argument	
x15	a5	Function argument	
x16	a6	Function argument	
x17	a7	Function argument	
x18	s2	Saved register	yes
x19	s3	Saved register	yes
x20	s4	Saved register	yes
x21	s5	Saved register	yes
x22	s6	Saved register	yes
x23	s7	Saved register	yes
x24	s8	Saved register	yes
x25	s9	Saved register	yes
x26	s10	Saved register	yes
x27	s11	Saved register	yes
x28	t3	Temporary	
x29	t4	Temporary	
x30	t5	Temporary	
x31	t6	Temporary	

5.5 memory

Note that RISC-V is a little-endian machine.

RV32I
Instruction!LUI

If a RISC-V processor implements the C (compressed) extension then instructions may be aligned to 2-byte boundaries. [1, p. 68]

Data alignment is not necessary but unaligned data can be inefficient. Accessing unaligned data using any of the load or store instructions can also prevent a memory access from operating atomically. [1, p.19] See also ??.

5.6 RV32I Base Instruction Set

RV32I refers to the basic 32-bit integer instructions.

5.6.1 LUI rd, imm


Load Upper Immediate.

$$\text{rd} \leftarrow \text{zr}(\text{imm})$$

Copy the immediate value into bits 31:12 of the destination register and place zeros into bits 11:0. When XLEN is 64 or 128, the immediate value is sign-extended to the left.


Instruction Format and Example:

LUI t0, 3



```
00010074: 000032b7 lui      x5, 0x3          // x5 = 0x3000
reg  0: 00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 00003000 f0f0f0f0 f0f0f0f0
reg  8: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
reg 16: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
reg 24: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
pc: 00010078
```

```
LUI t0, 0xffff
```



```
00010078: ffff2b7 lui      x5, 0xfffff // x5 = 0xfffff000
reg  0: 00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 ffff000 f0f0f0f0 f0f0f0f0
reg  8: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
reg 16: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
reg 24: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
      pc: 0001007c
```

5.6.2 AUIPC rd, imm

Instruction!AUIPC
Instruction!JAL


Add Upper Immediate to PC.

$$\text{rd} \leftarrow \text{pc} + \text{zr}(\text{imm})$$

Create a signed 32-bit value by zero-extending `imm[31:12]` to the right (see [subsection 5.2.4](#)) and add this value to the `pc` register, placing the result into `rd`.


When XLEN is 64 or 128, the immediate value is also sign-extended to the left prior to being added to the `pc` register.

AUIPC t0, 3



```
0001007c: 00003297  auipc    x5, 0x3           // x5 = 0x1307c = 0x1007c + 0x3000
reg  0: 00000000  f0f0f0f0  f0f0f0f0  f0f0f0f0-f0f0f0f0 0001307c f0f0f0f0 f0f0f0f0
reg  8: f0f0f0f0  f0f0f0f0  f0f0f0f0  f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
reg 16: f0f0f0f0  f0f0f0f0  f0f0f0f0  f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
reg 24: f0f0f0f0  f0f0f0f0  f0f0f0f0  f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
pc: 00010080
```

AUIPC t0, 0x81000



```
00010080: 81000297  auipc    x5, 0x81000    // x5 = 0x81010080 = 0x10080 + 0x81000000
reg  0: 00000000  f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 81010080 f0f0f0f0 f0f0f0f0
reg  8: f0f0f0f0  f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
reg 16: f0f0f0f0  f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
reg 24: f0f0f0f0  f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
      pc: 00010084
```

The AUIPC instruction supports two-instruction sequences to access arbitrary offsets from the PC for both control-flow transfers and data accesses. The combination of an AUIPC and the 12-bit immediate in a JALR can transfer control to any 32-bit PC-relative address, while an AUIPC plus the 12-bit immediate offset in regular load or store instructions can access any 32-bit PC-relative data address. [1, p. 14]

5.6.3 JAL rd, imm

Jump and link.


```
rd ← pc + 4
pc ← pc + sx(imm<<1)
```


This instruction saves the address of the next instruction that would otherwise execute (located at `pc+4`) into `rd` and then adds immediate value to the `pc` causing an unconditional branch to take place.

The standard software conventions for calling subroutines use `x1` as the return address (`rd` register in this case). [1, p. 16]

Encoding:

JAL `x7`, `+.16`



imm demultiplexed value = $0000000000000001000_2 \ll 1 = 16_{10}$

State of registers before execution:


`pc` = 0x11114444

State of registers after execution:

`pc` = 0x11114454 `x7` = 0x11114448

JAL provides a method to call a subroutine using a pc-relative address.

JAL `x7`, `-.16`



imm demultiplexed value = $1111111111111111000_2 \ll 1 = -16_{10}$

State of registers before execution:

`pc` = 0x11114444

State of registers after execution:

`pc` = 0x11114434 `x7` = 0x11114448

5.6.4 JALR `rd`, `rs1`, `imm`

Jump and link register.

`rd` \leftarrow `pc` + 4

`pc` \leftarrow (`rs1` + `sx(imm)`) & ~1

This instruction saves the address of the next instruction that would otherwise execute (located at `pc+4`) into `rd` and then adds the immediate value to the `rs1` register and stores the sum into the `pc` register causing an unconditional branch to take place.


Note that the branch target address is calculated by sign-extending the `imm[11:0]` bits from the instruc-

tion, adding it to the `rs1` register and *then* the LSB of the sum is to zero and the result is stored into the `pc` register. The discarding of the LSB allows the branch to refer to any even address. Instruction!BEQ

The standard software conventions for calling subroutines use `x1` as the return address (`rd` register in this case). [1, p. 16]

Encoding:

JALR `x1, x7, 4`



Before:


`pc` = 0x11114444
`x7` = 0x44444444

After

`pc` = 0x5555888c
`x1` = 0x11114448

JALR provides a method to call a subroutine using a base-displacement address.

JALR `x1, x0, 5`



Note that the least significant bit in the result of `rs1+imm` is discarded/set to zero before the result is saved in the `pc`.

`pc` = 0x11114444

After

`pc` = 0x00000004
`x1` = 0x11114448


5.6.5 BEQ `rs1, rs2, imm`

Branch if equal.

`pc` \leftarrow (`rs1` == `rs2`) ? `pc+sx(imm[12:1]<<1)` : `pc+4`

Encoding:

BEQ `x3, x15, 2064`



Instruction!BNE
Instruction!BLT

$\text{imm}[12:1] = 010000001000_2 = 1032_{10}$

$\text{imm} = 2064_{10}$

$\text{funct3} = 000_2$

$\text{rs1} = \text{x3}$

$\text{rs2} = \text{x15}$


5.6.6 BNE rs1, rs2, imm

Branch if Not Equal.

$\text{pc} \leftarrow (\text{rs1} \neq \text{rs2}) ? \text{pc} + \text{sx}(\text{imm}[12:1] \ll 1) : \text{pc} + 4$

Encoding:

BNE x3, x15, 2064



$\text{imm}[12:1] = 010000001000_2 = 1032_{10}$

$\text{imm} = 2064_{10}$

$\text{funct3} = 001_2$

$\text{rs1} = \text{x3}$

$\text{rs2} = \text{x15}$


5.6.7 BLT rs1, rs2, imm

Branch if Less Than.

$\text{pc} \leftarrow (\text{rs1} < \text{rs2}) ? \text{pc} + \text{sx}(\text{imm}[12:1] \ll 1) : \text{pc} + 4$

Encoding:

BLT x3, x15, 2064



$\text{imm}[12:1] = 010000001000_2 = 1032_{10}$

$\text{imm} = 2064_{10}$

$\text{funct3} = 100_2$

$\text{rs1} = \text{x3}$

$\text{rs2} = \text{x15}$

5.6.8 BGE rs1, rs2, imm

Instruction!BGE
Instruction!BLTU
Instruction!BGEU

Branch if Greater or Equal.

$pc \leftarrow (rs1 \geq rs2) ? pc + sx(imm[12:1] \ll 1) : pc + 4$

Encoding:

BGE x3, x15, 2064

imm[12:10:5]					rs2			rs1			funct3		imm[4:1:11]		opcode					
0	0	0	0	0	0	0	1	1	1	1	0	0	1	1	1	0	0	0	1	1
7					5			5			3		5		7					

B-type

$imm[12:1] = 010000001000_2 = 1032_{10}$

$imm = 2064_{10}$

$funct3 = 101_2$

$rs1 = x3$

$rs2 = x15$

5.6.9 BLTU rs1, rs2, imm

Branch if Less Than Unsigned.

$pc \leftarrow (rs1 < rs2) ? pc + sx(imm[12:1] \ll 1) : pc + 4$

Encoding:

BLTU x3, x15, 2064

312524201915141211760																				
imm[12:10:5]							rs2			rs1			funct3		imm[4:1:11]		opcode			
0000000							01111			00011			1110		100001		11100011			
7							5			5			3		5		7			

B-type

$imm[12:1] = 010000001000_2 = 1032_{10}$

$imm = 2064_{10}$

$funct3 = 110_2$

$rs1 = x3$

$rs2 = x15$


5.6.10 BGEU rs1, rs2, imm

Branch if Greater or Equal Unsigned.

$pc \leftarrow (rs1 \geq rs2) ? pc + sx(imm[12:1] \ll 1) : pc + 4$

Encoding:

BGEU x3, x15, 2064



Instruction!LB
Instruction!LH
Instruction!LW

► Fix Me:

use symbols in branch examples

$\text{imm}[12:1] = 010000001000_2 = 1032_{10}$

$\text{imm} = 2064_{10}$

$\text{funct3} = 111_2$

$\text{rs1} = \text{x3}$

$\text{rs2} = \text{x15}$

5.6.11 LB rd, imm(rs1)

Load byte.


$\text{rd} \leftarrow \text{sx}(\text{m8}(\text{rs1} + \text{sx}(\text{imm})))$

$\text{pc} \leftarrow \text{pc} + 4$

Load an 8-bit value from memory at address $\text{rs1} + \text{imm}$, then sign-extend it to 32 bits before storing it in rd

Encoding:

LB x7, 4(x3)



5.6.12 LH rd, imm(rs1)

Load halfword.


$\text{rd} \leftarrow \text{sx}(\text{m16}(\text{rs1} + \text{sx}(\text{imm})))$

$\text{pc} \leftarrow \text{pc} + 4$

Load a 16-bit value from memory at address $\text{rs1} + \text{imm}$, then sign-extend it to 32 bits before storing it in rd

Encoding:

LH x7, 4(x3)



5.6.13 LW rd, imm(rs1)

Load word.


$rd \leftarrow sx(m32(rs1+sx(imm)))$
 $pc \leftarrow pc+4$

Instruction!LBU
 Instruction!LHU

Load a 32-bit value from memory at address $rs1+imm$, then store it in rd

Encoding:

LW x7, 4(x3)



5.6.14 LBU rd, imm(rs1)


Load byte unsigned.

$rd \leftarrow zx(m8(rs1+sx(imm)))$
 $pc \leftarrow pc+4$

Load an 8-bit value from memory at address $rs1+imm$, then zero-extend it to 32 bits before storing it in rd

Encoding:

LBU x7, 4(x3)



5.6.15 LHU rd, imm(rs1)


Load halfword unsigned.

$rd \leftarrow zx(m16(rs1+sx(imm)))$
 $pc \leftarrow pc+4$

Load an 16-bit value from memory at address $rs1+imm$, then zero-extend it to 32 bits before storing it in rd

Encoding:

LHU x7, 4(x3)



5.6.16 SB rs2, imm(rs1)

Instruction!SB
 Instruction!SH
 Instruction!SW


Store Byte.

```
m8(rs1+sx(imm)) ← rs2[7:0]
pc ← pc+4
```

Store the 8-bit value in `rs2[7:0]` into memory at address `rs1+imm`.

Encoding:

SB x3, 19(x15)

**5.6.17 SH rs2, imm(rs1)**


Store Halfword.

```
m16(rs1+sx(imm)) ← rs2[15:0]
pc ← pc+4
```

Store the 16-bit value in `rs2[15:0]` into memory at address `rs1+imm`.

Encoding:

SH x3, 19(x15)

**5.6.18 SW rs2, imm(rs1)**


Store Word

```
m16(rs1+sx(imm)) ← rs2[31:0]
pc ← pc+4
```

Store the 32-bit value in `rs2` into memory at address `rs1+imm`.

Encoding:

SW x3, 19(x15)



Show pos & neg imm examples.

Instruction!ADDI
Instruction!SLTI

5.6.19 ADDI rd, rs1, imm


Add Immediate

$rd \leftarrow rs1 + sx(imm)$

$pc \leftarrow pc + 4$

Encoding:

ADDI x1, x7, 4



Before:

$x7 = 0x11111111$

After:

$x1 = 0x11111115$

5.6.20 SLTI rd, rs1, imm

Set LessThan Immediate


$rd \leftarrow (rs1 < sx(imm)) ? 1 : 0$

$pc \leftarrow pc + 4$

If the sign-extended immediate value is less than the value in the **rs1** register then the value 1 is stored in the **rd** register. Otherwise the value 0 is stored in the **rd** register.

Encoding:

SLTI x1, x7, 4



Before:

$x7 = 0x11111111$

After:

$x1 = 0x00000000$

5.6.21 SLTIU rd, rs1, immInstruction!SLTIU
Instruction!XORI

Set LessThan Immediate Unsigned

$$rd \leftarrow (rs1 < sx(imm)) ? 1 : 0$$

$$pc \leftarrow pc+4$$

If the sign-extended immediate value is less than the value in the **rs1** register then the value 1 is stored in the **rd** register. Otherwise the value 0 is stored in the **rd** register. Both the immediate and **rs1** register values are treated as unsigned numbers for the purposes of the comparison.³

Encoding:

SLTIU x1, x7, 4

31												20		19		15			14		12		11		7		6		0			
imm[11:0]												rs1					funct3			rd					opcode							
0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	1	1	1	0	1	1	0	0	0	0	1	0	0	1	0	0	1	1
12												5					3			5					7							

I-type

I-type

Before:

x7 = 0x81111111

After:

x1 = 0x00000001

5.6.22 XORI rd, rs1, imm

Exclusive Or Immediate

$$rd \leftarrow rs1 \wedge sx(imm)$$

$$pc \leftarrow pc+4$$

The logical XOR of the sign-extended immediate value and the value in the **rs1** register is stored in the **rd** register.

Encoding:

XORI x1, x7, 4

31												20		19		15			14		12		11		7		6		0		
imm[11:0]												rs1					funct3			rd					opcode						
0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1	0	0	0	0	0	0	1	0	0	1	0	0	1	1
12												5					3			5					7						

I-type

I-type

Before:

x7 = 0x81111111

After:

x1 = 0x81111115

³The immediate value is first sign-extended to XLEN bits then treated as an unsigned number.[1, p. 14]

5.6.23 ORI rd, rs1, immInstruction!ORI
Instruction!ANDI

Or Immediate


$$\text{rd} \leftarrow \text{rs1} \mid \text{sx}(\text{imm})$$

$$\text{pc} \leftarrow \text{pc} + 4$$

The logical OR of the sign-extended immediate value and the value in the **rs1** register is stored in the **rd** register.

Encoding:

ORI x1, x7, 4



Before:

x7 = 0x81111111

After:

x1 = 0x81111115

5.6.24 ANDI rd, rs1, imm

And Immediate


$$\text{rd} \leftarrow \text{rs1} \& \text{sx}(\text{imm})$$

$$\text{pc} \leftarrow \text{pc} + 4$$

The logical AND of the sign-extended immediate value and the value in the **rs1** register is stored in the **rd** register.

Encoding:

ANDI x1, x7, 4



Before:

x7 = 0x81111111

After:

x1 = 0x81111115

5.6.25 SLLI rd, rs1, shamt

Instruction!SLLI
 Instruction!SRLI
 Instruction!SRAI

Shift Left Logical Immediate

$rd \leftarrow rs1 \ll shamt$
 $pc \leftarrow pc+4$

SLLI is a logical left shift operation (zeros are shifted into the lower bits). The value in rs1 shifted left shamt number of bits and the result placed into rd. [1, p. 14]

Encoding:

SLLI x7, x3, 2

312524201915141211760																					
funct7							shamt			rs1			funct3		rd			opcode			
0000000							00010			00011			001		00111			0100011			
7							5			5			3		5			7			

R-type

R-type

x3 = 0x81111111

After:

x7 = 0x04444444

5.6.26 SRLI rd, rs1, shamt

Shift Right Logical Immediate

$rd \leftarrow rs1 \gg shamt$
 $pc \leftarrow pc+4$

SRLI is a logical right shift operation (zeros are shifted into the higher bits). The value in rs1 shifted right shamt number of bits and the result placed into rd. [1, p. 14]

Encoding:

SRLI x7, x3, 2

312524201915141211760																							
funct7							shamt			rs1			funct3			rd			opcode				
0000000							00010			00011			10100111			0010011							
7							5			5			3			5			7				R-type

R-type

x3 = 0x81111111

After:

x7 = 0x20444444

5.6.27 SRAI rd, rs1, shamt

Shift Right Arithmetic Immediate

$rd \leftarrow rs1 \gg shamt$
 $pc \leftarrow pc+4$

Instruction!ADD
 Instruction!SUB

SRAI is a logical right shift operation (zeros are shifted into the higher bits). The value in rs1 shifted right shamt number of bits and the result placed into rd. [1, p. 14]

Encoding:

SRAI x7, x3, 2

312524201915141211760							R-type
funct7shamtrs1funct3rdopcode							
010000000000001000001110011100100011							
755357							

R-type

x3 = 0x81111111

After:

x7 = 0xe0444444

5.6.28 ADD rd, rs1, rs2

Add

$rd \leftarrow rs1 + rs2$
 $pc \leftarrow pc+4$

ADD performs addition. Overflows are ignored and the low 32 bits of the result are written to rd. [1, p. 15]

Encoding:

ADD x7, x3, x31

312524201915141211760																							
funct7							rs2			rs1			funct3			rd			opcode				
0000000							11111			00011			000			00111			0110011				
7							5			5			3			5			7				
																							R-type

R-type

x3 = 0x81111111 x31 = 0x22222222

After:

x7 = 0xa3333333

5.6.29 SUB rd, rs1, rs2

Subtract

$rd \leftarrow rs1 - rs2$
 $pc \leftarrow pc+4$


SUB performs subtraction. Underflows are ignored and the low 32 bits of the result are written to rd. [1,

p. 15]

Instruction!SLL
Instruction!SLT

Encoding:

SUB x7, x3, x31



x3 = 0x83333333 x31 = 0x01111111

After:

x7 = 0x82222222

5.6.30 SLL rd, rs1, rs2


Shift Left Logical

 $rd \leftarrow rs1 \ll rs2$ $pc \leftarrow pc+4$

SLL performs a logical left shift on the value in register rs1 by the shift amount held in the lower 5 bits of register rs2. [1, p. 15]

Encoding:

SLL x7, x3, x31



x3 = 0x83333333

x31 = 0x00000002

After:

x7 = 0x0ccccccc

5.6.31 SLT rd, rs1, rs2

Set Less Than

 $rd \leftarrow (rs1 < rs2) ? 1 : 0$ $pc \leftarrow pc+4$ SLT performs a signed compare, writing 1 to rd if $rs1 < rs2$, 0 otherwise. [1, p. 15]

Encoding:

31	25	24	20	19	15	14	12	11	7	6	0										
funct7							rs2			rs1			funct3		rd			opcode			
0	0	0	0	0	0	0	1	1	1	1	0	0	0	1	1	1	0	0	1	1	1
7							5			5			3		5			7			

R-type

Instruction!SRL
Instruction!SRA

x3 = 0x83333333

x31 = 0x1888ffff

After:

x7 = 0x9bbccccc

5.6.34 SRL rd, rs1, rs2

Shift Right Logical

 $rd \leftarrow rs1 \gg rs2$ $pc \leftarrow pc+4$

SRL performs a logical right shift on the value in register rs1 by the shift amount held in the lower 5 bits of register rs2. [1, p. 15]

Encoding:

SRL x7, x3, x31

31	25	24	20	19	15	14	12	11	7	6	0										
funct7							rs2			rs1			funct3			rd			opcode		
0	0	0	0	0	0	0	1	1	1	1	0	0	0	1	1	1	0	0	1	1	1
7							5			5			3			5			7		

R-type

x3 = 0x83333333

x31 = 0x00000010

After:

x7 = 0x00008333

5.6.35 SRA rd, rs1, rs2

Shift Right Arithmetic

 $rd \leftarrow rs1 \gg rs2$ $pc \leftarrow pc+4$

SRA performs an arithmetic right shift (the original sign bit is copied into the vacated upper bits) on the value in register rs1 by the shift amount held in the lower 5 bits of register rs2. [1, p. 14, 15]

Encoding:


SRA x7, x3, x31

Instruction!OR
Instruction!AND

x7 = 0xffff8333

x7 = 0x83333773

AND x7, x3, x31



Instruction!FENCE
Instruction!FENCE.I

x3 = 0x83333333
x31 = 0x00000fe2

After:

x7 = 0x00000322

5.6.38 FENCE predecessor, successor

The FENCE instruction is used to order device I/O and memory accesses as viewed by other RISC-V harts and external devices or co-processors. Any combination of device input (I), device output (O), memory reads (R), and memory writes (W) may be ordered with respect to any combination of the same. Informally, no other RISC-V hart or external device can observe any operation in the successor set following a FENCE before any operation in the predecessor set preceding the FENCE. The execution environment will define what I/O operations are possible, and in particular, which load and store instructions might be treated and ordered as device input and device output operations respectively rather than memory reads and writes. For example, memory-mapped I/O devices will typically be accessed with uncached loads and stores that are ordered using the I and O bits rather than the R and W bits. Instruction-set extensions might also describe new coprocessor I/O instructions that will also be ordered using the I and O bits in a FENCE. [1, p. 21]

► Fix Me:


Which of the i, o, r and w goes into each bit? See what gas does.

Operation:

$pc \leftarrow pc+4$

Encoding:

FENCE iorw, iorw



5.6.39 FENCE.I

The FENCE.I instruction is used to synchronize the instruction and data streams. RISC-V does not guarantee that stores to instruction memory will be made visible to instruction fetches on the same RISC-V hart until a FENCE.I instruction is executed. A FENCE.I instruction only ensures that a subsequent instruction fetch on a RISC-V hart will see any previous data stores already visible to the same RISC-V hart. FENCE.I does not ensure that other RISC-V harts' instruction fetches will observe the local hart's stores in a multiprocessor system. To make a store to instruction memory visible to all RISC-V harts, the writing hart has to execute a data FENCE before requesting that all remote RISC-V harts execute a FENCE.I. [1, p. 21]


Operation:

$pc \leftarrow pc+4$

Encoding:

Instruction!ECALL
 Instruction!EBREAK
 Instruction!CSRW


FENCE.I



5.6.40 ECALL

The ECALL instruction is used to make a request to the supporting execution environment, which is usually an operating system. The ABI for the system will define how parameters for the environment request are passed, but usually these will be in defined locations in the integer register file. [1, p. 24]


ECALL



5.6.41 EBREAK

The EBREAK instruction is used by debuggers to cause control to be transferred back to a debugging environment. [1, p. 24]


EBREAK



5.6.42 CSRRW rd, csr, rs1

The CSRRW (Atomic Read/Write CSR) instruction atomically swaps values in the CSRs and integer registers. CSRRW reads the old value of the CSR, zero-extends the value to XLEN bits, then writes it to integer register rd. The initial value in rs1 is written to the CSR. If rd=x0, then the instruction shall not read the CSR and shall not cause any of the side-effects that might occur on a CSR read. [1, p. 22]

CSRRW x3, 2, x15



5.6.43 CSRRS rd, csr, rs1

Instruction!CSRWS
Instruction!CSRRC
Instruction!CSRWI

The CSRRS (Atomic Read and Set Bits in CSR) instruction reads the value of the CSR, zero-extends the value to XLEN bits, and writes it to integer register rd. The initial value in integer register rs1 is treated as a bit mask that specifies bit positions to be set in the CSR. Any bit that is high in rs1 will cause the corresponding bit to be set in the CSR, if that CSR bit is writable. Other bits in the CSR are unaffected (though CSRs might have side effects when written). [1, p. 22]

If rs1=x0, then the instruction will not write to the CSR at all, and so shall not cause any of the side effects that might otherwise occur on a CSR write, such as raising illegal instruction exceptions on accesses to read-only CSRs. Note that if rs1 specifies a register holding a zero value other than x0, the instruction will still attempt to write the unmodified value back to the CSR and will cause any attendant side effects. [1, p. 22]

CSRRS x3, 2, x15

31												20	19				15				14		12			11		7			6	0				
csr													rsl				funct3						rd					opcode								
0 0 0 0 0 0 0 0 0 0 1 0													0 1 1 1 1				0 1 0						0 0 0 1 1					1 1 1 0 0 1 1								
12													5				3						5					7								

I-type

5.6.44 CSRRC rd, csr, rs1

The CSRRC (Atomic Read and Clear Bits in CSR) instruction reads the value of the CSR, zero-extends the value to XLEN bits, and writes it to integer register rd. The initial value in integer register rs1 is treated as a bit mask that specifies bit positions to be cleared in the CSR. Any bit that is high in rs1 will cause the corresponding bit to be cleared in the CSR, if that CSR bit is writable. Other bits in the CSR are unaffected. [1, p. 22]

If rs1=x0, then the instruction will not write to the CSR at all, and so shall not cause any of the side effects that might otherwise occur on a CSR write, such as raising illegal instruction exceptions on accesses to read-only CSRs. Note that if rs1 specifies a register holding a zero value other than x0, the instruction will still attempt to write the unmodified value back to the CSR and will cause any attendant side effects. [1, p. 22]

CSRRC x3, 2, x15

31												20	19				15				14		12			11		7				6		0			
csr													rsl				funct3						rd					opcode									
0 0 0 0 0 0 0 0 0 0 0 0												0	0 1 1 1 1				0 1 1				0		0 0 0 1 1			1 1		1 1 1 0 0 1 1									
12													5				3						5					7									

I-type

5.6.45 CSRWRI rd, csr, imm

This instruction is the same as CSRRW except a 5-bit unsigned (zero-extended) immediate value is used rather than the value from a register.

CSRROI x3, 2, 7

31												20	19	15			14	12		11	7			6	0							
csr													zimm				funct3			rd				opcode								
0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1	0	1	0	0	0	1	1	1	1	1	0	0	1	1	
12													5				3			5				7								

I-type


5.6.46 CSRRSI rd, csr, rs1

Instruction!CSRRSI
 Instruction!CSRRCI
 RV32M
 Instruction!MUL

This instruction is the same as CSRRS except a 5-bit unsigned (zero-extended) immediate value is used rather than the value from a register.

If the uimm[4:0] field is zero, then this instruction will not write to the CSR, and shall not cause any of the side effects that might otherwise occur on a CSR write. For CSRRWI, if rd=x0, then the instruction shall not read the CSR and shall not cause any of the side-effects that might occur on a CSR read. [1, p. 22]


CSRRSI x3, 2, 7

**5.6.47 CSRRCI rd, csr, rs1**

This instruction is the same as CSRRC except a 5-bit unsigned (zero-extended) immediate value is used rather than the value from a register.

If the uimm[4:0] field is zero, then this instruction will not write to the CSR, and shall not cause any of the side effects that might otherwise occur on a CSR write. For CSRRWI, if rd=x0, then the instruction shall not read the CSR and shall not cause any of the side-effects that might occur on a CSR read. [1, p. 22]

CSRRCI x3, 2, 7


**5.7 RV32M Standard Extension**

32-bit integer multiply and divide instructions.

5.7.1 MUL rd, rs1, rs2

Multiply **rs1** by **rs2** and store the least significant 32-bits of the result in **rd**.

MUL x7, x3, x31



5.7.2 MULH rd, rs1, rs2

Instruction!MULH
 Instruction!MULHS
 Instruction!MULHU
 Instruction!DIV
 Instruction!DIVU

MULH x7, x3, x31

31							25	24	20				19	15				14	12		11	7				6	0									
funct7									rs2					rs1					funct3			rd			opcode											
0	0	0	0	0	0	1		1	1	1	1	1	0	0	1	1	1	0	0	1	0	0	1	1	1	0	1	1	0	0	1	1				
7									5					5					3			5			7											

R-type

R-type

5.7.3 MULHS rd, rs1, rs2

MULHS x7, x3, x31

31	25	24	20	19	15	14	12	11	7	6	0																				
funct7							rs2			rs1			funct3			rd			opcode												
0	0	0	0	0	0	1	1	1	1	1	1	0	0	1	1	1	0	1	0	0	0	1	1	1	0	1	1	0	0	1	1
7							5			5			3			5			7												

R-type

R-type

5.7.4 MULHU rd, rs1, rs2

MULHU x7, x3, x31

312524201915141211760																											
funct7							rs2					rs1				funct3			rd				opcode				
0000001							11111					00111				011			00111				0110011				
7							5					5				3			5				7				

R-type

R-type

5.7.5 DIV rd, rs1, rs2

DIV x7, x3, x31

312524201915141211760																										
funct7							rs2				rs1				funct3			rd				opcode				
0000001							11111				00111				11000			000111				0110011				
7							5				5				3			5				7				

R-type

R-type

5.7.6 DIVU rd, rs1, rs2

DIVU x7, x3, x31

312524201915141211760																											
funct7							rs2					rs1				funct3			rd				opcode				
0000001							11111					00111				1101			00111				0110011				
7							5					5				3			5				7				


R-type

R-type

5.7.7 REM rd, rs1, rs2


REM x7, x3, x31

Instruction!REM
Instruction!REMU
RV32A
RV32F
RV32D



5.7.8 REMU rd, rs1, rs2

REMU x7, x3, x31



5.8 RV32A Standard Extension

32-bit atomic operations.

5.9 RV32F Standard Extension

32-bit IEEE floating point instructions.

5.10 RV32D Standard Extension

64-bit IEEE floating point instructions.

Appendix A

Installing a RISC-V Toolchain

A.1 The GNU Toolchain

Discuss the GNU toolchain elements used to experiment with the material in this book.

The instructions and examples here were all implemented on Ubuntu 16.04 LTS.

Install custom code in a location that will not cause interference with other applications and allow for easy cleanup. These instructions install the toolchain in `/usr/local/riscv`. At any time you can remove the lot and start over by executing the following command:

```
rm -rf /usr/local/riscv/*
```

Tested on Ubuntu 16.04 LTS. 18.04 was just released... update accordingly.

These are the only commands that you should perform as root when installing the toolchain:

```
sudo apt-get install autoconf automake autotools-dev curl libmpc-dev \
libmpfr-dev libgmp-dev gawk build-essential bison flex texinfo gperf \
libtool patchutils bc zlib1g-dev libexpat-dev
sudo mkdir -p /usr/local/riscv/
sudo chmod 777 /usr/local/riscv/
```

All other commands should be executed as a regular user. This will eliminate the possibility of clobbering system files that should not be touched when tinkering with the toolchain applications.

To download, compile and “install” the toolchain:

```
# riscv toolchain:
#
# https://riscv.org/software-tools/risc-v-gnu-compiler-toolchain/

git clone --recursive https://github.com/riscv/riscv-gnu-toolchain
cd riscv-gnu-toolchain
```

► Fix Me:

It would be good to find some Mac and Windows users to write and test proper variations on this section to address those systems. Pull requests, welcome!

```
./configure --prefix=/usr/local/riscv/rv32i --with-arch=rv32i --with-abi=ilp32  
make  
make install
```

Need to discuss augmenting the PATH environment variable.

Discuss the choice of ilp32 as well as what the other variations would do.

Discuss rv32im and note that the details are found in [chapter 5](#).

A.2 rvddt

Discuss installing the rvddt simulator here.

Floating Point Numbers

This section provides an overview of the IEEE-754 32-bit binary floating point format.

- ... 128 64 32 16 8 4 2 1

- ... 128 64 32 16 8 4 2 1 . 1/2 1/4 1/8 1/16 1/32 1/64 1/128 ...

- We use scientific notation as in 2.7×10^{-47} to express either small fractions or large numbers when we are not concerned every last digit needed to represent the entire, exact, value of a number.
- The format of a number in scientific notation is $mantissa \times base^{exponent}$
- In binary we have $mantissa \times 2^{exponent}$
- IEEE-754 format requires binary numbers to be *normalized* to $1.significand \times 2^{exponent}$ where the *significand* is the portion of the *mantissa* that is to the right of the binary-point.

- The unnormalized binary value of -2.625 is 10.101
- The normalized value of -2.625 is 1.0101×2^1

- [illegible]

- $$\bullet -((1 + \frac{1}{4} + \frac{1}{16}) \times 2^{128-127}) = -((1 + \frac{1}{4} + \frac{1}{16}) \times 2^1) = -(2 + \frac{1}{2} + \frac{1}{8}) = -(2 + .5 + .125) = -2.625$$

- IEEE754 formats:

	IEEE754 32-bit	IEEE754 64-bit
sign	1 bit	1 bit
exponent	8 bits (excess-127)	11 bits (excess-1023)
mantissa	23 bits	52 bits
max exponent	127	1023
min exponent	-126	-1022

- When the exponent is all ones, the mantissa is all zeros, and the sign is zero, the number represents positive infinity.
- When the exponent is all ones, the mantissa is all zeros, and the sign is one, the number represents negative infinity.
- Note that the binary representation of an IEEE754 number in memory can be compared for magnitude with another one using the same logic as for comparing two's complement signed integers because the magnitude of an IEEE number grows upward and downward in the same fashion as signed integers. This is why we use excess notation and locate the significand's sign bit on the left of the exponent.
- Note that zero is a special case number. Recall that a normalized number has an implied 1-bit to the left of the significand. . . which means that there is no way to represent zero! Zero is represented by an exponent of all-zeros and a significand of all-zeros. This definition allows for a positive and a negative zero if we observe that the sign can be either 1 or 0.
- On the number-line, numbers between zero and the smallest fraction in either direction are in the *underflow* areas.
- On the number line, numbers greater than the mantissa of all-ones and the largest exponent allowed are in the *overflow* areas.
- Note that numbers have a higher resolution on the number line when the exponent is smaller.

► Fix Me:

Need to add the standard lecture number-line diagram showing where the over/under-flow areas are and why.

B.1.1 Floating Point Number Accuracy

Due to the finite number of bits used to store the value of a floating point number, it is not possible to represent every one of the infinite values on the real number line. The following C programs illustrate this point.

B.1.1.1 Powers Of Two

Just like the integer numbers, the powers of two that have bits to represent them can be represented perfectly... as can their sums (provided that the significand requires no more than 23 bits.)

Listing B.1: powersoftwo.c

Precise Powers of Two

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4
5 union floatbin
6 {
7     unsigned int    i;
8     float          f;
9 };

```

```

10 int main()
11 {
12     union floatbin x;
13     union floatbin y;
14     int i;
15     x.f = 1.0;
16     while (x.f > 1.0/1024.0)
17     {
18         y.f = -x.f;
19         printf("%25.10f = %08x      %25.10f = %08x\n", x.f, x.i, y.f, y.i);
20         x.f = x.f/2.0;
21     }
22 }

```

Listing B.2: powersoftwo.out

Output from powersoftwo.c

```

1  1.0000000000 = 3f800000          -1.0000000000 = bf800000
2  0.5000000000 = 3f000000          -0.5000000000 = bf000000
3  0.2500000000 = 3e800000          -0.2500000000 = be800000
4  0.1250000000 = 3e000000          -0.1250000000 = be000000
5  0.0625000000 = 3d800000          -0.0625000000 = bd800000
6  0.0312500000 = 3d000000          -0.0312500000 = bd000000
7  0.0156250000 = 3c800000          -0.0156250000 = bc800000
8  0.0078125000 = 3c000000          -0.0078125000 = bc000000
9  0.0039062500 = 3b800000          -0.0039062500 = bb800000
10 0.0019531250 = 3b000000          -0.0019531250 = bb000000

```

B.1.1.2 Clean Decimal Numbers

When dealing with decimal values, you will find that they don't map simply into binary floating point values.

Note how the decimal numbers are not accurately represented as they get larger. The decimal number on line 10 of [Listing B.4](#) can be perfectly represented in IEEE format. However, a problem arises in the 11th loop iteration. It is due to the fact that the binary number can not be represented accurately in IEEE format. Its least significant bits were truncated in a best-effort attempt at rounding the value off in order to fit the value into the bits provided. This is an example of *low order truncation*. Once this happens, the value of `x.f` is no longer as precise as it could be given more bits in which to save its value.

Listing B.3: cleandecimal.c

Print Clean Decimal Numbers

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4
5  union floatbin
6  {
7      unsigned int i;
8      float f;
9  };
10 int main()
11 {
12     union floatbin x, y;
13     int i;
14
15     x.f = 10;
16     while (x.f <= 10000000000000.0)
17     {
18         y.f = -x.f;
19         printf("%25.10f = %08x      %25.10f = %08x\n", x.f, x.i, y.f, y.i);

```

```

20     x.f = x.f*10.0;
21 }
22 }

```

Listing B.4: cleandecimal.out

Output from cleandecimal.c

```

1      10.0000000000 = 41200000      -10.0000000000 = c1200000
2      100.0000000000 = 42c80000      -100.0000000000 = c2c80000
3      1000.0000000000 = 447a0000      -1000.0000000000 = c47a0000
4      10000.0000000000 = 461c4000      -10000.0000000000 = c61c4000
5      100000.0000000000 = 47c35000      -100000.0000000000 = c7c35000
6      1000000.0000000000 = 49742400      -1000000.0000000000 = c9742400
7      10000000.0000000000 = 4b189680      -10000000.0000000000 = cb189680
8      100000000.0000000000 = 4cbebc20      -100000000.0000000000 = ccbebc20
9      1000000000.0000000000 = 4e6e6b28      -1000000000.0000000000 = ce6e6b28
10     10000000000.0000000000 = 501502f9      -10000000000.0000000000 = d01502f9
11     99999997952.0000000000 = 51ba43b7      -99999997952.0000000000 = d1ba43b7
12     999999995904.0000000000 = 5368d4a5      -999999995904.0000000000 = d368d4a5
13     9999999827968.0000000000 = 551184e7      -9999999827968.0000000000 = d51184e7

```

B.1.1.3 Accumulation of Error

These rounding errors can be exaggerated when the number we multiply the `x.f` value by is, itself, something that can not be accurately represented in IEEE form.¹

For example, if we multiply our `x.f` value by $\frac{1}{10}$ each time, we can never be accurate and we start accumulating errors immediately.

Listing B.5: erroraccumulation.c

Accumulation of Error

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4
5  union floatbin
6  {
7      unsigned int    i;
8      float          f;
9  };
10 int main()
11 {
12     union floatbin  x, y;
13     int             i;
14
15     x.f = .1;
16     while (x.f <= 2.0)
17     {
18         y.f = -x.f;
19         printf("%25.10f = %08x      %25.10f = %08x\n", x.f, x.i, y.f, y.i);
20         x.f += .1;
21     }
22 }

```

Listing B.6: erroraccumulation.out

Output from erroraccumulation.c

```

1  0.1000000015 = 3dcccccd      -0.1000000015 = bdcccccd

```

¹Applications requiring accurate decimal values, such as financial accounting systems, can use a packed-decimal numeric format to avoid unexpected oddities caused by the use of binary numbers.

➡ Fix Me:

In a lecture one would show that one tenth is a repeating non-terminating binary number that gets truncated. This discussion should be reproduced here in text form.

2	0.2000000030 = 3e4ccccd	-0.2000000030 = be4ccccd
3	0.3000000119 = 3e99999a	-0.3000000119 = be99999a
4	0.4000000060 = 3ecccccd	-0.4000000060 = becccccd
5	0.5000000000 = 3f000000	-0.5000000000 = bf000000
6	0.6000000238 = 3f19999a	-0.6000000238 = bf19999a
7	0.7000000477 = 3f333334	-0.7000000477 = bf333334
8	0.8000000715 = 3f4ccccce	-0.8000000715 = bf4ccccce
9	0.9000000954 = 3f666668	-0.9000000954 = bf666668
10	1.0000001192 = 3f800001	-1.0000001192 = bf800001
11	1.1000001431 = 3f8ccccce	-1.1000001431 = bf8ccccce
12	1.2000001669 = 3f99999b	-1.2000001669 = bf99999b
13	1.3000001907 = 3fa66668	-1.3000001907 = bfa66668
14	1.4000002146 = 3fb33335	-1.4000002146 = bfb33335
15	1.5000002384 = 3fc00002	-1.5000002384 = bfc00002
16	1.6000002623 = 3fcccccfc	-1.6000002623 = bfccccfc
17	1.7000002861 = 3fd9999c	-1.7000002861 = bfd9999c
18	1.8000003099 = 3fe66669	-1.8000003099 = bfe66669
19	1.9000003338 = 3ff33336	-1.9000003338 = bff33336

B.1.2 Reducing Error Accumulation

In order to use floating point numbers in a program without causing excessive rounding problems an algorithm can be redesigned such that the accumulation is eliminated. This example is similar to the previous one, but this time we recalculate the desired value from a known-accurate integer value. Some rounding errors remain present, but they can not accumulate.

Listing B.7: errorcompensation.c

Accumulation of Error

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4
5 union floatbin
6 {
7     unsigned int    i;
8     float          f;
9 };
10 int main()
11 {
12     union floatbin  x, y;
13     int             i;
14
15     i = 1;
16     while (i <= 20)
17     {
18         x.f = i/10.0;
19         y.f = -x.f;
20         printf("%25.10f = %08x      %25.10f = %08x\n", x.f, x.i, y.f, y.i);
21         i++;
22     }
23     return(0);
24 }
```

Listing B.8: errorcompensation.out

Output from erroraccumulation.c

1	0.1000000015 = 3dcccccd	-0.1000000015 = bdcccccd
2	0.2000000030 = 3e4ccccd	-0.2000000030 = be4ccccd
3	0.3000000119 = 3e99999a	-0.3000000119 = be99999a
4	0.4000000060 = 3ecccccd	-0.4000000060 = becccccd
5	0.5000000000 = 3f000000	-0.5000000000 = bf000000
6	0.6000000238 = 3f19999a	-0.6000000238 = bf19999a

7	0.6999999881 = 3f333333	-0.6999999881 = bf333333
8	0.8000000119 = 3f4ccccd	-0.8000000119 = bf4ccccd
9	0.8999999762 = 3f666666	-0.8999999762 = bf666666
10	1.0000000000 = 3f800000	-1.0000000000 = bf800000
11	1.1000000238 = 3f8ccccd	-1.1000000238 = bf8ccccd
12	1.2000000477 = 3f99999a	-1.2000000477 = bf99999a
13	1.2999999523 = 3fa66666	-1.2999999523 = bfa66666
14	1.3999999762 = 3fb33333	-1.3999999762 = bfb33333
15	1.5000000000 = 3fc00000	-1.5000000000 = bfc00000
16	1.6000000238 = 3fcccccd	-1.6000000238 = bfcccccd
17	1.7000000477 = 3fd9999a	-1.7000000477 = bfd9999a
18	1.7999999523 = 3fe66666	-1.7999999523 = bfe66666
19	1.8999999762 = 3ff33333	-1.8999999762 = bff33333
20	2.0000000000 = 40000000	-2.0000000000 = c0000000

Appendix C

Attribution 4.0 International

Creative Commons Corporation ("Creative Commons") is not a law firm and does not provide legal services or legal advice. Distribution of Creative Commons public licenses does not create a lawyer-client or other relationship. Creative Commons makes its licenses and related information available on an "as-is" basis. Creative Commons gives no warranties regarding its licenses, any material licensed under their terms and conditions, or any related information. Creative Commons disclaims all liability for damages resulting from their use to the fullest extent possible.

Using Creative Commons Public Licenses

Creative Commons public licenses provide a standard set of terms and conditions that creators and other rights holders may use to share original works of authorship and other material subject to copyright and certain other rights specified in the public license below. The following considerations are for informational purposes only, are not exhaustive, and do not form part of our licenses.

Considerations for licensors: Our public licenses are intended for use by those authorized to give the public permission to use material in ways otherwise restricted by copyright and certain other rights. Our licenses are irrevocable. Licensors should read and understand the terms and conditions of the license they choose before applying it. Licensors should also secure all rights necessary before applying our licenses so that the public can reuse the material as expected. Licensors should clearly mark any material not subject to the license. This includes other CC-licensed material, or material used under an exception or limitation to copyright. More considerations for licensors: http://wiki.creativecommons.org/Considerations_for_licensors

Considerations for the public: By using one of our public licenses, a licensor grants the public permission to use the licensed material under specified terms and conditions. If the licensor's permission is not necessary for any reason—for example, because of any applicable exception or limitation to copyright—then that use is not regulated by the license. Our licenses grant only permissions under copyright and certain other rights that a licensor has authority to grant. Use of the licensed material may still be restricted for other reasons, including because others have copyright or other rights in the material. A licensor may make special requests, such as asking that all changes be marked or described. Although not required by our licenses, you are encouraged to respect those requests where reasonable. More considerations for the public: http://wiki.creativecommons.org/Considerations_for_licensees

Creative Commons Attribution 4.0 International Public License

By exercising the Licensed Rights (defined below), You accept and agree to be bound by the terms and conditions of this Creative Commons Attribution 4.0 International Public License ("Public License"). To the extent this Public License may be interpreted as a contract, You are granted the Licensed Rights in consideration of Your acceptance of these terms and conditions, and the Licensor grants You such rights in consideration of benefits the Licensor receives from making the Licensed Material available under these terms and conditions.

Section 1. Definitions

- a. Adapted Material means material subject to Copyright and Similar Rights that is derived from or based upon the Licensed Material and in which the Licensed Material is translated, altered, arranged, transformed, or otherwise modified in a manner requiring permission under the Copyright and Similar Rights held by the Licensor. For purposes of this Public License, where the Licensed Material is a musical work, performance, or sound recording, Adapted Material is always produced where the Licensed Material is synched in timed relation with a moving image.
- b. Adapter's License means the license You apply to Your Copyright and Similar Rights in Your contributions to Adapted Material in accordance with the terms and conditions of this Public License.

-
- c. Copyright and Similar Rights means copyright and/or similar rights closely related to copyright including, without limitation, performance, broadcast, sound recording, and Sui Generis Database Rights, without regard to how the rights are labeled or categorized. For purposes of this Public License, the rights specified in Section 2(b)(1)-(2) are not Copyright and Similar Rights.
 - d. Effective Technological Measures means those measures that, in the absence of proper authority, may not be circumvented under laws fulfilling obligations under Article 11 of the WIPO Copyright Treaty adopted on December 20, 1996, and/or similar international agreements.
 - e. Exceptions and Limitations means fair use, fair dealing, and/or any other exception or limitation to Copyright and Similar Rights that applies to Your use of the Licensed Material.
 - f. Licensed Material means the artistic or literary work, database, or other material to which the Licensor applied this Public License.
 - g. Licensed Rights means the rights granted to You subject to the terms and conditions of this Public License, which are limited to all Copyright and Similar Rights that apply to Your use of the Licensed Material and that the Licensor has authority to license.
 - h. Licensor means the individual(s) or entity(ies) granting rights under this Public License.
 - i. Share means to provide material to the public by any means or process that requires permission under the Licensed Rights, such as reproduction, public display, public performance, distribution, dissemination, communication, or importation, and to make material available to the public including in ways that members of the public may access the material from a place and at a time individually chosen by them.
 - j. Sui Generis Database Rights means rights other than copyright resulting from Directive 96/9/EC of the European Parliament and of the Council of 11 March 1996 on the legal protection of databases, as amended and/or succeeded, as well as other essentially equivalent rights anywhere in the world.
 - k. You means the individual or entity exercising the Licensed Rights under this Public License. Your has a corresponding meaning.

Section 2. Scope

- a. License grant.
 - 1. Subject to the terms and conditions of this Public License, the Licensor hereby grants You a worldwide, royalty-free, non-sublicensable, non-exclusive, irrevocable license to exercise the Licensed Rights in the Licensed Material to:
 - a. reproduce and Share the Licensed Material, in whole or in part; and
 - b. produce, reproduce, and Share Adapted Material.
 - 2. Exceptions and Limitations. For the avoidance of doubt, where Exceptions and Limitations apply to Your use, this Public License does not apply, and You do not need to comply with its terms and conditions.
 - 3. Term. The term of this Public License is specified in Section 6(a).
 - 4. Media and formats; technical modifications allowed. The Licensor authorizes You to exercise the Licensed Rights in all media and formats whether now known or hereafter created, and to make technical modifications necessary to do so. The Licensor waives and/or agrees not to assert any right or authority to forbid You from making technical modifications necessary to exercise the Licensed Rights, including technical modifications necessary to circumvent Effective Technological Measures. For purposes of this Public License, simply making modifications authorized by this Section 2(a) (4) never produces Adapted Material.
 - 5. Downstream recipients.
 - a. Offer from the Licensor – Licensed Material. Every recipient of the Licensed Material automatically receives an offer from the Licensor to exercise the Licensed Rights under the terms and conditions of this Public License.
 - b. No downstream restrictions. You may not offer or impose any additional or different terms or conditions on, or apply any Effective Technological Measures to, the Licensed Material if doing so restricts exercise of the Licensed Rights by any recipient of the Licensed Material.
 - 6. No endorsement. Nothing in this Public License constitutes or may be construed as permission to assert or imply that You are, or that Your use of the Licensed Material is, connected with, or sponsored, endorsed, or granted official status by, the Licensor or others designated to receive attribution as provided in Section 3(a)(1)(A)(i).
- b. Other rights.
 - 1. Moral rights, such as the right of integrity, are not licensed under this Public License, nor are publicity, privacy, and/or other similar personality rights; however, to the extent possible, the Licensor waives and/or agrees not to assert any such rights held by the Licensor to the limited extent necessary to allow You to exercise the Licensed Rights, but not otherwise.
 - 2. Patent and trademark rights are not licensed under this Public License.
 - 3. To the extent possible, the Licensor waives any right to collect royalties from You for the exercise of the Licensed Rights, whether directly or through a collecting society under any voluntary or waivable statutory or compulsory licensing scheme. In all other cases the Licensor expressly reserves any right to collect such royalties.

Section 3. License Conditions

Your exercise of the Licensed Rights is expressly made subject to the following conditions.

- a. Attribution.
 1. If You Share the Licensed Material (including in modified form), You must:
 - a. retain the following if it is supplied by the Licensor with the Licensed Material:
 - i. identification of the creator(s) of the Licensed Material and any others designated to receive attribution, in any reasonable manner requested by the Licensor (including by pseudonym if designated);
 - ii. a copyright notice;
 - iii. a notice that refers to this Public License;
 - iv. a notice that refers to the disclaimer of warranties;
 - v. a URI or hyperlink to the Licensed Material to the extent reasonably practicable;
 - b. indicate if You modified the Licensed Material and retain an indication of any previous modifications; and
 - c. indicate the Licensed Material is licensed under this Public License, and include the text of, or the URI or hyperlink to, this Public License.
 2. You may satisfy the conditions in Section 3(a)(1) in any reasonable manner based on the medium, means, and context in which You Share the Licensed Material. For example, it may be reasonable to satisfy the conditions by providing a URI or hyperlink to a resource that includes the required information.
 3. If requested by the Licensor, You must remove any of the information required by Section 3(a)(1)(A) to the extent reasonably practicable.
 4. If You Share Adapted Material You produce, the Adapter's License You apply must not prevent recipients of the Adapted Material from complying with this Public License.

Section 4. Sui Generis Database Rights

Where the Licensed Rights include Sui Generis Database Rights that apply to Your use of the Licensed Material:

- a. for the avoidance of doubt, Section 2(a)(1) grants You the right to extract, reuse, reproduce, and Share all or a substantial portion of the contents of the database;
- b. if You include all or a substantial portion of the database contents in a database in which You have Sui Generis Database Rights, then the database in which You have Sui Generis Database Rights (but not its individual contents) is Adapted Material; and
- c. You must comply with the conditions in Section 3(a) if You Share all or a substantial portion of the contents of the database.

For the avoidance of doubt, this Section 4 supplements and does not replace Your obligations under this Public License where the Licensed Rights include other Copyright and Similar Rights.

Section 5. Disclaimer of Warranties and Limitation of Liability

- a. UNLESS OTHERWISE SEPARATELY UNDERTAKEN BY THE LICENSOR, TO THE EXTENT POSSIBLE, THE LICENSOR OFFERS THE LICENSED MATERIAL AS-IS AND AS-AVAILABLE, AND MAKES NO REPRESENTATIONS OR WARRANTIES OF ANY KIND CONCERNING THE LICENSED MATERIAL, WHETHER EXPRESS, IMPLIED, STATUTORY, OR OTHER. THIS INCLUDES, WITHOUT LIMITATION, WARRANTIES OF TITLE, MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, NON-INFRINGEMENT, ABSENCE OF LATENT OR OTHER DEFECTS, ACCURACY, OR THE PRESENCE OR ABSENCE OF ERRORS, WHETHER OR NOT KNOWN OR DISCOVERABLE. WHERE DISCLAIMERS OF WARRANTIES ARE NOT ALLOWED IN FULL OR IN PART, THIS DISCLAIMER MAY NOT APPLY TO YOU.
- b. TO THE EXTENT POSSIBLE, IN NO EVENT WILL THE LICENSOR BE LIABLE TO YOU ON ANY LEGAL THEORY (INCLUDING, WITHOUT LIMITATION, NEGLIGENCE) OR OTHERWISE FOR ANY DIRECT, SPECIAL, INDIRECT, INCIDENTAL, CONSEQUENTIAL, PUNITIVE, EXEMPLARY, OR OTHER LOSSES, COSTS, EXPENSES, OR DAMAGES ARISING OUT OF THIS PUBLIC LICENSE OR USE OF THE LICENSED MATERIAL, EVEN IF THE LICENSOR HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH LOSSES, COSTS, EXPENSES, OR DAMAGES. WHERE A LIMITATION OF LIABILITY IS NOT ALLOWED IN FULL OR IN PART, THIS LIMITATION MAY NOT APPLY TO YOU.
- c. The disclaimer of warranties and limitation of liability provided above shall be interpreted in a manner that, to the extent possible, most closely approximates an absolute disclaimer and waiver of all liability.

Section 6. Term and Termination

- a. This Public License applies for the term of the Copyright and Similar Rights licensed here. However, if You fail to comply with this Public License, then Your rights under this Public License terminate automatically.
- b. Where Your right to use the Licensed Material has terminated under Section 6(a), it reinstates:
 - 1. automatically as of the date the violation is cured, provided it is cured within 30 days of Your discovery of the violation; or
 - 2. upon express reinstatement by the Licensor.

For the avoidance of doubt, this Section 6(b) does not affect any right the Licensor may have to seek remedies for Your violations of this Public License.

- c. For the avoidance of doubt, the Licensor may also offer the Licensed Material under separate terms or conditions or stop distributing the Licensed Material at any time; however, doing so will not terminate this Public License.
- d. Sections 1, 5, 6, 7, and 8 survive termination of this Public License.

Section 7. Other Terms and Conditions

- a. The Licensor shall not be bound by any additional or different terms or conditions communicated by You unless expressly agreed.
- b. Any arrangements, understandings, or agreements regarding the Licensed Material not stated herein are separate from and independent of the terms and conditions of this Public License.

Section 8. Interpretation

- a. For the avoidance of doubt, this Public License does not, and shall not be interpreted to, reduce, limit, restrict, or impose conditions on any use of the Licensed Material that could lawfully be made without permission under this Public License.
- b. To the extent possible, if any provision of this Public License is deemed unenforceable, it shall be automatically reformed to the minimum extent necessary to make it enforceable. If the provision cannot be reformed, it shall be severed from this Public License without affecting the enforceability of the remaining terms and conditions.
- c. No term or condition of this Public License will be waived and no failure to comply consented to unless expressly agreed to by the Licensor.
- d. Nothing in this Public License constitutes or may be interpreted as a limitation upon, or waiver of, any privileges and immunities that apply to the Licensor or You, including from the legal processes of any jurisdiction or authority.

Creative Commons is not a party to its public licenses. Notwithstanding, Creative Commons may elect to apply one of its public licenses to material it publishes and in those instances will be considered the Licensor. The text of the Creative Commons public licenses is dedicated to the public domain under the CC0 Public Domain Dedication. Except for the limited purpose of indicating that material is shared under a Creative Commons public license or as otherwise permitted by the Creative Commons policies published at <http://creativecommons.org/policies>, Creative Commons does not authorize the use of the trademark “Creative Commons” or any other trademark or logo of Creative Commons without its prior written consent including, without limitation, in connection with any unauthorized modifications to any of its public licenses or any other arrangements, understandings, or agreements concerning use of licensed material. For the avoidance of doubt, this paragraph does not form part of the public licenses.

Creative Commons may be contacted at <http://creativecommons.org>.

Bibliography

- [1] RISC-V Foundation, *The RISC-V Instruction Set Manual, Volume I: User-Level ISA, Document Version 2.2*, 5 2017. Editors Andrew Waterman and Krste Asanović. [vii](#), [31](#), [32](#), [33](#), [34](#), [41](#), [43](#), [44](#), [45](#), [46](#), [47](#), [48](#), [49](#), [50](#), [51](#), [52](#)
- [2] D. Patterson and A. Waterman, *The RISC-V Reader: An Open Architecture Atlas*. Strawberry Canyon, 11 2017. ISBN: 978-0999249116. [vii](#)
- [3] D. Patterson and J. Hennessy, *Computer Organization and Design RISC-V Edition: The Hardware Software Interface*. Morgan Kaufmann, 4 2017. ISBN: 978-0128122754. [vii](#), [15](#)
- [4] W. F. Decker, “A modern approach to teaching computer organization and assembly language programming,” *SIGCSE Bull.*, vol. 17, pp. 38–44, 12 1985. [vii](#)
- [5] Texas Instruments, *SN54190, SN54191, SN54LS190, SN54LS191, SN74190, SN74191, SN74LS190, SN74LS191 Synchronous Up/Down Counters With Down/Up Mode Control*, 3 1988. [vii](#)
- [6] Texas Instruments, *SN54154, SN74154 4-line to 16-line Decoders/Demultiplexers*, 12 1972. [vii](#)
- [7] Intel, *MCS-85 User’s Manual*, 9 1978. [vii](#)
- [8] Radio Shack, *TRS-80 Editor/Assembler Operation and Reference Manual*, 1978. [vii](#)
- [9] Motorola, *MC68000 16-bit Microprocessor User’s Manual*, 2nd ed., 1 1980. MC68000UM(AD2). [vii](#)
- [10] R. A. Overbeek and W. E. Singletary, *Assembler Language With ASSIST*. Science Research Associates, Inc., 2nd ed., 1983. [vii](#)
- [11] IBM, *IBM System/370 Principals of Operation*, 7th ed., 3 1980. [vii](#)
- [12] IBM, *OS/VS-DOS/VSE-VM/370 Assembler Language*, 6th ed., 3 1979. [vii](#)
- [13] D. Cohen, “[IEN 137, On Holy Wars and a Plea for Peace](#),” Apr. 1980. This note discusses the Big-Endian/Little-Endian byte/bit-order controversy, but did not settle it. A decade later, David V. James in “Multiplexed Buses: The Endian Wars Continue”, *IEEE Micro*, **10**(3), 9–21 (1990) continued the discussion. [13](#)
- [14] R. M. Stallman and the GCC Developer Community, *Using the GNU Compiler Collection (For GCC version 7.3.0)*. Free Software Foundation, 51 Franklin Street, Fifth Floor, Boston, MA 02110-1301 USA: GNU Press, 2017. [20](#)
- [15] RISC-V Foundation, *The RISC-V Instruction Set Manual, Volume II: Privileged Architecture, Document Version 1.10*, 5 2017. Editors Andrew Waterman and Krste Asanović.
- [16] P. Dabbelt, S. O’Rear, K. Cheng, A. Waterman, M. Clark, A. Bradbury, D. Horner, M. Nordlund, and K. Merker, *RISC-V ELF psABI specification*, 2017.
- [17] National Semiconductor Corporation, *Series 32000 Databook*, 1986.

Index

A

ALU, 3

C

CPU, 2

H

hart, 3

I

Instruction

ADD, 44
ADDI, 40
AND, 48
ANDI, 42
AUIPC, 32
BEQ, 34
BGE, 36
BGEU, 36
BLT, 35
BLTU, 36
BNE, 35
CSRRC, 51
CSRRCI, 52
CSRRS, 51
CSRRSI, 52
CSRROW, 50
CSRROWI, 51
DIV, 53
DIVU, 53
EBREAK, 50
ECALL, 50
FENCE, 49
FENCE.I, 49
JAL, 32
JALR, 33
LB, 37
LBU, 38
LH, 37
LHU, 38
LUI, 31
LW, 37
MUL, 52
MULH, 53
MULHS, 53
MULHU, 53

OR, 48

ORI, 42

REM, 54

REMU, 54

SB, 39

SH, 39

SLL, 45

SLLI, 43

SLT, 45

SLTI, 40

SLTIU, 41

SLTU, 46

SRA, 47

SRAI, 43

SRL, 47

SRLI, 43

SUB, 44

SW, 39

XOR, 46

XORI, 41

instruction cycle, 4

instruction execute, 5

instruction fetch, 4

ISA, 3

R

register, 1, 3

RV32, 22

RV32A, 4, 54

RV32C, 4

RV32D, 4, 54

RV32F, 4, 54

RV32G, 4

RV32I, 4, 31

RV32M, 4, 52

RV32Q, 4

rvddt, 17

Glossary

address A numeric value used to uniquely identify each byte of main memory. [2](#), [69](#)

alignment Refers to a range of numeric values that begin at a multiple of some number. Primarily used when referring to a memory address. For example an alignment of two refers to one or more addresses starting at even address and continuing onto subsequent adjacent, increasing memory addresses. [14](#), [69](#)

big endian A number format where the most significant values are printed to the left of the lesser significant values. This is the method that everyone used to write decimal numbers every day. [18](#), [19](#), [69](#)

binary Something that has two parts or states. In computing these two states are represented by the numbers one and zero or by the conditions true and false and can be stored in one bit. [3](#), [69](#)

bit One binary digit. [3](#), [9](#), [69](#)

byte A binary value represented by 8 bits. [2](#), [69](#)

CPU Central Processing Unit. [1](#), [2](#), [69](#)

Doubleword A binary value represented by 64 bits. [69](#)

exception An error encountered by the CPU while executing an instruction that can not be completed. [14](#), [69](#)

Fullword A binary value represented by 32 bits. [69](#)

Halfword A binary value represented by 16 bits. [69](#)

hart Hardware Thread. [3](#), [69](#)

hexadecimal A base-16 numbering system whose digits are 0123456789abcdef. The hex digits (hits) are not case-sensitive. [18](#), [19](#), [69](#)

High order bits Some number of MSBs. [69](#)

hit One hex digit. [9](#), [10](#), [69](#)

ISA Instruction Set Architecture. [3](#), [69](#)

LaTeX Is a mark up language specially suited for scientific documents. [69](#)

little endian A number format where the least significant values are printed to the left of the more significant values. This is the opposite ordering that everyone learns in grade school when learning how to count. For example a big endian number written as “1234” would be written in little endian form as “4321”. [69](#)

Low order bits Some number of LSBs. [69](#)

LSB Least Significant Bit. [11](#), [23](#), [26](#), [28](#), [69](#)

machine language The instructions that are executed by a CPU that are expressed in the form of binary values. [1](#), [69](#)

mnemonic A method used to remember something. In the case of assembly language, each machine instruction is given a name so the programmer need not memorize the binary values of each machine instruction. [1](#), [69](#)

MSB Most Significant Bit. [11](#), [22](#), [23](#), [28](#), [69](#)

overflow The situation where the result of an addition or subtraction operation is approaching positive or negative infinity and exceeds the number of bits allotted to contain the result. This is typically caused by high-order truncation. [12](#), [58](#), [69](#)

program A ordered list of one or more instructions. [1](#), [69](#)

Quadword A binary value represented by 128 bits. [69](#)

register A unit of storage inside a CPU with the capacity of XLEN bits. [1](#), [69](#)

RV32 Short for RISC-V 32. The number 32 refers to the XLEN. [31](#), [69](#)

RV64 Short for RISC-V 64. The number 64 refers to the XLEN. [69](#)

rvddt A RV32I simulator and debugging tool inspired by the simplicity of the Dynamic Debugging Tool (ddt) that was part of the CP/M operating system. [17](#), [69](#)

thread An stream of instructions. When plural, it is used to refer to the ability of a CPU to execute multiple instruction streams at the same time. [3](#), [69](#)

underflow The situation where the result of an addition or subtraction operation is approaching zero and exceeds the number of bits allotted to contain the result. This is typically caused by low-order truncation. [58](#), [69](#)

XLEN The number of bits a RISC-V x integer register (such as x0). For RV32 XLEN=32, RV64 XLEN=64 etc. [27](#), [28](#), [69](#)