

1

2

3

RISC-V Assembly Language Programming

(Draft v0.1-36-g8be5378)

4

John Winans
jwinans@niu.edu

5

May 17, 2018

6 Copyright © 2018 John Winans

7 This document is made available under a Creative Commons Attribution 4.0 International License. See
8 [Appendix D](#) for more information.

9 Download your own copy of this book from github here: <https://github.com/johnwinans/rvalp>.

10 This document may contain inaccuracies or errors. The author provides no guarantee regarding the
11 accuracy of this document's contents. If you discover that this document contains errors, please notify
12 the author.

13

14 ARM® is a registered trademark of ARM Limited in the EU and other countries.

► Fix Me:

*Need to say something
about trademarks for things
mentioned in this text*

15

Contents

16	Preface	viii
17	1 Introduction	1
18	1.1 The Digital Computer	1
19	1.1.1 Storage Systems	1
20	1.1.1.1 Volatile Storage	1
21	1.1.1.2 Non-Volatile Storage	2
22	1.1.2 CPU	2
23	1.1.2.1 Execution Unit	2
24	1.1.2.2 Arithmetic and Logic Unit	3
25	1.1.2.3 Registers	3
26	1.1.2.4 Harts	3
27	1.1.3 Peripherals	3
28	1.2 Instruction Set Architecture	4
29	1.2.1 RV Base Modules	4
30	1.2.2 Extension Modules	4
31	1.3 How the CPU Executes a Program	4
32	1.3.1 Instruction Fetch	5
33	1.3.2 Instruction Decode	5
34	1.3.3 Instruction Execute	5
35	2 Numbers and Storage Systems	6
36	2.1 Boolean Functions	6
37	2.1.1 NOT	6
38	2.1.2 AND	7
39	2.1.3 OR	7
40	2.1.4 XOR	7
41	2.2 Integers and Counting	8
42	2.2.1 Converting Between Bases	9
43	2.2.1.1 From Binary to Decimal	9

44	2.2.1.2	From Binary to Hexadecimal	10
45	2.2.1.3	From Hexadecimal to Binary	10
46	2.2.1.4	From Decimal to Binary	11
47	2.2.1.5	From Decimal to Hex	11
48	2.2.2	Addition of Binary Numbers	11
49	2.2.3	Signed Numbers	12
50	2.2.3.1	Converting between Positive and Negative	12
51	2.2.4	Subtraction of Binary Numbers	13
52	2.2.5	Truncation and Overflow	14
53	2.3	Sign and Zero Extension	15
54	2.4	Shifting	15
55	2.5	Main Memory Storage	15
56	2.5.1	Memory Dump	15
57	2.5.2	Big Endian Representation	15
58	2.5.3	Little Endian Representation	16
59	2.5.4	Character Strings and Arrays	16
60	2.5.5	Context is Important!	16
61	2.5.6	Alignment	16
62	2.5.7	Instruction Alignment	16
63	3	The Elements of a Assembly Language Program	17
64	3.1	Assembly Language Statements	17
65	3.2	Memory Layout	17
66	3.3	A Sample Program Source Listing	17
67	3.4	Running a Program With rvddt	18
68	4	Using The RISC-V GNU Toolchain	21
69	5	Writing RISC-V Programs	23
70	5.1	Using <code>addi</code> to Set Register Values	23
71	5.2	Using <code>ebreak</code> to Stopping rvddt Execution	23
72	5.3	Other Instructions With Immediate Operands	23
73	5.4	Transferring Data Between Registers and Memory	24
74	5.5	RR operations	24
75	5.6	Setting values to large values using <code>lui</code> with <code>addi</code>	24
76	5.7	Labels and Branching	25
77	5.8	Relocation	25
78	5.9	Jumps	26
79	5.10	Pseudo Operations	26

80	5.11 The Linker and Relaxation	26
81	5.12 pic and nopic	27
82	6 RV32 Machine Instructions	28
83	6.1 Introduction	28
84	6.2 Conventions and Terminology	28
85	6.2.1 XLEN	28
86	6.2.2 sx(val)	28
87	6.2.3 zx(val)	29
88	6.2.4 zr(val)	29
89	6.2.5 Sign Extended Left and Zero Extend Right	30
90	6.2.6 m8(addr)	30
91	6.2.7 m16(addr)	31
92	6.2.8 m32(addr)	31
93	6.2.9 m64(addr)	31
94	6.2.10 m128(addr)	31
95	6.2.11 .+offset	31
96	6.2.12 .-offset	31
97	6.2.13 pc	31
98	6.2.14 rd	31
99	6.2.15 rs1	32
100	6.2.16 rs2	32
101	6.2.17 imm	32
102	6.2.18 rsN[h:l]	32
103	6.3 Addressing Modes	32
104	6.4 Instruction Encoding Formats	32
105	6.4.1 U Type	32
106	6.4.2 J Type	34
107	6.4.3 R Type	35
108	6.4.4 I Type	35
109	6.4.5 S Type	35
110	6.4.6 B Type	35
111	6.4.7 CPU Registers	35
112	6.5 memory	36
113	6.6 RV32I Base Instruction Set	37
114	6.6.1 LUI rd, imm	37
115	6.6.2 AUIPC rd, imm	37
116	6.6.3 JAL rd, imm	38

117	6.6.4	JALR rd, rs1, imm	39
118	6.6.5	BEQ rs1, rs2, imm	40
119	6.6.6	BNE rs1, rs2, imm	41
120	6.6.7	BLT rs1, rs2, imm	41
121	6.6.8	BGE rs1, rs2, imm	41
122	6.6.9	BLTU rs1, rs2, imm	42
123	6.6.10	BGEU rs1, rs2, imm	42
124	6.6.11	LB rd, imm(rs1)	43
125	6.6.12	LH rd, imm(rs1)	43
126	6.6.13	LW rd, imm(rs1)	43
127	6.6.14	LBU rd, imm(rs1)	44
128	6.6.15	LHU rd, imm(rs1)	44
129	6.6.16	SB rs2, imm(rs1)	44
130	6.6.17	SH rs2, imm(rs1)	45
131	6.6.18	SW rs2, imm(rs1)	45
132	6.6.19	ADDI rd, rs1, imm	45
133	6.6.20	SLTI rd, rs1, imm	46
134	6.6.21	SLTIU rd, rs1, imm	46
135	6.6.22	XORI rd, rs1, imm	47
136	6.6.23	ORI rd, rs1, imm	47
137	6.6.24	ANDI rd, rs1, imm	48
138	6.6.25	SLLI rd, rs1, shamt	48
139	6.6.26	SRLI rd, rs1, shamt	49
140	6.6.27	SRAI rd, rs1, shamt	49
141	6.6.28	ADD rd, rs1, rs2	50
142	6.6.29	SUB rd, rs1, rs2	50
143	6.6.30	SLL rd, rs1, rs2	51
144	6.6.31	SLT rd, rs1, rs2	51
145	6.6.32	SLTU rd, rs1, rs2	52
146	6.6.33	XOR rd, rs1, rs2	52
147	6.6.34	SRL rd, rs1, rs2	53
148	6.6.35	SRA rd, rs1, rs2	53
149	6.6.36	OR rd, rs1, rs2	54
150	6.6.37	AND rd, rs1, rs2	54
151	6.6.38	FENCE predecessor, successor	55
152	6.6.39	FENCE.I	55
153	6.6.40	ECALL	56
154	6.6.41	EBREAK	56

155	6.6.42 CSR _{RW} rd, csr, rs1	56
156	6.6.43 CSR _{RS} rd, csr, rs1	56
157	6.6.44 CSR _{RC} rd, csr, rs1	57
158	6.6.45 CSR _{RWI} rd, csr, imm	57
159	6.6.46 CSR _{RSI} rd, csr, rs1	57
160	6.6.47 CSR _{RCI} rd, csr, rs1	58
161	6.7 RV32M Standard Extension	58
162	6.7.1 MUL rd, rs1, rs2	58
163	6.7.2 MULH rd, rs1, rs2	58
164	6.7.3 MULHS rd, rs1, rs2	59
165	6.7.4 MULHU rd, rs1, rs2	59
166	6.7.5 DIV rd, rs1, rs2	59
167	6.7.6 DIVU rd, rs1, rs2	59
168	6.7.7 REM rd, rs1, rs2	59
169	6.7.8 REMU rd, rs1, rs2	60
170	6.8 RV32A Standard Extension	60
171	6.9 RV32F Standard Extension	60
172	6.10 RV32D Standard Extension	60
173	A Installing a RISC-V Toolchain	61
174	A.1 The GNU Toolchain	61
175	A.2 rvddt	62
176	B Floating Point Numbers	63
177	B.1 IEEE-754 Floating Point Number Representation	63
178	B.1.1 Floating Point Number Accuracy	64
179	B.1.1.1 Powers Of Two	64
180	B.1.1.2 Clean Decimal Numbers	65
181	B.1.1.3 Accumulation of Error	66
182	B.1.2 Reducing Error Accumulation	67
183	C The ASCII Character Set	69
184	C.1 NAME	69
185	C.2 DESCRIPTION	69
186	C.2.1 Tables	71
187	C.3 NOTES	71
188	C.3.1 History	71
189	C.4 COLOPHON	71
190	D Attribution 4.0 International	72

191	Bibliography	76
192	Index	77
193	Glossary	79

Preface

195 I set out to this book because I couldn't find it in a single volume elsewhere.

196 The closest thing to what I sought when deciding to collect my thoughts into this document would be
 197 select portions of *The RISC-V Instruction Set Manual, Volume I: User-Level ISA, Document Version*
 198 *2.2*[\[1\]](#), *The RISC-V Reader*[\[2\]](#), and *Computer Organization and Design RISC-V Edition: The Hardware*
 199 *Software Interface*[\[3\]](#).

200 There *are* some terse guides around the Internet that are suitable for those that already know an
 201 assembly language. With all the (deserved) excitement brewing over system organization (and the need
 202 to compress the time out of university courses targeting assembly language programming [\[4\]](#)), it is no
 203 surprise that RISC-V texts for the beginning assembly programmer are not (yet) available.

204 When I got started in computing I learned how to count in binary in a high school electronics course
 205 using data sheets for integrated circuits such as the 74191[\[5\]](#) and 74154[\[6\]](#) prior to knowing that assembly
 206 language even existed.

207 I learned assembler from data sheets and texts (that are still sitting on my shelves) such as:

- 208 • The MCS-85 User's Manual[\[7\]](#)
- 209 • The EDTASM Manual[\[8\]](#)
- 210 • The MC68000 User's Manual[\[9\]](#)
- 211 • Assembler Language With ASSIST[\[10\]](#)
- 212 • IBM System/370 Principals of Operation[\[11\]](#)
- 213 • OS/VS-DOS/VSE-VM/370 Assembler Language[\[12\]](#)
- 214 • ... and several others

215 One way or another all of them discuss each CPU instruction in excruciating detail with both a logical
 216 and narrative description. For RISC-V this is also the case for the *RISC-V Reader*[\[2\]](#) and the *Computer*
 217 *Organization and Design RISC-V Edition*[\[3\]](#) books and is also present in this text (I consider that to be
 218 the minimal level of responsibility.)

219 Where I hope this text will differentiate itself from the existing RISC-V titles is in its attempt to address
 220 the needs of those learning assembly language for the first time. To this end I have primed this project
 221 with some of the material from old handouts I used when teaching assembly language programming in
 222 the late '80s.

Chapter 1

Introduction

At its core, a digital computer has at least one **Central Processing Unit** (CPU). A CPU executes a continuous stream of instructions called a **program**. These program instructions are expressed in what is called **machine language**. Each machine language instruction is a **binary** value. In order to provide a method to simplify the management of machine language programs a symbolic mapping is provided where a **mnemonic** can be used to specify each machine instruction and any of its parameters... rather than require that programs be expressed as a series of binary values. A set of mnemonics, parameters and rules for specifying their use for the purpose of programming a CPU is called an *Assembly Language*.

1.1 The Digital Computer

There are different types of computers. A *digital* computer is the type that most people think of when they hear the word *computer*. Other varieties of computers include *analog* and *quantum*.

A digital computer is one that that processes data that are represented using numeric values (digits), most commonly expressed in binary (ones and zeros) form.

This text focuses on digital computing.

A typical digital computer is composed of storage systems (memory, disc drives, USB drives, etc.), a CPU (with one or more cores), input peripherals (a keyboard and mouse) and output peripherals (display, printer or speakers.)

1.1.1 Storage Systems

Computer storage systems are used to hold the data and instructions for the CPU.

Types of computer storage can be classified into two categories: volatile and non-volatile.

1.1.1.1 Volatile Storage

Volatile storage is characterized by the fact that it will lose its contents (forget) any time that it is powered off.

One type of volatile storage is provided inside the CPU itself in small blocks called [registers](#). These CPU registers are used to hold individual data values that can be manipulated by the instructions that are executed by the CPU.

Another type of volatile storage is main memory. Main memory is connected to a computer's CPU and is used to hold the data and instructions that can not fit into the CPU registers.

Typically, a CPU's registers can hold tens of data values while the main memory can contain many billions of data values.

To keep track of the data values, each register is assigned a number and the main memory is broken up into small blocks called [bytes](#) that are also each assigned number called an [address](#) (an address is often referred to as a *location*).

A CPU can process data in a register at a speed that can be an order of magnitude faster than the rate that it can process (specifically, transfer data and instructions to and from) the main memory.

Register storage costs an order of magnitude more to manufacture than main memory. While it is desirable to have many registers the economics dictate that the vast majority of volatile computer storage be provided in its main memory. As a result, optimizing the copying of data between the registers and main memory is a desirable trait of good programs.

1.1.1.2 Non-Volatile Storage

Non-volatile storage is characterized by the fact that it will *NOT* lose its contents when it is powered off.

Common types of non-volatile storage are disc drives, flash cards and USB drives. Prices can vary widely depending on size and transfer speeds.

It is typical for a computer system's non-volatile storage to operate more slowly than its main memory.

This text is not particularly concerned with non-volatile storage.

1.1.2 CPU

The [CPU](#) is a collection of registers and circuitry designed manipulate the register data and to exchange data and instructions with the storage system. The instructions that are read from the main memory tell the CPU to perform various mathematic and logical operations on the data in its registers and where to save the results of those operations.

► Fix Me:
Add a block diagram of the CPU components described here.

1.1.2.1 Execution Unit

The part of a CPU that coordinates all aspects of the operations of each instruction is called the *execution unit*. It is what performs the transfers of instructions and data between the CPU and the main memory and tells the registers when they are supposed to either store or recall data being transferred. The execution unit also controls the ALU (Arithmetic and Logic Unit).

1.1.2.2 Arithmetic and Logic Unit

ALU
register
hart

When an instruction manipulates data by performing things like an *addition*, *subtraction*, *comparison* or other similar operations, the ALU is what will calculate the sum, difference, and so on... under the control of the execution unit.

1.1.2.3 Registers

In the RV32 CPU there are 31 general purpose registers that each contain 32 *bits* (where each bit is one *binary* digit value of one or zero) and a number of special-purpose registers. Each of the general purpose registers is given a name such as *x1*, *x2*, ... on up to *x31* (*general purpose* refers to the fact that the CPU itself does not prescribe any particular function to any these registers.) Two important special-purpose registers are *x0* and *pc*.

Register *x0* will always represent the value zero or logical *false* no matter what. If any instruction tries to change the value in *x0* the operation will fail. The need for *zero* is so common that, other than the fact that it is hard-wired to zero, the *x0* register is made available as if it were otherwise a general purpose register.¹

The *pc* register is called the *program counter*. The CPU uses it to remember the memory address where its program instructions are located.

The number of bits in each register is defined by the *Instruction Set Architecture (ISA)*.

► Fix Me:
Say something about XLEN?

1.1.2.4 Harts

Analogous to a *core* in other types of CPUs, a *hart* (hardware *thread*) in a RISC-V CPU refers to the collection of 32 registers, instruction execution unit and ALU.

When more than one hart is present in a CPU, a different stream of instructions can be executed on each hart all at the same time. Programs that are written to take advantage of this are called *multithreaded*.

This text will primarily focus on CPUs that have only one hart.

1.1.3 Peripherals

A peripheral is a device that is not a CPU or main memory. They are typically used to transfer information/data into and out of the main memory.

This text is not particularly concerned with the peripherals of a computer system other than in those sections where instructions are discussed whose purpose is to address the needs of a peripheral device. Such instructions are used to initiate, execute and/or synchronize data transfers.

¹Having a special *zero* register allows the total set of instructions that the CPU can execute to be simplified. Thus reducing its complexity, power consumption and cost.

1.2 Instruction Set Architecture

The catalog of rules that describes the details of the instructions and features that a given CPU provides is called its [Instruction Set Architecture \(ISA\)](#).

An ISA is typically expressed in terms of the specific meaning of each binary instruction that a CPU can recognize and how it will process each one.

The RISC-V ISA is defined as a set of modules. The purpose of dividing the ISA into modules is to allow an implementer to select which features to incorporate into a CPU design.

Any given RISC-V implementation must provide one of the *base* modules and zero or more of the *extension* modules.

ISA
RV32I
RV32M
RV32A
RV32F
RV32D
RV32Q
RV32C
RV32G
instruction cycle

1.2.1 RV Base Modules

The base modules are RV32I (32-bit general purpose), RV32E (32-bit embedded), RV64I (64-bit general purpose) and RV128I (128-bit general purpose).

These base modules provide the minimal functional set of integer operations needed to execute a useful application. The differing bit-widths address the needs of different main-memory sizes.

This text primarily focuses on the RV32I base module and how to program it.

1.2.2 Extension Modules

RISC-V extension modules may be included by an implementer interested in optimizing a design for one or more purposes.

Available extension modules include M (integer math), A (atomic), F (32-bit floating point), D (64-bit floating point), Q (128-bit floating point), C (compressed size instructions) and others.

The extension name *G* is used to represent the combined set of IMAFD extensions as it is expected to be a common combination.

1.3 How the CPU Executes a Program

The process of executing a program is continuously repeating series of *instruction cycles* that are each comprised of a *fetch*, *decode* and *execute* phase.

The current status of a CPU hart is entirely embodied in the data values that are stored in its registers at any moment in time. Of particular interest to an executing a program is the *pc* register. The *pc* contains the memory address containing the instruction that the CPU is currently executing.²

For this to work, the instructions to be executed must have been previously stored in adjacent main memory locations and the address of the first instruction placed into the *pc* register.

²In the RISC-V ISA the *pc* register points to the *current* instruction where in most other designs, the *pc* register points to the *next* instruction.

1.3.1 Instruction Fetch

instruction fetch
instruction decode
instruction execute

In order to *fetch* an instruction from the main memory the CPU must have a method to identify which instruction should be fetched and a method to fetch it.

Given that the main memory is broken up and that each of its bytes is assigned an address, the `pc` is used to hold the address of the location where the next instruction to execute is located.

Given an instruction address, the CPU can request that the main memory locate and return the value of the data stored there using what is called a *memory read* operation and then the CPU can treat that *fetched* value as an instruction and execute it.³

1.3.2 Instruction Decode

Once an instruction has been fetched, it must be inspected to determine what operation(s) are to be performed. This primarily boils down to inspecting the portions of the instruction that dictate which registers are involved and, if the ALU is required, what it should do.

1.3.3 Instruction Execute

Typical instructions do things like add a number to the value currently stored in one of the registers or store the contents of a register into the main memory at some given address.

Also part of every instruction is a notion of what should be done next.

Most of the time an instruction will complete by indicating that the CPU should proceed to fetch and execute the instruction at the next larger main memory address. In these cases the `pc` is incremented to point to the memory address after the current instruction.

Any parameters that an instruction requires must either be part of the instruction itself or read from (or stored into) one or more of the general purpose registers.

Some instructions can specify that the CPU proceed to execute an instruction at an address other than the one that follows itself. This class of instructions have names like *jump* and *branch* and are available in a variety of different styles.

The RISC-V ISA uses the word *jump* to refer to an *unconditional* change in the sequential processing of instructions and the word *branch* to refer to a *conditional* change.

For example, a (conditional) branch instruction might instruct the CPU to proceed to the instruction at the next main memory address if the value in register number 8 is currently less than the value in register number 24 *but otherwise* proceed to an instruction at a different address when it is not. This type of instruction can therefore result in having one of two different actions pending the resulting *condition* of the comparison.⁴

Once the instruction execution phase has completed, the next instruction cycle will be performed using the new value in the `pc` register.

³RV32I instructions are more than one byte in size, but this general description is suitable for now.

⁴This is the fundamental method used by a CPU to make decisions.

Chapter 2

Numbers and Storage Systems

This chapter discusses how data are represented and stored in a computer.

In the context of computing, *boolean* refers to a condition that can be either true and false and *binary* refers to the use of a base-2 numeric system to represent numbers.

RISC-V assembly language uses binary to represent all values, be they boolean or numeric. It is the context within which they are used that determines whether they are boolean or numeric.

RISC-V assembly language uses zero to represent *false* and one to represent *true*. In general, however, it is useful to relax this and define zero **and only zero** to be *false* and anything that is not *false* is therefore *true*.¹

The reason for this relaxation is because, while a single binary digit (**bit**) can represent the two values zero and one, the vast majority of the time data is processed by the CPU in groups of bits. These groups have names like **byte**, **halfword** and **fullword**.

► Fix Me:

Add some diagrams here showing bits, bytes and the MSB, LSB, ... perhaps relocated from the RV32I chapter?

2.1 Boolean Functions

Boolean functions apply on a per-bit basis. When applied to multi-bit values, each bit position is operated upon independently of the other bits.

► Fix Me:

Probably should add basic truth table diagrams.

2.1.1 NOT

The *NOT* operator applies to a single operand and represents the opposite of the input.

If the input is 1 then the output is 0. If the input is 0 then the output is 1. In other words, the output value is *not* that of the input value.

► Fix Me:

Need to define unary, binary and ternary operators without confusing binary operators with binary numbers.

This text will use the operator used in the C language when discussing the *NOT* operator in symbolic form. Specifically the tilde: '~'.

```
~ 1 1 1 1 0 1 0 1 <== A
```

¹This is how *true* and *false* behave in C, C++, and many other languages as well as the common assembly language idioms discussed in this text.

```

395  -----
396      0 0 0 0 1 0 1 0  <== output

```

397 In a line of code the above might read like this: `output = ~A`

398 2.1.2 AND

399 The boolean *and* function has two or more inputs and the output is a single bit. The output is 1 if and
400 only if all of the input values are 1. Otherwise it is 0.

401 This text will use the operator used in the C language when discussing the *AND* operator in symbolic
402 form. Specifically the ampersand: '&'.
403 This function works like it does in spoken language. For example if A is 1 *AND* B is 1 then the output
404 is 1 (true). Otherwise the output is 0 (false). For example:

```

405      1 1 1 1 0 1 0 1  <== A
406      & 1 0 0 1 0 0 1 1  <== B
407  -----
408      1 0 0 1 0 0 0 1  <== output

```

409 In a line of code the above might read like this: `output = A & B`

410 2.1.3 OR

411 The boolean *or* function has two or more inputs and the output is a single bit. The output is 1 if at
412 least one of the input values are 1.

413 This text will use the operator used in the C language when discussing the *OR* operator in symbolic
414 form. Specifically the pipe: '|'.
415 This function works like it does in spoken language. For example if A is 1 *OR* B is 1 then the output is
416 1 (true). Otherwise the output is 0 (false). For example:

```

417      1 1 1 1 0 1 0 1  <== A
418      | 1 0 0 1 0 0 1 1  <== B
419  -----
420      1 1 1 1 0 1 1 1  <== output

```

421 In a line of code the above might read like this: `output = A | B`

422 2.1.4 XOR

423 The boolean *exclusive or* function has two or more inputs and the output is a single bit. The output is
424 1 if only an odd number of inputs are 1. Otherwise the output will be 0.

425 This text will use the operator used in the C language when discussing the *XOR* operator in symbolic
426 form. Specifically the carrot: '^'.

427 Note that when *XOR* is used with two inputs, the output is set to 1 (true) when the inputs have different
428 values and 0 (false) when the inputs both have the same value.

429 For example:

```
430      1 1 1 1 0 1 0 1  <== A
431      ^ 1 0 0 1 0 0 1 1  <== B
432      -----
433      0 1 1 0 0 1 1 0  <== output
```

434 In a line of code the above might read like this: `output = A ^ B`

435 **2.2 Integers and Counting**

436 A binary integer is constructed with only 1s and 0s in the same manner as decimal numbers are con-
437 structed with values from 0 to 9.

438 Counting in binary is the same as in decimal. For example, when adding 1 to 9, the carry is added to
439 the next place value. When subtracting 1 from 0, a borrow is required and so on.

440 Figure [Figure 2.1](#) shows an abridged table of the decimal, binary and hexadecimal values from 0 to 129.

Decimal			Binary								Hex	
10 ²	10 ¹	10 ⁰	2 ⁷	2 ⁶	2 ⁵	2 ⁴	2 ³	2 ²	2 ¹	2 ⁰	16 ¹	16 ⁰
100	10	1	128	64	32	16	8	4	2	1	16	1
0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	1	0	0	0	0	0	0	0	1	0	1
0	0	2	0	0	0	0	0	0	1	0	0	2
0	0	3	0	0	0	0	0	0	1	1	0	3
0	0	4	0	0	0	0	0	1	0	0	0	4
0	0	5	0	0	0	0	0	1	0	1	0	5
0	0	6	0	0	0	0	0	1	1	0	0	6
0	0	7	0	0	0	0	0	1	1	1	0	7
0	0	8	0	0	0	0	1	0	0	0	0	8
0	0	9	0	0	0	0	1	0	0	1	0	9
0	1	0	0	0	0	0	1	0	1	0	0	a
0	1	1	0	0	0	0	1	0	1	1	0	b
0	1	2	0	0	0	0	1	1	0	0	0	c
0	1	3	0	0	0	0	1	1	0	1	0	d
0	1	4	0	0	0	0	1	1	1	0	0	e
0	1	5	0	0	0	0	1	1	1	1	0	f
0	1	6	0	0	0	1	0	0	0	0	1	0
0	1	7	0	0	0	1	0	0	0	1	1	1
...			
1	2	5	0	1	1	1	1	1	0	1	7	d
1	2	6	0	1	1	1	1	1	1	0	7	e
1	2	7	0	1	1	1	1	1	1	1	7	f
1	2	8	1	0	0	0	0	0	0	0	8	0

Figure 2.1: Counting in decimal, binary and hexadecimal.

One way to look at this table is on a per-row basis where each place value is represented by the base raised to the power of the place value position (shown in the column headings.) This is useful when converting arbitrary values between bases. For example to interpret the decimal value on the fourth row:

$$0 \times 10^2 + 0 \times 10^1 + 3 \times 10^0 = 3_{10}$$

And to interpret binary value on the same row by converting it to decimal:

$$0 \times 2^7 + 0 \times 2^6 + 0 \times 2^5 + 0 \times 2^4 + 0 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 1 \times 2^0 = 3_{10}$$

And the same for the hexadecimal value:

$$0 \times 16^1 + 3 \times 16^0 = 3_{10}$$

Another way to look at this table is on a per-column basis. When tasked with drawing such a table by hand, it might be useful to observe that, just as in decimal, the right-most column will cycle through all of the values represented in the chosen base then cycle back to zero and repeat. (For example, in binary this pattern is 0-1-0-1-0-1-0-...) The next column in each base will cycle in the same manner except each of the values is repeated as many times as is represented by the place value (in the case of decimal, 10^1 times, binary 2^1 times, hex 16^1 times. Again, the for binary numbers this pattern is 0-0-1-1-0-0-1-1-...) This continues for as many columns as are needed to represent the magnitude of the desired number.

Another item worth noting is that any even binary number will always have a 0 LSB and odd numbers will always have a 1 LSB.

As is customary in decimal, leading zeros are sometimes not shown for readability.

The relationship between binary and hex values is also worth taking note. Because $2^4 = 16$, there is a clean and simple grouping of 4 bits to 1 hit. There is no such relationship between binary and decimal.

Writing and reading numbers in binary that are longer than 8 bits is cumbersome and prone to error. The simple conversion between binary and hex makes hex a convenient shorthand for expressing binary values in many situations.

For example, consider the following value expressed in binary, hexadecimal and decimal (spaced to show the relationship between binary and hex):

Binary value:	0010	0111	1011	1010	1100	1100	1111	0101
Hex Value:	2	7	B	A	C	C	F	5
Decimal Value:	666553589							

Empirically we can see that grouping the bits into sets of four allows an easy conversion to hex and expressing it as such is $\frac{1}{4}$ as long as in binary while at the same time allowing for easy conversion back to binary.

The decimal value in this example does not easily convey a sense of the binary value.

2.2.1 Converting Between Bases

2.2.1.1 From Binary to Decimal

Alas, it is occasionally necessary to convert between decimal, binary and/or hex.

To convert from binary to decimal, put the decimal value of the place values ... 8 4 2 1 over the binary digits like this:

	128	64	32	16	8	4	2	1
	0	0	0	1	1	0	1	1

Now sum the place-values that are expressed in decimal for each bit with the value of 1: $16 + 8 + 2 + 1$. The integer binary value 00011011_2 represents the decimal value 27_{10} .

2.2.1.2 From Binary to Hexadecimal

Conversion from binary to hex involves grouping the bits into sets of four and then performing the same summing process as shown above. If there is not a multiple of four bits then extend the binary to the left with zeros to make it so.

Grouping the bits into sets of four and summing:

Place:	8 4 2 1	8 4 2 1	8 4 2 1	8 4 2 1
Binary:	0 1 1 0	1 1 0 1	1 0 1 0	1 1 1 0
Decimal:	$4+2 = 6$	$8+4+1 = 13$	$8+2 = 10$	$8+4+2 = 14$

After the summing, convert each decimal value to hex. The decimal values from 0–9 are the same values in hex. Because we don't have any more numerals to represent the values from 10–15, we use the first 6 letters (See the right-most column of [Figure 2.1](#).) Fortunately there are only six hex mappings involving letters. Thus it is reasonable to memorize them.

Continuing this example:

Decimal:	6	13	10	14
Hex:	6	D	A	E

2.2.1.3 From Hexadecimal to Binary

Again, the four-bit mapping between binary and hex makes this task as straight forward as using a look-up table.

For each [hit](#) (Hex digIT), translate it to its unique four-bit pattern. Perform this task either by memorizing each of the 16 patterns or by converting each hit to decimal first and then converting each four-bit binary value to decimal using the place-value summing method discussed in [subsection 2.2.1.1](#).

For example:

Hex:	4	C
Binary:	0 1 0 0	1 1 0 0
Decimal:	128 64 32 16	8 4 2 1
Sum:	$64+$	$8+4 = 76$

2.2.1.4 From Decimal to Binary

To convert arbitrary decimal numbers to binary, extend the list of binary place values until it exceeds the value of the decimal number being converted. Then make successive subtractions of each of the place values that would yield a non-negative result.

For example, to convert 1234_{10} to binary:

Place values: 2048-1024-512-256-128-64-32-16-8-4-2-1

0	2048	(too big)
1	$1234 - 1024 = 210$	
0	512	(too big)
0	256	(too big)
1	$210 - 128 = 82$	
1	$82 - 64 = 18$	
0	32	(too big)
1	$18 - 16 = 2$	
0	8	(too big)
0	4	(too big)
1	$2 - 2 = 0$	
0	1	(too big)

The answer using this notation is listed vertically in the left column with the **MSB** on the top and the **LSB** on the bottom line: 010011010010_2 .

2.2.1.5 From Decimal to Hex

Conversion from decimal to hex can be done by using the place values for base-16 and the same math as from decimal to binary or by first converting the decimal value to binary and then from binary to hex by using the methods discussed above.

Because binary and hex are so closely related, performing a conversion by way of binary is quite straight forward.

2.2.2 Addition of Binary Numbers

The addition of binary numbers can be performed long-hand the same way decimal addition is taught in grade school. In fact binary addition is easier since it only involves adding 0 or 1.

The first thing to note that in any number base $0 + 0 = 0$, $0 + 1 = 1$, and $1 + 0 = 1$. Since there is no “two” in binary (just like there is no “ten” decimal) adding $1 + 1$ results in a zero with a carry as in: $1 + 1 = 10_2$ and in: $1 + 1 + 1 = 11_2$. Using these five sums, any two binary integers can be added.

For example:

	111111	1111	<== carries
	0110101111001111		<== addend
+	0000011101100011		<== addend

	0111001100110010		<== sum

2.2.3 Signed Numbers

There are multiple methods used to represent signed binary integers. The method used by most modern computers is called “two’s complement.”

A two’s complement number is encoded in such a manner as to simplify the hardware used to add, subtract and compare integers.

A simple method of thinking about two’s complement numbers is to negate the place value of the [MSB](#). For example, the number one is represented the same as discussed before:

```
-128 64 32 16  8  4  2  1
      0  0  0  0  0  0  0  1
```

The [MSB](#) of any negative number in this format will always be 1. For example the value -1_{10} is:

```
-128 64 32 16  8  4  2  1
      1  1  1  1  1  1  1  1
```

... because: $-128 + 64 + 32 + 16 + 8 + 4 + 2 + 1 = -1$.

This format has the virtue of allowing the same addition logic discussed above to be used to calculate $-1 + 1 = 0$.

```
-128 64 32 16  8  4  2  1 <== place value
  1  1  1  1  1  1  1  0 <== carries
    1  1  1  1  1  1  1  1 <== addend (-1)
+  0  0  0  0  0  0  0  1 <== addend (1)
-----
  1  0  0  0  0  0  0  0 <== sum (0 with an overflow)
```

In order for this to work, the [overflow](#) carry out of the sum of the MSBs is ignored.

2.2.3.1 Converting between Positive and Negative

Changing the sign on two’s complement numbers can be described as inverting all of the bits (which is also known as the one’s complement) and then add one.

For example, inverting the number *four*:

```
-128 64 32 16  8  4  2  1
      0  0  0  0  0  1  0  0 <== 4
                                     1  1 <== carries
      1  1  1  1  1  0  1  1 <== one's complement of 4
+  0  0  0  0  0  0  0  1 <== plus 1
-----
      1  1  1  1  1  1  0  0 <== -4
```

574 This can be verified by adding 5 to the result and observe that the sum is 1:

```

575 -128 64 32 16 8 4 2 1
576   1  1  1  1  1          <== carries
577   1  1  1  1  1  1  0  0 <== -4
578 + 0  0  0  0  0  1  0  1 <== 5
579 -----
580  1  0  0  0  0  0  0  1

```

581 Note that the changing of the sign using this method is symmetric in that it is identical when converting
 582 from negative to positive and when converting from positive to negative: flip the bits and add 1.

583 For example, changing the value -4 to 4 to illustrate the reverse of the conversion above:

```

584 -128 64 32 16 8 4 2 1
585   1  1  1  1  1  1  0  0 <== -4
586
587           1  1          <== carries
588   0  0  0  0  0  0  1  1 <== one's complement of -4
589 + 0  0  0  0  0  0  0  1 <== plus 1
590 -----
591   0  0  0  0  0  1  0  0 <== 4

```

592 2.2.4 Subtraction of Binary Numbers

593 Subtraction of binary numbers is performed by first negating the subtrahend and then adding the two
 594 numbers. Due to the nature of two's complement numbers this will work for both signed and unsigned
 595 numbers.

596 To calculate $-4 - 8 = -12$

```

597 -128 64 32 16 8 4 2 1
598   1  1  1  1  1  1  0  0 <== -4
599 - 0  0  0  0  1  0  0  0 <== 8
600
601
602           1  1  1          <== carries
603   1  1  1  1  0  1  1  1 <== one's complement of -8
604 + 0  0  0  0  0  0  0  1 <== plus 1
605 -----
606   1  1  1  1  1  0  0  0 <== -8
607
608
609   1  1  1  1          <== carries
610   1  1  1  1  1  1  0  0 <== -4
611 + 1  1  1  1  1  0  0  0 <== -8
612 -----
613  1  1  1  1  1  0  1  0 < == -12

```

Fix Me:

This section needs more examples of subtracting signed and unsigned numbers and a discussion on how signedness is not relevant until the results are interpreted. For example adding $-4 + -8 = -12$ using two 8-bit numbers is the same as adding $252 + 248 = 500$ and truncating the result to 244.

2.2.5 Truncation and Overflow

Discuss the details of truncation and overflow here.

I prefer to define *truncation* as the loss of data as result of the bit-length of the destination being too small to hold result of an operation and *overflow* as when the carry into a sign bit is not the same as the carry out of the sign bit.

Where addition and subtraction on the RV32 is concerned, the sum or difference of two unsigned 32-bit numbers will be *truncated* when the operation results in a carry out of bit 31. Unsigned operations can not overflow (as defined above).

(show a truncation picture here)

An Overflow occurs with signed numbers when the two addends are positive and sum is negative or the addends are both negative and the sum is positive.

(show an overflow picture here)

(show mixed overflow and truncation situations here to drive home the need to ignore truncation when dealing with signed numbers.)

$0xffffffff + 0x00000002$ has truncation but not overflow (OK for signed, not OK for unsigned).

$0xffffffff + 0xffffffe$ also has truncation but not overflow.

$0x40000000 + 0x40000000$ has overflow but not truncation. (We care if are signed numbers.)

$0x80000000 + 0x80000000$ has both overflow and truncation. (we care regardless of signedness)

Where subtraction is concerned the notion of a borrow is the same as carry.

Page 13 of [1] mixes these two notions of (and never mentions the word *truncate*) like this:

We did not include special instruction set support for overflow checks on integer arithmetic operations in the base instruction set, as many overflow checks can be cheaply implemented using RISC-V branches. Overflow checking for unsigned addition requires only a single additional branch instruction after the addition: `add t0, t1, t2; bltu t0, t1, overflow`.

For signed addition, if one operand's sign is known, overflow checking requires only a single branch after the addition: `addi t0, t1, +imm; blt t0, t1, overflow`. This covers the common case of addition with an immediate operand.

For general signed addition, three additional instructions after the addition are required, leveraging the observation that the sum should be less than one of the operands if and only if the other operand is negative.

```
add t0, t1, t2
slti t3, t2, 0
slt t4, t0, t1
bne t3, t4, overflow
```

In RV64, checks of 32-bit signed additions can be optimized further by comparing the results of ADD and ADDW on the operands.

► Fix Me:

This chapter should be made consistent in its use of truncation and overflow as occur with signed and unsigned addition and subtraction.

► Fix Me:

I think that overloading the word overflow like this can be is confusing to new programmers.

2.3 Sign and Zero Extension

Seems like a good place to discuss extension.

► Fix Me:

Refactor the `sx()` and `zx()` discussion in the RV32I chapter and locate the details here.

2.4 Shifting

Seems like a good place to discuss logical and arithmetic shifting.

shift left logical

shift right logical

shift right arithmetic

2.5 Main Memory Storage

When transferring data between its registers registers and main memory a RISC-V system uses the little-endian byte order.²

► Fix Me:

Consider refactoring the memory discussion in RV32I reference chapter and placing some of it in this section.

2.5.1 Memory Dump

Introduce the memory dump and how to read them here.

Listing 2.1: `rvddt_memdump.out`

`rvddt` memory dump

```
ddt> d 0x00002600
00002600: 93 05 00 00 13 06 00 00-93 06 00 00 13 07 00 00 *.....*
00002610: 93 07 00 00 93 08 d0 05-73 00 00 63 54 05 02 *.....s...cT..*
00002620: 13 01 01 ff 23 24 81 00-13 04 05 00 23 26 11 00 *...#$....#&..*
00002630: 33 04 80 40 97 00 00 00-e7 80 40 01 23 20 85 00 *3..@.....@.# ..*
00002640: 6f 00 00 00 6f 00 00 00-b7 87 00 00 03 a5 07 43 *o...o.....C*
00002650: 67 80 00 00 00 00 00 00-76 61 6c 3d 00 00 00 00 *g.....val=....*
00002660: 00 00 00 00 80 84 2e 41-1f 85 45 41 80 40 9a 44 *.....A..EA.@.D*
00002670: 4f 11 f3 c3 6e 8a 67 41-20 1b 00 00 20 1b 00 00 *0...n.gA ... ..*
00002680: 44 1b 00 00 14 1b 00 00-14 1b 00 00 04 1c 00 00 *D.....*
00002690: 44 1b 00 00 14 1b 00 00-04 1c 00 00 14 1b 00 00 *D.....*
000026a0: 44 1b 00 00 10 1b 00 00-10 1b 00 00 10 1b 00 00 *D.....*
000026b0: 04 1c 00 00 54 1f 00 00-54 1f 00 00 d4 1f 00 00 *....T...T.....*
000026c0: 4c 1f 00 00 4c 1f 00 00-34 20 00 00 d4 1f 00 00 *L...L...4 .....*
000026d0: 4c 1f 00 00 34 20 00 00-4c 1f 00 00 d4 1f 00 00 *L...4 ..L.....*
000026e0: 48 1f 00 00 48 1f 00 00-48 1f 00 00 34 20 00 00 *H...H...H...4 ..*
000026f0: 00 01 02 02 03 03 03 03-04 04 04 04 04 04 04 04 *.....*
```

2.5.2 Big Endian Representation

Using the memory dump contents in prior section, discuss how big endian values are stored.

² See[13] for some history of the big/little-endian “controversy.”

2.5.3 Little Endian Representation

Using the memory dump contents in prior section, discuss how little endian values are stored.

2.5.4 Character Strings and Arrays

Define character strings and arrays.

Using the prior memory dump, discuss how and where things are stored and retrieved.

2.5.5 Context is Important!

Data values can be interpreted differently depending on the context in which they are used. Assuming what a set of bytes is used for based on their contents can be very misleading! For example, there is a 0x76 at address 0x00002658. This is a 'v' if you use it as an ASCII (see [Appendix C](#)) character, a 118₁₀ if it is an integer value and TRUE if it is a conditional.

2.5.6 Alignment

Draw a diagram showing the overlapping data types when they are all aligned.

2.5.7 Instruction Alignment

Every possible instruction that an RV32I CPU can execute contains exactly 32 bits. Therefore each one must be stored in four bytes of the main memory.

To simplify the hardware, each instruction must be placed into four adjacent bytes whose numeric address sequence begins with a multiple four. For example, an instruction might be located in bytes 4, 5, 6 and 7 (but not in 5, 6, 7 and 8 nor in 9, 3, 1, and 0...).

This sort of addressing requirement is common and is referred to as [alignment](#). An aligned instruction begins at a memory address that is a multiple of four. An *unaligned* instruction would be one beginning at any other address and is *illegal*.

An attempt to fetch an instruction from an unaligned address will result in an error referred to as an alignment [exception](#). This and other exceptions cause the CPU to stop executing the current instruction and start executing a different set of instructions that are prepared to handle the problem. Often an exception is handled by completely stopping the program in a way that is commonly referred to as a system or application *crash*.

Given a properly aligned instruction address, the CPU can request that the main memory locate and deliver the values of the four bytes in the address sequence to the CPU using what is called a memory read operation. Some systems can deliver four (or more) bytes at the same time while others might only be capable of delivering one or two bytes at a time. These differences in hardware typically impact the cost and performance of a system.³

➡ Fix Me:

Rewrite this section for data rather than instructions and then note here that instructions must be naturally aligned. For RV32 that is on a 4-byte boundary

³The design and implementation choices that determine how any given system operates are part of what is called a system's *organization* and is beyond the scope of this text. See [\[3\]](#) for more information on computer organization.

Chapter 3

The Elements of a Assembly Language Program

3.1 Assembly Language Statements

Introduce the assembly language grammar. Statement = 1 line of text containing an instruction or directive.

Instruction = label, mnemonic, operands, comment.

Directive = Used to control the operation of the assembler.

3.2 Memory Layout

Is this a good place to introduce the text, data, bss, heap and stack regions?

Or does that belong in a new section/chapter that discusses addressing modes?

3.3 A Sample Program Source Listing

A simple program that illustrates how this text presents program source code is seen in [Listing 3.1](#). This program will place a zero in each of the 4 registers named x28, x29, x30 and x31.

Listing 3.1: `zero4regs.S`
Setting four registers to zero.

```
.text                # put this into the text section
.align 2             # align to 2^2
.globl _start
_start:
    addi    x28, x0, 0    # set register x28 to zero
    addi    x29, x0, 0    # set register x29 to zero
    addi    x30, x0, 0    # set register x30 to zero
    addi    x31, x0, 0    # set register x31 to zero
```

This program listing illustrates a number of things:

- Listings are identified by the name of the file within which they are stored. This listing is from a file named: `zero4regs.S`.
- The assembly language programs discussed in this text will be saved in files that end with: `.S` (Alternately you can use `.sx` on systems that don't understand the difference between upper and lowercase letters.¹)
- A description of the listing's purpose appears under the name of the file. The description of [Listing 3.1](#) is *Setting four registers to zero*.
- The lines of the listing are numbered on the left margin for easy reference.
- An assembly program consists of lines of plain text.
- The RISC-V ISA does not provide an operation that will simply set a register to a numeric value. To accomplish our goal this program will add zero to zero and place the sum in each of the four registers.
- The lines that start with a dot `'.'` (on lines 1, 2 and 3) are called *assembler directives* as they tell the assembler itself how we want it to translate the following *assembly language instructions* into *machine language instructions*.
- Line 4 shows a *label* named `_start`. The colon at the end is the indicator to the assembler that causes it to recognize the preceding characters as a label.
- Lines 5-8 are the four assembly language instructions that make up the program. Each instruction in this program consists of four *fields*. (Different instructions can have a different number of fields.) The fields on line 5 are:
 - `addi` The instruction mnemonic. It indicates the operation that the CPU will perform.
 - `x28` The *destination* register that will receive the sum when the *addi* instruction is finished. The names of the 32 registers are expressed as `x0 – x31`.
 - `x0` One of the addends of the sum operation. (The `x0` register will always contain the value zero. It can never be changed.)
 - `0` The second addend is the number zero.
- `# set ...` Any text anywhere in a RISC-V assembly language program that starts with the pound-sign is ignored by the assembler. They are used to place a *comment* in the program to help the reader better understand the motive of the programmer.

3.4 Running a Program With rvddt

To illustrate what a CPU does when it executes instructions this text will use the [rvddt](#) simulator to display shows sequence of events and the binary values involved. This simulator supports the RV32I ISA and has a configurable amount of memory.²

[Listing 3.2](#) shows the operation of the four *addi* instructions from [Listing 3.1](#) when it is executed in trace-mode.

¹The author of this text prefers to avoid using such systems.

²The *rvddt* simulator was written to generate the listings for this text. It is similar to the fancier *spike* simulator. Given the simplicity of the RV32I ISA, *rvddt* is less than 1700 lines of C++ and was written in one (long) afternoon.

Listing 3.2: zero4regs.out

Running a program with the rvddt simulator

```

774 [winans@w510 src]$ ./rvddt -f ../t1/load4regs.bin
775 Loading '../t1/load4regs.bin' to 0x0
776
777 ddt> t4
778   x0: 00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
779   x8: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
780  x16: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
781  x24: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
782   pc: 00000000
783 00000000: 00000e13 addi    x28, x0, 0    # x28 = 0x00000000 = 0x00000000 + 0x00000000
784   x0: 00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
785   x8: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
786  x16: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
787  x24: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0
788   pc: 00000004
789 00000004: 00000e93 addi    x29, x0, 0    # x29 = 0x00000000 = 0x00000000 + 0x00000000
790   x0: 00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
791   x8: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
792  x16: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
793  x24: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-00000000 00000000 f0f0f0f0 f0f0f0f0
794   pc: 00000008
795 00000008: 00000f13 addi    x30, x0, 0    # x30 = 0x00000000 = 0x00000000 + 0x00000000
796   x0: 00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
797   x8: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
798  x16: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
799  x24: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-00000000 00000000 00000000 f0f0f0f0
800   pc: 0000000c
801 0000000c: 00000f93 addi    x31, x0, 0    # x31 = 0x00000000 = 0x00000000 + 0x00000000
802 ddt> r
803   x0: 00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
804   x8: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
805  x16: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
806  x24: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-00000000 00000000 00000000 00000000
807   pc: 00000010
808 ddt> x
809 [winans@w510 src]$
810

```

- ℓ 1 This listing includes the command-line that shows how the simulator was executed to load a file containing the machine instructions (aka machine code) from the assembler.
- ℓ 2 A message from the simulator indicating that it loaded the machine code into simulated memory at address 0.
- ℓ 3 This line shows the prompt from the debugger and the command `t4` that the user entered to request that the simulator trace the execution of four instructions.
- ℓ 4-8 Prior to executing the first instruction, the state of the CPU registers is displayed.
- ℓ 4 The values in registers 0, 1, 2, 3, 4, 5, 6 and 7 are printed from left to right in [big endian](#), [hexadecimal](#) form. The dash ‘-’ character in the middle of the line is a reference to make it easier to visually navigate across the line without being forced to count the values from the far left when seeking the value of, say, `x5`.
- ℓ 5-7 The values of registers 8–31 are printed.
- ℓ 8 The *program counter* (`pc`) register is printed. It contains the address of the instruction that the CPU will execute. After each instruction, the `pc` will either advance four bytes ahead or be set to another value by a branch instruction as discussed above.
- ℓ 9 A four-byte instruction is fetched from memory at the address in the `pc` register, is decoded and printed. From left to right the fields shown on this line are:

```

828 00000000 The memory address from which the instruction was fetched. This address is displayed in big
829 endian, hexadecimal form.
830 00000e13 The machine code of the instruction displayed in big endian, hexadecimal form.
831     addi The mnemonic for the machine instruction.
832     x28 The rd field of the addi instruction.
833     x0 The rs1 field of the addi instruction that holds one of the two addends of the operation.
834     0 The imm field of the addi instruction that holds the second of the two addends of the operation.
835     # ... A simulator-generated comment that explains what the instruction is doing. For this instruc-
836         tion it indicates that x28 will have the value zero stored into it as a result of performing the
837         addition: 0 + 0.

838ℓ 10-14 These lines are printed as the prelude while tracing the second instruction. Lines 7 and 13 show
839     that x28 has changed from f0f0f0f0 to 00000000 as a result of executing the first instruction and
840     lines 8 and 14 show that the pc has advanced from zero (the location of the first instruction) to
841     four, where the second instruction will be fetched. None of the rest of the registers have changed
842     values.

843 ℓ 15 The second instruction decoded executed and described. This time register x29 will be assigned a
844     value.

845ℓ 16-27 The third and fourth instructions are traced.

846 ℓ 28 Tracing has completed. The simulator prints its prompt and the user enters the ‘r’ command to
847     see the register state after the fourth instruction has completed executing.

848ℓ 29-33 Following the fourth instruction it can be observed that registers x28, x29, x30 and x31 have been
849     set to zero and that the pc has advanced from zero to four, then eight, then 12 (the hex value for
850     12 is c) and then to 16 (which, in hex, is 10).

851 ℓ 34 The simulator exit command ‘x’ is entered by the user and the terminal displays the shell prompt.

```

Chapter 4

Using The RISC-V GNU Toolchain

This chapter discusses using the GNU toolchain elements to experiment with the material in this book.

See [Appendix A](#) if you do not already have the GNU crosscompiler toolchain available on your system.

Discuss the choice of ilp32 as well as what the other variations would do.

Discuss rv32im and note that the details are found in [chapter 6](#).

Discuss installing and using one of the RISC-V simulators here.

Describe the pre-processor, compiler, assembler and linker.

Source, object, and binary files

Assembly syntax (label: mnemonic op1, op2, op3 # comment).

text, data, bss, stack

Labels and scope.

Forward & backward references to throw-away labels.

The entry address of an application.

.s file contain assembler code. .S (or .sx) files contain assembler code that must be preprocessed. [[14](#), p. 29]

Pre-processing conditional assembly using #if.

Building with -mabi=ilp32 -march=rv32i -mno-fdiv -mno-div to match the config options on the toolchain.

Linker scripts.

Makefiles

objdump

nm

Chapter 5

Writing RISC-V Programs

Ideas for order of introducing operations and instructions.

► Fix Me:
Introduce the register names and aliases here. Should we use ISA names or actual names here?

5.1 Using addi to Set Register Values

```
addi    t0, zero, 4      # t0 = 4
addi    t1, t1, 100      # t1 = 104

addi    t0, zero, 0x123   # t0 = 0x123
addi    t0, t0, 0xfff     # t0 = 0x122 (subtract 1)

addi    t0, zero, 0xfff   # t0 = 0xffffffff (-1) (diagram out the chaining carry)
                                # refer back to the overflow/truncation discussion in binary chapter
```

Demonstrate various addi instructions.

5.2 Using ebreak to Stopping rvddt Execution

Introduce ebreak & demonstrate.

5.3 Other Instructions With Immediate Operands

```
addi
andi
ori
xori

slti
sltiu
```



```
900     srai
901     slli
902     srli
```

903 5.4 Transferring Data Between Registers and Memory

904 RV is a load-store architecture. This means that the only way that the CPU can interact with the
905 memory is via the *load* and *store* instructions. All other data manipulation must be performed on
906 register values.

907 Copying values from memory to a register (first examples using regs set with addi):

```
908     lb
909     lh
910     lw
911     lbu
912     lhu
```

913 Copying values from a register to memory:

```
914     sb
915     sh
916     sw
```

917

➡ Fix Me:

*Mention the rvddt UART
I/O address for writing to
the console here?*

918 5.5 RR operations

```
919     add
920     sub
921     and
922     or
923     sra
924     srl
925     sll
926     xor
927     sltu
928     slt
```

929 5.6 Setting values to large values using lui with addi

```
930     addi    // useful for values from -2048 to 2047
931     lui      // useful for loading any multiple of 0x1000
```

932

933 Setting a register to any other value must be done using a combo of insns:

934

```

935     auipc      // Load an address relative the the current PC (see la pseudo)
936     addi
937
938
939     lui        // Load constant into into bits 31:12 (see li pseudo)
940     addi        // add a constant to fill in bits 11:0
941                if bit 11 is set then need to +1 the lui value to compensate

```

5.7 Labels and Branching

Start to introduce addressing here?

```

944     beq
945     bne
946     blt
947     bge
948     bltu
949     bgeu
950
951     bgt rs, rt, offset    # pseudo for: blt rt, rs, offset    (reverse the operands)
952     ble rs, rt, offset    # pseudo for: bge rt, rs, offset    (reverse the operands)
953     bgtu rs, rt, offset   # pseudo for: bltu rt, rs, offset   (reverse the operands)
954     bleu rs, rt, offset   # pseudo for: bgeu rt, rs, offset   (reverse the operands)
955
956     beqz beqz rs, offset  # pseudo for: beq rs, x0, offset
957     bnez rs, offset       # pseudo for: bne rs, x0, offset
958     blez rs, offset       # pseudo for: bge x0, rs, offset
959     bgez rs, offset       # pseudo for: bge rs, x0, offset
960     bltz rs, offset       # pseudo for: blt rs, x0, offset
961     bgtz rs, offset       # pseudo for: blt x0, rs, offset

```

5.8 Relocation

```

963 Absolute:
964     %hi(symbol)
965     %lo(symbol)
966
967 PC-relative:
968     %pcrel_hi(symbol)
969     %pcrel_lo(label)
970
971 Using the auipc & addi pair with label references:
972     The %pcrel_lo() uses the label to find the associated %pcrel_hi()
973     The label MUST be on a line that used a %pcrel_hi() or get an error.
974     This is needed to calculate the proper offset.
975     Things like this are legal (though not sure of the value):
976     label: auipc    t1, %pcrel_hi(symbol)
977            addi     t2, t1, %pcrel_lo(label)
978            addi     t3, t1, %pcrel_lo(label)

```

```

979         lw      t4, %pcrel_lo(label)(t1)
980         sw      t5, %pcrel_lo(label)(t1)
981
982 Discuss how relaxation works.
983 see: https://github.com/riscv/riscv-elf-psabi-doc/blob/master/riscv-elf.md

```

5.9 Jumps

Introduce and present subroutines but not nesting until introduce stack operations.

```

986     jal
987     jalr

```

5.10 Pseudo Operations

```

989
990     la rd, symbol
991         auipc rd, symbol[31:12]
992         addi rd, rd, symbol[11:0]
993
994     l{b|h|w|d} rd, symbol
995         auipc rd, symbol[31:12]
996         l{b|h|w|d} rd, symbol[11:0](rd)
997
998     s{b|h|w|d} rd, symbol, rt           # rt is the temp reg to use for the operation
999         auipc rt, symbol[31:12]
1000         s{b|h|w|d} rd, symbol[11:0](rt)
1001
1002
1003     j offset        jal x0, offset
1004     jal offset      jal x1, offset
1005     jr rs           jalr x0, rs, 0
1006     jalr rs         jalr x1, rs, 0
1007     ret             jalr x0, x1, 0
1008
1009     call offset     auipc x6, offset[31:12]
1010                   jalr x1, x6, offset[11:0]
1011
1012     tail offset     auipc x6, offset[31:12]   # same as call but no x1
1013                   jalr x0, x6, offset[11:0]

```

5.11 The Linker and Relaxation

I don't know where this should go just yet.

► Fix Me:

Explain why we have pseudo ops. These mappings are lifted from the ISM, Vol 1, V2.2

► Fix Me:

Needs research. I'm not sure if/how the linker alone can relax the AUIPC+JALR pair since the assembler could have used a pcrel branch across one of these pairs.

1016 5.12 pic and nopic

1017 pic is *needed* for shared libs. Should discuss it but probably best to leave the topic for a later chapter.

Chapter 6

RV32 Machine Instructions

6.1 Introduction

6.2 Conventions and Terminology

When discussing instructions, the following abbreviations/notations are used:

6.2.1 XLEN

XLEN represents the bit-length of an `x` register in the machine architecture. Possible values are 32, 64 and 128.

6.2.2 `sx(val)`

Sign extend *val* to the left.

This is used to convert a signed integer value expressed using some number of bits to a larger number of bits by adding more bits to the left. In doing so, the sign will be preserved. In this case *val* represents the least MSBs of the value. For more on binary numbers see [Appendix B](#).

[Figure 6.1](#) illustrates extending the negative sign bit of *val* to the left by replicating it. When *val* is negative, its MSB (bit 19 in this example) will be set to 1. Extending this value to the left will set all the new bits to the left of it to 1 as well.




Figure 6.1: Sign-extending a negative integer from 20 bits to 32 bits.

[Figure 6.2](#) illustrates extending the positive sign bit of *val* to the left by replicating it. When *val* is

positive, its **MSB** will be set to 0. Extending this value to the left will set all the new bits to the left of it to 0 as well.




Figure 6.2: Sign-extending a positive integer from 20 bits to 32 bits.

6.2.3 `zx(val)`

Zero extend *val* to the left.

This is used to convert an unsigned integer value expressed using some number of bits to a larger number of bits by adding more bits to the left. In doing so, the new bits added will all be set to zero. As is the case with `sx(val)`, *val* represents the **LSBs** of the final value. Figure 6.3 illustrates zero-extending a 20-bit *val* to the left to form a 32-bit fullword.

For more on binary numbers see [Appendix B](#).




Figure 6.3: Zero-extending an unsigned integer from 20 bits to 32 bits.

6.2.4 `zr(val)`

Zero extend *val* to the right.

Some times a binary value is encoded such that a set of bits represented by *val* are used to represent the **MSBs** of some longer (more bits) value. In this case it is necessary to append zeros to the right to convert *val* to the longer value.

Figure 6.4 illustrates converting a 20-bit *val* to a 32-bit fullword.




Figure 6.4: Zero-extending an integer to the right from 20 bits to 32 bits.

6.2.5 Sign Extended Left and Zero Extend Right

Some instructions such as the J-type (see [subsection 6.4.2](#)) include immediate operands that are extended in both directions.

[Figure 6.5](#) and [Figure 6.6](#) illustrates zero-extending a 20-bit negative number one bit to the right and sign-extending it 11 bits to the left:




Figure 6.5: Sign-extending a positive 20-bit number 11 bits to the left and one bit to the right.




Figure 6.6: Sign-extending a negative 20-bit number 11 bits to the left and one bit to the right.

6.2.6 $\text{m8}(\text{addr})$

The contents of an 8-bit value in memory at address *addr*.

Given the contents of the memory dump shown in [Figure 6.7](#), $\text{m8}(42)$ refers to the memory location at address 42_{16} that currently contains the 8-bit value fc_{16} .

The $\text{mn}(\text{addr})$ notation can be used to refer to memory that is being read or written depending on the context.

When memory is being written, the following notation is used to indicate that the least significant 8 bits of *source* will be written into memory at the address *addr*:

$\text{m8}(\text{addr}) \leftarrow \text{source}$

When memory is being read, the following notation is used to indicate that the 8 bit value at the address *addr* will be read and stored into *dest*:

$\text{dest} \leftarrow \text{m8}(\text{addr})$

Note that *source* and *dest* are typically registers.

```
00000030  2f 20 72 65 61 64 20 61  20 62 69 6e 61 72 79 20
00000040  66 69 fc 65 20 66 69 6c  6c 65 64 20 77 69 74 68
00000050  20 72 76 33 32 49 20 69  6e 73 74 72 75 63 74 69
00000060  6f 6e 73 20 61 6e 64 20  66 65 65 64 20 74 68 65
```

Figure 6.7: Sample memory contents.

1068 **6.2.7 m16(addr)**

1069 The contents of an 16-bit little-endian value in memory at address *addr*.

1070 Given the contents of the memory dump shown in [Figure 6.7](#), `m16(42)` refers to the memory location at
1071 address `4216` that currently contains `65fc16`. See also [subsection 6.2.6](#).

1072 **6.2.8 m32(addr)**

1073 The contents of an 32-bit little-endian value in memory at address *addr*.

1074 Given the contents of the memory dump shown in [Figure 6.7](#), `m32(42)` refers to the memory location at
1075 address `4216` that currently contains `662065fc16`. See also [subsection 6.2.6](#).

1076 **6.2.9 m64(addr)**

1077 The contents of an 64-bit little-endian value in memory at address *addr*.

1078 Given the contents of the memory dump shown in [Figure 6.7](#), `m64(42)` refers to the memory location at
1079 address `4216` that currently contains `656c6c69662065fc16`. See also [subsection 6.2.6](#).

1080 **6.2.10 m128(addr)**

1081 The contents of an 128-bit little-endian value in memory at address *addr*.

1082 Given the contents of the memory dump shown in [Figure 6.7](#), `m128(42)` refers to the memory location at
1083 address `4216` that currently contains `7220687469772064656c6c69662065fc16`. See also [subsection 6.2.6](#).

1084 **6.2.11 .+offset**

1085 The address of the current instruction plus a numeric offset.

1086 **6.2.12 .-offset**

1087 The address of the current instruction minus a numeric offset.

1088 **6.2.13 pc**

1089 The current value of the program counter.

1090 **6.2.14 rd**

1091 An x-register used to store the result of instruction.

6.2.15 rs1

An x-register value used as a source operand for an instruction.

6.2.16 rs2

An x-register value used as a source operand for an instruction.

6.2.17 imm

An immediate numeric operand. The word *immediate* refers to the fact that the operand is stored within an instruction.

6.2.18 rsN[h:l]

The value of bits from h through l of x-register rsN. For example: rs1[15:0] refers to the contents of the 16 [LSBs](#) of rs1.

6.3 Addressing Modes

immediate, register, base-displacement, pc-relative

► Fix Me:

Write this section.

6.4 Instruction Encoding Formats

This document concerns itself with the following RISC-V instruction formats.

XXX Show and discuss a stack of formats explaining how the unnatural ordering of the *imm* fields reduces the number of possible locations that the hardware has to be prepared to *look* for various bits. For example, the opcode, rd, rs1, rs1, func3 and the sign bit (when used) are all always in the same position. Also note that imm[19:12] and imm[10:5] can only be found in one place. imm[4:0] can only be found in one of two places...

► Fix Me:

Should discuss types and sizes beyond the fundamentals. Will add if/when instruction details are added in the future.

The point to all this is that it is easier to build a machine if it does not have to accommodate many different ways to perform the same task. This simplification can also allow it operate faster.

[Figure 6.8](#) Shows the RISC-V instruction formats.

6.4.1 U Type

The U-Type format is used for instructions that use a 20-bit immediate operand and a destination register.

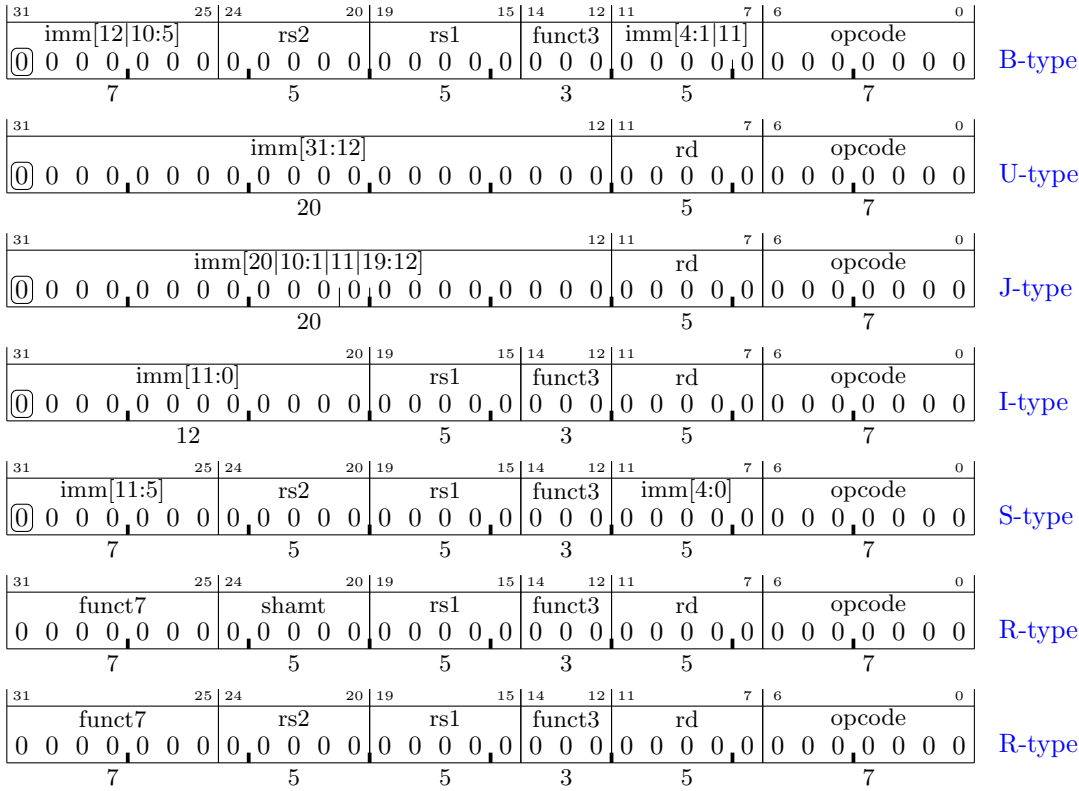



Figure 6.8: RISC-V instruction formats.




1117

1118 The **rd** field contains an **x** register number to be set to a value that depends on the instruction.

1119 The **imm** field contains a 20-bit value that will be converted into **XLEN** bits by using the *imm* operand
 1120 for bits 31:12 and then sign-extending it to the left¹ and zero-extending the LSBs as discussed in
 1121 subsection 6.2.4.

1122 If **XLEN**=32 then the **imm** value in this example will be converted as shown below.



1123

1124 Notice that the 20-bits of the **imm** field are mapped in the same order and in the same relative position
 1125 that they appear in the instruction when they are used to create the value of the immediate operand.
 1126 Shifting the **imm** value to the left, into the “upper bits” of the immediate value suggests a rationale for
 1127 the name of this format.


1128 If **XLEN**=64 then the **imm** value in this example will be converted to the same two’s complement integer

¹When **XLEN** is larger than 32.

value by extending the sign to the left.

6.4.2 J Type

The J-type format is used for instructions that use a 20-bit immediate operand and a destination register. It is similar to the U-type. However, the immediate operand is constructed by arranging the *imm* bits in a different manner.




The *rd* field contains an *x* register number to be set to a value that depends on the instruction.

In the J-type format the 20 *imm* bits are arranged such that they represent the “lower” portion of the immediate value. Unlike the U-type instructions, the J-type requires the bits to be re-ordered and shifted to the right before they are used.²


The example above shows that the bit positions in the *imm* field description. We see that the 20 *imm* bits are re-ordered according to: [20:10:11:19:12]. This means that the *MSB* of the *imm* field is to be placed into bit 20 of the immediate integer value ultimately used by the instruction when it is converted into *XLEN* bits. The next bit to the right in the *imm* field is to be placed into bit 10 of the immediate value and so on.

After the *imm* bits are re-positioned into bits 20:1 of the immediate value being constructed, a zero-bit will be added to the *LSB* and the value in bit-position 20 will be replicated to sign-extend the value to *XLEN* bits as discussed in subsection 6.2.5.


If *XLEN*=32 then the *imm* value in this example will be converted as shown below.



A J-type example with a negative *imm* field:




If *XLEN*=32 then the *imm* field in this example will be converted as shown below.




²The reason that the J-type bits are reordered like this is because it simplifies the implementation of hardware as discussed in section 6.4.

The J-type format is used by the Jump And Link instruction that calculates a target address by adding a signed immediate value to the current program counter. Since no instruction can be placed at an odd address the 20-bit imm value is zero-extended to the right to represent a 21-bit signed offset capable of representing numbers twice the magnitude of the 20-bit imm value.


6.4.3 R Type




A special case of the R-type used for shift-immediate instructions where the *rs2* field is used as an immediate value named *shamt* representing the number of bit positions to shift:




6.4.4 I Type



6.4.5 S Type



6.4.6 B Type



6.4.7 CPU Registers

The registers are names x0 through x31 and have aliases suited to their conventional use. The following table describes each register.

Note that the calling convention specifies that only some of the registers are to be saved by functions if they alter their contents. The idea being that accessing memory is time-consuming and that

► Fix Me:
Need to add a section that discusses the calling conventions

by classifying some registers as “temporary” (not saved by any function that alter its contents) it is possible to carefully implement a function with less need to store register values on the stack in order to use them to perform the operations of the function.

The lack of grouping the temporary and saved registers is due to the fact that the C extension provides access to only the first 16 registers when executing instructions in the compressed format.

Reg	Alias	Description	Saved
x0	zero	Hard-wired zero	yes
x1	ra	Return address	
x2	sp	Stack pointer	
x3	gp	Global pointer	
x4	tp	Thread pointer	
x5	t0	Temporary/alternate link register	
x6	t1	Temporary	
x7	t2	Temporary	yes
x8	s0/fp	Saved register/frame pointer	
x9	s1	Saved register	yes
x10	a0	Function argument/return value	yes
x11	a1	Function argument/return value	
x12	a2	Function argument	
x13	a3	Function argument	
x14	a4	Function argument	
x15	a5	Function argument	
x16	a6	Function argument	
x17	a7	Function argument	
x18	s2	Saved register	yes
x19	s3	Saved register	yes
x20	s4	Saved register	yes
x21	s5	Saved register	yes
x22	s6	Saved register	yes
x23	s7	Saved register	yes
x24	s8	Saved register	yes
x25	s9	Saved register	yes
x26	s10	Saved register	yes
x27	s11	Saved register	yes
x28	t3	Temporary	yes
x29	t4	Temporary	
x30	t5	Temporary	
x31	t6	Temporary	

6.5 memory

Note that RISC-V is a little-endian machine.

All instructions must be naturally aligned to their 4-byte boundaries. [1, p. 5]

If a RISC-V processor implements the C (compressed) extension then instructions may be aligned to 2-byte boundaries. [1, p. 68]

Data alignment is not necessary but unaligned data can be inefficient. Accessing unaligned data using any of the load or store instructions can also prevent a memory access from operating atomically. [1,

1186 p.19] See also ??.

RV32I
Instruction!LUI
Instruction!AUIPC

1187 6.6 RV32I Base Instruction Set

1188 [RV32I](#) refers to the basic 32-bit integer instructions.

1189 6.6.1 LUI rd, imm


1190 Load Upper Immediate.

1191 $rd \leftarrow zr(imm)$

1192 Copy the immediate value into bits 31:12 of the destination register and place zeros into bits 11:0. When
1193 XLEN is 64 or 128, the immediate value is sign-extended to the left.


1194 Instruction Format and Example:

1195 LUI t0, 3



1198 00010074: 000032b7 lui x5, 0x3 // x5 = 0x3000
 1199 reg 0: 00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 00003000 f0f0f0f0 f0f0f0f0
 1200 reg 8: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
 1201 reg 16: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
 1202 reg 24: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
 1203 pc: 00010078

1204 LUI t0, 0xfffff



1207 00010078: fffff2b7 lui x5, 0xfffff // x5 = 0xfffff000
 1208 reg 0: 00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 fffff000 f0f0f0f0 f0f0f0f0
 1209 reg 8: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
 1210 reg 16: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
 1211 reg 24: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
 1212 pc: 0001007c

1213 6.6.2 AUIPC rd, imm

1214 Add Upper Immediate to PC.

Instruction!JAL

When XLEN is 64 or 128, the immediate value is also sign-extended to the left prior to being added to the pc register.

1221



1230




1243 6.6.3 JAL rd, imm

The standard software conventions for calling subroutines use `x1` as the return address (`rd` register in this case). [1, p. 16]

Instruction!JALR

JAL x7, .+16



imm demultiplexed value = 000000000000000001000₂ \ll 1 = 16₁₀

State of registers before execution:


```
pc = 0x11114444
```

State of registers after execution:

```
pc = 0x11114454 x7 = 0x11114448
```

JAL provides a method to call a subroutine using a pc-relative address.

JAL x7, .-16



imm demultiplexed value = 1111111111111111000₂ \ll 1 = -16₁₀

State of registers before execution:

```
pc = 0x1114444
```

State of registers after execution:

```
pc = 0x11114434 x7 = 0x11114448
```

6.6.4 JALR rd, rs1, imm

Jump and link register.

$$\text{rd} \leftarrow \text{pc} + 4$$
$$pc \leftarrow (rs1 + sx(imm)) \& \sim 1$$


This instruction saves the address of the next instruction that would otherwise execute (located at **pc**+4) into **rd** and then adds the immediate value to the **rs1** register and stores the sum into the **pc** register causing an unconditional branch to take place.

Note that the branch target address is calculated by sign-extending the imm[11:0] bits from the instruction, adding it to the **rs1** register and *then* the LSB of the sum is to zero and the result is stored into the **pc** register. The discarding of the LSB allows the branch to refer to any even address.

The standard software conventions for calling subroutines use `x1` as the return address (`rd` register in this case). [1, p. 16]

1281 Encoding:

1282 JALR x1, x7, 4



1285 Before:

1286 pc = 0x11114444

1287 x7 = 0x44444444


1288 After

1289 pc = 0x5555888c

1290 x1 = 0x11114448

1291 JALR provides a method to call a subroutine using a base-displacement address.

1292 JALR x1, x0, 5



1295 Note that the least significant bit in the result of rs1+imm is discarded/set to zero before the result is
 1296 saved in the pc.

1297 pc = 0x11114444

1298 After

1299 pc = 0x00000004

1300 x1 = 0x11114448


1301 **6.6.5 BEQ rs1, rs2, imm**

1302 Branch if equal.

1303 $pc \leftarrow (rs1 == rs2) ? pc + sx(imm[12:1] \ll 1) : pc + 4$

1304 Encoding:

1305 BEQ x3, x15, 2064

1308 $imm[12:1] = 010000001000_2 = 1032_{10}$ 1309 $imm = 2064_{10}$ 1310 $funct3 = 000_2$

1311 $rs1 = x3$ Instruction!BNE
 1312 $rs2 = x15$ Instruction!BLT
 Instruction!BGE

1313 6.6.6 BNE $rs1, rs2, imm$

1314 Branch if Not Equal.

1315 $pc \leftarrow (rs1 \neq rs2) ? pc + sx(imm[12:1] \ll 1) : pc + 4$

1316 Encoding:

1317 BNE $x3, x15, 2064$

1318

31	25	24	20	19	15	14	12	11	7	6	0	
imm[12:10:5]			rs2		rs1		funct3		imm[4:1:11]		opcode	
0 0 0 0 0 0			0 1 1 1 1		0 0 0 1 1		0 0 1		1 0 0 0 1		1 1 0 0 1 1	
7			5		5		3		5		7	

B-type

1319

1320 $imm[12:1] = 010000001000_2 = 1032_{10}$

1321 $imm = 2064_{10}$

1322 $funct3 = 001_2$

1323 $rs1 = x3$

1324 $rs2 = x15$

1325 6.6.7 BLT $rs1, rs2, imm$

1326 Branch if Less Than.

1327 $pc \leftarrow (rs1 < rs2) ? pc + sx(imm[12:1] \ll 1) : pc + 4$

1328 Encoding:

1329 BLT $x3, x15, 2064$

1330

31	25	24	20	19	15	14	12	11	7	6	0	
imm[12:10:5]			rs2		rs1		funct3		imm[4:1:11]		opcode	
0 0 0 0 0 0			0 1 1 1 1		0 0 0 1 1		1 0 0		1 0 0 0 1		1 1 0 0 1 1	
7			5		5		3		5		7	

B-type

1331

1332 $imm[12:1] = 010000001000_2 = 1032_{10}$

1333 $imm = 2064_{10}$

1334 $funct3 = 100_2$

1335 $rs1 = x3$

1336 $rs2 = x15$


1337 6.6.8 BGE $rs1, rs2, imm$

1338 Branch if Greater or Equal.

1339 $pc \leftarrow (rs1 \geq rs2) ? pc + sx(imm[12:1] \ll 1) : pc + 4$

1340 Encoding:

1341 BGE x3, x15, 2064


Instruction!BLTU
Instruction!BGEU1343
1344 $\text{imm}[12:1] = 010000001000_2 = 1032_{10}$ 1345 $\text{imm} = 2064_{10}$ 1346 $\text{funct3} = 101_2$ 1347 $\text{rs1} = \text{x3}$ 1348 $\text{rs2} = \text{x15}$ 1349 **6.6.9 BLTU rs1, rs2, imm**

1350 Branch if Less Than Unsigned.

1351 $\text{pc} \leftarrow (\text{rs1} < \text{rs2}) ? \text{pc} + \text{sx}(\text{imm}[12:1] \ll 1) : \text{pc} + 4$

1352 Encoding:

1353 BLTU x3, x15, 2064


1355
1356 $\text{imm}[12:1] = 010000001000_2 = 1032_{10}$ 1357 $\text{imm} = 2064_{10}$ 1358 $\text{funct3} = 110_2$ 1359 $\text{rs1} = \text{x3}$ 1360 $\text{rs2} = \text{x15}$ 1361 **6.6.10 BGEU rs1, rs2, imm**

1362 Branch if Greater or Equal Unsigned.

1363 $\text{pc} \leftarrow (\text{rs1} \geq \text{rs2}) ? \text{pc} + \text{sx}(\text{imm}[12:1] \ll 1) : \text{pc} + 4$

1364 Encoding:

1365 BGEU x3, x15, 2064

1367
1368 $\text{imm}[12:1] = 010000001000_2 = 1032_{10}$ 1369 $\text{imm} = 2064_{10}$ 1370 $\text{funct3} = 111_2$ 1371 $\text{rs1} = \text{x3}$ 1372 $\text{rs2} = \text{x15}$

► Fix Me:

use symbols in branch examples


1373 **6.6.11 LB rd, imm(rs1)**Instruction!LB
Instruction!LH
Instruction!LW

1374 Load byte.

1375 $rd \leftarrow sx(m8(rs1+sx(imm)))$ 1376 $pc \leftarrow pc+4$ 1377 Load an 8-bit value from memory at address $rs1+imm$, then sign-extend it to 32 bits before storing it in
1378 rd

1379 Encoding:

1380 LB x7, 4(x3)


1383 **6.6.12 LH rd, imm(rs1)**

1384 Load halfword.

1385 $rd \leftarrow sx(m16(rs1+sx(imm)))$ 1386 $pc \leftarrow pc+4$ 1387 Load a 16-bit value from memory at address $rs1+imm$, then sign-extend it to 32 bits before storing it in
1388 rd

1389 Encoding:

1390 LH x7, 4(x3)

1393 **6.6.13 LW rd, imm(rs1)**


1394 Load word.

1395 $rd \leftarrow sx(m32(rs1+sx(imm)))$ 1396 $pc \leftarrow pc+4$ 1397 Load a 32-bit value from memory at address $rs1+imm$, then store it in rd

1398 Encoding:

1399 LW x7, 4(x3)

1400



Instruction!LBU
Instruction!LHU
Instruction!SB

1401

1402 6.6.14 LBU rd, imm(rs1)

1403 Load byte unsigned.

1404 $rd \leftarrow zx(m8(rs1+sx(imm)))$


1405 $pc \leftarrow pc+4$

1406 Load an 8-bit value from memory at address **rs1+imm**, then zero-extend it to 32 bits before storing it in
1407 **rd**

1408 Encoding:

1409 LBU x7, 4(x3)

1410



1411

1412 6.6.15 LHU rd, imm(rs1)

1413 Load halfword unsigned.

1414 $rd \leftarrow zx(m16(rs1+sx(imm)))$


1415 $pc \leftarrow pc+4$

1416 Load an 16-bit value from memory at address **rs1+imm**, then zero-extend it to 32 bits before storing it
1417 in **rd**

1418 Encoding:

1419 LHU x7, 4(x3)

1420



1421

1422 6.6.16 SB rs2, imm(rs1)

1423 Store Byte.

1424 $m8(rs1+sx(imm)) \leftarrow rs2[7:0]$


1425 $pc \leftarrow pc+4$

1426 Store the 8-bit value in **rs2[7:0]** into memory at address **rs1+imm**.

1427 Encoding:

Instruction!SH
Instruction!SW
Instruction!ADDI

1428 SB x3, 19(x15)


1431 **6.6.17 SH rs2, imm(rs1)**

1432 Store Halfword.

1433 $m16(rs1+sx(imm)) \leftarrow rs2[15:0]$ 1434 $pc \leftarrow pc+4$ 1435 Store the 16-bit value in `rs2[15:0]` into memory at address `rs1+imm`.

1436 Encoding:

1437 SH x3, 19(x15)


1440 **6.6.18 SW rs2, imm(rs1)**

1441 Store Word

1442 $m16(rs1+sx(imm)) \leftarrow rs2[31:0]$ 1443 $pc \leftarrow pc+4$ 1444 Store the 32-bit value in `rs2` into memory at address `rs1+imm`.

1445 Encoding:

1446 SW x3, 19(x15)



1449 Show pos & neg imm examples.

1450 **6.6.19 ADDI rd, rs1, imm**


1451 Add Immediate

1452 $rd \leftarrow rs1+sx(imm)$ 1453 $pc \leftarrow pc+4$

1454 Encoding:

Instruction!SLTI
Instruction!SLTIU

1455 ADDI x1, x7, 4



1458 Before:

1459 x7 = 0x11111111

1460 After:

1461 x1 = 0x11111115

1462 6.6.20 SLTI rd, rs1, imm

1463 Set LessThan Immediate


1464 $rd \leftarrow (rs1 < sx(imm)) ? 1 : 0$

1465 $pc \leftarrow pc+4$

1466 If the sign-extended immediate value is less than the value in the **rs1** register then the value 1 is stored
1467 in the **rd** register. Otherwise the value 0 is stored in the **rd** register.

1468 Encoding:

1469 SLTI x1, x7, 4



1472 Before:

1473 x7 = 0x11111111

1474 After:

1475 x1 = 0x00000000

1476 6.6.21 SLTIU rd, rs1, imm

1477 Set LessThan Immediate Unsigned

1478 $rd \leftarrow (rs1 < sx(imm)) ? 1 : 0$

1479 $pc \leftarrow pc+4$


1480 If the sign-extended immediate value is less than the value in the **rs1** register then the value 1 is stored
1481 in the **rd** register. Otherwise the value 0 is stored in the **rd** register. Both the immediate and **rs1**

1482 register values are treated as unsigned numbers for the purposes of the comparison.³

Instruction!XORI
Instruction!ORI

1483 Encoding:

1484 SLTIU x1, x7, 4



1486 Before:

1488 x7 = 0x81111111

1489 After:

1490 x1 = 0x00000001

1491 6.6.22 XORI rd, rs1, imm

1492 Exclusive Or Immediate


1493 $rd \leftarrow rs1 \oplus sx(imm)$

1494 $pc \leftarrow pc+4$

1495 The logical XOR of the sign-extended immediate value and the value in the **rs1** register is stored in the
1496 **rd** register.

1497 Encoding:

1498 XORI x1, x7, 4



1500 Before:

1502 x7 = 0x81111111

1503 After:

1504 x1 = 0x81111115

1505 6.6.23 ORI rd, rs1, imm

1506 Or Immediate

1507 $rd \leftarrow rs1 \mid sx(imm)$


1508 $pc \leftarrow pc+4$

³The immediate value is first sign-extended to XLEN bits then treated as an unsigned number.[1, p. 14]

The logical OR of the sign-extended immediate value and the value in the **rs1** register is stored in the **rd** register. Instruction!ANDI
Instruction!SLLI

Encoding:

ORI x1, x7, 4



Before:

x7 = 0x81111111

After:

x1 = 0x81111115

6.6.24 ANDI rd, rs1, imm

And Immediate


$rd \leftarrow rs1 \ \& \ sx(imm)$

$pc \leftarrow pc+4$

The logical AND of the sign-extended immediate value and the value in the **rs1** register is stored in the **rd** register.

Encoding:

ANDI x1, x7, 4



Before:

x7 = 0x81111111

After:

x1 = 0x81111115

6.6.25 SLLI rd, rs1, shamt

Shift Left Logical Immediate


$rd \leftarrow rs1 \ll shamt$

$pc \leftarrow pc+4$

1537 SLLI is a logical left shift operation (zeros are shifted into the lower bits). The value in rs1 shifted left
 1538 shamt number of bits and the result placed into rd. [1, p. 14] Instruction!SRLI
Instruction!SRAI

1539 Encoding:

1540 SLLI x7, x3, 2



1542

1543 x3 = 0x81111111

1544 After:

1545 x7 = 0x04444444

1546 6.6.26 SRLI rd, rs1, shamt

1547 Shift Right Logical Immediate


1548 $rd \leftarrow rs1 \gg shamt$

1549 $pc \leftarrow pc+4$

1550 SRLI is a logical right shift operation (zeros are shifted into the higher bits). The value in rs1 shifted
 1551 right shamt number of bits and the result placed into rd. [1, p. 14]

1552 Encoding:

1553 SRLI x7, x3, 2



1555

1556 x3 = 0x81111111

1557 After:

1558 x7 = 0x20444444

1559 6.6.27 SRAI rd, rs1, shamt

1560 Shift Right Arithmetic Immediate

1561 $rd \leftarrow rs1 \gg shamt$


1562 $pc \leftarrow pc+4$

1563 SRAI is a logical right shift operation (zeros are shifted into the higher bits). The value in rs1 shifted
 1564 right shamt number of bits and the result placed into rd. [1, p. 14]

1565 Encoding:

1566 SRAI x7, x3, 2
 1567

Instruction!ADD
 Instruction!SUB



1568
 1569 x3 = 0x81111111

1570 After:

1571 x7 = 0xe0444444

1572 6.6.28 ADD rd, rs1, rs2

1573 Add


1574 $rd \leftarrow rs1 + rs2$

1575 $pc \leftarrow pc+4$

1576 ADD performs addition. Overflows are ignored and the low 32 bits of the result are written to rd. [1,
 1577 p. 15]

1578 Encoding:

1579 ADD x7, x3, x31



1580
 1581
 1582 x3 = 0x81111111 x31 = 0x22222222

1583 After:

1584 x7 = 0xa3333333

1585 6.6.29 SUB rd, rs1, rs2

1586 Subtract

1587 $rd \leftarrow rs1 - rs2$


1588 $pc \leftarrow pc+4$

1589 SUB performs subtraction. Underflows are ignored and the low 32 bits of the result are written to rd. [1,
 1590 p. 15]

1591 Encoding:

1592 SUB x7, x3, x31

1593



R-type

Instruction!SLL
Instruction!SLT

x3 = 0x83333333 x31 = 0x01111111

After:

x7 = 0x82222222

6.6.30 SLL rd, rs1, rs2


Shift Left Logical

 $rd \leftarrow rs1 \ll rs2$ $pc \leftarrow pc+4$

SLL performs a logical left shift on the value in register rs1 by the shift amount held in the lower 5 bits of register rs2. [1, p. 15]

Encoding:

SLL x7, x3, x31



R-type

x3 = 0x83333333

x31 = 0x00000002

After:

x7 = 0x0ccccccc

6.6.31 SLT rd, rs1, rs2


Set Less Than

 $rd \leftarrow (rs1 < rs2) ? 1 : 0$ $pc \leftarrow pc+4$

SLT performs a signed compare, writing 1 to rd if $rs1 < rs2$, 0 otherwise. [1, p. 15]

Encoding:

SLT x7, x3, x31



R-type

1621 x3 = 0x83333333
 1622 x31 = 0x00000002

Instruction!SLTU
 Instruction!XOR

1623 After:

1624 x7 = 0x00000001

1625 6.6.32 SLTU rd, rs1, rs2


1626 Set Less Than Unsigned

1627 $rd \leftarrow (rs1 < rs2) ? 1 : 0$
 1628 $pc \leftarrow pc+4$

1629 SLTU performs an unsigned compare, writing 1 to **rd** if **rs1** < **rs2**, 0 otherwise. Note, SLTU rd, x0, rs2
 1630 sets **rd** to 1 if **rs2** is not equal to zero, otherwise sets **rd** to zero (assembler pseudo-op SNEZ rd, rs). [1,
 1631 p. 15]

1632 Encoding:

1633 SLTU x7, x3, x31



1636 x3 = 0x83333333
 1637 x31 = 0x00000002

1638 After:

1639 x7 = 0x00000000

1640 6.6.33 XOR rd, rs1, rs2


1641 Exclusive Or

1642 $rd \leftarrow rs1 \oplus rs2$
 1643 $pc \leftarrow pc+4$

1644 XOR performs a bit-wise exclusive or on rs1 and rs2. The result is stored on rd.

1645 Encoding:

1646 XOR x7, x3, x31



1649 x3 = 0x83333333
 1650 x31 = 0x1888ffff

1651 After:

Instruction!SRL
Instruction!SRA

1652 $x7 = 0x9bbbcxxx$

1653 6.6.34 SRL rd, rs1, rs2

1654 Shift Right Logical


1655 $rd \leftarrow rs1 \gg rs2$

1656 $pc \leftarrow pc+4$

1657 SRL performs a logical right shift on the value in register rs1 by the shift amount held in the lower 5
1658 bits of register rs2. [1, p. 15]

1659 Encoding:

1660 SRL x7, x3, x31



1663 $x3 = 0x83333333$

1664 $x31 = 0x00000010$

1665 After:

1666 $x7 = 0x00008333$

1667 6.6.35 SRA rd, rs1, rs2

1668 Shift Right Arithmetic


1669 $rd \leftarrow rs1 \gg rs2$

1670 $pc \leftarrow pc+4$

1671 SRA performs an arithmetic right shift (the original sign bit is copied into the vacated upper bits) on
1672 the value in register rs1 by the shift amount held in the lower 5 bits of register rs2. [1, p. 14, 15]

1673 Encoding:

1674 SRA x7, x3, x31



1677 $x3 = 0x83333333$

1678 $x31 = 0x00000010$

1679 After:

1680 $x7 = 0xffff8333$

Instruction!OR
Instruction!AND

1681 6.6.36 OR rd, rs1, rs2

1682 Or


1683 $rd \leftarrow rs1 \mid rs2$

1684 $pc \leftarrow pc+4$

1685 OR is a logical operation that performs a bit-wise OR on register rs1 and rs2 and then places the result
1686 in rd. [1, p. 14]

1687 Encoding:

1688 OR x7, x3, x31



1691 $x3 = 0x83333333$

1692 $x31 = 0x00000440$

1693 After:

1694 $x7 = 0x83333773$

1695 6.6.37 AND rd, rs1, rs2

1696 And


1697 $rd \leftarrow rs1 \& rs2$

1698 $pc \leftarrow pc+4$

1699 AND is a logical operation that performs a bit-wise AND on register rs1 and rs2 and then places the
1700 result in rd. [1, p. 14]

1701 Encoding:

1702 AND x7, x3, x31



1705 $x3 = 0x83333333$

1706 $x31 = 0x00000fe2$

1707 After:

1708 $x7 = 0x00000322$

6.6.38 FENCE predecessor, successor

Instruction!FENCE
Instruction!FENCE.I

The FENCE instruction is used to order device I/O and memory accesses as viewed by other RISC-V harts and external devices or co-processors. Any combination of device input (I), device output (O), memory reads (R), and memory writes (W) may be ordered with respect to any combination of the same. Informally, no other RISC-V hart or external device can observe any operation in the successor set following a FENCE before any operation in the predecessor set preceding the FENCE. The execution environment will define what I/O operations are possible, and in particular, which load and store instructions might be treated and ordered as device input and device output operations respectively rather than memory reads and writes. For example, memory-mapped I/O devices will typically be accessed with uncached loads and stores that are ordered using the I and O bits rather than the R and W bits. Instruction-set extensions might also describe new coprocessor I/O instructions that will also be ordered using the I and O bits in a FENCE. [1, p. 21]

► Fix Me:


Which of the i, o, r and w goes into each bit? See what gas does.

Operation:

$pc \leftarrow pc+4$

Encoding:

FENCE iorw, iorw



6.6.39 FENCE.I


The FENCE.I instruction is used to synchronize the instruction and data streams. RISC-V does not guarantee that stores to instruction memory will be made visible to instruction fetches on the same RISC-V hart until a FENCE.I instruction is executed. A FENCE.I instruction only ensures that a subsequent instruction fetch on a RISC-V hart will see any previous data stores already visible to the same RISC-V hart. FENCE.I does not ensure that other RISC-V harts' instruction fetches will observe the local hart's stores in a multiprocessor system. To make a store to instruction memory visible to all RISC-V harts, the writing hart has to execute a data FENCE before requesting that all remote RISC-V harts execute a FENCE.I. [1, p. 21]

Operation:

$pc \leftarrow pc+4$

Encoding:

FENCE.I




6.6.40 ECALL

Instruction!ECALL
Instruction!EBREAK
Instruction!CSRRW
Instruction!CSRWS

The ECALL instruction is used to make a request to the supporting execution environment, which is usually an operating system. The ABI for the system will define how parameters for the environment request are passed, but usually these will be in defined locations in the integer register file. [1, p. 24]


ECALL



6.6.41 EBREAK

The EBREAK instruction is used by debuggers to cause control to be transferred back to a debugging environment. [1, p. 24]


EBREAK



6.6.42 CSRRW rd, csr, rs1

The CSRRW (Atomic Read/Write CSR) instruction atomically swaps values in the CSRs and integer registers. CSRRW reads the old value of the CSR, zero-extends the value to XLEN bits, then writes it to integer register rd. The initial value in rs1 is written to the CSR. If rd=x0, then the instruction shall not read the CSR and shall not cause any of the side-effects that might occur on a CSR read. [1, p. 22]

CSRRW x3, 2, x15



6.6.43 CSRWS rd, csr, rs1


The CSRWS (Atomic Read and Set Bits in CSR) instruction reads the value of the CSR, zero-extends the value to XLEN bits, and writes it to integer register rd. The initial value in integer register rs1 is treated as a bit mask that specifies bit positions to be set in the CSR. Any bit that is high in rs1 will cause the corresponding bit to be set in the CSR, if that CSR bit is writable. Other bits in the CSR are unaffected (though CSRs might have side effects when written). [1, p. 22]

If rs1=x0, then the instruction will not write to the CSR at all, and so shall not cause any of the side effects that might otherwise occur on a CSR write, such as raising illegal instruction exceptions on accesses to read-only CSRs. Note that if rs1 specifies a register holding a zero value other than x0, the

instruction will still attempt to write the unmodified value back to the CSR and will cause any attendant side effects. [1, p. 22]

Instruction!CSRRC
Instruction!CSRRWI
Instruction!CSRRSI

CSRRS x3, 2, x15




6.6.44 CSRRC rd, csr, rs1

The CSRRC (Atomic Read and Clear Bits in CSR) instruction reads the value of the CSR, zero-extends the value to XLEN bits, and writes it to integer register rd. The initial value in integer register rs1 is treated as a bit mask that specifies bit positions to be cleared in the CSR. Any bit that is high in rs1 will cause the corresponding bit to be cleared in the CSR, if that CSR bit is writable. Other bits in the CSR are unaffected. [1, p. 22]

If rs1=x0, then the instruction will not write to the CSR at all, and so shall not cause any of the side effects that might otherwise occur on a CSR write, such as raising illegal instruction exceptions on accesses to read-only CSRs. Note that if rs1 specifies a register holding a zero value other than x0, the instruction will still attempt to write the unmodified value back to the CSR and will cause any attendant side effects. [1, p. 22]


CSRRC x3, 2, x15



6.6.45 CSRRWI rd, csr, imm

This instruction is the same as CSRRW except a 5-bit unsigned (zero-extended) immediate value is used rather than the value from a register.

CSRRWI x3, 2, 7




6.6.46 CSRRSI rd, csr, rs1

This instruction is the same as CSRRS except a 5-bit unsigned (zero-extended) immediate value is used rather than the value from a register.

If the uimm[4:0] field is zero, then this instruction will not write to the CSR, and shall not cause any of the side effects that might otherwise occur on a CSR write. For CSRRWI, if rd=x0, then the instruction shall not read the CSR and shall not cause any of the side-effects that might occur on a CSR read. [1, p. 22]

1804 CSRRCI x3, 2, 7

1805



1806

Instruction!CSRRCI
RV32M
Instruction!MUL
Instruction!MULH


1807 **6.6.47 CSRRCI rd, csr, rs1**

1808 This instruction is the same as CSRRC except a 5-bit unsigned (zero-extended) immediate value is used
1809 rather than the value from a register.

1810 If the uimm[4:0] field is zero, then this instruction will not write to the CSR, and shall not cause any of
1811 the side effects that might otherwise occur on a CSR write. For CSRRWI, if rd=x0, then the instruction
1812 shall not read the CSR and shall not cause any of the side-effects that might occur on a CSR read. [1,
1813 p. 22]

1814 CSRRCI x3, 2, 7

1815



1816

1817 **6.7 RV32M Standard Extension**


1818 32-bit integer multiply and divide instructions.

1819 **6.7.1 MUL rd, rs1, rs2**

1820 Multiply rs1 by rs2 and store the least significant 32-bits of the result in rd.

1821 MUL x7, x3, x31

1822




1823

1824 **6.7.2 MULH rd, rs1, rs2**

1825 MULH x7, x3, x31


1826




1827

1828 **6.7.3 MULHS rd, rs1, rs2**


1829 MULHS x7, x3, x31

1832 **6.7.4 MULHU rd, rs1, rs2**


1833 MULHU x7, x3, x31

1836 **6.7.5 DIV rd, rs1, rs2**


1837 DIV x7, x3, x31

1840 **6.7.6 DIVU rd, rs1, rs2**

1841 DIVU x7, x3, x31

1844 **6.7.7 REM rd, rs1, rs2**


1845 REM x7, x3, x31



1848 **6.7.8 REMU rd, rs1, rs2**

Instruction!REMURV32ARV32F RV32D

1849 REMU x7, x3, x31



1852 **6.8 RV32A Standard Extension**

1853 32-bit atomic operations.

1854 **6.9 RV32F Standard Extension**

1855 32-bit IEEE floating point instructions.

1856 **6.10 RV32D Standard Extension**

1857 64-bit IEEE floating point instructions.

Appendix A

Installing a RISC-V Toolchain

A.1 The GNU Toolchain

Discuss the GNU toolchain elements used to experiment with the material in this book.

The instructions and examples here were all implemented on Ubuntu 16.04 LTS.

Install custom code in a location that will not cause interference with other applications and allow for easy cleanup. These instructions install the toolchain in `/usr/local/riscv`. At any time you can remove the lot and start over by executing the following command:

```
rm -rf /usr/local/riscv/*
```

Tested on Ubuntu 16.04 LTS. 18.04 was just released... update accordingly.

These are the only commands that you should perform as root when installing the toolchain:

```
sudo apt-get install autoconf automake autotools-dev curl libmpc-dev \
libmpfr-dev libgmp-dev gawk build-essential bison flex texinfo gperf \
libtool patchutils bc zlib1g-dev libexpat-dev
sudo mkdir -p /usr/local/riscv/
sudo chmod 777 /usr/local/riscv/
```

All other commands should be executed as a regular user. This will eliminate the possibility of clobbering system files that should not be touched when tinkering with the toolchain applications.

To download, compile and “install” the toolchain:

```
# riscv toolchain:
#
# https://riscv.org/software-tools/risc-v-gnu-compiler-toolchain/
git clone --recursive https://github.com/riscv/riscv-gnu-toolchain
cd riscv-gnu-toolchain
```

► Fix Me:

It would be good to find some Mac and Windows users to write and test proper variations on this section to address those systems. Pull requests, welcome!

```
1883 ./configure --prefix=/usr/local/riscv/rv32i --with-arch=rv32i --with-abi=ilp32
1884 make
1885 make install
```

1886 Need to discuss augmenting the PATH environment variable.

1887 Discuss the choice of ilp32 as well as what the other variations would do.

1888 Discuss rv32im and note that the details are found in [chapter 6](#).

1889 A.2 rvddt

1890 Discuss installing the rvddt simulator here.

1891

1892

1893

1894

- 1895

1896

- 1897

1898

1899

- 1900

1901

1002

1000

1906

1907

1908

1909

1910

- 1911

- IEEE754 formats:

	IEEE754 32-bit	IEEE754 64-bit
sign	1 bit	1 bit
exponent	8 bits (excess-127)	11 bits (excess-1023)
mantissa	23 bits	52 bits
max exponent	127	1023
min exponent	-126	-1022

- When the exponent is all ones, the mantissa is all zeros, and the sign is zero, the number represents positive infinity.
- When the exponent is all ones, the mantissa is all zeros, and the sign is one, the number represents negative infinity.
- Note that the binary representation of an IEEE754 number in memory can be compared for magnitude with another one using the same logic as for comparing two's complement signed integers because the magnitude of an IEEE number grows upward and downward in the same fashion as signed integers. This is why we use excess notation and locate the significand's sign bit on the left of the exponent.
- Note that zero is a special case number. Recall that a normalized number has an implied 1-bit to the left of the significand. . . which means that there is no way to represent zero! Zero is represented by an exponent of all-zeros and a significand of all-zeros. This definition allows for a positive and a negative zero if we observe that the sign can be either 1 or 0.
- On the number-line, numbers between zero and the smallest fraction in either direction are in the *underflow* areas.
- On the number line, numbers greater than the mantissa of all-ones and the largest exponent allowed are in the *overflow* areas.
- Note that numbers have a higher resolution on the number line when the exponent is smaller.

► Fix Me:

Need to add the standard lecture number-line diagram showing where the over/under-flow areas are and why.

B.1.1 Floating Point Number Accuracy

Due to the finite number of bits used to store the value of a floating point number, it is not possible to represent every one of the infinite values on the real number line. The following C programs illustrate this point.

B.1.1.1 Powers Of Two

Just like the integer numbers, the powers of two that have bits to represent them can be represented perfectly... as can their sums (provided that the significand requires no more than 23 bits.)

Listing B.1: powersoftwo.c

Precise Powers of Two

```

#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

union floatbin
{
    unsigned int    i;
    float          f;
};

```

```

19400 int main()
19401 {
19402     union floatbin x;
19403     union floatbin y;
19404     int i;
19405     x.f = 1.0;
19406     while (x.f > 1.0/1024.0)
19407     {
19408         y.f = -x.f;
19409         printf("%25.10f = %08x      %25.10f = %08x\n", x.f, x.i, y.f, y.i);
19410         x.f = x.f/2.0;
19411     }
19412 }
19413

```

Listing B.2: powersoftwo.out

Output from powersoftwo.c

```

1963 1.0000000000 = 3f800000      -1.0000000000 = bf800000
1964 0.5000000000 = 3f000000      -0.5000000000 = bf000000
1965 0.2500000000 = 3e800000      -0.2500000000 = be800000
1966 0.1250000000 = 3e000000      -0.1250000000 = be000000
1967 0.0625000000 = 3d800000      -0.0625000000 = bd800000
1968 0.0312500000 = 3d000000      -0.0312500000 = bd000000
1969 0.0156250000 = 3c800000      -0.0156250000 = bc800000
1970 0.0078125000 = 3c000000      -0.0078125000 = bc000000
1971 0.0039062500 = 3b800000      -0.0039062500 = bb800000
1972 0.0019531250 = 3b000000      -0.0019531250 = bb000000
1973

```

B.1.1.2 Clean Decimal Numbers

When dealing with decimal values, you will find that they don't map simply into binary floating point values.

Note how the decimal numbers are not accurately represented as they get larger. The decimal number on line 10 of [Listing B.4](#) can be perfectly represented in IEEE format. However, a problem arises in the 11th loop iteration. It is due to the fact that the binary number can not be represented accurately in IEEE format. Its least significant bits were truncated in a best-effort attempt at rounding the value off in order to fit the value into the bits provided. This is an example of *low order truncation*. Once this happens, the value of `x.f` is no longer as precise as it could be given more bits in which to save its value.

Listing B.3: cleandecimal.c

Print Clean Decimal Numbers

```

1984 #include <stdio.h>
1985 #include <stdlib.h>
1986 #include <unistd.h>
1987
1988 union floatbin
1989 {
1990     unsigned int i;
1991     float f;
1992 };
1993
1994 int main()
1995 {
1996     union floatbin x, y;
1997     int i;
1998
1999     x.f = 10;
2000     while (x.f <= 10000000000000.0)
2001     {
2002         y.f = -x.f;
2003         printf("%25.10f = %08x      %25.10f = %08x\n", x.f, x.i, y.f, y.i);

```

```

20040         x.f = x.f*10.0;
20051     }
20062 }
20067

```

Listing B.4: cleandecimal.out

Output from cleandecimal.c

```

2008      10.0000000000 = 41200000          -10.0000000000 = c1200000
20091      100.0000000000 = 42c80000         -100.0000000000 = c2c80000
20102      1000.0000000000 = 447a0000        -1000.0000000000 = c47a0000
20113      10000.0000000000 = 461c4000       -10000.0000000000 = c61c4000
20124      100000.0000000000 = 47c35000      -100000.0000000000 = c7c35000
20135      1000000.0000000000 = 49742400     -1000000.0000000000 = c9742400
20146      10000000.0000000000 = 4b189680    -10000000.0000000000 = cb189680
20157      100000000.0000000000 = 4cbebc20   -100000000.0000000000 = ccbebc20
20168      1000000000.0000000000 = 4e6e6b28  -1000000000.0000000000 = ce6e6b28
20179      10000000000.0000000000 = 501502f9 -10000000000.0000000000 = d01502f9
20180      99999997952.0000000000 = 51ba43b7  -99999997952.0000000000 = d1ba43b7
20191      999999995904.0000000000 = 5368d4a5 -999999995904.0000000000 = d368d4a5
20202      9999999827968.0000000000 = 551184e7 -9999999827968.0000000000 = d51184e7
20213
20224
20235

```

2023 B.1.1.3 Accumulation of Error

2024 These rounding errors can be exaggerated when the number we multiply the `x.f` value by is, itself,
 2025 something that can not be accurately represented in IEEE form.¹

2026 For example, if we multiply our `x.f` value by $\frac{1}{10}$ each time, we can never be accurate and we start
 2027 accumulating errors immediately.

Listing B.5: erroraccumulation.c

Accumulation of Error

```

2028 #include <stdio.h>
20291 #include <stdlib.h>
20302 #include <unistd.h>
20313
20324
20335 union floatbin
20346 {
20357     unsigned int    i;
20368     float          f;
20379 };
20380 int main()
20391 {
20402     union floatbin  x, y;
20413     int             i;
20424
20435     x.f = .1;
20446     while (x.f <= 2.0)
20457     {
20468         y.f = -x.f;
20479         printf("%25.10f = %08x      %25.10f = %08x\n", x.f, x.i, y.f, y.i);
20480         x.f += .1;
20491     }
20502 }
20513

```

➡ Fix Me:

In a lecture one would show that one tenth is a repeating non-terminating binary number that gets truncated. This discussion should be reproduced here in text form.

Listing B.6: erroraccumulation.out

Output from erroraccumulation.c

```

2052 0.1000000015 = 3dcccccd          -0.1000000015 = bdcccccd
20531

```

¹Applications requiring accurate decimal values, such as financial accounting systems, can use a packed-decimal numeric format to avoid unexpected oddities caused by the use of binary numbers.

20542	0.2000000030 = 3e4ccccd	-0.2000000030 = be4ccccd
20553	0.3000000119 = 3e99999a	-0.3000000119 = be99999a
20564	0.4000000060 = 3ecccccd	-0.4000000060 = becccccd
20575	0.5000000000 = 3f000000	-0.5000000000 = bf000000
20586	0.6000000238 = 3f19999a	-0.6000000238 = bf19999a
20597	0.7000000477 = 3f333334	-0.7000000477 = bf333334
20608	0.8000000715 = 3f4ccccce	-0.8000000715 = bf4ccccce
20619	0.9000000954 = 3f666668	-0.9000000954 = bf666668
20620	1.0000001192 = 3f800001	-1.0000001192 = bf800001
20631	1.1000001431 = 3f8ccccce	-1.1000001431 = bf8ccccce
20642	1.2000001669 = 3f99999b	-1.2000001669 = bf99999b
20653	1.3000001907 = 3fa66668	-1.3000001907 = bfa66668
20664	1.4000002146 = 3fb33335	-1.4000002146 = bfb33335
20675	1.5000002384 = 3fc00002	-1.5000002384 = bfc00002
20686	1.6000002623 = 3fcccccfc	-1.6000002623 = bfccccfc
20697	1.7000002861 = 3fd9999c	-1.7000002861 = bfd9999c
20708	1.8000003099 = 3fe66669	-1.8000003099 = bfe66669
20719	1.9000003338 = 3ff33336	-1.9000003338 = bff33336

B.1.2 Reducing Error Accumulation

In order to use floating point numbers in a program without causing excessive rounding problems an algorithm can be redesigned such that the accumulation is eliminated. This example is similar to the previous one, but this time we recalculate the desired value from a known-accurate integer value. Some rounding errors remain present, but they can not accumulate.

Listing B.7: errorcompensation.c
Accumulation of Error

```

2078 #include <stdio.h>
2079 #include <stdlib.h>
2080 #include <unistd.h>
2081
2082 union floatbin
2083 {
2084     unsigned int    i;
2085     float          f;
2086 };
2087
2088 int main()
2089 {
2090     union floatbin  x, y;
2091     int             i;
2092
2093     i = 1;
2094     while (i <= 20)
2095     {
2096         x.f = i/10.0;
2097         y.f = -x.f;
2098         printf("%25.10f = %08x      %25.10f = %08x\n", x.f, x.i, y.f, y.i);
2099         i++;
2100     }
2101     return(0);
2102 }

```

Listing B.8: errorcompensation.out
Output from erroraccumulation.c

2104	0.1000000015 = 3dcccccd	-0.1000000015 = bdcccccd
21051	0.2000000030 = 3e4ccccd	-0.2000000030 = be4ccccd
21062	0.3000000119 = 3e99999a	-0.3000000119 = be99999a
21073	0.4000000060 = 3ecccccd	-0.4000000060 = becccccd
21084	0.5000000000 = 3f000000	-0.5000000000 = bf000000
21095	0.6000000238 = 3f19999a	-0.6000000238 = bf19999a

21117	0.6999999881 = 3f333333	-0.6999999881 = bf333333
21128	0.8000000119 = 3f4ccccd	-0.8000000119 = bf4ccccd
21139	0.8999999762 = 3f666666	-0.8999999762 = bf666666
21140	1.0000000000 = 3f800000	-1.0000000000 = bf800000
21151	1.1000000238 = 3f8ccccd	-1.1000000238 = bf8ccccd
21162	1.2000000477 = 3f99999a	-1.2000000477 = bf99999a
21173	1.2999999523 = 3fa66666	-1.2999999523 = bfa66666
21184	1.3999999762 = 3fb33333	-1.3999999762 = bfb33333
21195	1.5000000000 = 3fc00000	-1.5000000000 = bfc00000
21206	1.6000000238 = 3fcccccd	-1.6000000238 = bfcccccd
21217	1.7000000477 = 3fd9999a	-1.7000000477 = bfd9999a
21228	1.7999999523 = 3fe66666	-1.7999999523 = bfe66666
21239	1.8999999762 = 3ff33333	-1.8999999762 = bff33333
21250	2.0000000000 = 40000000	-2.0000000000 = c0000000

Appendix C

The ASCII Character Set

A slightly abridged version of the Linux “ASCII” man(1) page.

C.1 NAME

ascii - ASCII character set encoded in octal, decimal, and hexadecimal

C.2 DESCRIPTION

ASCII is the American Standard Code for Information Interchange. It is a 7-bit code. Many 8-bit codes (e.g., ISO 8859-1) contain ASCII as their lower half. The international counterpart of ASCII is known as ISO 646-IRV.

The following table contains the 128 ASCII characters.

C program ‘\X’ escapes are noted.

Oct	Dec	Hex	Char	Oct	Dec	Hex	Char
000	0	00	NUL ‘\0’ (null character)	100	64	40	@
001	1	01	SOH (start of heading)	101	65	41	A
002	2	02	STX (start of text)	102	66	42	B
003	3	03	ETX (end of text)	103	67	43	C
004	4	04	EOT (end of transmission)	104	68	44	D
005	5	05	ENQ (enquiry)	105	69	45	E
006	6	06	ACK (acknowledge)	106	70	46	F
007	7	07	BEL ‘\a’ (bell)	107	71	47	G
010	8	08	BS ‘\b’ (backspace)	110	72	48	H
011	9	09	HT ‘\t’ (horizontal tab)	111	73	49	I
012	10	0A	LF ‘\n’ (new line)	112	74	4A	J
013	11	0B	VT ‘\v’ (vertical tab)	113	75	4B	K
014	12	0C	FF ‘\f’ (form feed)	114	76	4C	L
015	13	0D	CR ‘\r’ (carriage ret)	115	77	4D	M

2153	016	14	0E	S0 (shift out)	116	78	4E	N	
2154	017	15	0F	SI (shift in)	117	79	4F	O	
2155	020	16	10	DLE (data link escape)	120	80	50	P	
2156	021	17	11	DC1 (device control 1)	121	81	51	Q	
2157	022	18	12	DC2 (device control 2)	122	82	52	R	
2158	023	19	13	DC3 (device control 3)	123	83	53	S	
2159	024	20	14	DC4 (device control 4)	124	84	54	T	
2160	025	21	15	NAK (negative ack.)	125	85	55	U	
2161	026	22	16	SYN (synchronous idle)	126	86	56	V	
2162	027	23	17	ETB (end of trans. blk)	127	87	57	W	
2163	030	24	18	CAN (cancel)	130	88	58	X	
2164	031	25	19	EM (end of medium)	131	89	59	Y	
2165	032	26	1A	SUB (substitute)	132	90	5A	Z	
2166	033	27	1B	ESC (escape)	133	91	5B	[
2167	034	28	1C	FS (file separator)	134	92	5C	\	'\'
2168	035	29	1D	GS (group separator)	135	93	5D]	
2169	036	30	1E	RS (record separator)	136	94	5E	^	
2170	037	31	1F	US (unit separator)	137	95	5F	_	
2171	040	32	20	SPACE	140	96	60	`	
2172	041	33	21	!	141	97	61	a	
2173	042	34	22	"	142	98	62	b	
2174	043	35	23	#	143	99	63	c	
2175	044	36	24	\$	144	100	64	d	
2176	045	37	25	%	145	101	65	e	
2177	046	38	26	&	146	102	66	f	
2178	047	39	27	'	147	103	67	g	
2179	050	40	28	(150	104	68	h	
2180	051	41	29)	151	105	69	i	
2181	052	42	2A	*	152	106	6A	j	
2182	053	43	2B	+	153	107	6B	k	
2183	054	44	2C	,	154	108	6C	l	
2184	055	45	2D	-	155	109	6D	m	
2185	056	46	2E	.	156	110	6E	n	
2186	057	47	2F	/	157	111	6F	o	
2187	060	48	30	0	160	112	70	p	
2188	061	49	31	1	161	113	71	q	
2189	062	50	32	2	162	114	72	r	
2190	063	51	33	3	163	115	73	s	
2191	064	52	34	4	164	116	74	t	
2192	065	53	35	5	165	117	75	u	
2193	066	54	36	6	166	118	76	v	
2194	067	55	37	7	167	119	77	w	
2195	070	56	38	8	170	120	78	x	
2196	071	57	39	9	171	121	79	y	
2197	072	58	3A	:	172	122	7A	z	
2198	073	59	3B	;	173	123	7B	{	
2199	074	60	3C	<	174	124	7C		
2200	075	61	3D	=	175	125	7D	}	
2201	076	62	3E	>	176	126	7E	~	
2202	077	63	3F	?	177	127	7F	DEL	

C.2.1 Tables

For convenience, below are more compact tables in hex and decimal.

2 3 4 5 6 7	30 40 50 60 70 80 90 100 110 120
-----	-----
0: 0 @ P ' p	0: (2 < F P Z d n x
1: ! 1 A Q a q	1:) 3 = G Q [e o y
2: " 2 B R b r	2: * 4 > H R \ f p z
3: # 3 C S c s	3: ! + 5 ? I S] g q {
4: \$ 4 D T d t	4: " , 6 @ J T ^ h r
5: % 5 E U e u	5: # - 7 A K U _ i s }
6: & 6 F V f v	6: \$. 8 B L V ' j t ~
7: ' 7 G W g w	7: % / 9 C M W a k u DEL
8: (8 H X h x	8: & 0 : D N X b l v
9:) 9 I Y i y	9: ' 1 ; E O Y c m w
A: * : J Z j z	
B: + ; K [k {	
C: , < L \ l	
D: - = M] m }	
E: . > N ^ n ~	
F: / ? 0 _ o DEL	

C.3 NOTES

C.3.1 History

An ascii manual page appeared in Version 7 of AT&T UNIX.

On older terminals, the underscore code is displayed as a left arrow, called backarrow, the caret is displayed as an up-arrow and the vertical bar has a hole in the middle.

Uppercase and lowercase characters differ by just one bit and the ASCII character 2 differs from the double quote by just one bit, too. That made it much easier to encode characters mechanically or with a non-microcontroller-based electronic keyboard and that pairing was found on old teletypes.

The ASCII standard was published by the United States of America Standards Institute (USASI) in 1968.

C.4 COLOPHON

This page is part of release 4.04 of the Linux man-pages project. A description of the project, information about reporting bugs, and the latest version of this page, can be found at <http://www.kernel.org/doc/man-pages/>.

Appendix D

Attribution 4.0 International

Creative Commons Corporation ("Creative Commons") is not a law firm and does not provide legal services or legal advice. Distribution of Creative Commons public licenses does not create a lawyer-client or other relationship. Creative Commons makes its licenses and related information available on an "as-is" basis. Creative Commons gives no warranties regarding its licenses, any material licensed under their terms and conditions, or any related information. Creative Commons disclaims all liability for damages resulting from their use to the fullest extent possible.

Using Creative Commons Public Licenses

Creative Commons public licenses provide a standard set of terms and conditions that creators and other rights holders may use to share original works of authorship and other material subject to copyright and certain other rights specified in the public license below. The following considerations are for informational purposes only, are not exhaustive, and do not form part of our licenses.

Considerations for licensors: Our public licenses are intended for use by those authorized to give the public permission to use material in ways otherwise restricted by copyright and certain other rights. Our licenses are irrevocable. Licensors should read and understand the terms and conditions of the license they choose before applying it. Licensors should also secure all rights necessary before applying our licenses so that the public can reuse the material as expected. Licensors should clearly mark any material not subject to the license. This includes other CC-licensed material, or material used under an exception or limitation to copyright. More considerations for licensors: http://wiki.creativecommons.org/Considerations_for_licensors

Considerations for the public: By using one of our public licenses, a licensor grants the public permission to use the licensed material under specified terms and conditions. If the licensor's permission is not necessary for any reason—for example, because of any applicable exception or limitation to copyright—then that use is not regulated by the license. Our licenses grant only permissions under copyright and certain other rights that a licensor has authority to grant. Use of the licensed material may still be restricted for other reasons, including because others have copyright or other rights in the material. A licensor may make special requests, such as asking that all changes be marked or described. Although not required by our licenses, you are encouraged to respect those requests where reasonable. More considerations for the public: http://wiki.creativecommons.org/Considerations_for_licensees

Creative Commons Attribution 4.0 International Public License

By exercising the Licensed Rights (defined below), You accept and agree to be bound by the terms and conditions of this Creative Commons Attribution 4.0 International Public License ("Public License"). To the extent this Public License may be interpreted as a contract, You are granted the Licensed Rights in consideration of Your acceptance of these terms and conditions, and the Licensor grants You such rights in consideration of benefits the Licensor receives from making the Licensed Material available under these terms and conditions.

Section 1. Definitions

- a. Adapted Material means material subject to Copyright and Similar Rights that is derived from or based upon the Licensed Material and in which the Licensed Material is translated, altered, arranged, transformed, or otherwise modified in a manner requiring permission under the Copyright and Similar Rights held by the Licensor. For purposes of this Public License, where the Licensed Material is a musical work, performance, or sound recording, Adapted Material is always produced where the Licensed Material is synched in timed relation with a moving image.
- b. Adapter's License means the license You apply to Your Copyright and Similar Rights in Your contributions to Adapted Material in accordance with the terms and conditions of this Public License.

-
- 2276 c. Copyright and Similar Rights means copyright and/or similar rights closely related to copyright including, without limi-
2277 tation, performance, broadcast, sound recording, and Sui Generis Database Rights, without regard to how the rights are
2278 labeled or categorized. For purposes of this Public License, the rights specified in Section 2(b)(1)-(2) are not Copyright
2279 and Similar Rights.
- 2280 d. Effective Technological Measures means those measures that, in the absence of proper authority, may not be circumvented
2281 under laws fulfilling obligations under Article 11 of the WIPO Copyright Treaty adopted on December 20, 1996, and/or
2282 similar international agreements.
- 2283 e. Exceptions and Limitations means fair use, fair dealing, and/or any other exception or limitation to Copyright and Similar
2284 Rights that applies to Your use of the Licensed Material.
- 2285 f. Licensed Material means the artistic or literary work, database, or other material to which the Licensor applied this Public
2286 License.
- 2287 g. Licensed Rights means the rights granted to You subject to the terms and conditions of this Public License, which are
2288 limited to all Copyright and Similar Rights that apply to Your use of the Licensed Material and that the Licensor has
2289 authority to license.
- 2290 h. Licensor means the individual(s) or entity(ies) granting rights under this Public License.
- 2291 i. Share means to provide material to the public by any means or process that requires permission under the Licensed Rights,
2292 such as reproduction, public display, public performance, distribution, dissemination, communication, or importation, and
2293 to make material available to the public including in ways that members of the public may access the material from a place
2294 and at a time individually chosen by them.
- 2295 j. Sui Generis Database Rights means rights other than copyright resulting from Directive 96/9/EC of the European Parlia-
2296 ment and of the Council of 11 March 1996 on the legal protection of databases, as amended and/or succeeded, as well as
2297 other essentially equivalent rights anywhere in the world.
- 2298 k. You means the individual or entity exercising the Licensed Rights under this Public License. Your has a corresponding
2299 meaning.

2300 Section 2. Scope

- 2301 a. License grant.
- 2302 1. Subject to the terms and conditions of this Public License, the Licensor hereby grants You a worldwide, royalty-free,
2303 non-sublicensable, non-exclusive, irrevocable license to exercise the Licensed Rights in the Licensed Material to:
- 2304 a. reproduce and Share the Licensed Material, in whole or in part; and
2305 b. produce, reproduce, and Share Adapted Material.
- 2306 2. Exceptions and Limitations. For the avoidance of doubt, where Exceptions and Limitations apply to Your use, this
2307 Public License does not apply, and You do not need to comply with its terms and conditions.
- 2308 3. Term. The term of this Public License is specified in Section 6(a).
- 2309 4. Media and formats; technical modifications allowed. The Licensor authorizes You to exercise the Licensed Rights in
2310 all media and formats whether now known or hereafter created, and to make technical modifications necessary to
2311 do so. The Licensor waives and/or agrees not to assert any right or authority to forbid You from making technical
2312 modifications necessary to exercise the Licensed Rights, including technical modifications necessary to circumvent
2313 Effective Technological Measures. For purposes of this Public License, simply making modifications authorized by
2314 this Section 2(a) (4) never produces Adapted Material.
- 2315 5. Downstream recipients.
- 2316 a. Offer from the Licensor – Licensed Material. Every recipient of the Licensed Material automatically receives an
2317 offer from the Licensor to exercise the Licensed Rights under the terms and conditions of this Public License.
- 2318 b. No downstream restrictions. You may not offer or impose any additional or different terms or conditions on,
2319 or apply any Effective Technological Measures to, the Licensed Material if doing so restricts exercise of the
2320 Licensed Rights by any recipient of the Licensed Material.
- 2321 6. No endorsement. Nothing in this Public License constitutes or may be construed as permission to assert or imply
2322 that You are, or that Your use of the Licensed Material is, connected with, or sponsored, endorsed, or granted official
2323 status by, the Licensor or others designated to receive attribution as provided in Section 3(a)(1)(A)(i).
- 2324 b. Other rights.
- 2325 1. Moral rights, such as the right of integrity, are not licensed under this Public License, nor are publicity, privacy,
2326 and/or other similar personality rights; however, to the extent possible, the Licensor waives and/or agrees not to
2327 assert any such rights held by the Licensor to the limited extent necessary to allow You to exercise the Licensed
2328 Rights, but not otherwise.
- 2329 2. Patent and trademark rights are not licensed under this Public License.
- 2330 3. To the extent possible, the Licensor waives any right to collect royalties from You for the exercise of the Licensed
2331 Rights, whether directly or through a collecting society under any voluntary or waivable statutory or compulsory
2332 licensing scheme. In all other cases the Licensor expressly reserves any right to collect such royalties.

2333 Section 3. License Conditions

2334 Your exercise of the Licensed Rights is expressly made subject to the following conditions.

2335 a. Attribution.

2336 1. If You Share the Licensed Material (including in modified form), You must:

2337 a. retain the following if it is supplied by the Licensor with the Licensed Material:

2338 i. identification of the creator(s) of the Licensed Material and any others designated to receive attribution,
2339 in any reasonable manner requested by the Licensor (including by pseudonym if designated);

2340 ii. a copyright notice;

2341 iii. a notice that refers to this Public License;

2342 iv. a notice that refers to the disclaimer of warranties;

2343 v. a URI or hyperlink to the Licensed Material to the extent reasonably practicable;

2344 b. indicate if You modified the Licensed Material and retain an indication of any previous modifications; and

2345 c. indicate the Licensed Material is licensed under this Public License, and include the text of, or the URI or
2346 hyperlink to, this Public License.

2347 2. You may satisfy the conditions in Section 3(a)(1) in any reasonable manner based on the medium, means, and context
2348 in which You Share the Licensed Material. For example, it may be reasonable to satisfy the conditions by providing
2349 a URI or hyperlink to a resource that includes the required information.

2350 3. If requested by the Licensor, You must remove any of the information required by Section 3(a)(1)(A) to the extent
2351 reasonably practicable.

2352 4. If You Share Adapted Material You produce, the Adapter's License You apply must not prevent recipients of the
2353 Adapted Material from complying with this Public License.

2354 Section 4. Sui Generis Database Rights

2355 Where the Licensed Rights include Sui Generis Database Rights that apply to Your use of the Licensed Material:

2356 a. for the avoidance of doubt, Section 2(a)(1) grants You the right to extract, reuse, reproduce, and Share all or a substantial
2357 portion of the contents of the database;

2358 b. if You include all or a substantial portion of the database contents in a database in which You have Sui Generis Database
2359 Rights, then the database in which You have Sui Generis Database Rights (but not its individual contents) is Adapted
2360 Material; and

2361 c. You must comply with the conditions in Section 3(a) if You Share all or a substantial portion of the contents of the
2362 database.

2363 For the avoidance of doubt, this Section 4 supplements and does not replace Your obligations under this Public License where the
2364 Licensed Rights include other Copyright and Similar Rights.

2365 Section 5. Disclaimer of Warranties and Limitation of Liability

2366 a. UNLESS OTHERWISE SEPARATELY UNDERTAKEN BY THE LICENSOR, TO THE EXTENT POSSIBLE, THE
2367 LICENSOR OFFERS THE LICENSED MATERIAL AS-IS AND AS-AVAILABLE, AND MAKES NO REPRESENTA-
2368 TIONS OR WARRANTIES OF ANY KIND CONCERNING THE LICENSED MATERIAL, WHETHER EXPRESS, IM-
2369 PLIED, STATUTORY, OR OTHER. THIS INCLUDES, WITHOUT LIMITATION, WARRANTIES OF TITLE, MER-
2370 CHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, NON-INFRINGEMENT, ABSENCE OF LATENT OR
2371 OTHER DEFECTS, ACCURACY, OR THE PRESENCE OR ABSENCE OF ERRORS, WHETHER OR NOT KNOWN
2372 OR DISCOVERABLE. WHERE DISCLAIMERS OF WARRANTIES ARE NOT ALLOWED IN FULL OR IN PART,
2373 THIS DISCLAIMER MAY NOT APPLY TO YOU.

2374 b. TO THE EXTENT POSSIBLE, IN NO EVENT WILL THE LICENSOR BE LIABLE TO YOU ON ANY LEGAL THE-
2375 ORY (INCLUDING, WITHOUT LIMITATION, NEGLIGENCE) OR OTHERWISE FOR ANY DIRECT, SPECIAL, IN-
2376 DIRECT, INCIDENTAL, CONSEQUENTIAL, PUNITIVE, EXEMPLARY, OR OTHER LOSSES, COSTS, EXPENSES,
2377 OR DAMAGES ARISING OUT OF THIS PUBLIC LICENSE OR USE OF THE LICENSED MATERIAL, EVEN IF THE
2378 LICENSOR HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH LOSSES, COSTS, EXPENSES, OR DAMAGES.
2379 WHERE A LIMITATION OF LIABILITY IS NOT ALLOWED IN FULL OR IN PART, THIS LIMITATION MAY NOT
2380 APPLY TO YOU.

2381 c. The disclaimer of warranties and limitation of liability provided above shall be interpreted in a manner that, to the extent
2382 possible, most closely approximates an absolute disclaimer and waiver of all liability.

2383 Section 6. Term and Termination

- 2384 a. This Public License applies for the term of the Copyright and Similar Rights licensed here. However, if You fail to comply
2385 with this Public License, then Your rights under this Public License terminate automatically.
- 2386 b. Where Your right to use the Licensed Material has terminated under Section 6(a), it reinstates:
- 2387 1. automatically as of the date the violation is cured, provided it is cured within 30 days of Your discovery of the
2388 violation; or
 - 2389 2. upon express reinstatement by the Licensor.
- 2390 For the avoidance of doubt, this Section 6(b) does not affect any right the Licensor may have to seek remedies for Your
2391 violations of this Public License.
- 2392 c. For the avoidance of doubt, the Licensor may also offer the Licensed Material under separate terms or conditions or stop
2393 distributing the Licensed Material at any time; however, doing so will not terminate this Public License.
- 2394 d. Sections 1, 5, 6, 7, and 8 survive termination of this Public License.

2395 Section 7. Other Terms and Conditions

- 2396 a. The Licensor shall not be bound by any additional or different terms or conditions communicated by You unless expressly
2397 agreed.
- 2398 b. Any arrangements, understandings, or agreements regarding the Licensed Material not stated herein are separate from and
2399 independent of the terms and conditions of this Public License.

2400 Section 8. Interpretation

- 2401 a. For the avoidance of doubt, this Public License does not, and shall not be interpreted to, reduce, limit, restrict, or impose
2402 conditions on any use of the Licensed Material that could lawfully be made without permission under this Public License.
- 2403 b. To the extent possible, if any provision of this Public License is deemed unenforceable, it shall be automatically reformed
2404 to the minimum extent necessary to make it enforceable. If the provision cannot be reformed, it shall be severed from this
2405 Public License without affecting the enforceability of the remaining terms and conditions.
- 2406 c. No term or condition of this Public License will be waived and no failure to comply consented to unless expressly agreed
2407 to by the Licensor.
- 2408 d. Nothing in this Public License constitutes or may be interpreted as a limitation upon, or waiver of, any privileges and
2409 immunities that apply to the Licensor or You, including from the legal processes of any jurisdiction or authority.
2410

2411 Creative Commons is not a party to its public licenses. Notwithstanding, Creative Commons may elect to apply one of its public
2412 licenses to material it publishes and in those instances will be considered the Licensor. The text of the Creative Commons public
2413 licenses is dedicated to the public domain under the CC0 Public Domain Dedication. Except for the limited purpose of indicating
2414 that material is shared under a Creative Commons public license or as otherwise permitted by the Creative Commons policies
2415 published at <http://creativecommons.org/policies>, Creative Commons does not authorize the use of the trademark “Creative
2416 Commons” or any other trademark or logo of Creative Commons without its prior written consent including, without limitation,
2417 in connection with any unauthorized modifications to any of its public licenses or any other arrangements, understandings, or
2418 agreements concerning use of licensed material. For the avoidance of doubt, this paragraph does not form part of the public
2419 licenses.

2420 Creative Commons may be contacted at <http://creativecommons.org>.

Bibliography

- [1] RISC-V Foundation, *The RISC-V Instruction Set Manual, Volume I: User-Level ISA, Document Version 2.2*, 5 2017. Editors Andrew Waterman and Krste Asanović. [viii](#), [14](#), [36](#), [37](#), [38](#), [39](#), [47](#), [49](#), [50](#), [51](#), [52](#), [53](#), [54](#), [55](#), [56](#), [57](#), [58](#)
- [2] D. Patterson and A. Waterman, *The RISC-V Reader: An Open Architecture Atlas*. Strawberry Canyon, 11 2017. ISBN: 978-0999249116. [viii](#)
- [3] D. Patterson and J. Hennessy, *Computer Organization and Design RISC-V Edition: The Hardware Software Interface*. Morgan Kaufmann, 4 2017. ISBN: 978-0128122754. [viii](#), [16](#)
- [4] W. F. Decker, “A modern approach to teaching computer organization and assembly language programming,” *SIGCSE Bull.*, vol. 17, pp. 38–44, 12 1985. [viii](#)
- [5] Texas Instruments, *SN54190, SN54191, SN54LS190, SN54LS191, SN74190, SN74191, SN74LS190, SN74LS191 Synchronous Up/Down Counters With Down/Up Mode Control*, 3 1988. [viii](#)
- [6] Texas Instruments, *SN54154, SN74154 4-line to 16-line Decoders/Demultiplexers*, 12 1972. [viii](#)
- [7] Intel, *MCS-85 User’s Manual*, 9 1978. [viii](#)
- [8] Radio Shack, *TRS-80 Editor/Assembler Operation and Reference Manual*, 1978. [viii](#)
- [9] Motorola, *MC68000 16-bit Microprocessor User’s Manual*, 2nd ed., 1 1980. MC68000UM(AD2). [viii](#)
- [10] R. A. Overbeek and W. E. Singletary, *Assembler Language With ASSIST*. Science Research Associates, Inc., 2nd ed., 1983. [viii](#)
- [11] IBM, *IBM System/370 Principals of Operation*, 7th ed., 3 1980. [viii](#)
- [12] IBM, *OS/VS-DOS/VSE-VM/370 Assembler Language*, 6th ed., 3 1979. [viii](#)
- [13] D. Cohen, “*IEN 137, On Holy Wars and a Plea for Peace*,” Apr. 1980. This note discusses the Big-Endian/Little-Endian byte/bit-order controversy, but did not settle it. A decade later, David V. James in “Multiplexed Buses: The Endian Wars Continue”, *IEEE Micro*, **10**(3), 9–21 (1990) continued the discussion. [15](#)
- [14] R. M. Stallman and the GCC Developer Community, *Using the GNU Compiler Collection (For GCC version 7.3.0)*. Free Software Foundation, 51 Franklin Street, Fifth Floor, Boston, MA 02110-1301 USA: GNU Press, 2017. [21](#)
- [15] RISC-V Foundation, *The RISC-V Instruction Set Manual, Volume II: Privileged Architecture, Document Version 1.10*, 5 2017. Editors Andrew Waterman and Krste Asanović.
- [16] P. Dabbelt, S. O’Rear, K. Cheng, A. Waterman, M. Clark, A. Bradbury, D. Horner, M. Nordlund, and K. Merker, *RISC-V ELF psABI specification*, 2017.
- [17] National Semiconductor Corporation, *Series 32000 Databook*, 1986.

Index

2454	A	2497	MULHU, 59
2455	ALU, 3	2498	OR, 54
2456	ASCII, 69	2499	ORI, 47
		2500	REM, 59
2457	C	2501	REMU, 60
2458	CPU, 2	2502	SB, 44
		2503	SH, 45
2459	H	2504	SLL, 51
2460	hart, 3	2505	SLLI, 48
		2506	SLT, 51
2461	I	2507	SLTI, 46
2462	Instruction	2508	SLTIU, 46
2463	ADD, 50	2509	SLTU, 52
2464	ADDI, 45	2510	SRA, 53
2465	AND, 54	2511	SRAI, 49
2466	ANDI, 48	2512	SRL, 53
2467	AUIPC, 37	2513	SRLI, 49
2468	BEQ, 40	2514	SUB, 50
2469	BGE, 41	2515	SW, 45
2470	BGEU, 42	2516	XOR, 52
2471	BLT, 41	2517	XORI, 47
2472	BLTU, 42	2518	instruction cycle, 4
2473	BNE, 41	2519	instruction decode, 5
2474	CSRRC, 57	2520	instruction execute, 5
2475	CSRRCI, 58	2521	instruction fetch, 5
2476	CSRRS, 56	2522	ISA, 4
2477	CSRRSI, 57		
2478	CSRROW, 56	2523	R
2479	CSRROWI, 57	2524	register, 1, 3
2480	DIV, 59	2525	RV32, 28
2481	DIVU, 59	2526	RV32A, 4, 60
2482	EBREAK, 56	2527	RV32C, 4
2483	ECALL, 56	2528	RV32D, 4, 60
2484	FENCE, 55	2529	RV32F, 4, 60
2485	FENCE.I, 55	2530	RV32G, 4
2486	JAL, 38	2531	RV32I, 4, 37
2487	JALR, 39	2532	RV32M, 4, 58
2488	LB, 43	2533	RV32Q, 4
2489	LBU, 44	2534	rvddt, 18
2490	LH, 43		
2491	LHU, 44		
2492	LUI, 37		
2493	LW, 43		
2494	MUL, 58		
2495	MULH, 58		
2496	MULHS, 59		

2536

Glossary

- 2537 **address** A numeric value used to uniquely identify each byte of main memory. [2](#), [78](#)
- 2538 **alignment** Refers to a range of numeric values that begin at a multiple of some number. Primarily
2539 used when referring to a memory address. For example an alignment of two refers to one or more
2540 addresses starting at even address and continuing onto subsequent adjacent, increasing memory
2541 addresses. [16](#), [78](#)
- 2542 **big endian** A number format where the most significant values are printed to the left of the lesser
2543 significant values. This is the method that everyone used to write decimal numbers every day. [19](#),
2544 [20](#), [78](#)
- 2545 **binary** Something that has two parts or states. In computing these two states are represented by the
2546 numbers one and zero or by the conditions true and false and can be stored in one bit. [1](#), [3](#), [78](#)
- 2547 **bit** One binary digit. [3](#), [6](#), [9](#), [78](#)
- 2548 **byte** A binary value represented by 8 bits. [2](#), [6](#), [78](#)
- 2549 **CPU** Central Processing Unit. [1](#), [2](#), [78](#)
- 2550 **doubleword** A binary value represented by 64 bits. [78](#)
- 2551 **exception** An error encountered by the CPU while executing an instruction that can not be completed.
2552 [16](#), [78](#)
- 2553 **fullword** A binary value represented by 32 bits. [6](#), [78](#)
- 2554 **halfword** A binary value represented by 16 bits. [6](#), [78](#)
- 2555 **hart** Hardware Thread. [3](#), [78](#)
- 2556 **hexadecimal** A base-16 numbering system whose digits are 0123456789abcdef. The hex digits (hits)
2557 are not case-sensitive. [19](#), [20](#), [78](#)
- 2558 **high order bits** Some number of MSBs. [78](#)
- 2559 **hit** One hex digit. [9](#), [10](#), [78](#)
- 2560 **ISA** Instruction Set Architecture. [3](#), [78](#)
- 2561 **LaTeX** Is a mark up language specially suited for scientific documents. [78](#)

- little endian** A number format where the least significant values are printed to the left of the more significant values. This is the opposite ordering that everyone learns in grade school when learning how to count. For example a big endian number written as “1234” would be written in little endian form as “4321”. [78](#)
- low order bits** Some number of LSBs. [78](#)
- LSB** Least Significant Bit. [11](#), [29](#), [32](#), [34](#), [78](#)
- machine language** The instructions that are executed by a CPU that are expressed in the form of binary values. [1](#), [78](#)
- mnemonic** A method used to remember something. In the case of assembly language, each machine instruction is given a name so the programmer need not memorize the binary values of each machine instruction. [1](#), [78](#)
- MSB** Most Significant Bit. [11](#), [12](#), [28](#), [29](#), [34](#), [78](#)
- overflow** The situation where the result of an addition or subtraction operation is approaching positive or negative infinity and exceeds the number of bits allotted to contain the result. This is typically caused by high-order truncation. [12](#), [64](#), [78](#)
- program** A ordered list of one or more instructions. [1](#), [78](#)
- quadword** A binary value represented by 128 bits. [78](#)
- register** A unit of storage inside a CPU with the capacity of XLEN bits. [1](#), [78](#)
- RV32** Short for RISC-V 32. The number 32 refers to the XLEN. [37](#), [78](#)
- RV64** Short for RISC-V 64. The number 64 refers to the XLEN. [78](#)
- rvddt** A RV32I simulator and debugging tool inspired by the simplicity of the Dynamic Debugging Tool (ddt) that was part of the CP/M operating system. [18](#), [78](#)
- thread** An stream of instructions. When plural, it is used to refer to the ability of a CPU to execute multiple instruction streams at the same time. [3](#), [78](#)
- underflow** The situation where the result of an addition or subtraction operation is approaching zero and exceeds the number of bits allotted to contain the result. This is typically caused by low-order truncation. [64](#), [78](#)
- XLEN** The number of bits a RISC-V x integer register (such as x0). For RV32 XLEN=32, RV64 XLEN=64 etc. [33](#), [34](#), [78](#)