

1

2

RISC-V Assembly Language Programming

3 (Draft v0.1-61-gcaaa6b3)

4 John Winans
jwinans@niu.edu

5 May 24, 2018

6 Copyright © 2018 John Winans
7 This document is made available under a Creative Commons Attribution 4.0 International License.
8 See [Appendix E](#) for more information.
9 Download your own copy of this book from github here: <https://github.com/johnwinans/rvalp>.
10 This document may contain inaccuracies or errors. The author provides no guarantee regarding the
11 accuracy of this document's contents. If you discover that this document contains errors, please notify
12 the author.
13
14 ARM® is a registered trademark of ARM Limited in the EU and other countries.

► Fix Me:
*Need to say something
about trademarks for things
mentioned in this text*

Contents

Preface	viii
----------------	-------------

1 Introduction	1
-----------------------	----------

1.1 The Digital Computer	1
------------------------------------	---

1.1.1 Storage Systems	1
---------------------------------	---

1.1.1.1 Volatile Storage	2
------------------------------------	---

1.1.1.2 Non-Volatile Storage	2
--	---

1.1.2 CPU	2
---------------------	---

1.1.2.1 Execution Unit	2
----------------------------------	---

1.1.2.2 Arithmetic and Logic Unit	3
---	---

1.1.2.3 Registers	3
-----------------------------	---

1.1.2.4 Harts	3
-------------------------	---

1.1.3 Peripherals	3
-----------------------------	---

1.2 Instruction Set Architecture	4
--	---

1.2.1 RV Base Modules	4
---------------------------------	---

1.2.2 Extension Modules	4
-----------------------------------	---

1.3 How the CPU Executes a Program	4
--	---

1.3.1 Instruction Fetch	5
-----------------------------------	---

1.3.2 Instruction Decode	5
------------------------------------	---

1.3.3 Instruction Execute	5
-------------------------------------	---

2 Numbers and Storage Systems	6
--------------------------------------	----------

2.1 Boolean Functions	6
---------------------------------	---

2.1.1 NOT	6
---------------------	---

2.1.2 AND	7
---------------------	---

2.1.3 OR	8
--------------------	---

2.1.4 XOR	8
---------------------	---

2.2 Integers and Counting	9
-------------------------------------	---

2.2.1 Converting Between Bases	11
--	----

2.2.1.1 From Binary to Decimal	11
--	----

44	2.2.1.2	From Binary to Hexadecimal	11
45	2.2.1.3	From Hexadecimal to Binary	11
46	2.2.1.4	From Decimal to Binary	12
47	2.2.1.5	From Decimal to Hex	12
48	2.2.2	Addition of Binary Numbers	12
49	2.2.3	Signed Numbers	13
50	2.2.3.1	Converting between Positive and Negative	14
51	2.2.4	Subtraction of Binary Numbers	15
52	2.2.5	Truncation and Overflow	15
53	2.3	Sign and Zero Extension	16
54	2.4	Shifting	17
55	2.5	Main Memory Storage	18
56	2.5.1	Memory Dump	18
57	2.5.2	Big Endian Representation	18
58	2.5.3	Little Endian Representation	18
59	2.5.4	Character Strings and Arrays	18
60	2.5.5	Context is Important!	19
61	2.5.6	Alignment	19
62	2.5.7	Instruction Alignment	19
63	3	The Elements of a Assembly Language Program	20
64	3.1	Assembly Language Statements	20
65	3.2	Memory Layout	20
66	3.3	A Sample Program Source Listing	20
67	3.4	Running a Program With rvddt	21
68	4	Writing RISC-V Programs	24
69	4.1	Use ebreak to Stop rvddt Execution	24
70	4.2	Using the addi Instruction	25
71	4.2.1	No Operation	25
72	4.2.2	Copying the Contents of One Register to Another	27
73	4.2.3	Setting a Register to Zero	28
74	4.2.4	Adding a 12-bit Signed Value	29
75	4.3	todo	29
76	4.4	Other Instructions With Immediate Operands	29
77	4.5	Transferring Data Between Registers and Memory	29
78	4.6	RR operations	30
79	4.7	Setting registers to large values using lui with addi	30
80	4.8	Labels and Branching	30

81	4.9 Relocation	31
82	4.10 Jumps	31
83	4.11 Pseudo Operations	32
84	4.12 The Linker and Relaxation	32
85	4.13 pic and nopic	32
86	5 RV32 Machine Instructions	33
87	5.1 Introduction	33
88	5.2 Conventions and Terminology	33
89	5.2.1 XLEN	33
90	5.2.2 sx(val)	33
91	5.2.3 zx(val)	33
92	5.2.4 zr(val)	34
93	5.2.5 Sign Extended Left and Zero Extend Right	34
94	5.2.6 m8(addr)	34
95	5.2.7 m16(addr)	35
96	5.2.8 m32(addr)	35
97	5.2.9 m64(addr)	35
98	5.2.10 m128(addr)	35
99	5.2.11 .+offset	36
100	5.2.12 .-offset	36
101	5.2.13 pc	36
102	5.2.14 rd	36
103	5.2.15 rs1	36
104	5.2.16 rs2	36
105	5.2.17 imm	36
106	5.2.18 rsN[h:l]	36
107	5.3 Addressing Modes	36
108	5.4 Instruction Encoding Formats	37
109	5.4.1 U Type	37
110	5.4.2 J Type	38
111	5.4.3 R Type	39
112	5.4.4 I Type	39
113	5.4.5 S Type	40
114	5.4.6 B Type	40
115	5.4.7 CPU Registers	40
116	5.5 memory	41
117	5.6 RV32I Base Instruction Set	41

118	5.6.1	LUI rd, imm	42
119	5.6.2	AUIPC rd, imm	42
120	5.6.3	JAL rd, imm	43
121	5.6.4	JALR rd, rs1, imm	44
122	5.6.5	BEQ rs1, rs2, imm	45
123	5.6.6	BNE rs1, rs2, imm	45
124	5.6.7	BLT rs1, rs2, imm	46
125	5.6.8	BGE rs1, rs2, imm	46
126	5.6.9	BLTU rs1, rs2, imm	47
127	5.6.10	BGEU rs1, rs2, imm	47
128	5.6.11	LB rd, imm(rs1)	47
129	5.6.12	LH rd, imm(rs1)	48
130	5.6.13	LW rd, imm(rs1)	48
131	5.6.14	LBU rd, imm(rs1)	48
132	5.6.15	LHU rd, imm(rs1)	49
133	5.6.16	SB rs2, imm(rs1)	49
134	5.6.17	SH rs2, imm(rs1)	49
135	5.6.18	SW rs2, imm(rs1)	50
136	5.6.19	ADDI rd, rs1, imm	50
137	5.6.20	SLTI rd, rs1, imm	50
138	5.6.21	SLTIU rd, rs1, imm	51
139	5.6.22	XORI rd, rs1, imm	51
140	5.6.23	ORI rd, rs1, imm	52
141	5.6.24	ANDI rd, rs1, imm	52
142	5.6.25	SLLI rd, rs1, shamt	53
143	5.6.26	SRLI rd, rs1, shamt	53
144	5.6.27	SRAI rd, rs1, shamt	54
145	5.6.28	ADD rd, rs1, rs2	54
146	5.6.29	SUB rd, rs1, rs2	55
147	5.6.30	SLL rd, rs1, rs2	55
148	5.6.31	SLT rd, rs1, rs2	56
149	5.6.32	SLTU rd, rs1, rs2	56
150	5.6.33	XOR rd, rs1, rs2	56
151	5.6.34	SRL rd, rs1, rs2	57
152	5.6.35	SRA rd, rs1, rs2	57
153	5.6.36	OR rd, rs1, rs2	58
154	5.6.37	AND rd, rs1, rs2	58
155	5.6.38	FENCE predecessor, successor	59

156	5.6.39 FENCE.I	59
157	5.6.40 ECALL	60
158	5.6.41 EBREAK	60
159	5.6.42 CSRRW rd, csr, rs1	60
160	5.6.43 CSRRS rd, csr, rs1	60
161	5.6.44 CSRRRC rd, csr, rs1	61
162	5.6.45 CSRRWI rd, csr, imm	61
163	5.6.46 CSRRSI rd, csr, rs1	61
164	5.6.47 CSRRCI rd, csr, rs1	62
165	5.7 RV32M Standard Extension	62
166	5.7.1 MUL rd, rs1, rs2	62
167	5.7.2 MULH rd, rs1, rs2	62
168	5.7.3 MULHS rd, rs1, rs2	63
169	5.7.4 MULHU rd, rs1, rs2	63
170	5.7.5 DIV rd, rs1, rs2	63
171	5.7.6 DIVU rd, rs1, rs2	63
172	5.7.7 REM rd, rs1, rs2	63
173	5.7.8 REMU rd, rs1, rs2	63
174	5.8 RV32A Standard Extension	64
175	5.9 RV32F Standard Extension	64
176	5.10 RV32D Standard Extension	64
177	A Installing a RISC-V Toolchain	65
178	A.1 The GNU Toolchain	65
179	A.2 rvddt	66
180	B Using The RISC-V GNU Toolchain	67
181	C Floating Point Numbers	69
182	C.1 IEEE-754 Floating Point Number Representation	69
183	C.1.1 Floating Point Number Accuracy	70
184	C.1.1.1 Powers Of Two	70
185	C.1.1.2 Clean Decimal Numbers	71
186	C.1.1.3 Accumulation of Error	72
187	C.1.2 Reducing Error Accumulation	73
188	D The ASCII Character Set	75
189	D.1 NAME	75
190	D.2 DESCRIPTION	75
191	D.2.1 Tables	77

192	D.3 NOTES	77
193	D.3.1 History	77
194	D.4 COLOPHON	77
195	E Attribution 4.0 International	78
196	Bibliography	82
197	Glossary	83
198	Index	84

Preface

200 I set out to this book because I couldn't find it in a single volume elsewhere.

201 The closest thing to what I sought when deciding to collect my thoughts into this document would
 202 be select portions of *The RISC-V Instruction Set Manual, Volume I: User-Level ISA, Document*
 203 *Version 2.2*[1], The RISC-V Reader[2], and Computer Organization and Design RISC-V Edition: The
 204 Hardware Software Interface[3].

205 There *are* some terse guides around the Internet that are suitable for those that already know an
 206 assembly language. With all the (deserved) excitement brewing over system organization (and the
 207 need to compress the time out of university courses targeting assembly language programming [4]), it
 208 is no surprise that RISC-V texts for the beginning assembly programmer are not (yet) available.

209 When I got started in computing I learned how to count in binary in a high school electronics course
 210 using data sheets for integrated circuits such as the 74191[5] and 74154[6] prior to knowing that
 211 assembly language even existed.

212 I learned assembler from data sheets and texts (that are still sitting on my shelves) such as:

- 213 • The MCS-85 User's Manual[7]
- 214 • The EDTASM Manual[8]
- 215 • The MC68000 User's Manual[9]
- 216 • Assembler Language With ASSIST[10]
- 217 • IBM System/370 Principals of Operation[11]
- 218 • OS/VS-DOS/VSE-VM/370 Assembler Language[12]
- 219 • ... and several others

220 One way or another all of them discuss each CPU instruction in excruciating detail with both a
 221 logical and narrative description. For RISC-V this is also the case for the *RISC-V Reader*[2] and
 222 the *Computer Organization and Design RISC-V Edition*[3] books and is also present in this text (I
 223 consider that to be the minimal level of responsibility.)

224 Where I hope this text will differentiate itself from the existing RISC-V titles is in its attempt to
 225 address the needs of those learning assembly language for the first time. To this end I have primed
 226 this project with some of the material from old handouts I used when teaching assembly language
 227 programming in the late '80s.

Chapter 1

Introduction

At its core, a digital computer has at least one **Central Processing Unit (CPU)**. A CPU executes a continuous stream of instructions called a **program**. These program instructions are expressed in what is called **machine language**. Each machine language instruction is a **binary** value. In order to provide a method to simplify the management of machine language programs a symbolic mapping is provided where a **mnemonic** can be used to specify each machine instruction and any of its parameters... rather than require that programs be expressed as a series of binary values. A set of mnemonics, parameters and rules for specifying their use for the purpose of programming a CPU is called an *Assembly Language*.

1.1 The Digital Computer

There are different types of computers. A *digital* computer is the type that most people think of when they hear the word *computer*. Other varieties of computers include *analog* and *quantum*.

A digital computer is one that processes data that are represented using numeric values (digits), most commonly expressed in binary (ones and zeros) form.

This text focuses on digital computing.

A typical digital computer is composed of storage systems (memory, disc drives, USB drives, etc.), a CPU (with one or more cores), input peripherals (a keyboard and mouse) and output peripherals (display, printer or speakers.)

1.1.1 Storage Systems

Computer storage systems are used to hold the data and instructions for the CPU.

Types of computer storage can be classified into two categories: volatile and non-volatile.

register
CPU

1.1.1.1 Volatile Storage

Volatile storage is characterized by the fact that it will lose its contents (forget) any time that it is powered off.

One type of volatile storage is provided inside the CPU itself in small blocks called [registers](#). These registers are used to hold individual data values that can be manipulated by the instructions that are executed by the CPU.

Another type of volatile storage is main memory (sometimes called [RAM](#)) Main memory is connected to a computer's CPU and is used to hold the data and instructions that can not fit into the CPU registers.

Typically, a CPU's registers can hold tens of data values while the main memory can contain many billions of data values.

To keep track of the data values, each register is assigned a number and the main memory is broken up into small blocks called [bytes](#) that are also each assigned number called an [address](#) (an address is often referred to as a *location*).

A CPU can process data in a register at a speed that can be an order of magnitude faster than the rate that it can process (specifically, transfer data and instructions to and from) the main memory.

Register storage costs an order of magnitude more to manufacture than main memory. While it is desirable to have many registers the economics dictate that the vast majority of volatile computer storage be provided in its main memory. As a result, optimizing the copying of data between the registers and main memory is a desirable trait of good programs.

1.1.1.2 Non-Volatile Storage

Non-volatile storage is characterized by the fact that it will *NOT* lose its contents when it is powered off.

Common types of non-volatile storage are disc drives, [ROM](#) flash cards and USB drives. Prices can vary widely depending on size and transfer speeds.

It is typical for a computer system's non-volatile storage to operate more slowly than its main memory.

This text is not particularly concerned with non-volatile storage.

1.1.2 CPU

The [CPU](#) is a collection of registers and circuitry designed manipulate the register data and to exchange data and instructions with the storage system. The instructions that are read from the main memory tell the CPU to perform various mathematic and logical operations on the data in its registers and where to save the results of those operations.

► Fix Me:
Add a block diagram of the
CPU components described
here.

1.1.2.1 Execution Unit

The part of a CPU that coordinates all aspects of the operations of each instruction is called the *execution unit*. It is what performs the transfers of instructions and data between the CPU and

ALU
register
hart

the main memory and tells the registers when they are supposed to either store or recall data being transferred. The execution unit also controls the ALU (Arithmetic and Logic Unit).

1.1.2.2 Arithmetic and Logic Unit

When an instruction manipulates data by performing things like an *addition*, *subtraction*, *comparison* or other similar operations, the ALU is what will calculate the sum, difference, and so on... under the control of the execution unit.

1.1.2.3 Registers

In the RV32 CPU there are 31 general purpose registers that each contain 32 *bits* (where each bit is one *binary* digit value of one or zero) and a number of special-purpose registers. Each of the general purpose registers is given a name such as *x1*, *x2*, ... on up to *x31* (*general purpose* refers to the fact that the CPU itself does not prescribe any particular function to any these registers.) Two important special-purpose registers are *x0* and *pc*.

Register *x0* will always represent the value zero or logical *false* no matter what. If any instruction tries to change the value in *x0* the operation will fail. The need for *zero* is so common that, other than the fact that it is hard-wired to zero, the *x0* register is made available as if it were otherwise a general purpose register.¹

The *pc* register is called the *program counter*. The CPU uses it to remember the memory address where its program instructions are located.

The number of bits in each register is defined by the *Instruction Set Architecture* (ISA).

► Fix Me:
Say something about XLEN?

1.1.2.4 Harts

Analogous to a *core* in other types of CPUs, a *hart* (hardware *thread*) in a RISC-V CPU refers to the collection of 32 registers, instruction execution unit and ALU.^[1, p. 20]

When more than one hart is present in a CPU, a different stream of instructions can be executed on each hart all at the same time. Programs that are written to take advantage of this are called *multithreaded*.

This text will primarily focus on CPUs that have only one hart.

1.1.3 Peripherals

A peripheral is a device that is not a CPU or main memory. They are typically used to transfer information/data into and out of the main memory.

This text is not particularly concerned with the peripherals of a computer system other than in those sections where instructions are discussed whose purpose is to address the needs of a peripheral device. Such instructions are used to initiate, execute and/or synchronize data transfers.

¹Having a special *zero* register allows the total set of instructions that the CPU can execute to be simplified. Thus reducing its complexity, power consumption and cost.

ISA
RV32I
RV32M
RV32A
RV32F
RV32D
RV32Q
RV32C
RV32G
instruction cycle

1.2 Instruction Set Architecture

The catalog of rules that describes the details of the instructions and features that a given CPU provides is called its [Instruction Set Architecture \(ISA\)](#).

An ISA is typically expressed in terms of the specific meaning of each binary instruction that a CPU can recognize and how it will process each one.

The RISC-V ISA is defined as a set of modules. The purpose of dividing the ISA into modules is to allow an implementer to select which features to incorporate into a CPU design.^[1, p. 4]

Any given RISC-V implementation must provide one of the *base* modules and zero or more of the *extension* modules.^[1, p. 4]

1.2.1 RV Base Modules

The base modules are RV32I (32-bit general purpose), RV32E (32-bit embedded), RV64I (64-bit general purpose) and RV128I (128-bit general purpose).^[1, p. 4]

These base modules provide the minimal functional set of integer operations needed to execute a useful application. The differing bit-widths address the needs of different main-memory sizes.

This text primarily focuses on the RV32I base module and how to program it.

1.2.2 Extension Modules

RISC-V extension modules may be included by an implementer interested in optimizing a design for one or more purposes.^[1, p. 4]

Available extension modules include M (integer math), A (atomic), F (32-bit floating point), D (64-bit floating point), Q (128-bit floating point), C (compressed size instructions) and others.

The extension name *G* is used to represent the combined set of IMAFD extensions as it is expected to be a common combination.

1.3 How the CPU Executes a Program

The process of executing a program is continuously repeating series of *instruction cycles* that are each comprised of a *fetch*, *decode* and *execute* phase.

The current status of a CPU hart is entirely embodied in the data values that are stored in its registers at any moment in time. Of particular interest to an executing a program is the `pc` register. The `pc` contains the memory address containing the instruction that the CPU is currently executing.²

For this to work, the instructions to be executed must have been previously stored in adjacent main memory locations and the address of the first instruction placed into the `pc` register.

²In the RISC-V ISA the `pc` register points to the *current* instruction where in most other designs, the `pc` register points to the *next* instruction.

instruction fetch
instruction decode
instruction execute

1.3.1 Instruction Fetch

In order to *fetch* an instruction from the main memory the CPU must have a method to identify which instruction should be fetched and a method to fetch it.

Given that the main memory is broken up and that each of its bytes is assigned an address, the `pc` is used to hold the address of the location where the next instruction to execute is located.

Given an instruction address, the CPU can request that the main memory locate and return the value of the data stored there using what is called a *memory read* operation and then the CPU can treat that *fetch*ed value as an instruction and execute it.³

1.3.2 Instruction Decode

Once an instruction has been fetched, it must be inspected to determine what operation(s) are to be performed. This primarily boils down to inspecting the portions of the instruction that dictate which registers are involved and, if the ALU is required, what it should do.

1.3.3 Instruction Execute

Typical instructions do things like add a number to the value currently stored in one of the registers or store the contents of a register into the main memory at some given address.

Also part of every instruction is a notion of what should be done next.

Most of the time an instruction will complete by indicating that the CPU should proceed to fetch and execute the instruction at the next larger main memory address. In these cases the `pc` is incremented to point to the memory address after the current instruction.

Any parameters that an instruction requires must either be part of the instruction itself or read from (or stored into) one or more of the general purpose registers.

Some instructions can specify that the CPU proceed to execute an instruction at an address other than the one that follows itself. This class of instructions have names like *jump* and *branch* and are available in a variety of different styles.

The RISC-V ISA uses the word *jump* to refer to an *unconditional* change in the sequential processing of instructions and the word *branch* to refer to a *conditional* change.

For example, a (conditional) branch instruction might instruct the CPU to proceed to the instruction at the next main memory address if the value in register number 8 is currently less than the value in register number 24 *but otherwise* proceed to an instruction at a different address when it is not. This type of instruction can therefore result in having one of two different actions pending the resulting *condition* of the comparison.⁴

Once the instruction execution phase has completed, the next instruction cycle will be performed using the new value in the `pc` register.

³RV32I instructions are more than one byte in size, but this general description is suitable for now.

⁴This is the fundamental method used by a CPU to make decisions.

Chapter 2

Numbers and Storage Systems

This chapter discusses how data are represented and stored in a computer.

In the context of computing, *boolean* refers to a condition that can be either true and false and *binary* refers to the use of a base-2 numeric system to represent numbers.

RISC-V assembly language uses binary to represent all values, be they boolean or numeric. It is the context within which they are used that determines whether they are boolean or numeric.

► Fix Me:

Add some diagrams here showing bits, bytes and the MSB, LSB, ... perhaps relocated from the RV32I chapter?

2.1 Boolean Functions

Boolean functions apply on a per-bit basis. When applied to multi-bit values, each bit position is operated upon independently of the other bits.

RISC-V assembly language uses zero to represent *false* and one to represent *true*. In general, however, it is useful to relax this and define zero **and only zero** to be *false* and anything that is not *false* is therefore *true*.¹

The reason for this relaxation is because, while a single binary digit (*bit*) can represent the two values zero and one, the vast majority of the time data is processed by the CPU in groups of bits. These groups have names like *byte* (8 bits), *halfword* (16 bits) and *fullword* (32 bits).

2.1.1 NOT

The *NOT* operator applies to a single operand and represents the opposite of the input.

If the input is 1 then the output is 0. If the input is 0 then the output is 1. In other words, the output value is *not* that of the input value.

Expressing the *not* function in the form of a truth table:

► Fix Me:

Need to define unary, binary and ternary operators without confusing binary operators with binary numbers.

¹This is how *true* and *false* behave in C, C++, and many other languages as well as the common assembly language idioms discussed in this text.

A	\overline{A}
0	1
1	0

A truth table is drawn by indicating all of the possible input values on the left of the vertical bar with each row displaying the output values that correspond to the input for that row. The column headings are used to define the illustrated operation expressed using a mathematical notation. The *not* operation is indicated by the presense of an *overline*.

In computer programming languages, things like an overline can not be efficiently expressed using a standard keyboard. Therefore it is common to use a notation such as that used by the C language when discussing the *NOT* operator in symbolic form. Specifically the tilde: ‘~’.

It is also uncommon to for programming languages to express boolean operations on single-bit input(s). A more generalized operation is used that applies to a set of bits all at once. For example, performing a *not* operation of eight bits at once can be illustrated as:

```

~ 1 1 1 1 0 1 0 1  <== A
-----
0 0 0 0 1 0 1 0  <== output

```

In a line of code the above might read like this: `output = ~A`

2.1.2 AND

The boolean *and* function has two or more inputs and the output is a single bit. The output is 1 if and only if all of the input values are 1. Otherwise it is 0.

This function works like it does in spoken language. For example if A is 1 *AND* B is 1 then the output is 1 (true). Otherwise the output is 0 (false).

In mathamatical notion, the *and* operator is expressed the same way as is *multiplication*. That is by a raised dot between, or by juxtaposition of, two variable names. It is also worth noting that, in base-2, the *and* operation actually *is* multiplication!

A	B	AB
0	0	0
0	1	0
1	0	0
1	1	1

This text will use the operator used in the C language when discussing the *AND* operator in symbolic form. Specifically the ampersand: ‘&’.

An eight-bit example:

```

1 1 1 1 0 1 0 1  <== A
& 1 0 0 1 0 0 1 1  <== B
-----
1 0 0 1 0 0 0 1  <== output

```

In a line of code the above might read like this: `output = A & B`

2.1.3 OR

The boolean *or* function has two or more inputs and the output is a single bit. The output is 1 if at least one of the input values are 1.

This function works like it does in spoken language. For example if A is 1 *OR* B is 1 then the output is 1 (true). Otherwise the output is 0 (false).

In mathematical notion, the *or* operator is expressed using the plus (+).

A	B	A+B
0	0	0
0	1	1
1	0	1
1	1	1

This text will use the operator used in the C language when discussing the *OR* operator in symbolic form. Specifically the pipe: '|'.

An eight-bit example:

```

  1 1 1 1 0 1 0 1  <== A
| 1 0 0 1 0 0 1 1  <== B
-----
  1 1 1 1 0 1 1 1  <== output

```

In a line of code the above might read like this: `output = A | B`

2.1.4 XOR

The boolean *exclusive or* function has two or more inputs and the output is a single bit. The output is 1 if only an odd number of inputs are 1. Otherwise the output will be 0.

Note that when *XOR* is used with two inputs, the output is set to 1 (true) when the inputs have different values and 0 (false) when the inputs both have the same value.

In mathematical notion, the *xor* operator is expressed using the plus in a circle (\oplus).

A	B	A \oplus B
0	0	0
0	1	1
1	0	1
1	1	0

This text will use the operator used in the C language when discussing the *XOR* operator in symbolic form. Specifically the carrot: '^'.

An eight-bit example:

Decimal			Binary								Hex	
10^2	10^1	10^0	2^7	2^6	2^5	2^4	2^3	2^2	2^1	2^0	16^1	16^0
100	10	1	128	64	32	16	8	4	2	1	16	1
0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	1	0	0	0	0	0	0	0	1	0	1
0	0	2	0	0	0	0	0	0	1	0	0	2
0	0	3	0	0	0	0	0	0	1	1	0	3
0	0	4	0	0	0	0	0	1	0	0	0	4
0	0	5	0	0	0	0	0	1	0	1	0	5
0	0	6	0	0	0	0	0	1	1	0	0	6
0	0	7	0	0	0	0	0	1	1	1	0	7
0	0	8	0	0	0	0	1	0	0	0	0	8
0	0	9	0	0	0	0	1	0	0	1	0	9
0	1	0	0	0	0	0	1	0	1	0	0	a
0	1	1	0	0	0	0	1	0	1	1	0	b
0	1	2	0	0	0	0	1	1	0	0	0	c
0	1	3	0	0	0	0	1	1	0	1	0	d
0	1	4	0	0	0	0	1	1	1	0	0	e
0	1	5	0	0	0	0	1	1	1	1	0	f
0	1	6	0	0	0	1	0	0	0	0	1	0
0	1	7	0	0	0	1	0	0	0	1	1	1
...			
1	2	5	0	1	1	1	1	1	0	1	7	d
1	2	6	0	1	1	1	1	1	1	0	7	e
1	2	7	0	1	1	1	1	1	1	1	7	f
1	2	8	1	0	0	0	0	0	0	0	8	0

Figure 2.1: Counting in decimal, binary and hexadecimal.

```

459   1 1 1 1 0 1 0 1 <== A
460 ^ 1 0 0 1 0 0 1 1 <== B
461 -----
462   0 1 1 0 0 1 1 0 <== output

```

463 In a line of code the above might read like this: `output = A ^ B`

464 2.2 Integers and Counting

465 A binary integer is constructed with only 1s and 0s in the same manner as decimal numbers are
466 constructed with values from 0 to 9.

467 Counting in binary (base-2) uses the same basic rules as decimal (base-10). The difference comes in
468 when we consider that there are ten decimal digits and only two binary digits. Therefore, in base-10,
469 we must carry when adding one to nine (because there is no digit representing a ten) and, in base-2,
470 we must carry when adding one to one (because there is no digit representing a two.)

471 Figure 2.1 shows an abridged table of the decimal, binary and hexadecimal values ranging from 0_{10}
472 to 129_{10} .

473 One way to look at this table is on a per-row basis where each place value is represented by the

base raised to the power of the place value position (shown in the column headings.) For example to interpret the decimal value on the fourth row:

Most significant bit
MSB—see Most
significant bit
Least significant bit
LSB—see Least
significant bit

$$0 \times 10^2 + 0 \times 10^1 + 3 \times 10^0 = 3_{10} \quad (2.2.1)$$

Interpreting the binary value on the fourth row by converting it to decimal:

$$0 \times 2^7 + 0 \times 2^6 + 0 \times 2^5 + 0 \times 2^4 + 0 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 1 \times 2^0 = 3_{10} \quad (2.2.2)$$

Interpreting the hexadecimal value on the fourth row by converting it to decimal:

$$0 \times 16^1 + 3 \times 16^0 = 3_{10} \quad (2.2.3)$$

We refer to the place values with the largest exponent (the one furthest to the left for any given base) as the *most significant* digit and the place value with the lowest exponent as the *least significant* digit. For binary numbers these are the **Most Significant Bit (MSB)** and **Least Significant Bit (LSB)** respectively.²

Another way to look at this table is on a per-column basis. When tasked with drawing such a table by hand, it might be useful to observe that, just as in decimal, the right-most column will cycle through all of the values represented in the chosen base then cycle back to zero and repeat. (For example, in binary this pattern is 0-1-0-1-0-1-0-...) The next column in each base will cycle in the same manner except each of the values is repeated as many times as is represented by the place value (in the case of decimal, 10^1 times, binary 2^1 times, hex 16^1 times. Again, the for binary numbers this pattern is 0-0-1-1-0-0-1-1-...) This continues for as many columns as are needed to represent the magnitude of the desired number.

Another item worth noting is that any even binary number will always have a 0 LSB and odd numbers will always have a 1 LSB.

As is customary in decimal, leading zeros are sometimes not shown for readability.

The relationship between binary and hex values is also worth taking note. Because $2^4 = 16$, there is a clean and simple grouping of 4 bits to 1 hit (aka **nybble**). There is no such relationship between binary and decimal.

Writing and reading numbers in binary that are longer than 8 bits is cumbersome and prone to error. The simple conversion between binary and hex makes hex a convenient shorthand for expressing binary values in many situations.

For example, consider the following value expressed in binary, hexadecimal and decimal (spaced to show the relationship between binary and hex):

Binary value:	0010	0111	1011	1010	1100	1100	1111	0101
Hex Value:	2	7	B	A	C	C	F	5
Decimal Value:							666553589	

Empirically we can see that grouping the bits into sets of four allows an easy conversion to hex and

²Changing the value of the MSB will have a more *significant* impact on the numeric value than changing the value of the LSB.

expressing it as such is $\frac{1}{4}$ as long as in binary while at the same time allowing for easy conversion back to binary.

The decimal value in this example does not easily convey a sense of the binary value.

2.2.1 Converting Between Bases

2.2.1.1 From Binary to Decimal

Alas, it is occasionally necessary to convert between decimal, binary and/or hex.

To convert from binary to decimal, put the decimal value of the place values ...8 4 2 1 over the binary digits like this:

Base-2 place values:	128	64	32	16	8	4	2	1
Binary:	0	0	0	1	1	0	1	1
Decimal:				16	+8		+2	+1 = 27

Now sum the place-values that are expressed in decimal for each bit with the value of 1: $16 + 8 + 2 + 1$. The integer binary value 00011011_2 represents the decimal value 27_{10} .

2.2.1.2 From Binary to Hexadecimal

Conversion from binary to hex involves grouping the bits into sets of four and then performing the same summing process as shown above. If there is not a multiple of four bits then extend the binary to the left with zeros to make it so.

Grouping the bits into sets of four and summing:

Base-2 place values:	8	4	2	1	8	4	2	1	8	4	2	1	8	4	2	1
Binary:	0	1	1	0	1	1	0	1	1	0	1	0	1	1	1	0
Decimal:				4+2 =6		8+4+ 1=13		8+ 2 =10		8+4+2 =14						

After the summing, convert each decimal value to hex. The decimal values from 0–9 are the same values in hex. Because we don't have any more numerals to represent the values from 10–15, we use the first 6 letters (See the right-most column of [Figure 2.1](#).) Fortunately there are only six hex mappings involving letters. Thus it is reasonable to memorize them.

Continuing this example:

Decimal:	6	13	10	14
Hex:	6	D	A	E

2.2.1.3 From Hexadecimal to Binary

Again, the four-bit mapping between binary and hex makes this task as straight forward as using a look-up table.

For each [hit](#) (Hex digIT), translate it to its unique four-bit pattern. Perform this task either by memorizing each of the 16 patterns or by converting each hit to decimal first and then converting each four-bit binary value to decimal using the place-value summing method discussed in [subsection 2.2.1.1](#).

For example:

Hex:		7		C
Decimal Sum:		4+2+1=7	8+4	=12
Binary:	0	1	1	1
			1	1
			0	0

2.2.1.4 From Decimal to Binary

To convert arbitrary decimal numbers to binary, extend the list of binary place values until it exceeds the value of the decimal number being converted. Then make successive subtractions of each of the place values that would yield a non-negative result.

For example, to convert 1234_{10} to binary:

Base-2 place values: 2048-1024-512-256-128-64-32-16-8-4-2-1

0		2048	(too big)
1	1234 - 1024 =	210	
0		512	(too big)
0		256	(too big)
1	210 - 128 =	82	
1	82 - 64 =	18	
0		32	(too big)
1	18 - 16 =	2	
0		8	(too big)
0		4	(too big)
1	2 - 2 =	0	
0		1	(too big)

The answer using this notation is listed vertically in the left column with the [MSB](#) on the top and the [LSB](#) on the bottom line: 010011010010₂.

2.2.1.5 From Decimal to Hex

Conversion from decimal to hex can be done by using the place values for base-16 and the same math as from decimal to binary or by first converting the decimal value to binary and then from binary to hex by using the methods discussed above.

Because binary and hex are so closely related, performing a conversion by way of binary is quite straight forward.

2.2.2 Addition of Binary Numbers

The addition of binary numbers can be performed long-hand the same way decimal addition is taught in grade school. In fact binary addition is easier since it only involves adding 0 or 1.

The first thing to note that in any number base $0 + 0 = 0$, $0 + 1 = 1$, and $1 + 0 = 1$. Since there is no “two” in binary (just like there is no “ten” decimal) adding $1 + 1$ results in a zero with a carry as in: $1 + 1 = 10_2$ and in: $1 + 1 + 1 = 11_2$. Using these five sums, any two binary integers can be added.

For example:

```

      111111 1111 <== carries
    0110101111001111 <== addend
+ 0000011101100011 <== addend
-----
    0111001100110010 <== sum

```

2.2.3 Signed Numbers

There are multiple methods used to represent signed binary integers. The method used by most modern computers is called “two’s complement.”

A two’s complement number is encoded in such a manner as to simplify the hardware used to add, subtract and compare integers.

A simple method of thinking about two’s complement numbers is to negate the place value of the **MSB**. For example, the number one is represented the same as discussed before:

```

Base-2 place values:  -128 64 32 16  8  4  2  1
Binary:                0  0  0  0  0  0  0  1

```

The **MSB** of any negative number in this format will always be 1. For example the value -1_{10} is:

```

Base-2 place values:  -128 64 32 16  8  4  2  1
Binary:                1  1  1  1  1  1  1  1

```

... because: $-128 + 64 + 32 + 16 + 8 + 4 + 2 + 1 = -1$.

Calculating $4 + 5 = 9$

```

      1      <== carries
    000100 <== 4
+000101 <== 5
-----
    001001 <== 9

```

Calculating $-4 + -5 = -9$

```

    1 11      <== carries
    111100 <== -4
+111011 <== -5
-----
    1 110111 <== -9 (with a truncation)

```

```

608
609   -32 16 8 4 2 1
610     1  1 0 1 1 1
611   -32 + 16 + 4 + 2 + 1 = -9

```

612 This format has the virtue of allowing the same addition logic discussed above to be used to calculate
 613 $-1 + 1 = 0$.

```

614   -128 64 32 16  8  4  2  1 <== place value
615     1  1  1  1  1  1  1  1  0 <== carries
616       1  1  1  1  1  1  1  1 <== addend (-1)
617     + 0  0  0  0  0  0  0  1 <== addend (1)
618     -----
619     1  0  0  0  0  0  0  0  0 <== sum (0 with a truncation)

```

620 In order for this to work, the carry out of the sum of the MSBs is ignored.

621 2.2.3.1 Converting between Positive and Negative

622 Changing the sign on two's complement numbers can be described as inverting all of the bits (which
 623 is also known as the one's complement) and then add one.

624 For example, inverting the number *four*:

```

625   -128 64 32 16  8  4  2  1
626     0  0  0  0  0  1  0  0 <== 4
627
628           1  1      <== carries
629     1  1  1  1  1  0  1  1 <== one's complement of 4
630   + 0  0  0  0  0  0  0  1 <== plus 1
631   -----
632     1  1  1  1  1  1  0  0 <== -4

```

633 This can be verified by adding 5 to the result and observe that the sum is 1:

```

634   -128 64 32 16  8  4  2  1
635     1  1  1  1  1      <== carries
636     1  1  1  1  1  1  0  0 <== -4
637   + 0  0  0  0  0  1  0  1 <== 5
638   -----
639     1  0  0  0  0  0  0  1

```

640 Note that the changing of the sign using this method is symmetric in that it is identical when converting
 641 from negative to positive and when converting from positive to negative: flip the bits and add 1.

642 For example, changing the value -4 to 4 to illustrate the reverse of the conversion above:

```

643   -128 64 32 16  8  4  2  1

```

```

1 1 1 1 1 1 0 0 <== -4
          1 1 <== carries
0 0 0 0 0 0 1 1 <== one's complement of -4
+ 0 0 0 0 0 0 0 1 <== plus 1
-----
0 0 0 0 0 1 0 0 <== 4

```

2.2.4 Subtraction of Binary Numbers

Subtraction of binary numbers is performed by first negating the subtrahend and then adding the two numbers. Due to the nature of two's complement numbers this will work for both signed and unsigned numbers.

To calculate $-4 - 8 = -12$

```

-128 64 32 16 8 4 2 1
 1 1 1 1 1 1 0 0 <== -4
- 0 0 0 0 1 0 0 0 <== 8
          1 1 1 <== carries
 1 1 1 1 0 1 1 1 <== one's complement of -8
+ 0 0 0 0 0 0 0 1 <== plus 1
-----
 1 1 1 1 1 0 0 0 <== -8
          1 1 1 1 <== carries
 1 1 1 1 1 1 0 0 <== -4
+ 1 1 1 1 1 0 0 0 <== -8
-----
 1 1 1 1 1 0 1 0 <== -12

```

2.2.5 Truncation and Overflow

Discuss the details of truncation and overflow here.

I prefer to define *truncation* as the loss of data as result of the bit-length of the destination being too small to hold result of an operation and *overflow* as when the carry into a sign bit is not the same as the carry out of the sign bit.

Where addition and subtraction on the RV32 is concerned, the sum or difference of two unsigned 32-bit numbers will be *truncated* when the operation results in a carry out of bit 31. Unsigned operations can not overflow (as defined above).

(show a truncation picture here)

An Overflow occurs with signed numbers when the two addends are positive and sum is negative or the addends are both negative and the sum is positive.

(show an overflow picture here)

► Fix Me:

This section needs more examples of subtracting signed and unsigned numbers and a discussion on how signedness is not relevant until the results are interpreted. For example adding $-4 + -8 = -12$ using two 8-bit numbers is the same as adding $252 + 248 = 500$ and truncating the result to 244.

► Fix Me:

This chapter should be made consistent in its use of truncation and overflow as occur with signed and unsigned addition and subtraction.

(show mixed overflow and truncation situations here to drive home the need to ignore truncation when sign extension dealing with signed numbers.)

0xffffffff + 0x00000002 has truncation but not overflow (OK for signed, not OK for unsigned).

0xffffffff + 0xffffffe also has truncation but not overflow.

0x40000000 + 0x40000000 has overflow but not truncation. (We care if are signed numbers.)

0x80000000 + 0x80000000 has both overflow and truncation. (we care regardless of signedness)

Where subtraction is concerned the notion of a borrow is the same as carry.

Page 13 of [1] mixes these two notions of (and never mentions the word *truncate*) like this:

➡ Fix Me:

I think that overloading the word overflow like this can be is confusing to new programmers.

We did not include special instruction set support for overflow checks on integer arithmetic operations in the base instruction set, as many overflow checks can be cheaply implemented using RISC-V branches. Overflow checking for unsigned addition requires only a single additional branch instruction after the addition: `add t0, t1, t2; bltu t0, t1, overflow`.

For signed addition, if one operand's sign is known, overflow checking requires only a single branch after the addition: `addi t0, t1, +imm; blt t0, t1, overflow`. This covers the common case of addition with an immediate operand.

For general signed addition, three additional instructions after the addition are required, leveraging the observation that the sum should be less than one of the operands if and only if the other operand is negative.

```
add t0, t1, t2
slti t3, t2, 0
slt t4, t0, t1
bne t3, t4, overflow
```

In RV64, checks of 32-bit signed additions can be optimized further by comparing the results of ADD and ADDW on the operands.

2.3 Sign and Zero Extension

Due to the nature of the two's complement encoding scheme, the following numbers all represent the same value:

```

                                1111 <== -1
                        11111111 <== -1
                11111111111111111111 <== -1
        11111111111111111111111111111111 <== -1
```

As do these:

[illegible]

The phenomenon illustrated here is called *sign extension*. That is any signed number can have any quantity of additional MSBs added to it, provided that they repeat the value of the sign bit.

Figure 2.2 illustrates extending the negative sign bit of *val* to the left by replicating it. When *val* is negative, its MSB (bit 19 in this example) will be set to 1. Extending this value to the left will set all the new bits to the left of it to 1 as well.




Figure 2.2: Sign-extending a negative integer from 20 bits to 32 bits.

Figure 2.3 illustrates extending the positive sign bit of *val* to the left by replicating it. When *val* is positive, its MSB will be set to 0. Extending this value to the left will set all the new bits to the left of it to 0 as well.




Figure 2.3: Sign-extending a positive integer from 20 bits to 32 bits.

In a similar vein, any *unsigned* number also may have any quantity of additional MSBs added to it provided that they are all zero. For example, the following all represent the same value:

```

1111 <== 15
01111 <== 15
0000000000000000000000001111 <== 15

```

The observation here is that any *unsigned* number may be *zero extended* to any size.

Figure 2.4 illustrates zero-extending a 20-bit *val* to the left to form a 32-bit fullword.




Figure 2.4: Zero-extending an unsigned integer from 20 bits to 32 bits.

2.4 Shifting

Seems like a good place to discuss logical and arithmetic shifting.

shift left logical

shift right logical

shift right arithmetic

2.5 Main Memory Storage

When transferring data between its registers registers and main memory a RISC-V system uses the little-endian byte order.³

► Fix Me:

Consider refactoring the memory discussion in RV32 reference chapter and placing some of it in this section.

2.5.1 Memory Dump

Introduce the memory dump and how to read them here.

Listing 2.1: rvddt_memdump.out

rvddt memory dump

```
ddt> d 0x00002600
00002600: 93 05 00 00 13 06 00 00-93 06 00 00 13 07 00 00 *.*****
00002610: 93 07 00 00 93 08 d0 05-73 00 00 00 63 54 05 02 *.*****s...cT..*
00002620: 13 01 01 ff 23 24 81 00-13 04 05 00 23 26 11 00 *....#$.*****#&...*
00002630: 33 04 80 40 97 00 00 00-e7 80 40 01 23 20 85 00 *3..@.....@.# ...*
00002640: 6f 00 00 00 6f 00 00 00-b7 87 00 00 03 a5 07 43 *o...o.....C*
00002650: 67 80 00 00 00 00 00 00-76 61 6c 3d 00 00 00 00 *g.....val=....*
00002660: 00 00 00 00 80 84 2e 41-1f 85 45 41 80 40 9a 44 *.*****A..EA.@.D*
00002670: 4f 11 f3 c3 6e 8a 67 41-20 1b 00 00 20 1b 00 00 *0...n.gA ... ..*
00002680: 44 1b 00 00 14 1b 00 00-14 1b 00 00 04 1c 00 00 *D.....*
00002690: 44 1b 00 00 14 1b 00 00-04 1c 00 00 14 1b 00 00 *D.....*
000026a0: 44 1b 00 00 10 1b 00 00-10 1b 00 00 10 1b 00 00 *D.....*
000026b0: 04 1c 00 00 54 1f 00 00-54 1f 00 00 d4 1f 00 00 *....T...T.....*
000026c0: 4c 1f 00 00 4c 1f 00 00-34 20 00 00 d4 1f 00 00 *L...L...4 .....*
000026d0: 4c 1f 00 00 34 20 00 00-4c 1f 00 00 d4 1f 00 00 *L...4 ..L.....*
000026e0: 48 1f 00 00 48 1f 00 00-48 1f 00 00 34 20 00 00 *H...H...H...4 ...*
000026f0: 00 01 02 02 03 03 03 03-04 04 04 04 04 04 04 *.*****
```

2.5.2 Big Endian Representation

Using the memory dump contents in prior section, discuss how big endian values are stored.

2.5.3 Little Endian Representation

Using the memory dump contents in prior section, discuss how little endian values are stored.

2.5.4 Character Strings and Arrays

Define character strings and arrays.

Using the prior memory dump, discuss how and where things are stored and retrieved.

³ See[13] for some history of the big/little-endian “controversy.”

2.5.5 Context is Important!

Data values can be interpreted differently depending on the context in which they are used. Assuming what a set of bytes is used for based on their contents can be very misleading! For example, there is a 0x76 at address 0x00002658. This is a ‘v’ if you use it as an ASCII (see [Appendix D](#)) character, a 118₁₀ if it is an integer value and TRUE if it is a conditional.

2.5.6 Alignment

With respect to memory and storage, *alignment* refers to the *location* of a data element when the address that it is stored is a precise multiple of a power-of-2.

The primary alignments of concern are typically 2 (a halfword), 4 (a fullword), 8 (a double word) and 16 (a quad-word) bytes.

For example, any data element that is aligned to 2-byte boundary must have an (hex) address that ends in any of: 0, 2, 4, 6, 8, A, C or E. Any 4-byte aligned element must be located at an address ending in 0, 4, 8 or C. An 8-byte aligned element at an address ending with 0 or 8, and 16-byte aligned elements must be located at addresses ending in zero.

Such alignments are important when exchanging data between the CPU and memory because the hardware implementations are optimized to transfer aligned data. Therefore, aligning data used by any program will reap the benefit of running faster.

An element of data is considered to be *aligned to its natural size* when its address is an exact multiple of the number of bytes used to represent the data. Note that the ISA we are concerned with *only* operates on elements that have sizes that are powers of two.

For example, a 32-bit integer consumes one full word. If the four bytes are stored in main memory at an address that is a multiple of 4 then the integer is considered to be naturally aligned.

The same would apply to 16-bit, 64-bit, 128-bit and other such values as they fit into 2, 8 and 16 byte elements respectively.

Some CPUs can deliver four (or more) bytes at the same time while others might only be capable of delivering one or two bytes at a time. Such differences in hardware typically impact the cost and performance of a system.⁴

2.5.7 Instruction Alignment

The RISC-V ISA requires that all instructions be aligned to their natural boundaries.

Every possible instruction that an RV32I CPU can execute contains exactly 32 bits. Therefore they are always stored on a full word boundary. Any *unaligned* instruction would be *illegal*.

An attempt to fetch an instruction from an unaligned address will result in an error referred to as an alignment *exception*. This and other exceptions cause the CPU to stop executing the current instruction and start executing a different set of instructions that are prepared to handle the problem. Often an exception is handled by completely stopping the program in a way that is commonly referred to as a system or application *crash*.

⁴The design and implementation choices that determine how any given system operates are part of what is called a system’s *organization* and is beyond the scope of this text. See [3] for more information on computer organization.

► Fix Me:
Include the obligatory
diagram showing the
overlapping data types when
they are all aligned.

Chapter 3

The Elements of a Assembly Language Program

3.1 Assembly Language Statements

Introduce the assembly language grammar. Statement = 1 line of text containing an instruction or directive.

Instruction = label, mnemonic, operands, comment.

Directive = Used to control the operation of the assembler.

3.2 Memory Layout

Is this a good place to introduce the text, data, bss, heap and stack regions?

Or does that belong in a new section/chapter that discusses addressing modes?

3.3 A Sample Program Source Listing

A simple program that illustrates how this text presents program source code is seen in [Listing 3.1](#). This program will place a zero in each of the 4 registers named x28, x29, x30 and x31.

Listing 3.1: `zero4regs.S`
Setting four registers to zero.

```
1  .text                # put this into the text section
2  .align 2             # align to 2^2
3  .globl _start
4  _start:
5      addi    x28, x0, 0    # set register x28 to zero
6      addi    x29, x0, 0    # set register x29 to zero
7      addi    x30, x0, 0    # set register x30 to zero
8      addi    x31, x0, 0    # set register x31 to zero
```

This program listing illustrates a number of things:

- Listings are identified by the name of the file within which they are stored. This listing is from a file named: `zero4regs.S`.
- The assembly language programs discussed in this text will be saved in files that end with: `.S` (Alternately you can use `.sx` on systems that don't understand the difference between upper and lowercase letters.¹)
- A description of the listing's purpose appears under the name of the file. The description of [Listing 3.1](#) is *Setting four registers to zero*.
- The lines of the listing are numbered on the left margin for easy reference.
- An assembly program consists of lines of plain text.
- The RISC-V ISA does not provide an operation that will simply set a register to a numeric value. To accomplish our goal this program will add zero to zero and place the sum in each of the four registers.
- The lines that start with a dot `'.'` (on lines 1, 2 and 3) are called *assembler directives* as they tell the assembler itself how we want it to translate the following *assembly language instructions* into *machine language instructions*.
- Line 4 shows a *label* named `_start`. The colon at the end is the indicator to the assembler that causes it to recognize the preceding characters as a label.
- Lines 5-8 are the four assembly language instructions that make up the program. Each instruction in this program consists of four *fields*. (Different instructions can have a different number of fields.) The fields on line 5 are:
 - `addi` The instruction mnemonic. It indicates the operation that the CPU will perform.
 - `x28` The *destination* register that will receive the sum when the *addi* instruction is finished. The names of the 32 registers are expressed as `x0 – x31`.
 - `x0` One of the addends of the sum operation. (The `x0` register will always contain the value zero. It can never be changed.)
 - `0` The second addend is the number zero.
- `# set ...` Any text anywhere in a RISC-V assembly language program that starts with the pound-sign is ignored by the assembler. They are used to place a *comment* in the program to help the reader better understand the motive of the programmer.

3.4 Running a Program With rvddt

To illustrate what a CPU does when it executes instructions this text will use the `rvddt` simulator to display shows sequence of events and the binary values involved. This simulator supports the RV32I ISA and has a configurable amount of memory.²

[Listing 3.2](#) shows the operation of the four *addi* instructions from [Listing 3.1](#) when it is executed in trace-mode.

¹The author of this text prefers to avoid using such systems.

²The `rvddt` simulator was written to generate the listings for this text. It is similar to the fancier *spike* simulator. Given the simplicity of the RV32I ISA, `rvddt` is less than 1700 lines of C++ and was written in one (long) afternoon.

Listing 3.2: zero4regs.out

Running a program with the rvdtdt simulator

```

867 [winans@w510 src]$ ./rvdtdt -f ../examples/load4regs.bin
868 1 Loading '../examples/load4regs.bin' to 0x0
869 2 ddt> t4
870 3
871 4   x0: 00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
872 5   x8: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
873 6  x16: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
874 7  x24: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
875 8   pc: 00000000
876 9 00000000: 00000e13 addi    x28, x0, 0    # x28 = 0x00000000 = 0x00000000 + 0x00000000
877 10  x0: 00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
878 11  x8: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
879 12  x16: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
880 13  x24: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0
881 14  pc: 00000004
882 15 00000004: 00000e93 addi    x29, x0, 0    # x29 = 0x00000000 = 0x00000000 + 0x00000000
883 16  x0: 00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
884 17  x8: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
885 18  x16: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
886 19  x24: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-00000000 00000000 f0f0f0f0 f0f0f0f0
887 20  pc: 00000008
888 21 00000008: 00000f13 addi    x30, x0, 0    # x30 = 0x00000000 = 0x00000000 + 0x00000000
889 22  x0: 00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
890 23  x8: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
891 24  x16: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
892 25  x24: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-00000000 00000000 00000000 f0f0f0f0
893 26  pc: 0000000c
894 27 0000000c: 00000f93 addi    x31, x0, 0    # x31 = 0x00000000 = 0x00000000 + 0x00000000
895 28 ddt> r
896 29   x0: 00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
897 30   x8: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
898 31  x16: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
899 32  x24: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-00000000 00000000 00000000 00000000
900 33  pc: 00000010
901 34 ddt> x
902 35 [winans@w510 src]$

```

ℓ 1 This listing includes the command-line that shows how the simulator was executed to load a file containing the machine instructions (aka machine code) from the assembler.

ℓ 2 A message from the simulator indicating that it loaded the machine code into simulated memory at address 0.

ℓ 3 This line shows the prompt from the debugger and the command `t4` that the user entered to request that the simulator trace the execution of four instructions.

ℓ 4-8 Prior to executing the first instruction, the state of the CPU registers is displayed.

ℓ 4 The values in registers 0, 1, 2, 3, 4, 5, 6 and 7 are printed from left to right in [big endian](#), [hexadecimal](#) form. The dash ‘-’ character in the middle of the line is a reference to make it easier to visually navigate across the line without being forced to count the values from the far left when seeking the value of, say, `x5`.

ℓ 5-7 The values of registers 8–31 are printed.

ℓ 8 The *program counter* (`pc`) register is printed. It contains the address of the instruction that the CPU will execute. After each instruction, the `pc` will either advance four bytes ahead or be set to another value by a branch instruction as discussed above.

ℓ 9 A four-byte instruction is fetched from memory at the address in the `pc` register, is decoded and printed. From left to right the fields shown on this line are:

```

921 00000000 The memory address from which the instruction was fetched. This address is displayed in
922         big endian, hexadecimal form.
923 00000e13 The machine code of the instruction displayed in big endian, hexadecimal form.
924     addi The mnemonic for the machine instruction.
925     x28 The rd field of the addi instruction.
926     x0 The rs1 field of the addi instruction that holds one of the two addends of the operation.
927     0 The imm field of the addi instruction that holds the second of the two addends of the
928       operation.
929     # ... A simulator-generated comment that explains what the instruction is doing. For this in-
930           struction it indicates that x28 will have the value zero stored into it as a result of performing
931           the addition: 0 + 0.

932 ℓ 10-14 These lines are printed as the prelude while tracing the second instruction. Lines 7 and 13 show
933         that x28 has changed from f0f0f0f0 to 00000000 as a result of executing the first instruction and
934         lines 8 and 14 show that the pc has advanced from zero (the location of the first instruction) to
935         four, where the second instruction will be fetched. None of the rest of the registers have changed
936         values.

937 ℓ 15 The second instruction decoded executed and described. This time register x29 will be assigned
938       a value.

939 ℓ 16-27 The third and fourth instructions are traced.

940 ℓ 28 Tracing has completed. The simulator prints its prompt and the user enters the 'r' command
941       to see the register state after the fourth instruction has completed executing.

942 ℓ 29-33 Following the fourth instruction it can be observed that registers x28, x29, x30 and x31 have
943         been set to zero and that the pc has advanced from zero to four, then eight, then 12 (the hex
944         value for 12 is c) and then to 16 (which, in hex, is 10).

945 ℓ 34 The simulator exit command 'x' is entered by the user and the terminal displays the shell prompt.

```


Chapter 4

Writing RISC-V Programs

This chapter introduces each of the RV32I instructions by developing programs that demonstrate their usefulness.


► Fix Me:
Introduce the ISA register names and aliases in here?

4.1 Use ebreak to Stop rvddt Execution

The **ebreak** instruction exists for the sole purpose of transferring control back to a debugging environment.^[1, p. 24]

When **rvddt** executes an **ebreak** instruction, it will immediately terminate any executing *trace* or *go* command currently executing and return to the command prompt without advancing the **pc** register.

The machine language encoding shows that **ebreak** has no operands.



Listing 4.2 demonstrates that since **rvddt** does not advance the **pc** when it encounters an **ebreak** instruction, subsequent *trace* and/or *go* commands will re-execute the same **ebreak** and halt the simulation again (and again). This feature is intended to help prevent overzealous users from accidentally running past the end of a code fragment.¹

Listing 4.1: **ebreak/ebreak.S**
A one-line **ebreak** program.

```
1 .text          # put this into the text section
2 .align 2       # align to a multiple of 4
3 .globl _start
4
5 _start:
6     ebreak
```

Listing 4.2: **ebreak/ebreak.out**
ebreak stopps **rvddt** without advancing **pc**.

¹This was one of the first *enhancements* I needed for myself :-)

```

970
971 1 $ rvddt -f ebreak.bin
972 2 sp initialized to top of memory: 0x0000fff0
973 3 Loading 'ebreak.bin' to 0x0
974 4 This is rvddt. Enter ? for help.
975 5 ddt> d 0 16
976 6 00000000: 73 00 10 00 a5 a5 a5 a5-a5 a5 a5 a5 a5 a5 a5 *s.....*
977 7 ddt> r
978 8 x0 00000000 f0f0f0f0 0000fff0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
979 9 x8 f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
980 10 x16 f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
981 11 x24 f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
982 12 pc 00000000
983 13 ddt> ti 0 1000
984 14 00000000: ebreak
985 15 ddt> ti
986 16 00000000: ebreak
987 17 ddt> g 0
988 18 00000000: ebreak
989 19 ddt> r
990 20 x0 00000000 f0f0f0f0 0000fff0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
991 21 x8 f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
992 22 x16 f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
993 23 x24 f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
994 24 pc 00000000
995 25 ddt> x

```

Instruction!addi
Instruction!nop

997 4.2 Using the addi Instruction


998 The detailed description of how the `addi` instruction is executed is that it:

► Fix Me:
Define what constant and
immediate values are
somewhere.

- 999 1. Sign-extends the immediate operand.
- 1000 2. Add the sign-extended immediate operand to the contents of the `rs1` register.
- 1001 3. Store the sum in the `rd` register.
- 1002 4. Add four to the `pc` register (point to the next instruction.)

1003 In the following example `rs1 = x28`, `rd = x29` and the immediate operand is -1.

1004 `addi x29, x28, -1`



1006 Depending on the values of the fields in this instruction a number of different operations can be
 1007 performed. The most obvious is that it can add things. But it can also be used to copy registers, set
 1008 a register to zero and even, when you need to, accomplish nothing.

1009 4.2.1 No Operation

1010 It might seem odd but it is sometimes important to be able to execute an instruction that accomplishes
 1011 nothing while simply advancing the `pc` to the next instruction. One reason for this is to fill unused


objdump

memory between two instructions in a program.²

An instruction that accomplishes nothing is called a **nop** (some times systems call these **noop**). The name means *no operation*. The intent of a **nop** is to execute without having any side effects other than to advance the pc register.

The **addi** instruction can serve as a **nop** by coding it like this:

addi x0, x0, 0



The result will be to add zero to zero and discard the result (because you can never store a value into the x0 register.)

The RISC-V assembler provides a pseudoinstruction specifically for this purpose that you can use to improve the readability of your code. Note that the **addi** and **nop** instructions in Listing 4.3 are assembled into the exact same binary machine instruction (The 0x00000013 you can see are stored at addresses 0x0 and 0x4) as seen by looking at the objdump listing in Listing 4.4. In fact, you can see that objdump shows both instructions as a **nop** while Listing 4.5 shows that **rvddt** displays both as **addi x0, x0, 0**.

Listing 4.3: **nop/nop.S**

Demonstrate that an **addi** can be the same as **nop**.

```

1      .text                # put this into the text section
2      .align 2             # align to a multiple of 4
3      .globl _start
4
5      _start:
6      addi    x0, x0, 0     # these two instructions assemble into the same thing!
7      nop
8
9      ebreak

```

Listing 4.4: **nop/nop.lst**

Using **addi** to perform a **nop**

```

1      nop:      file format elf32-littleriscv
2      Disassembly of section .text:
3      00000000 <_start>:
4      0:      00000013          nop
5      4:      00000013          nop
6      8:      00100073          ebreak

```

Listing 4.5: **nop/nop.out**

Using **addi** to perform a **nop**

```

1      $ rvddt -f nop.bin
2      sp initialized to top of memory: 0x0000fff0
3      Loading 'nop.bin' to 0x0
4      This is rvddt. Enter ? for help.
5      ddt> d 0 16
6      00000000: 13 00 00 00 13 00 00 00-73 00 10 00 a5 a5 a5 a5 *.....s.....*
7      ddt> r
8      x0 00000000 f0f0f0f0 0000fff0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0

```

²This can happen during the evolution of one portion of code that reduces in size but has to continue to fit into a system without altering any other code... or some times you just need to waste a small amount of time in a device driver.

```

1055 9      x8 f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
1056 10      x16 f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
1057 11      x24 f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
1058 12      pc 00000000
1059 13 ddt> ti 0 1000
1060 14 00000000: 00000013 addi    x0, x0, 0      # x0 = 0x00000000 = 0x00000000 + 0x00000000
1061 15 00000004: 00000013 addi    x0, x0, 0      # x0 = 0x00000000 = 0x00000000 + 0x00000000
1062 16 00000008: ebreak
1063 17 ddt> r
1064 18      x0 00000000 f0f0f0f0 0000ffff f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
1065 19      x8 f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
1066 20      x16 f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
1067 21      x24 f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
1068 22      pc 00000008
1069 23 ddt> x


```

Instruction!mv

1071 4.2.2 Copying the Contents of One Register to Another

1072 By adding zero to one register and storing the sum in another register the **addi** instruction can be
 1073 used to copy the value stored in one register to another register. The following instruction will copy
 1074 the contents of **t4** into **t3**.

1075 **addi t3, t4, 0**



1076 This is a commonly required operation. To make your intent clear you may use the **mv** pseudoinstruction for this purpose.

1079 [Listing 4.6](#) shows the source of a program that is dumped in [Listing 4.7](#) illustrating that the assembler
 1080 has generated the same machine instruction (0x000e8e13 at addresses 0x0 and 0x4) for both of the
 1081 instructions.

Listing 4.6: mv/mv.S
Comparing addi to mv

```

1082 1      .text                # put this into the text section
1083 2      .align 2             # align to a multiple of 4
1084 3      .globl _start
1085 4
1086 5 _start:
1087 6     addi    t3, t4, 0      # t3 = t4
1088 7     mv     t3, t4         # t3 = t4
1089 8
1090 9     ebreak
1091

```

Listing 4.7: mv/mv.lst

An objdump of an addi and mv Instruction.

```

1093 1 mv:      file format elf32-littleriscv
1094 2 Disassembly of section .text:
1095 3 00000000 <_start>:
1096 4 0: 000e8e13      mv     t3,t4
1097 5 4: 000e8e13      mv     t3,t4
1098 6 8: 00100073      ebreak
1099


```

4.2.3 Setting a Register to Zero

Recall that `x0` always contains the value zero. Any register can be set to zero by copying the contents of `x0` using `mv` (aka `addi`).³

For example, to set `t3` to zero:

```
addi t3, x0, 0
```



Listing 4.8: `mvzero/mv.S`

Using `mv` (aka `addi`) to zero-out a register.

```

1  .text                # put this into the text section
2  .align 2            # align to a multiple of 4
3  .globl _start
4
5  _start:
6      mv      t3, x0    # t3 = 0
7
8      ebreak

```

Listing 4.9 traces the execution of the program in Listing 4.8 showing how `t3` is changed from `0xf0f0f0f0` (seen on `ℓ16`) to `0x00000000` (seen on `ℓ26`.)

Listing 4.9: `mvzero/mv.out`

Setting `t3` to zero.

```


1  $ rvddt -f mv.bin
2  sp initialized to top of memory: 0x0000fff0
3  Loading 'mv.bin' to 0x0
4  This is rvddt. Enter ? for help.
5  ddt> a
6  ddt> d 0 16
7  00000000: 13 0e 00 00 73 00 10 00-a5 a5 a5 a5 a5 a5 a5 *....s.....*
8  ddt> t 0 1000
9  zero x0 00000000 ra x1 f0f0f0f0 sp x2 0000fff0 gp x3 f0f0f0f0
10 tp x4 f0f0f0f0 t0 x5 f0f0f0f0 t1 x6 f0f0f0f0 t2 x7 f0f0f0f0
11 s0 x8 f0f0f0f0 s1 x9 f0f0f0f0 a0 x10 f0f0f0f0 a1 x11 f0f0f0f0
12 a2 x12 f0f0f0f0 a3 x13 f0f0f0f0 a4 x14 f0f0f0f0 a5 x15 f0f0f0f0
13 a6 x16 f0f0f0f0 a7 x17 f0f0f0f0 s2 x18 f0f0f0f0 s3 x19 f0f0f0f0
14 s4 x20 f0f0f0f0 s5 x21 f0f0f0f0 s6 x22 f0f0f0f0 s7 x23 f0f0f0f0
15 s8 x24 f0f0f0f0 s9 x25 f0f0f0f0 s10 x26 f0f0f0f0 S11 x27 f0f0f0f0
16 t3 x28 f0f0f0f0 t4 x29 f0f0f0f0 t5 x30 f0f0f0f0 t6 x31 f0f0f0f0
17 pc 00000000
18 00000000: 00000e13 addi t3, zero, 0 # t3 = 0x00000000 = 0x00000000 + 0x00000000
19 zero x0 00000000 ra x1 f0f0f0f0 sp x2 0000fff0 gp x3 f0f0f0f0
20 tp x4 f0f0f0f0 t0 x5 f0f0f0f0 t1 x6 f0f0f0f0 t2 x7 f0f0f0f0
21 s0 x8 f0f0f0f0 s1 x9 f0f0f0f0 a0 x10 f0f0f0f0 a1 x11 f0f0f0f0
22 a2 x12 f0f0f0f0 a3 x13 f0f0f0f0 a4 x14 f0f0f0f0 a5 x15 f0f0f0f0
23 a6 x16 f0f0f0f0 a7 x17 f0f0f0f0 s2 x18 f0f0f0f0 s3 x19 f0f0f0f0
24 s4 x20 f0f0f0f0 s5 x21 f0f0f0f0 s6 x22 f0f0f0f0 s7 x23 f0f0f0f0
25 s8 x24 f0f0f0f0 s9 x25 f0f0f0f0 s10 x26 f0f0f0f0 S11 x27 f0f0f0f0
26 t3 x28 00000000 t4 x29 f0f0f0f0 t5 x30 f0f0f0f0 t6 x31 f0f0f0f0
27 pc 00000004
28 00000004: ebreak
29 ddt> x

```

³There are other pseudoinstructions (such as `li`) that can also turn into an `addi` instruction. Objdump might display `'addi t3,x0,0'` as `'mv t3,x0'` or `'li t3,0'`.

4.2.4 Adding a 12-bit Signed Value

addi x1, x7, 4



```

addi    t0, zero, 4      # t0 = 4
addi    t1, t1, 100      # t1 = 104

addi    t0, zero, 0x123   # t0 = 0x123
addi    t0, t0, 0xffff    # t0 = 0x122 (subtract 1)

addi    t0, zero, 0xffff  # t0 = 0xffffffff (-1) (diagram out the chaining carry)
                                # refer back to the overflow/truncation discussion in binary chapter

addi x0, x0, 0 # no operation (pseudo: nop)
addi rd, rs, 0 # copy reg rs to rd (pseudo: mv rd, rs)

```

Demonstrate various addi instructions.

4.3 todo

Ideas for the order of introducing instructions.

4.4 Other Instructions With Immediate Operands

```

addi
andi
ori
xori

slti
sltiu
srai
slli
srli

```

4.5 Transferring Data Between Registers and Memory

RV is a load-store architecture. This means that the only way that the CPU can interact with the memory is via the *load* and *store* instructions. All other data manipulation must be performed on register values.

Copying values from memory to a register (first examples using regs set with addi):

```

1183     lb
1184     lh
1185     lw
1186     lbu
1187     lhu

```

1188 Copying values from a register to memory:

```

1189     sb
1190     sh
1191     sw

```

1192

► Fix Me:

*Mention the rvddt UART
I/O address for writing to
the console here?*

1193 4.6 RR operations

```

1194     add
1195     sub
1196     and
1197     or
1198     sra
1199     srl
1200     sll
1201     xor
1202     sltu
1203     slt

```

1204 4.7 Setting registers to large values using lui with addi

```

1205     addi    // useful for values from -2048 to 2047
1206     lui     // useful for loading any multiple of 0x1000
1207
1208     Setting a register to any other value must be done using a combo of insns:
1209
1210     auipc    // Load an address relative the the current PC (see la pseudo)
1211     addi
1212
1213
1214     lui     // Load constant into into bits 31:12 (see li pseudo)
1215     addi    // add a constant to fill in bits 11:0
1216             if bit 11 is set then need to +1 the lui value to compensate

```

1217 4.8 Labels and Branching

1218 Start to introduce addressing here?

```

1219     beq
1220     bne
1221     blt
1222     bge
1223     bltu
1224     bgeu
1225
1226     bgt rs, rt, offset    # pseudo for: blt rt, rs, offset    (reverse the operands)
1227     ble rs, rt, offset    # pseudo for: bge rt, rs, offset    (reverse the operands)
1228     bgtu rs, rt, offset   # pseudo for: bltu rt, rs, offset   (reverse the operands)
1229     bleu rs, rt, offset   # pseudo for: bgeu rt, rs, offset   (reverse the operands)
1230
1231     beqz beqz rs, offset  # pseudo for: beq rs, x0, offset
1232     bnez rs, offset       # pseudo for: bne rs, x0, offset
1233     blez rs, offset       # pseudo for: bge x0, rs, offset
1234     bgez rs, offset       # pseudo for: bge rs, x0, offset
1235     bltz rs, offset       # pseudo for: blt rs, x0, offset
1236     bgtz rs, offset       # pseudo for: blt x0, rs, offset

```

1237 4.9 Relocation

```

1238 Absolute:
1239     %hi(symbol)
1240     %lo(symbol)
1241
1242 PC-relative:
1243     %pcrel_hi(symbol)
1244     %pcrel_lo(label)
1245
1246 Using the auipc & addi pair with label references:
1247     The %pcrel_lo() uses the label to find the associated %pcrel_hi()
1248     The label MUST be on a line that used a %pcrel_hi() or get an error.
1249     This is needed to calculate the proper offset.
1250     Things like this are legal (though not sure of the value):
1251     label: auipc    t1, %pcrel_hi(symbol)
1252            addi     t2, t1, %pcrel_lo(label)
1253            addi     t3, t1, %pcrel_lo(label)
1254            lw       t4, %pcrel_lo(label)(t1)
1255            sw       t5, %pcrel_lo(label)(t1)
1256
1257 Discuss how relaxation works.
1258 see: https://github.com/riscv/riscv-elf-psabi-doc/blob/master/riscv-elf.md

```

1259 4.10 Jumps

1260 Introduce and present subroutines but not nesting until introduce stack operations.

```

1261     jal
1262     jalr

```


4.11 Pseudo Operations

► Fix Me:

Explain why we have pseudo ops. These mappings are lifted from the ISM, Vol 1, V2.2

```

1263
1264
1265     la rd, symbol
1266             auipc rd, symbol[31:12]
1267             addi rd, rd, symbol[11:0]
1268
1269     l{b|h|w|d} rd, symbol
1270             auipc rd, symbol[31:12]
1271             l{b|h|w|d} rd, symbol[11:0](rd)
1272
1273     s{b|h|w|d} rd, symbol, rt           # rt is the temp reg to use for the operation
1274             auipc rt, symbol[31:12]
1275             s{b|h|w|d} rd, symbol[11:0](rt)
1276
1277
1278     j offset      jal x0, offset
1279     jal offset    jal x1, offset
1280     jr rs         jalr x0, rs, 0
1281     jalr rs       jalr x1, rs, 0
1282     ret           jalr x0, x1, 0
1283
1284     call offset   auipc x6, offset[31:12]
1285                 jalr x1, x6, offset[11:0]
1286
1287     tail offset   auipc x6, offset[31:12]   # same as call but no x1
1288                 jalr x0, x6, offset[11:0]

```

4.12 The Linker and Relaxation

I don't know where this should go just yet.

► Fix Me:

Needs research. I'm not sure if/how the linker alone can relax the AUIPC+JALR pair since the assembler could have used a pcrel branch across one of these pairs.

4.13 pic and nopic

pic is *needed* for shared libs. Should discuss it but probably best to leave the topic for a later chapter.

1293

Chapter 5

1294

RV32 Machine Instructions

1295

5.1 Introduction

1296

5.2 Conventions and Terminology

1297

When discussing instructions, the following abbreviations/notations are used:

1298

5.2.1 XLEN

1299

XLEN represents the bit-length of an `x` register in the machine architecture. Possible values are 32, 64 and 128.

1300

1301

5.2.2 `sx(val)`

1302

Sign extend *val* to the left.

1303

This is used to convert a signed integer value expressed using some number of bits to a larger number of bits by adding more bits to the left. In doing so, the sign will be preserved. In this case *val* represents the least [MSBs](#) of the value.

1304

1305

1306

For more on sign-extension see [section 2.3](#).

1307

5.2.3 `zx(val)`

1308

Zero extend *val* to the left.

1309

This is used to convert an unsigned integer value expressed using some number of bits to a larger number of bits by adding more bits to the left. In doing so, the new bits added will all be set to zero. As is the case with `sx(val)`, *val* represents the [LSBs](#) of the final value.

1310

1311

1312

For more on zero-extension see [Figure 2.3](#).

5.2.4 `zr(val)`

Zero extend *val* to the right.

Some times a binary value is encoded such that a set of bits represented by *val* are used to represent the **MSBs** of some longer (more bits) value. In this case it is necessary to append zeros to the right to convert *val* to the longer value.

Figure 5.1 illustrates converting a 20-bit *val* to a 32-bit fullword.




Figure 5.1: Zero-extending an integer to the right from 20 bits to 32 bits.

5.2.5 Sign Extended Left and Zero Extend Right

Some instructions such as the J-type (see subsection 5.4.2) include immediate operands that are extended in both directions.

Figure 5.2 and Figure 5.3 illustrates zero-extending a 20-bit negative number one bit to the right and sign-extending it 11 bits to the left:




Figure 5.2: Sign-extending a positive 20-bit number 11 bits to the left and one bit to the right.




Figure 5.3: Sign-extending a negative 20-bit number 11 bits to the left and one bit to the right.

5.2.6 `m8(addr)`

The contents of an 8-bit value in memory at address *addr*.

Given the contents of the memory dump shown in Figure 5.4, `m8(42)` refers to the memory location at address 42_{16} that currently contains the 8-bit value fc_{16} .

The `mn(addr)` notation can be used to refer to memory that is being read or written depending on the context.

1330 When memory is being written, the following notation is used to indicate that the least significant 8
1331 bis of *source* will be written into memory at the address *addr*:

1332 $\text{m8(addr)} \leftarrow \text{source}$

1333 When memory is being read, the following notation is used to indicate that the 8 bit value at the
1334 address *addr* will be read and stored into *dest*:

1335 $\text{dest} \leftarrow \text{m8(addr)}$

1336 Note that *source* and *dest* are typically registers.

00000030	2f 20 72 65 61 64 20 61	20 62 69 6e 61 72 79 20
00000040	66 69 fc 65 20 66 69 6c	6c 65 64 20 77 69 74 68
00000050	20 72 76 33 32 49 20 69	6e 73 74 72 75 63 74 69
00000060	6f 6e 73 20 61 6e 64 20	66 65 65 64 20 74 68 65

Figure 5.4: Sample memory contents.

1337 5.2.7 m16(addr)

1338 The contents of an 16-bit little-endian value in memory at address *addr*.

1339 Given the contents of the memory dump shown in [Figure 5.4](#), $\text{m16}(42)$ refers to the memory location
1340 at address 42_{16} that currently contains $65fc_{16}$. See also [subsection 5.2.6](#).

1341 5.2.8 m32(addr)

1342 The contents of an 32-bit little-endian value in memory at address *addr*.

1343 Given the contents of the memory dump shown in [Figure 5.4](#), $\text{m32}(42)$ refers to the memory location
1344 at address 42_{16} that currently contains $662065fc_{16}$. See also [subsection 5.2.6](#).

1345 5.2.9 m64(addr)

1346 The contents of an 64-bit little-endian value in memory at address *addr*.

1347 Given the contents of the memory dump shown in [Figure 5.4](#), $\text{m64}(42)$ refers to the memory location
1348 at address 42_{16} that currently contains $656c6c69662065fc_{16}$. See also [subsection 5.2.6](#).

1349 5.2.10 m128(addr)

1350 The contents of an 128-bit little-endian value in memory at address *addr*.

1351 Given the contents of the memory dump shown in [Figure 5.4](#), $\text{m128}(42)$ refers to the memory location
1352 at address 42_{16} that currently contains $7220687469772064656c6c69662065fc_{16}$. See also [subsec-](#)
1353 [tion 5.2.6](#).

1354 5.2.11 .+offset

1355 The address of the current instruction plus a numeric offset.

1356 5.2.12 .-offset

1357 The address of the current instruction minus a numeric offset.

1358 5.2.13 pc

1359 The current value of the program counter.

1360 5.2.14 rd

1361 An x-register used to store the result of instruction.

1362 5.2.15 rs1

1363 An x-register value used as a source operand for an instruction.

1364 5.2.16 rs2

1365 An x-register value used as a source operand for an instruction.

1366 5.2.17 imm

1367 An immediate numeric operand. The word *immediate* refers to the fact that the operand is stored
1368 within an instruction.

1369 5.2.18 rsN[h:l]

1370 The value of bits from *h* through *l* of x-register rsN. For example: rs1[15:0] refers to the contents of
1371 the 16 [LSBs](#) of rs1.

1372 5.3 Addressing Modes

1373 immediate, register, base-displacement, pc-relative

► Fix Me:
Write this section.

5.4 Instruction Encoding Formats

This document concerns itself with the following RISC-V instruction formats.

XXX Show and discuss a stack of formats explaining how the unnatural ordering of the *imm* fields reduces the number of possible locations that the hardware has to be prepared to *look* for various bits. For example, the opcode, rd, rs1, rs1, funct3 and the sign bit (when used) are all always in the same position. Also note that imm[19:12] and imm[10:5] can only be found in one place. imm[4:0] can only be found in one of two places...

The point to all this is that it is easier to build a machine if it does not have to accommodate many different ways to perform the same task. This simplification can also allow it operate faster.

Figure 5.5 Shows the RISC-V instruction formats.

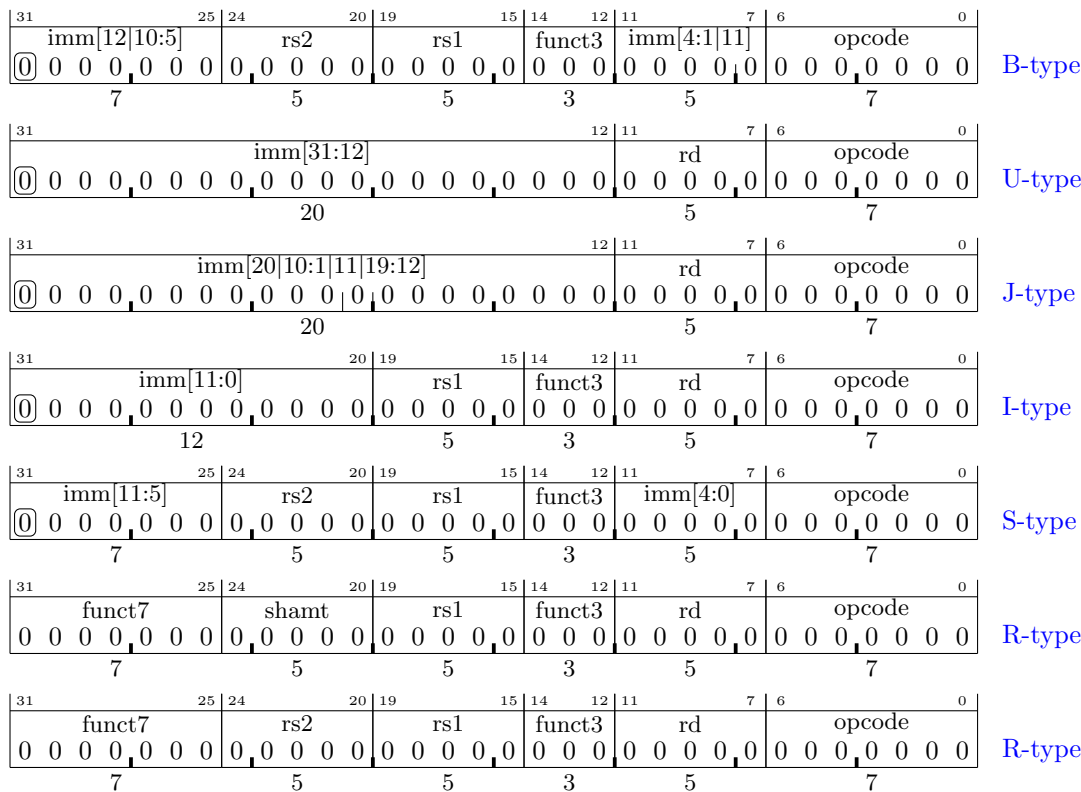



Figure 5.5: RISC-V instruction formats.

5.4.1 U Type


The U-Type format is used for instructions that use a 20-bit immediate operand and a destination register.



The **rd** field contains an **x** register number to be set to a value that depends on the instruction.

The **imm** field contains a 20-bit value that will be converted into **XLEN** bits by using the *imm* operand for bits 31:12 and then sign-extending it to the left¹ and zero-extending the LSBs as discussed in subsection 5.2.4.

If **XLEN**=32 then the **imm** value in this example will be converted as shown below.




Notice that the 20-bits of the **imm** field are mapped in the same order and in the same relative position that they appear in the instruction when they are used to create the value of the immediate operand. Shifting the **imm** value to the left, into the “upper bits” of the immediate value suggests a rationale for the name of this format.

If **XLEN**=64 then the **imm** value in this example will be converted to the same two’s complement integer value by extending the sign to the left.

5.4.2 J Type

The J-type format is used for instructions that use a 20-bit immediate operand and a destination register. It is similar to the U-type. However, the immediate operand is constructed by arranging the *imm* bits in a different manner.



The **rd** field contains an **x** register number to be set to a value that depends on the instruction.

In the J-type format the 20 *imm* bits are arranged such that they represent the “lower” portion of the immediate value. Unlike the U-type instructions, the J-type requires the bits to be re-ordered and shifted to the right before they are used.²


The example above shows that the bit positions in the *imm* field description. We see that the 20 *imm* bits are re-ordered according to: [20|10:1|11|19:12]. This means that the **MSB** of the *imm* field is to be placed into bit 20 of the immediate integer value ultimately used by the instruction when it is converted into **XLEN** bits. The next bit to the right in the *imm* field is to be placed into bit 10 of the immediate value and so on.

After the *imm* bits are re-positioned into bits 20:1 of the immediate value being constructed, a zero-bit will be added to the **LSB** and the value in bit-position 20 will be replicated to sign-extend the value to **XLEN** bits as discussed in subsection 5.2.5.


If **XLEN**=32 then the *imm* value in this example will be converted as shown below.

¹When **XLEN** is larger than 32.


²The reason that the J-type bits are reordered like this is because it simplifies the implementation of hardware as discussed in section 5.4.



A J-type example with a negative imm field:




If $XLEN=32$ then the *imm* field in this example will be converted as shown below.




The J-type format is used by the Jump And Link instruction that calculates a target address by adding a signed immediate value to the current program counter. Since no instruction can be placed at an odd address the 20-bit imm value is zero-extended to the right to represent a 21-bit signed offset capable of representing numbers twice the magnitude of the 20-bit imm value.


5.4.3 R Type




A special case of the R-type used for shift-immediate instructions where the *rs2* field is used as an immediate value named *shamt* representing the number of bit positions to shift:




5.4.4 I Type



5.4.5 S Type



5.4.6 B Type



5.4.7 CPU Registers

The registers are names x0 through x31 and have aliases suited to their conventional use. The following table describes each register.

Note that the calling convention specifies that only some of the registers are to be saved by functions if they alter their contents. The idea being that accessing memory is time-consuming and that by classifying some registers as “temporary” (not saved by any function that alter its contents) it is possible to carefully implement a function with less need to store register values on the stack in order to use them to perform the operations of the function.

The lack of grouping the temporary and saved registers is due to the fact that the C extension provides access to only the first 16 registers when executing instructions in the compressed format.

► Fix Me:
Need to add a section that discusses the calling conventions

Reg	Alias	Description	Saved
x0	zero	Hard-wired zero	
x1	ra	Return address	
x2	sp	Stack pointer	yes
x3	gp	Global pointer	
x4	tp	Thread pointer	
x5	t0	Temporary/alternate link register	
x6	t1	Temporary	
x7	t2	Temporary	
x8	s0/fp	Saved register/frame pointer	yes
x9	s1	Saved register	yes
x10	a0	Function argument/return value	
x11	a1	Function argument/return value	
x12	a2	Function argument	
x13	a3	Function argument	
x14	a4	Function argument	
x15	a5	Function argument	
x16	a6	Function argument	
x17	a7	Function argument	
x18	s2	Saved register	yes
x19	s3	Saved register	yes
x20	s4	Saved register	yes
x21	s5	Saved register	yes
x22	s6	Saved register	yes
x23	s7	Saved register	yes
x24	s8	Saved register	yes
x25	s9	Saved register	yes
x26	s10	Saved register	yes
x27	s11	Saved register	yes
x28	t3	Temporary	
x29	t4	Temporary	
x30	t5	Temporary	
x31	t6	Temporary	

5.5 memory

Note that RISC-V is a little-endian machine.

All instructions must be naturally aligned to their 4-byte boundaries. [1, p. 5]

If a RISC-V processor implements the C (compressed) extension then instructions may be aligned to 2-byte boundaries. [1, p. 68]

Data alignment is not necessary but unaligned data can be inefficient. Accessing unaligned data using any of the load or store instructions can also prevent a memory access from operating atomically. [1, p.19] See also ??.

5.6 RV32I Base Instruction Set

RV32I refers to the basic 32-bit integer instructions.

Instruction!LUI
Instruction!AUIPC**5.6.1 LUI rd, imm**


Load Upper Immediate.

 $rd \leftarrow zr(imm)$


Copy the immediate value into bits 31:12 of the destination register and place zeros into bits 11:0.
When XLEN is 64 or 128, the immediate value is sign-extended to the left.

Instruction Format and Example:

LUI t0, 3



Instruction!JAL




```

1489
1490 0001007c: 00003297 auipc   x5, 0x3          // x5 = 0x1307c = 0x1007c + 0x3000
1491 reg 0: 00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 0001307c f0f0f0f0 f0f0f0f0
1492 reg 8: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
1493 reg 16: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
1494 reg 24: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
1495 pc: 00010080

```

AUIPC t0, 0x81000



```

1496
1498 00010080: 81000297 auipc   x5, 0x81000          // x5 = 0x81010080 = 0x10080 + 0x81000000
1499 reg 0: 00000000 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 81010080 f0f0f0f0 f0f0f0f0
1500 reg 8: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
1501 reg 16: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
1502 reg 24: f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0-f0f0f0f0 f0f0f0f0 f0f0f0f0 f0f0f0f0
1503 pc: 00010084

```

The AUIPC instruction supports two-instruction sequences to access arbitrary offsets from the PC for both control-flow transfers and data accesses. The combination of an AUIPC and the 12-bit immediate in a JALR can transfer control to any 32-bit PC-relative address, while an AUIPC plus the 12-bit immediate offset in regular load or store instructions can access any 32-bit PC-relative data address. [1, p. 14]

5.6.3 JAL rd, imm

Jump and link.

```

1511 rd ← pc + 4
1512 pc ← pc + sx(imm<<1)


```

This instruction saves the address of the next instruction that would otherwise execute (located at pc+4) into rd and then adds immediate value to the pc causing an unconditional branch to take place.

The standard software conventions for calling subroutines use x1 as the return address (rd register in this case). [1, p. 16]

Encoding:

JAL x7, .+16



imm demultiplexed value = 00000000000000001000₂ << 1 = 16₁₀

Instruction!JALR

State of registers before execution:


pc = 0x11114444

State of registers after execution:

pc = 0x11114454 x7 = 0x11114448

JAL provides a method to call a subroutine using a pc-relative address.

JAL x7, -16

imm demultiplexed value = $1111111111111111000_2 \ll 1 = -16_{10}$

State of registers before execution:

pc = 0x11114444

State of registers after execution:

pc = 0x11114434 x7 = 0x11114448

5.6.4 JALR rd, rs1, imm

Jump and link register.

 $rd \leftarrow pc + 4$ $pc \leftarrow (rs1 + sx(imm)) \& \sim 1$


This instruction saves the address of the next instruction that would otherwise execute (located at pc+4) into rd and then adds the immediate value to the rs1 register and stores the sum into the pc register causing an unconditional branch to take place.

Note that the branch target address is calculated by sign-extending the imm[11:0] bits from the instruction, adding it to the rs1 register and *then* the LSB of the sum is to zero and the result is stored into the pc register. The discarding of the LSB allows the branch to refer to any even address.

The standard software conventions for calling subroutines use x1 as the return address (rd register in this case). [1, p. 16]

Encoding:

JALR x1, x7, 4



Before:

pc = 0x11114444

x7 = 0x44444444

Instruction!BEQ
Instruction!BNE


After

pc = 0x5555888c

x1 = 0x11114448

JALR provides a method to call a subroutine using a base-displacement address.

JALR x1, x0, 5



Note that the least significant bit in the result of rs1+imm is discarded/set to zero before the result is saved in the pc.

pc = 0x11114444

After

pc = 0x00000004

x1 = 0x11114448


5.6.5 BEQ rs1, rs2, imm

Branch if equal.

$pc \leftarrow (rs1 == rs2) ? pc + sx(imm[12:1] \ll 1) : pc + 4$

Encoding:

BEQ x3, x15, 2064



$imm[12:1] = 010000001000_2 = 1032_{10}$

$imm = 2064_{10}$

$funct3 = 000_2$

rs1 = x3

rs2 = x15


5.6.6 BNE rs1, rs2, imm

Branch if Not Equal.

$pc \leftarrow (rs1 != rs2) ? pc + sx(imm[12:1] \ll 1) : pc + 4$

Encoding:

BNE x3, x15, 2064

Instruction!BLT
Instruction!BGEimm[12:1] = 010000001000₂ = 1032₁₀imm = 2064₁₀funct3 = 001₂

rs1 = x3

rs2 = x15


5.6.7 BLT rs1, rs2, imm

Branch if Less Than.

 $pc \leftarrow (rs1 < rs2) ? pc + sx(imm[12:1] \ll 1) : pc + 4$

Encoding:

BLT x3, x15, 2064

imm[12:1] = 010000001000₂ = 1032₁₀imm = 2064₁₀funct3 = 100₂

rs1 = x3

rs2 = x15


5.6.8 BGE rs1, rs2, imm

Branch if Greater or Equal.

 $pc \leftarrow (rs1 \geq rs2) ? pc + sx(imm[12:1] \ll 1) : pc + 4$

Encoding:

BGE x3, x15, 2064

imm[12:1] = 010000001000₂ = 1032₁₀imm = 2064₁₀funct3 = 101₂

rs1 = x3

rs2 = x15

Instruction!BLTU
Instruction!BGEU
Instruction!LB

5.6.9 BLTU rs1, rs2, imm

Branch if Less Than Unsigned.

$pc \leftarrow (rs1 < rs2) ? pc + sx(imm[12:1] \ll 1) : pc + 4$

Encoding:

BLTU x3, x15, 2064

31	25	24	20	19	15	14	12	11	7	6	0	
imm[12:10:5]							rs2			rs1		
0 0 0 0 0 0 0							0 1 1 1 1			0 0 0 1 1		
7							5			5		
							funct3			imm[4:1 11]		
							1 1 0			1 0 0 0 1		
							3			5		
										opcode		
										1 1 0 0 0 1 1		
										7		

B-type

$imm[12:1] = 010000001000_2 = 1032_{10}$

$imm = 2064_{10}$

$funct3 = 110_2$

$rs1 = x3$

$rs2 = x15$

5.6.10 BGEU rs1, rs2, imm

Branch if Greater or Equal Unsigned.

$pc \leftarrow (rs1 \geq rs2) ? pc + sx(imm[12:1] \ll 1) : pc + 4$

Encoding:

BGEU x3, x15, 2064

31	25	24	20	19	15	14	12	11	7	6	0	
imm[12:10:5]							rs2			rs1		
0 0 0 0 0 0 0							0 1 1 1 1			0 0 0 1 1		
7							5			5		
							funct3			imm[4:1 11]		
							1 1 1			1 0 0 0 1		
							3			5		
										opcode		
										1 1 0 0 0 1 1		
										7		

B-type

$imm[12:1] = 010000001000_2 = 1032_{10}$

$imm = 2064_{10}$

$funct3 = 111_2$

$rs1 = x3$

$rs2 = x15$

► Fix Me:

use symbols in branch
examples

5.6.11 LB rd, imm(rs1)

Load byte.

$rd \leftarrow sx(m8(rs1 + sx(imm)))$


$pc \leftarrow pc + 4$

Load an 8-bit value from memory at address $rs1 + imm$, then sign-extend it to 32 bits before storing it in rd

Encoding:

LB x7, 4(x3)

Instruction!LH
Instruction!LW
Instruction!LBU



5.6.12 LH rd, imm(rs1)

Load halfword.


$rd \leftarrow sx(m16(rs1+sx(imm)))$

$pc \leftarrow pc+4$

Load a 16-bit value from memory at address $rs1+imm$, then sign-extend it to 32 bits before storing it in rd

Encoding:

LH x7, 4(x3)



5.6.13 LW rd, imm(rs1)

Load word.


$rd \leftarrow sx(m32(rs1+sx(imm)))$

$pc \leftarrow pc+4$

Load a 32-bit value from memory at address $rs1+imm$, then store it in rd

Encoding:

LW x7, 4(x3)



5.6.14 LBU rd, imm(rs1)

Load byte unsigned.


$rd \leftarrow zx(m8(rs1+sx(imm)))$

$pc \leftarrow pc+4$

Load an 8-bit value from memory at address $rs1+imm$, then zero-extend it to 32 bits before storing it in rd

Encoding:

Instruction!LHU
Instruction!SB
Instruction!SH



5.6.15 LHU rd, imm(rs1)

Load halfword unsigned.


```
rd ← zx(m16(rs1+sx(imm)))
```

$$pc \leftarrow pc+4$$

Load an 16-bit value from memory at address **rs1+imm**, then zero-extend it to 32 bits before storing it in **rd**

Encoding:

LHU x7, 4(x3)



5.6.16 SB rs2, imm(rs1)


Store Byte.

$$m8(rs1+sx(imm)) \leftarrow rs2[7:0]$$
$$pc \leftarrow pc+4$$

Store the 8-bit value in `rs2[7:0]` into memory at address `rs1+imm`.

Encoding:

SB x3, 19(x15)



5.6.17 SH rs2, imm(rs1)

Store Halfword.


$$\text{m16}(\text{rs1} + \text{sx}(\text{imm})) \leftarrow \text{rs2}[15:0]$$
$$pc \leftarrow pc+4$$

Store the 16-bit value in `rs2[15:0]` into memory at address `rs1+imm`.

Encoding:

Instruction!SW
Instruction!ADDI
Instruction!SLTI

SH x3, 19(x15)



5.6.18 SW rs2, imm(rs1)


Store Word

$m16(rs1+sx(imm)) \leftarrow rs2[31:0]$
 $pc \leftarrow pc+4$

Store the 32-bit value in `rs2` into memory at address `rs1+imm`.

Encoding:

SW x3, 19(x15)



Show pos & neg imm examples.


5.6.19 ADDI rd, rs1, imm

Add Immediate

$rd \leftarrow rs1+sx(imm)$
 $pc \leftarrow pc+4$

Encoding:

ADDI x1, x7, 4



Before:

$x7 = 0x11111111$

After:

$x1 = 0x11111115$


5.6.20 SLTI rd, rs1, imm

Set LessThan Immediate

The logical AND of the sign-extended immediate value and the value in the **rs1** register is stored in the **rd** register. Instruction!SLLI
Instruction!SRLI

Encoding:

ANDI x1, x7, 4



Before:

x7 = 0x81111111

After:

x1 = 0x81111115

5.6.25 SLLI rd, rs1, shamt

Shift Left Logical Immediate


$rd \leftarrow rs1 \ll shamt$

$pc \leftarrow pc+4$

SLLI is a logical left shift operation (zeros are shifted into the lower bits). The value in rs1 shifted left shamt number of bits and the result placed into rd. [1, p. 14]

Encoding:

SLLI x7, x3, 2



x3 = 0x81111111

After:

x7 = 0x04444444

5.6.26 SRLI rd, rs1, shamt

Shift Right Logical Immediate

$rd \leftarrow rs1 \gg shamt$

$pc \leftarrow pc+4$

SRLI is a logical right shift operation (zeros are shifted into the higher bits). The value in rs1 shifted right shamt number of bits and the result placed into rd. [1, p. 14]

Encoding:

Instruction!SRAI
Instruction!ADD

SRLI x7, x3, 2

312524201915141211760																						
funct7							shamt			rs1			funct3			rd			opcode			
0000000							00010			00011			101			00111			0010011			
7							5			5			3			5			7			

R-type

x3 = 0x81111111

After:

x7 = 0x20444444

5.6.27 SRAI rd, rs1, shamt

Shift Right Arithmetic Immediate

$rd \leftarrow rs1 \gg shamt$

$pc \leftarrow pc+4$

SRAI is a logical right shift operation (zeros are shifted into the higher bits). The value in rs1 shifted right shamt number of bits and the result placed into rd. [1, p. 14]

Encoding:

SRAI x7, x3, 2

25							24	20				19	15				14	12		11	7	6	0							
funct7							shamt					rs1				funct3			rd					opcode						
0	1	0	0	0	0	0	0	0	0	1	0	0	0	1	1	1	0	1	0	0	1	1	1	0	0	1	0	0	1	1
7							5					5				3			5					7						

R-type

x3 = 0x81111111

After:

x7 = 0xe0444444

5.6.28 ADD rd, rs1, rs2

Add

$rd \leftarrow rs1 + rs2$

$pc \leftarrow pc+4$

ADD performs addition. Overflows are ignored and the low 32 bits of the result are written to rd. [1, p. 15]

Encoding:

ADD x7, x3, x31

312524201915141211760																											
funct7							rs2					rs1				funct3			rd				opcode				
0000000							11111					00011				0000			00111				0110011				
7							5					5				3			5				7				

R-type

Instruction!SUB
Instruction!SLL

x3 = 0x81111111 x31 = 0x22222222

After:

x7 = 0xa3333333

5.6.29 SUB rd, rs1, rs2

Subtract

$rd \leftarrow rs1 - rs2$

$pc \leftarrow pc+4$

SUB performs subtraction. Underflows are ignored and the low 32 bits of the result are written to rd. [1, p. 15]

Encoding:

SUB x7, x3, x31

312524201915141211760																								
funct7							rs2				rs1				funct3		rd				opcode			
0100000							11111				00011				0000		00111				0110011			
7							5				5				3		5				7			

R-type

x3 = 0x83333333 x31 = 0x01111111

After:

x7 = 0x82222222

5.6.30 SLL rd, rs1, rs2

Shift Left Logical

$rd \leftarrow rs1 \ll rs2$

$pc \leftarrow pc+4$

SLL performs a logical left shift on the value in register rs1 by the shift amount held in the lower 5 bits of register rs2. [1, p. 15]

Encoding:

SLL x7, x3, x31

312524201915141211760																											
funct7							rs2					rs1				funct3			rd				opcode				
0000000							11111					00011				001			00111				0110011				
7							5					5				3			5				7				

R-type

x3 = 0x83333333

x31 = 0x00000002

After:

x7 = 0x0cccccc

Instruction!SLT
Instruction!SLTU
Instruction!XOR

5.6.31 SLT rd, rs1, rs2

Set Less Than


$rd \leftarrow (rs1 < rs2) ? 1 : 0$

$pc \leftarrow pc+4$

SLT performs a signed compare, writing 1 to **rd** if **rs1** < **rs2**, 0 otherwise. [1, p. 15]

Encoding:

SLT x7, x3, x31



x3 = 0x83333333

x31 = 0x00000002

After:

x7 = 0x00000001

5.6.32 SLTU rd, rs1, rs2

Set Less Than Unsigned


$rd \leftarrow (rs1 < rs2) ? 1 : 0$

$pc \leftarrow pc+4$

SLTU performs an unsigned compare, writing 1 to **rd** if **rs1** < **rs2**, 0 otherwise. Note, SLTU rd, x0, rs2 sets **rd** to 1 if **rs2** is not equal to zero, otherwise sets **rd** to zero (assembler pseudo-op SNEZ **rd**, **rs**). [1, p. 15]

Encoding:

SLTU x7, x3, x31



x3 = 0x83333333

x31 = 0x00000002

After:

x7 = 0x00000000

5.6.33 XOR rd, rs1, rs2


Exclusive Or

1876 $rd \leftarrow rs1 \wedge rs2$ Instruction!SRL
 1877 $pc \leftarrow pc+4$ Instruction!SRA

1878 XOR performs a bit-wise exclusive or on rs1 and rs2. The result is stored on rd.

1879 Encoding:

1880 XOR x7, x3, x31



1882 $x3 = 0x83333333$

1883 $x31 = 0x1888ffff$

1884 After:

1885 $x7 = 0x9bbbcccc$

1886 5.6.34 SRL rd, rs1, rs2

1887 Shift Right Logical


1888 $rd \leftarrow rs1 \gg rs2$

1889 $pc \leftarrow pc+4$

1890 SRL performs a logical right shift on the value in register rs1 by the shift amount held in the lower 5
 1891 bits of register rs2. [1, p. 15]

1892 Encoding:

1893 SRL x7, x3, x31



1895 $x3 = 0x83333333$

1896 $x31 = 0x00000010$

1897 After:

1898 $x7 = 0x00008333$

1899 5.6.35 SRA rd, rs1, rs2

1900 Shift Right Arithmetic

1901 $rd \leftarrow rs1 \ggg rs2$

1902 $pc \leftarrow pc+4$

1903 SRA performs an arithmetic right shift (the original sign bit is copied into the vacated upper bits) on
 1904 the value in register rs1 by the shift amount held in the lower 5 bits of register rs2. [1, p. 14, 15]

Instruction!OR
Instruction!AND

Encoding:

SLA x7, x3, x31

312524201915141211760																									
funct7							rs2				rs1				funct3			rd				opcode			
0100000							11111				00011				101			00111				0110011			
7							5				5				3			5				7			

R-type

x3 = 0x83333333

x31 = 0x00000010

After:

x7 = 0xffff8333

5.6.36 OR rd, rs1, rs2

Or

$rd \leftarrow rs1 \mid rs2$

$pc \leftarrow pc+4$

OR is a logical operation that performs a bit-wise OR on register rs1 and rs2 and then places the result in rd. [1, p. 14]

Encoding:

OR x7, x3, x31

312524201915141211760																									
funct7							rs2				rs1				funct3			rd				opcode			
0000000							11111				00011				1001			000111				0110011			
7							5				5				3			5				7			

R-type

x3 = 0x83333333

x31 = 0x00000440

After:

x7 = 0x83333773

5.6.37 AND rd, rs1, rs2

And

$rd \leftarrow rs1 \& rs2$


$pc \leftarrow pc+4$

AND is a logical operation that performs a bit-wise AND on register rs1 and rs2 and then places the result in rd. [1, p. 14]

Encoding:

AND x7, x3, x31

Instruction!FENCE
Instruction!FENCE.I



x3 = 0x83333333

x31 = 0x00000fe2

After:

x7 = 0x00000322

5.6.38 FENCE predecessor, successor

The FENCE instruction is used to order device I/O and memory accesses as viewed by other RISC-V harts and external devices or co-processors. Any combination of device input (I), device output (O), memory reads (R), and memory writes (W) may be ordered with respect to any combination of the same. Informally, no other RISC-V hart or external device can observe any operation in the successor set following a FENCE before any operation in the predecessor set preceding the FENCE. The execution environment will define what I/O operations are possible, and in particular, which load and store instructions might be treated and ordered as device input and device output operations respectively rather than memory reads and writes. For example, memory-mapped I/O devices will typically be accessed with uncached loads and stores that are ordered using the I and O bits rather than the R and W bits. Instruction-set extensions might also describe new coprocessor I/O instructions that will also be ordered using the I and O bits in a FENCE. [1, p. 21]

► Fix Me:


Which of the i, o, r and w goes into each bit? See what gas does.

Operation:

$pc \leftarrow pc+4$

Encoding:

FENCE iorw, iorw



5.6.39 FENCE.I

The FENCE.I instruction is used to synchronize the instruction and data streams. RISC-V does not guarantee that stores to instruction memory will be made visible to instruction fetches on the same RISC-V hart until a FENCE.I instruction is executed. A FENCE.I instruction only ensures that a subsequent instruction fetch on a RISC-V hart will see any previous data stores already visible to the same RISC-V hart. FENCE.I does not ensure that other RISC-V harts' instruction fetches will observe the local hart's stores in a multiprocessor system. To make a store to instruction memory visible to all RISC-V harts, the writing hart has to execute a data FENCE before requesting that all remote RISC-V harts execute a FENCE.I. [1, p. 21]

Operation:

$pc \leftarrow pc+4$

Instruction!ECALL
 Instruction!EBREAK
 Instruction!CSRRW
 Instruction!CSRWS

Encoding:

FENCE.I

31	28	27	24	23	20	19	15	14	12	11	7	6	0
pred				succ				funct3			opcode		
0	0	0	0	0	0	0	0	0	0	1	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	1	1
0	0	0	0	0	0	0	0	0	0	0	0	1	1
4				4				3			7		

FENCE

5.6.40 ECALL

The ECALL instruction is used to make a request to the supporting execution environment, which is usually an operating system. The ABI for the system will define how parameters for the environment request are passed, but usually these will be in defined locations in the integer register file. [1, p. 24]

ECALL

31												20				19				15				14		12		11		7				6		0																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																									
funct3												opcode																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																	
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

I-type

5.6.41 EBREAK

The EBREAK instruction is used by debuggers to cause control to be transferred back to a debugging environment. [1, p. 24]

EBREAK

31												20			19		15			14		12		11		7			6		0		
funct3												opcode																					
0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	0	0	1	1		
12												5					3			5					7								

I-type

I-type

5.6.42 CSRRW rd, csr, rs1

The CSRRW (Atomic Read/Write CSR) instruction atomically swaps values in the CSRs and integer registers. CSRRW reads the old value of the CSR, zero-extends the value to XLEN bits, then writes it to integer register rd. The initial value in rs1 is written to the CSR. If rd=x0, then the instruction shall not read the CSR and shall not cause any of the side-effects that might occur on a CSR read. [1, p. 22]

CSRRW x3, 2, x15

31												20				19		15				14		12		11		7				6		0			
csr																rs1				funct3				rd				opcode									
0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	0	0	0	1	0	0	0	1	1	1	1	1	1	0	0	1	1		
12												5				3				5				7													

I-type

I-type

5.6.43 CSRWS rd, csr, rs1


The CSRWS (Atomic Read and Set Bits in CSR) instruction reads the value of the CSR, zero-extends the value to XLEN bits, and writes it to integer register rd. The initial value in integer register rs1 is treated as a bit mask that specifies bit positions to be set in the CSR. Any bit that is high in rs1 will

Instruction!CSRRC
Instruction!CSRRIWI
Instruction!CSRRSI

cause the corresponding bit to be set in the CSR, if that CSR bit is writable. Other bits in the CSR are unaffected (though CSRs might have side effects when written). [1, p. 22]

If $rs1=x0$, then the instruction will not write to the CSR at all, and so shall not cause any of the side effects that might otherwise occur on a CSR write, such as raising illegal instruction exceptions on accesses to read-only CSRs. Note that if $rs1$ specifies a register holding a zero value other than $x0$, the instruction will still attempt to write the unmodified value back to the CSR and will cause any attendant side effects. [1, p. 22]

CSRRS $x3, 2, x15$




5.6.44 CSRRC $rd, csr, rs1$

The CSRRC (Atomic Read and Clear Bits in CSR) instruction reads the value of the CSR, zero-extends the value to XLEN bits, and writes it to integer register rd . The initial value in integer register $rs1$ is treated as a bit mask that specifies bit positions to be cleared in the CSR. Any bit that is high in $rs1$ will cause the corresponding bit to be cleared in the CSR, if that CSR bit is writable. Other bits in the CSR are unaffected. [1, p. 22]

If $rs1=x0$, then the instruction will not write to the CSR at all, and so shall not cause any of the side effects that might otherwise occur on a CSR write, such as raising illegal instruction exceptions on accesses to read-only CSRs. Note that if $rs1$ specifies a register holding a zero value other than $x0$, the instruction will still attempt to write the unmodified value back to the CSR and will cause any attendant side effects. [1, p. 22]


CSRRC $x3, 2, x15$



5.6.45 CSRRWI rd, csr, imm

This instruction is the same as CSRRW except a 5-bit unsigned (zero-extended) immediate value is used rather than the value from a register.

CSRRWI $x3, 2, 7$




5.6.46 CSRRSI $rd, csr, rs1$

This instruction is the same as CSRRS except a 5-bit unsigned (zero-extended) immediate value is used rather than the value from a register.

If the `uimm[4:0]` field is zero, then this instruction will not write to the CSR, and shall not cause any of the side effects that might otherwise occur on a CSR write. For CSRRWI, if `rd=x0`, then the instruction shall not read the CSR and shall not cause any of the side-effects that might occur on a CSR read. [1, p. 22]

Instruction!CSRRCI
RV32M
Instruction!MUL
Instruction!MULH

CSRRSI `x3, 2, 7`




5.6.47 CSRRCI `rd, csr, rs1`

This instruction is the same as CSRRC except a 5-bit unsigned (zero-extended) immediate value is used rather than the value from a register.

If the `uimm[4:0]` field is zero, then this instruction will not write to the CSR, and shall not cause any of the side effects that might otherwise occur on a CSR write. For CSRRWI, if `rd=x0`, then the instruction shall not read the CSR and shall not cause any of the side-effects that might occur on a CSR read. [1, p. 22]

CSRRCI `x3, 2, 7`




5.7 RV32M Standard Extension

32-bit integer multiply and divide instructions.

5.7.1 MUL `rd, rs1, rs2`


Multiply `rs1` by `rs2` and store the least significant 32-bits of the result in `rd`.

MUL `x7, x3, x31`



5.7.2 MULH `rd, rs1, rs2`

MULH `x7, x3, x31`



5.7.3 MULHS rd, rs1, rs2

MULHS x7, x3, x31

31	25	24	20	19	15	14	12	11	7	6	0												
funct7							rs2			rs1			funct3			rd			opcode				
0	0	0	0	0	0	1	1	1	1	1	0	0	1	1	1	0	0	1	1	0	0	1	1
7							5			5			3			5			7				

R-type

Instruction!MULHS
 Instruction!MULHU
 Instruction!DIV
 Instruction!DIVU
 Instruction!REM
 Instruction!REMU

5.7.4 MULHU rd, rs1, rs2

MULHU x7, x3, x31

31	25	24	20	19	15	14	12	11	7	6	0												
funct7							rs2			rs1			funct3			rd			opcode				
0	0	0	0	0	0	1	1	1	1	1	0	0	1	1	1	0	0	1	1	0	0	1	1
7							5			5			3			5			7				

R-type

5.7.5 DIV rd, rs1, rs2

DIV x7, x3, x31

31	25	24	20	19	15	14	12	11	7	6	0												
funct7							rs2			rs1			funct3			rd			opcode				
0	0	0	0	0	0	1	1	1	1	1	0	0	0	1	1	1	0	1	1	0	0	1	1
7							5			5			3			5			7				

R-type

5.7.6 DIVU rd, rs1, rs2

DIVU x7, x3, x31

31	25	24	20	19	15	14	12	11	7	6	0													
funct7							rs2			rs1			funct3			rd			opcode					
0	0	0	0	0	0	1	1	1	1	1	0	1	0	0	1	1	1	0	1	1	0	0	1	1
7							5			5			3			5			7					

R-type

5.7.7 REM rd, rs1, rs2

REM x7, x3, x31

31	25	24	20	19	15	14	12	11	7	6	0												
funct7							rs2			rs1			funct3			rd			opcode				
0	0	0	0	0	0	1	1	1	1	1	0	0	0	1	1	1	0	1	1	0	0	1	1
7							5			5			3			5			7				

R-type

5.7.8 REMU rd, rs1, rs2

REMU x7, x3, x31

31	25	24	20	19	15	14	12	11	7	6	0												
funct7							rs2			rs1			funct3			rd			opcode				
0	0	0	0	0	0	1	1	1	1	1	0	0	1	1	1	1	0	1	1	0	0	1	1
7							5			5			3			5			7				

R-type

2064

5.8 RV32A Standard Extension

RV32A
RV32F
RV32D

2065

32-bit atomic operations.

2066

5.9 RV32F Standard Extension

2067

32-bit IEEE floating point instructions.

2068

5.10 RV32D Standard Extension

2069

64-bit IEEE floating point instructions.

2070

Appendix A

2071

Installing a RISC-V Toolchain

2072

A.1 The GNU Toolchain

2073

Discuss the GNU toolchain elements used to experiment with the material in this book.

2074

The instructions and examples here were all implemented on Ubuntu 16.04 LTS.

2075

Install custom code in a location that will not cause interference with other applications and allow for easy cleanup. These instructions install the toolchain in `/usr/local/riscv`. At any time you can remove the lot and start over by executing the following command:

2076

2077

```
rm -rf /usr/local/riscv/*
```

2078

2079

Tested on Ubuntu 16.04 LTS. 18.04 was just released... update accordingly.

2080

These are the only commands that you should perform as root when installing the toolchain:

2081

2082

2083

2084

2085

```
sudo apt-get install autoconf automake autotools-dev curl libmpc-dev \
libmpfr-dev libgmp-dev gawk build-essential bison flex texinfo gperf \
libtool patchutils bc zlib1g-dev libexpat-dev
sudo mkdir -p /usr/local/riscv/
sudo chmod 777 /usr/local/riscv/
```

2086

All other commands should be executed as a regular user. This will eliminate the possibility of clobbering system files that should not be touched when tinkering with the toolchain applications.

2087

2088

To download, compile and “install” the toolchain:

2089

2090

2091

2092

2093

2094

```
# riscv toolchain:
#
# https://riscv.org/software-tools/risc-v-gnu-compiler-toolchain/

git clone --recursive https://github.com/riscv/riscv-gnu-toolchain
cd riscv-gnu-toolchain
```

► Fix Me:

It would be good to find some Mac and Windows users to write and test proper variations on this section to address those systems. Pull requests, welcome!

```
2095 ./configure --prefix=/usr/local/riscv/rv32i --with-arch=rv32i --with-abi=ilp32
2096 make
2097 make install
```

2098 Need to discuss augmenting the PATH environment variable.

2099 Discuss the choice of ilp32 as well as what the other variations would do.

2100 Discuss rv32im and note that the details are found in [chapter 5](#).

2101 A.2 rvddt

2102 Discuss installing the rvddt simulator here.

2103

Appendix B

2104

Using The RISC-V GNU Toolchain

2105

This chapter discusses using the GNU toolchain elements to experiment with the material in this book.

2106

2107

See [Appendix A](#) if you do not already have the GNU crosscompiler toolchain available on your system.

2108

Discuss the choice of ilp32 as well as what the other variations would do.

2109

Discuss rv32im and note that the details are found in [chapter 5](#).

2110

Discuss installing and using one of the RISC-V simulators here.

2111

Describe the pre-processor, compiler, assembler and linker.

2112

Source, object, and binary files

2113

Assembly syntax (label: mnemonic op1, op2, op3 # comment).

2114

text, data, bss, stack

2115

Labels and scope.

2116

Forward & backward references to throw-away labels.

2117

The entry address of an application.

2118

.s file contain assembler code. .S (or .sx) files contain assembler code that must be preprocessed. [[14](#), p. 29]

2119

2120

Pre-processing conditional assembly using #if.

2121

Building with `-mabi=ilp32 -march=rv32i -mno-fdiv -mno-div` to match the config options on the toolchain.

2122

2123

Linker scripts.

2124

Makefiles

2125

objdump

2126

nm

2127 hexdump -C

2128

Appendix C

2129

Floating Point Numbers

2130

C.1 IEEE-754 Floating Point Number Representation

2131

This section provides an overview of the IEEE-754 32-bit binary floating point format.

2132

- Recall that the place values for integer binary numbers are:

2133

... 128 64 32 16 8 4 2 1

2134

- We can extend this to the right in binary similar to the way we do for decimal numbers:

2135

... 128 64 32 16 8 4 2 1 . 1/2 1/4 1/8 1/16 1/32 1/64 1/128 ...

2136

The ‘.’ in a binary number is a binary point, not a decimal point.

2137

- We use scientific notation as in 2.7×10^{-47} to express either small fractions or large numbers when we are not concerned every last digit needed to represent the entire, exact, value of a number.

2138

2139

2140

- The format of a number in scientific notation is $\text{mantissa} \times \text{base}^{\text{exponent}}$

2141

- In binary we have $\text{mantissa} \times 2^{\text{exponent}}$

2142

2143

- IEEE-754 format requires binary numbers to be *normalized* to $1.\text{significand} \times 2^{\text{exponent}}$ where the *significand* is the portion of the *mantissa* that is to the right of the binary-point.

2144

– The unnormalized binary value of -2.625 is 10.101

2145

– The normalized value of -2.625 is 1.0101×2^1

2146

2147

- We need not store the ‘1.’ because *all* normalized floating point numbers will start that way. Thus we can save memory when storing normalized values by adding 1 to the significand.

2148

31	30	23	22	0
1	1 0 0 0 0 0 0 0	0	1 0 1 0	
sign	exponent		significand	

2149

- $-((1 + \frac{1}{4} + \frac{1}{16}) \times 2^{128-127}) = -((1 + \frac{1}{4} + \frac{1}{16}) \times 2^1) = -(2 + \frac{1}{2} + \frac{1}{8}) = -(2 + .5 + .125) = -2.625$

- IEEE754 formats:

	IEEE754 32-bit	IEEE754 64-bit
sign	1 bit	1 bit
exponent	8 bits (excess-127)	11 bits (excess-1023)
mantissa	23 bits	52 bits
max exponent	127	1023
min exponent	-126	-1022

- When the exponent is all ones, the mantissa is all zeros, and the sign is zero, the number represents positive infinity.
- When the exponent is all ones, the mantissa is all zeros, and the sign is one, the number represents negative infinity.
- Note that the binary representation of an IEEE754 number in memory can be compared for magnitude with another one using the same logic as for comparing two's complement signed integers because the magnitude of an IEEE number grows upward and downward in the same fashion as signed integers. This is why we use excess notation and locate the significand's sign bit on the left of the exponent.
- Note that zero is a special case number. Recall that a normalized number has an implied 1-bit to the left of the significand... which means that there is no way to represent zero! Zero is represented by an exponent of all-zeros and a significand of all-zeros. This definition allows for a positive and a negative zero if we observe that the sign can be either 1 or 0.
- On the number-line, numbers between zero and the smallest fraction in either direction are in the *underflow* areas.
- On the number line, numbers greater than the mantissa of all-ones and the largest exponent allowed are in the *overflow* areas.
- Note that numbers have a higher resolution on the number line when the exponent is smaller.

► Fix Me:

Need to add the standard lecture number-line diagram showing where the over/under-flow areas are and why.

C.1.1 Floating Point Number Accuracy

Due to the finite number of bits used to store the value of a floating point number, it is not possible to represent every one of the infinite values on the real number line. The following C programs illustrate this point.

C.1.1.1 Powers Of Two

Just like the integer numbers, the powers of two that have bits to represent them can be represented perfectly... as can their sums (provided that the significand requires no more than 23 bits.)

Listing C.1: `powersoftwo.c`
Precise Powers of Two

```

2177 1 #include <stdio.h>
2178 2 #include <stdlib.h>
2179 3 #include <unistd.h>
2180 4
2181 5 union floatbin
2182 6 {
2183 7     unsigned int    i;
2184 8     float          f;
2185 9 };
2186

```

```

2187 10 int main()
2188 11 {
2189 12     union floatbin x;
2190 13     union floatbin y;
2191 14     int i;
2192 15     x.f = 1.0;
2193 16     while (x.f > 1.0/1024.0)
2194 17     {
2195 18         y.f = -x.f;
2196 19         printf("%25.10f = %08x    %25.10f = %08x\n", x.f, x.i, y.f, y.i);
2197 20         x.f = x.f/2.0;
2198 21     }
2199 22 }

```

Listing C.2: powersoftwo.out

Output from powersoftwo.c

```

2201
2202 1 1.0000000000 = 3f800000          -1.0000000000 = bf800000
2203 2 0.5000000000 = 3f000000          -0.5000000000 = bf000000
2204 3 0.2500000000 = 3e800000          -0.2500000000 = be800000
2205 4 0.1250000000 = 3e000000          -0.1250000000 = be000000
2206 5 0.0625000000 = 3d800000          -0.0625000000 = bd800000
2207 6 0.0312500000 = 3d000000          -0.0312500000 = bd000000
2208 7 0.0156250000 = 3c800000          -0.0156250000 = bc800000
2209 8 0.0078125000 = 3c000000          -0.0078125000 = bc000000
2210 9 0.0039062500 = 3b800000          -0.0039062500 = bb800000
2211 10 0.0019531250 = 3b000000         -0.0019531250 = bb000000

```

C.1.1.2 Clean Decimal Numbers

When dealing with decimal values, you will find that they don't map simply into binary floating point values.

Note how the decimal numbers are not accurately represented as they get larger. The decimal number on line 10 of [Listing C.4](#) can be perfectly represented in IEEE format. However, a problem arises in the 11th loop iteration. It is due to the fact that the binary number can not be represented accurately in IEEE format. Its least significant bits were truncated in a best-effort attempt at rounding the value off in order to fit the value into the bits provided. This is an example of *low order truncation*. Once this happens, the value of `x.f` is no longer as precise as it could be given more bits in which to save its value.

Listing C.3: cleandecimal.c

Print Clean Decimal Numbers

```

2223
2224 1 #include <stdio.h>
2225 2 #include <stdlib.h>
2226 3 #include <unistd.h>
2227 4
2228 5 union floatbin
2229 6 {
2230 7     unsigned int i;
2231 8     float f;
2232 9 };
2233 10 int main()
2234 11 {
2235 12     union floatbin x, y;
2236 13     int i;
2237 14
2238 15     x.f = 10;
2239 16     while (x.f <= 10000000000000.0)
2240 17     {
2241 18         y.f = -x.f;

```



```

2242 19     printf("%25.10f = %08x      %25.10f = %08x\n", x.f, x.i, y.f, y.i);
2243 20     x.f = x.f*10.0;
2244 21 }
2245 22 }

```

Listing C.4: cleandecimal.out

Output from cleandecimal.c

```

2247      10.0000000000 = 41200000      -10.0000000000 = c1200000
2248 1      100.0000000000 = 42c80000     -100.0000000000 = c2c80000
2249 2      1000.0000000000 = 447a0000    -1000.0000000000 = c47a0000
2250 3      10000.0000000000 = 461c4000   -10000.0000000000 = c61c4000
2251 4      100000.0000000000 = 47c35000  -100000.0000000000 = c7c35000
2252 5      1000000.0000000000 = 49742400 -1000000.0000000000 = c9742400
2253 6      10000000.0000000000 = 4b189680 -10000000.0000000000 = cb189680
2254 7      100000000.0000000000 = 4cbebc20 -100000000.0000000000 = ccbebc20
2255 8      1000000000.0000000000 = 4e6e6b28 -1000000000.0000000000 = ce6e6b28
2256 9      10000000000.0000000000 = 501502f9 -10000000000.0000000000 = d01502f9
2257 10     99999997952.0000000000 = 51ba43b7 -99999997952.0000000000 = d1ba43b7
2258 11     999999995904.0000000000 = 5368d4a5 -999999995904.0000000000 = d368d4a5
2259 12     9999999827968.0000000000 = 551184e7 -9999999827968.0000000000 = d51184e7
2260 13

```

C.1.1.3 Accumulation of Error

These rounding errors can be exaggerated when the number we multiply the `x.f` value by is, itself, something that can not be accurately represented in IEEE form.¹

For example, if we multiply our `x.f` value by $\frac{1}{10}$ each time, we can never be accurate and we start accumulating errors immediately.

Listing C.5: erroraccumulation.c

Accumulation of Error

```

2267 1 #include <stdio.h>
2268 2 #include <stdlib.h>
2269 3 #include <unistd.h>
2270 4
2271 5 union floatbin
2272 6 {
2273 7     unsigned int    i;
2274 8     float          f;
2275 9 };
2276 10 int main()
2277 11 {
2278 12     union floatbin  x, y;
2279 13     int             i;
2280 14
2281 15     x.f = .1;
2282 16     while (x.f <= 2.0)
2283 17     {
2284 18         y.f = -x.f;
2285 19         printf("%25.10f = %08x      %25.10f = %08x\n", x.f, x.i, y.f, y.i);
2286 20         x.f += .1;
2287 21     }
2288 22 }
2289

```

Listing C.6: erroraccumulation.out

Output from erroraccumulation.c

¹Applications requiring accurate decimal values, such as financial accounting systems, can use a packed-decimal numeric format to avoid unexpected oddities caused by the use of binary numbers.

► Fix Me:

In a lecture one would show that one tenth is a repeating non-terminating binary number that gets truncated. This discussion should be reproduced here in text form.

2292 1	0.1000000015 = 3dcccccd	-0.1000000015 = bdcccccd
2293 2	0.2000000030 = 3e4ccccd	-0.2000000030 = be4ccccd
2294 3	0.3000000119 = 3e99999a	-0.3000000119 = be99999a
2295 4	0.4000000060 = 3ecccccd	-0.4000000060 = becccccd
2296 5	0.5000000000 = 3f000000	-0.5000000000 = bf000000
2297 6	0.6000000238 = 3f19999a	-0.6000000238 = bf19999a
2298 7	0.7000000477 = 3f333334	-0.7000000477 = bf333334
2299 8	0.8000000715 = 3f4ccccce	-0.8000000715 = bf4ccccce
2300 9	0.9000000954 = 3f666668	-0.9000000954 = bf666668
2301 10	1.0000001192 = 3f800001	-1.0000001192 = bf800001
2302 11	1.1000001431 = 3f8ccccce	-1.1000001431 = bf8ccccce
2303 12	1.2000001669 = 3f99999b	-1.2000001669 = bf99999b
2304 13	1.3000001907 = 3fa66668	-1.3000001907 = bfa66668
2305 14	1.4000002146 = 3fb33335	-1.4000002146 = bfb33335
2306 15	1.5000002384 = 3fc00002	-1.5000002384 = bfc00002
2307 16	1.6000002623 = 3fcccccf	-1.6000002623 = bfcccccf
2308 17	1.7000002861 = 3fd9999c	-1.7000002861 = bfd9999c
2309 18	1.8000003099 = 3fe66669	-1.8000003099 = bfe66669
2310 19	1.9000003338 = 3ff33336	-1.9000003338 = bff33336

C.1.2 Reducing Error Accumulation

In order to use floating point numbers in a program without causing excessive rounding problems an algorithm can be redesigned such that the accumulation is eliminated. This example is similar to the previous one, but this time we recalculate the desired value from a known-accurate integer value. Some rounding errors remain present, but they can not accumulate.

Listing C.7: errorcompensation.c

Accumulation of Error

```

2317 1 #include <stdio.h>
2318 2 #include <stdlib.h>
2319 3 #include <unistd.h>
2320 4
2321 5 union floatbin
2322 6 {
2323 7     unsigned int    i;
2324 8     float          f;
2325 9 };
2326 10 int main()
2327 11 {
2328 12     union floatbin  x, y;
2329 13     int             i;
2330 14
2331 15     i = 1;
2332 16     while (i <= 20)
2333 17     {
2334 18         x.f = i/10.0;
2335 19         y.f = -x.f;
2336 20         printf("%25.10f = %08x    %25.10f = %08x\n", x.f, x.i, y.f, y.i);
2337 21         i++;
2338 22     }
2339 23     return(0);
2340 24 }
2341 25

```

Listing C.8: errorcompensation.out

Output from erroraccumulation.c

2343 1	0.1000000015 = 3dcccccd	-0.1000000015 = bdcccccd
2344 2	0.2000000030 = 3e4ccccd	-0.2000000030 = be4ccccd
2345 3	0.3000000119 = 3e99999a	-0.3000000119 = be99999a
2346 4	0.4000000060 = 3ecccccd	-0.4000000060 = becccccd
2347 5	0.5000000000 = 3f000000	-0.5000000000 = bf000000

2349	6	0.6000000238 = 3f19999a	-0.6000000238 = bf19999a
2350	7	0.6999999881 = 3f333333	-0.6999999881 = bf333333
2351	8	0.8000000119 = 3f4ccccd	-0.8000000119 = bf4ccccd
2352	9	0.8999999762 = 3f666666	-0.8999999762 = bf666666
2353	10	1.0000000000 = 3f800000	-1.0000000000 = bf800000
2354	11	1.1000000238 = 3f8ccccd	-1.1000000238 = bf8ccccd
2355	12	1.2000000477 = 3f99999a	-1.2000000477 = bf99999a
2356	13	1.2999999523 = 3fa66666	-1.2999999523 = bfa66666
2357	14	1.3999999762 = 3fb33333	-1.3999999762 = bfb33333
2358	15	1.5000000000 = 3fc00000	-1.5000000000 = bfc00000
2359	16	1.6000000238 = 3fcccccd	-1.6000000238 = bfcccccd
2360	17	1.7000000477 = 3fd9999a	-1.7000000477 = bfd9999a
2361	18	1.7999999523 = 3fe66666	-1.7999999523 = bfe66666
2362	19	1.8999999762 = 3ff33333	-1.8999999762 = bff33333
2363	20	2.0000000000 = 40000000	-2.0000000000 = c0000000
2364			

Appendix D

The ASCII Character Set

A slightly abridged version of the Linux “ASCII” man(1) page.

D.1 NAME

ascii - ASCII character set encoded in octal, decimal, and hexadecimal

D.2 DESCRIPTION

ASCII is the American Standard Code for Information Interchange. It is a 7-bit code. Many 8-bit codes (e.g., ISO 8859-1) contain ASCII as their lower half. The international counterpart of ASCII is known as ISO 646-IRV.

The following table contains the 128 ASCII characters.

C program '\X' escapes are noted.

Oct	Dec	Hex	Char	Oct	Dec	Hex	Char
000	0	00	NUL '\0' (null character)	100	64	40	@
001	1	01	SOH (start of heading)	101	65	41	A
002	2	02	STX (start of text)	102	66	42	B
003	3	03	ETX (end of text)	103	67	43	C
004	4	04	EOT (end of transmission)	104	68	44	D
005	5	05	ENQ (enquiry)	105	69	45	E
006	6	06	ACK (acknowledge)	106	70	46	F
007	7	07	BEL '\a' (bell)	107	71	47	G
010	8	08	BS '\b' (backspace)	110	72	48	H
011	9	09	HT '\t' (horizontal tab)	111	73	49	I
012	10	0A	LF '\n' (new line)	112	74	4A	J
013	11	0B	VT '\v' (vertical tab)	113	75	4B	K
014	12	0C	FF '\f' (form feed)	114	76	4C	L
015	13	0D	CR '\r' (carriage ret)	115	77	4D	M

2392	016	14	0E	S0 (shift out)	116	78	4E	N	
2393	017	15	0F	SI (shift in)	117	79	4F	O	
2394	020	16	10	DLE (data link escape)	120	80	50	P	
2395	021	17	11	DC1 (device control 1)	121	81	51	Q	
2396	022	18	12	DC2 (device control 2)	122	82	52	R	
2397	023	19	13	DC3 (device control 3)	123	83	53	S	
2398	024	20	14	DC4 (device control 4)	124	84	54	T	
2399	025	21	15	NAK (negative ack.)	125	85	55	U	
2400	026	22	16	SYN (synchronous idle)	126	86	56	V	
2401	027	23	17	ETB (end of trans. blk)	127	87	57	W	
2402	030	24	18	CAN (cancel)	130	88	58	X	
2403	031	25	19	EM (end of medium)	131	89	59	Y	
2404	032	26	1A	SUB (substitute)	132	90	5A	Z	
2405	033	27	1B	ESC (escape)	133	91	5B	[
2406	034	28	1C	FS (file separator)	134	92	5C	\	'\'
2407	035	29	1D	GS (group separator)	135	93	5D]	
2408	036	30	1E	RS (record separator)	136	94	5E	^	
2409	037	31	1F	US (unit separator)	137	95	5F	_	
2410	040	32	20	SPACE	140	96	60	`	
2411	041	33	21	!	141	97	61	a	
2412	042	34	22	"	142	98	62	b	
2413	043	35	23	#	143	99	63	c	
2414	044	36	24	\$	144	100	64	d	
2415	045	37	25	%	145	101	65	e	
2416	046	38	26	&	146	102	66	f	
2417	047	39	27	'	147	103	67	g	
2418	050	40	28	(150	104	68	h	
2419	051	41	29)	151	105	69	i	
2420	052	42	2A	*	152	106	6A	j	
2421	053	43	2B	+	153	107	6B	k	
2422	054	44	2C	,	154	108	6C	l	
2423	055	45	2D	-	155	109	6D	m	
2424	056	46	2E	.	156	110	6E	n	
2425	057	47	2F	/	157	111	6F	o	
2426	060	48	30	0	160	112	70	p	
2427	061	49	31	1	161	113	71	q	
2428	062	50	32	2	162	114	72	r	
2429	063	51	33	3	163	115	73	s	
2430	064	52	34	4	164	116	74	t	
2431	065	53	35	5	165	117	75	u	
2432	066	54	36	6	166	118	76	v	
2433	067	55	37	7	167	119	77	w	
2434	070	56	38	8	170	120	78	x	
2435	071	57	39	9	171	121	79	y	
2436	072	58	3A	:	172	122	7A	z	
2437	073	59	3B	;	173	123	7B	{	
2438	074	60	3C	<	174	124	7C		
2439	075	61	3D	=	175	125	7D	}	
2440	076	62	3E	>	176	126	7E	~	
2441	077	63	3F	?	177	127	7F	DEL	

D.2.1 Tables

For convenience, below are more compact tables in hex and decimal.

2 3 4 5 6 7	30 40 50 60 70 80 90 100 110 120
-----	-----
0: 0 @ P ' p	0: (2 < F P Z d n x
1: ! 1 A Q a q	1:) 3 = G Q [e o y
2: " 2 B R b r	2: * 4 > H R \ f p z
3: # 3 C S c s	3: ! + 5 ? I S] g q {
4: \$ 4 D T d t	4: " , 6 @ J T ^ h r
5: % 5 E U e u	5: # - 7 A K U _ i s }
6: & 6 F V f v	6: \$. 8 B L V ' j t ~
7: ' 7 G W g w	7: % / 9 C M W a k u DEL
8: (8 H X h x	8: & 0 : D N X b l v
9:) 9 I Y i y	9: ' 1 ; E O Y c m w
A: * : J Z j z	
B: + ; K [k {	
C: , < L \ l	
D: - = M] m }	
E: . > N ^ n ~	
F: / ? 0 _ o DEL	

D.3 NOTES

D.3.1 History

An ascii manual page appeared in Version 7 of AT&T UNIX.

On older terminals, the underscore code is displayed as a left arrow, called backarrow, the caret is displayed as an up-arrow and the vertical bar has a hole in the middle.

Uppercase and lowercase characters differ by just one bit and the ASCII character 2 differs from the double quote by just one bit, too. That made it much easier to encode characters mechanically or with a non-microcontroller-based electronic keyboard and that pairing was found on old teletypes.

The ASCII standard was published by the United States of America Standards Institute (USASI) in 1968.

D.4 COLOPHON

This page is part of release 4.04 of the Linux man-pages project. A description of the project, information about reporting bugs, and the latest version of this page, can be found at <http://www.kernel.org/doc/man-pages/>.

2476

Appendix E

2477

Attribution 4.0 International

2478
2479
2480
2481
2482

Creative Commons Corporation ("Creative Commons") is not a law firm and does not provide legal services or legal advice. Distribution of Creative Commons public licenses does not create a lawyer-client or other relationship. Creative Commons makes its licenses and related information available on an "as-is" basis. Creative Commons gives no warranties regarding its licenses, any material licensed under their terms and conditions, or any related information. Creative Commons disclaims all liability for damages resulting from their use to the fullest extent possible.

2483

Using Creative Commons Public Licenses

2484
2485
2486
2487

Creative Commons public licenses provide a standard set of terms and conditions that creators and other rights holders may use to share original works of authorship and other material subject to copyright and certain other rights specified in the public license below. The following considerations are for informational purposes only, are not exhaustive, and do not form part of our licenses.

2488
2489
2490
2491
2492
2493

Considerations for licensors: Our public licenses are intended for use by those authorized to give the public permission to use material in ways otherwise restricted by copyright and certain other rights. Our licenses are irrevocable. Licensors should read and understand the terms and conditions of the license they choose before applying it. Licensors should also secure all rights necessary before applying our licenses so that the public can reuse the material as expected. Licensors should clearly mark any material not subject to the license. This includes other CC-licensed material, or material used under an exception or limitation to copyright. More considerations for licensors: http://wiki.creativecommons.org/Considerations_for_licensors

2494
2495
2496
2497
2498
2499
2500
2501
2502

Considerations for the public: By using one of our public licenses, a licensor grants the public permission to use the licensed material under specified terms and conditions. If the licensor's permission is not necessary for any reason—for example, because of any applicable exception or limitation to copyright—then that use is not regulated by the license. Our licenses grant only permissions under copyright and certain other rights that a licensor has authority to grant. Use of the licensed material may still be restricted for other reasons, including because others have copyright or other rights in the material. A licensor may make special requests, such as asking that all changes be marked or described. Although not required by our licenses, you are encouraged to respect those requests where reasonable. More considerations for the public: http://wiki.creativecommons.org/Considerations_for_licensees

2503

Creative Commons Attribution 4.0 International Public License

2504
2505
2506
2507
2508

By exercising the Licensed Rights (defined below), You accept and agree to be bound by the terms and conditions of this Creative Commons Attribution 4.0 International Public License ("Public License"). To the extent this Public License may be interpreted as a contract, You are granted the Licensed Rights in consideration of Your acceptance of these terms and conditions, and the Licensor grants You such rights in consideration of benefits the Licensor receives from making the Licensed Material available under these terms and conditions.

2509

Section 1. Definitions

2510
2511
2512
2513
2514

- a. Adapted Material means material subject to Copyright and Similar Rights that is derived from or based upon the Licensed Material and in which the Licensed Material is translated, altered, arranged, transformed, or otherwise modified in a manner requiring permission under the Copyright and Similar Rights held by the Licensor. For purposes of this Public License, where the Licensed Material is a musical work, performance, or sound recording, Adapted Material is always produced where the Licensed Material is synched in timed relation with a moving image.

2515

- b. Adapter's License means the license You apply to Your Copyright and Similar Rights in Your contributions to Adapted

Material in accordance with the terms and conditions of this Public License.

- c. Copyright and Similar Rights means copyright and/or similar rights closely related to copyright including, without limitation, performance, broadcast, sound recording, and Sui Generis Database Rights, without regard to how the rights are labeled or categorized. For purposes of this Public License, the rights specified in Section 2(b)(1)-(2) are not Copyright and Similar Rights.
- d. Effective Technological Measures means those measures that, in the absence of proper authority, may not be circumvented under laws fulfilling obligations under Article 11 of the WIPO Copyright Treaty adopted on December 20, 1996, and/or similar international agreements.
- e. Exceptions and Limitations means fair use, fair dealing, and/or any other exception or limitation to Copyright and Similar Rights that applies to Your use of the Licensed Material.
- f. Licensed Material means the artistic or literary work, database, or other material to which the Licensor applied this Public License.
- g. Licensed Rights means the rights granted to You subject to the terms and conditions of this Public License, which are limited to all Copyright and Similar Rights that apply to Your use of the Licensed Material and that the Licensor has authority to license.
- h. Licensor means the individual(s) or entity(ies) granting rights under this Public License.
- i. Share means to provide material to the public by any means or process that requires permission under the Licensed Rights, such as reproduction, public display, public performance, distribution, dissemination, communication, or importation, and to make material available to the public including in ways that members of the public may access the material from a place and at a time individually chosen by them.
- j. Sui Generis Database Rights means rights other than copyright resulting from Directive 96/9/EC of the European Parliament and of the Council of 11 March 1996 on the legal protection of databases, as amended and/or succeeded, as well as other essentially equivalent rights anywhere in the world.
- k. You means the individual or entity exercising the Licensed Rights under this Public License. Your has a corresponding meaning.

Section 2. Scope

- a. License grant.
 - 1. Subject to the terms and conditions of this Public License, the Licensor hereby grants You a worldwide, royalty-free, non-sublicensable, non-exclusive, irrevocable license to exercise the Licensed Rights in the Licensed Material to:
 - a. reproduce and Share the Licensed Material, in whole or in part; and
 - b. produce, reproduce, and Share Adapted Material.
 - 2. Exceptions and Limitations. For the avoidance of doubt, where Exceptions and Limitations apply to Your use, this Public License does not apply, and You do not need to comply with its terms and conditions.
 - 3. Term. The term of this Public License is specified in Section 6(a).
 - 4. Media and formats; technical modifications allowed. The Licensor authorizes You to exercise the Licensed Rights in all media and formats whether now known or hereafter created, and to make technical modifications necessary to do so. The Licensor waives and/or agrees not to assert any right or authority to forbid You from making technical modifications necessary to exercise the Licensed Rights, including technical modifications necessary to circumvent Effective Technological Measures. For purposes of this Public License, simply making modifications authorized by this Section 2(a) (4) never produces Adapted Material.
 - 5. Downstream recipients.
 - a. Offer from the Licensor – Licensed Material. Every recipient of the Licensed Material automatically receives an offer from the Licensor to exercise the Licensed Rights under the terms and conditions of this Public License.
 - b. No downstream restrictions. You may not offer or impose any additional or different terms or conditions on, or apply any Effective Technological Measures to, the Licensed Material if doing so restricts exercise of the Licensed Rights by any recipient of the Licensed Material.
 - 6. No endorsement. Nothing in this Public License constitutes or may be construed as permission to assert or imply that You are, or that Your use of the Licensed Material is, connected with, or sponsored, endorsed, or granted official status by, the Licensor or others designated to receive attribution as provided in Section 3(a)(1)(A)(i).
- b. Other rights.
 - 1. Moral rights, such as the right of integrity, are not licensed under this Public License, nor are publicity, privacy, and/or other similar personality rights; however, to the extent possible, the Licensor waives and/or agrees not to assert any such rights held by the Licensor to the limited extent necessary to allow You to exercise the Licensed Rights, but not otherwise.
 - 2. Patent and trademark rights are not licensed under this Public License.
 - 3. To the extent possible, the Licensor waives any right to collect royalties from You for the exercise of the Licensed Rights, whether directly or through a collecting society under any voluntary or waivable statutory or compulsory licensing scheme. In all other cases the Licensor expressly reserves any right to collect such royalties.

2576
2577
2578
2579
2580
2581
2582
2583
2584
2585
2586
2587
2588
2589
2590
2591
2592
2593
2594
2595
2596

2597

2598

2599
2600

2601
2602
2603

2604
2605

2606
2607

2608

2609
2610
2611
2612
2613
2614
2615
2616

2617
2618
2619
2620
2621
2622
2623

2624
2625

Section 3. License Conditions

- Your exercise of the Licensed Rights is expressly made subject to the following conditions.
- a. Attribution.
 1. If You Share the Licensed Material (including in modified form), You must:
 - a. retain the following if it is supplied by the Licensor with the Licensed Material:
 - i. identification of the creator(s) of the Licensed Material and any others designated to receive attribution, in any reasonable manner requested by the Licensor (including by pseudonym if designated);
 - ii. a copyright notice;
 - iii. a notice that refers to this Public License;
 - iv. a notice that refers to the disclaimer of warranties;
 - v. a URI or hyperlink to the Licensed Material to the extent reasonably practicable;
 - b. indicate if You modified the Licensed Material and retain an indication of any previous modifications; and
 - c. indicate the Licensed Material is licensed under this Public License, and include the text of, or the URI or hyperlink to, this Public License.
 2. You may satisfy the conditions in Section 3(a)(1) in any reasonable manner based on the medium, means, and context in which You Share the Licensed Material. For example, it may be reasonable to satisfy the conditions by providing a URI or hyperlink to a resource that includes the required information.
 3. If requested by the Licensor, You must remove any of the information required by Section 3(a)(1)(A) to the extent reasonably practicable.
 4. If You Share Adapted Material You produce, the Adapter's License You apply must not prevent recipients of the Adapted Material from complying with this Public License.

Section 4. Sui Generis Database Rights

- Where the Licensed Rights include Sui Generis Database Rights that apply to Your use of the Licensed Material:
- a. for the avoidance of doubt, Section 2(a)(1) grants You the right to extract, reuse, reproduce, and Share all or a substantial portion of the contents of the database;
 - b. if You include all or a substantial portion of the database contents in a database in which You have Sui Generis Database Rights, then the database in which You have Sui Generis Database Rights (but not its individual contents) is Adapted Material; and
 - c. You must comply with the conditions in Section 3(a) if You Share all or a substantial portion of the contents of the database.
- For the avoidance of doubt, this Section 4 supplements and does not replace Your obligations under this Public License where the Licensed Rights include other Copyright and Similar Rights.

Section 5. Disclaimer of Warranties and Limitation of Liability

- a. UNLESS OTHERWISE SEPARATELY UNDERTAKEN BY THE LICENSOR, TO THE EXTENT POSSIBLE, THE LICENSOR OFFERS THE LICENSED MATERIAL AS-IS AND AS-AVAILABLE, AND MAKES NO REPRESENTATIONS OR WARRANTIES OF ANY KIND CONCERNING THE LICENSED MATERIAL, WHETHER EXPRESS, IMPLIED, STATUTORY, OR OTHER. THIS INCLUDES, WITHOUT LIMITATION, WARRANTIES OF TITLE, MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, NON-INFRINGEMENT, ABSENCE OF LATENT OR OTHER DEFECTS, ACCURACY, OR THE PRESENCE OR ABSENCE OF ERRORS, WHETHER OR NOT KNOWN OR DISCOVERABLE. WHERE DISCLAIMERS OF WARRANTIES ARE NOT ALLOWED IN FULL OR IN PART, THIS DISCLAIMER MAY NOT APPLY TO YOU.
- b. TO THE EXTENT POSSIBLE, IN NO EVENT WILL THE LICENSOR BE LIABLE TO YOU ON ANY LEGAL THEORY (INCLUDING, WITHOUT LIMITATION, NEGLIGENCE) OR OTHERWISE FOR ANY DIRECT, SPECIAL, INDIRECT, INCIDENTAL, CONSEQUENTIAL, PUNITIVE, EXEMPLARY, OR OTHER LOSSES, COSTS, EXPENSES, OR DAMAGES ARISING OUT OF THIS PUBLIC LICENSE OR USE OF THE LICENSED MATERIAL, EVEN IF THE LICENSOR HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH LOSSES, COSTS, EXPENSES, OR DAMAGES. WHERE A LIMITATION OF LIABILITY IS NOT ALLOWED IN FULL OR IN PART, THIS LIMITATION MAY NOT APPLY TO YOU.
- c. The disclaimer of warranties and limitation of liability provided above shall be interpreted in a manner that, to the extent possible, most closely approximates an absolute disclaimer and waiver of all liability.

2626

Section 6. Term and Termination

2627

a. This Public License applies for the term of the Copyright and Similar Rights licensed here. However, if You fail to comply with this Public License, then Your rights under this Public License terminate automatically.

2628

2629

b. Where Your right to use the Licensed Material has terminated under Section 6(a), it reinstates:

2630

1. automatically as of the date the violation is cured, provided it is cured within 30 days of Your discovery of the violation; or

2631

2632

2. upon express reinstatement by the Licensor.

2633

For the avoidance of doubt, this Section 6(b) does not affect any right the Licensor may have to seek remedies for Your violations of this Public License.

2634

2635

c. For the avoidance of doubt, the Licensor may also offer the Licensed Material under separate terms or conditions or stop distributing the Licensed Material at any time; however, doing so will not terminate this Public License.

2636

2637

d. Sections 1, 5, 6, 7, and 8 survive termination of this Public License.

2638

Section 7. Other Terms and Conditions

2639

a. The Licensor shall not be bound by any additional or different terms or conditions communicated by You unless expressly agreed.

2640

2641

b. Any arrangements, understandings, or agreements regarding the Licensed Material not stated herein are separate from and independent of the terms and conditions of this Public License.

2642

2643

Section 8. Interpretation

2644

a. For the avoidance of doubt, this Public License does not, and shall not be interpreted to, reduce, limit, restrict, or impose conditions on any use of the Licensed Material that could lawfully be made without permission under this Public License.

2645

2646

2647

b. To the extent possible, if any provision of this Public License is deemed unenforceable, it shall be automatically reformed to the minimum extent necessary to make it enforceable. If the provision cannot be reformed, it shall be severed from this Public License without affecting the enforceability of the remaining terms and conditions.

2648

2649

2650

c. No term or condition of this Public License will be waived and no failure to comply consented to unless expressly agreed to by the Licensor.

2651

2652

d. Nothing in this Public License constitutes or may be interpreted as a limitation upon, or waiver of, any privileges and immunities that apply to the Licensor or You, including from the legal processes of any jurisdiction or authority.

2653

2654

2655

Creative Commons is not a party to its public licenses. Notwithstanding, Creative Commons may elect to apply one of its public licenses to material it publishes and in those instances will be considered the Licensor. The text of the Creative Commons public licenses is dedicated to the public domain under the CC0 Public Domain Dedication. Except for the limited purpose of indicating that material is shared under a Creative Commons public license or as otherwise permitted by the Creative Commons policies published at <http://creativecommons.org/policies>, Creative Commons does not authorize the use of the trademark “Creative Commons” or any other trademark or logo of Creative Commons without its prior written consent including, without limitation, in connection with any unauthorized modifications to any of its public licenses or any other arrangements, understandings, or agreements concerning use of licensed material. For the avoidance of doubt, this paragraph does not form part of the public licenses.

2656

2657

2658

2659

2660

2661

2662

2663

2664

Creative Commons may be contacted at <http://creativecommons.org>.

2665

Bibliography

- 2666 [1] RISC-V Foundation, *The RISC-V Instruction Set Manual, Volume I: User-Level ISA, Document*
 2667 *Version 2.2*, 5 2017. Editors Andrew Waterman and Krste Asanović. [viii](#), [3](#), [4](#), [16](#), [24](#), [41](#), [43](#),
 2668 [44](#), [51](#), [53](#), [54](#), [55](#), [56](#), [57](#), [58](#), [59](#), [60](#), [61](#), [62](#)
- 2669 [2] D. Patterson and A. Waterman, *The RISC-V Reader: An Open Architecture Atlas*. Strawberry
 2670 Canyon, 11 2017. ISBN: 978-0999249116. [viii](#)
- 2671 [3] D. Patterson and J. Hennessy, *Computer Organization and Design RISC-V Edition: The Hard-*
 2672 *ware Software Interface*. Morgan Kaufmann, 4 2017. ISBN: 978-0128122754. [viii](#), [19](#)
- 2673 [4] W. F. Decker, “A modern approach to teaching computer organization and assembly language
 2674 programming,” *SIGCSE Bull.*, vol. 17, pp. 38–44, 12 1985. [viii](#)
- 2675 [5] Texas Instruments, *SN54190, SN54191, SN54LS190, SN54LS191, SN74190, SN74191,*
 2676 *SN74LS190, SN74LS191 Synchronous Up/Down Counters With Down/Up Mode Control*, 3 1988.
 2677 [viii](#)
- 2678 [6] Texas Instruments, *SN54154, SN74154 4-line to 16-line Decoders/Demultiplexers*, 12 1972. [viii](#)
- 2679 [7] Intel, *MCS-85 User’s Manual*, 9 1978. [viii](#)
- 2680 [8] Radio Shack, *TRS-80 Editor/Assembler Operation and Reference Manual*, 1978. [viii](#)
- 2681 [9] Motorola, *MC68000 16-bit Microprocessor User’s Manual*, 2nd ed., 1 1980. MC68000UM(AD2).
 2682 [viii](#)
- 2683 [10] R. A. Overbeek and W. E. Singletary, *Assembler Language With ASSIST*. Science Research
 2684 Associates, Inc., 2nd ed., 1983. [viii](#)
- 2685 [11] IBM, *IBM System/370 Principals of Operation*, 7th ed., 3 1980. [viii](#)
- 2686 [12] IBM, *OS/VS-DOS/VSE-VM/370 Assembler Language*, 6th ed., 3 1979. [viii](#)
- 2687 [13] D. Cohen, “*IEN 137, On Holy Wars and a Plea for Peace*,” Apr. 1980. This note discusses the
 2688 Big-Endian/Little-Endian byte/bit-order controversy, but did not settle it. A decade later, David
 2689 V. James in “Multiplexed Buses: The Endian Wars Continue”, *IEEE Micro*, **10**(3), 9–21 (1990)
 2690 continued the discussion. [18](#)
- 2691 [14] R. M. Stallman and the GCC Developer Community, *Using the GNU Compiler Collection (For*
 2692 *GCC version 7.3.0)*. Free Software Foundation, 51 Franklin Street, Fifth Floor, Boston, MA
 2693 02110-1301 USA: GNU Press, 2017. [67](#)
- 2694 [15] RISC-V Foundation, *The RISC-V Instruction Set Manual, Volume II: Privileged Architecture,*
 2695 *Document Version 1.10*, 5 2017. Editors Andrew Waterman and Krste Asanović.
- 2696 [16] P. Dabbelt, S. O’Rear, K. Cheng, A. Waterman, M. Clark, A. Bradbury, D. Horner, M. Nordlund,
 2697 and K. Merker, *RISC-V ELF psABI specification*, 2017.

- 2698 [17] National Semiconductor Corporation, *Series 32000 Databook*, 1986.

2699

2700

Glossary

2701

address A numeric value used to uniquely identify each [byte](#) of main memory. [2](#), [82](#)

2702

alignment Refers to a range of numeric values that begin at a multiple of some number. Primarily used when referring to a memory address. For example an alignment of two refers to one or more addresses starting at even address and continuing onto subsequent adjacent, increasing memory addresses. [19](#), [82](#)

2703

2704

2705

2706

big endian A number format where the most significant values are printed to the left of the lesser significant values. This is the method that everyone uses to write decimal numbers every day. [22](#), [23](#), [82](#), [83](#)

2707

2708

2709

binary Something that has two parts or states. In computing these two states are represented by the numbers one and zero or by the conditions true and false and can be stored in one [bit](#). [1](#), [3](#), [82](#), [83](#), [84](#)

2710

2711

2712

bit One binary digit. [3](#), [6](#), [10](#), [82](#), [83](#), [84](#)

2713

byte A [binary](#) value represented by 8 [bits](#). [2](#), [6](#), [82](#), [83](#), [84](#)

2714

CPU Central Processing Unit. [1](#), [2](#), [82](#)

2715

doubleword A [binary](#) value represented by 64 [bits](#). [82](#)

2716

exception An error encountered by the CPU while executing an instruction that can not be completed. [19](#), [82](#)

2717

2718

fullword A [binary](#) value represented by 32 [bits](#). [6](#), [82](#)

2719

halfword A [binary](#) value represented by 16 [bits](#). [6](#), [82](#)

2720

hart Hardware Thread. [3](#), [82](#)

2721

hexadecimal A base-16 numbering system whose digits are 0123456789abcdef. The hex digits ([hits](#)) are not case-sensitive. [22](#), [23](#), [82](#), [83](#)

2722

2723

high order bits Some number of [MSBs](#). [82](#)

2724

hit One [hexadecimal](#) digit. [10](#), [11](#), [82](#), [83](#), [84](#)

2725

ISA Instruction Set Architecture. [3](#), [4](#), [82](#)

2726

LaTeX Is a mark up language specially suited for scientific documents. [82](#)

- little endian** A number format where the least significant values are printed to the left of the more significant values. This is the opposite ordering that everyone learns in grade school when learning how to count. For example a **big endian** number written as “1234” would be written in little endian form as “4321”. [82](#)
- low order bits** Some number of **LSBs**. [82](#)
- LSB** Least Significant Bit. [10](#), [12](#), [33](#), [36](#), [38](#), [82](#), [83](#)
- machine language** The instructions that are executed by a CPU that are expressed in the form of **binary** values. [1](#), [82](#)
- mnemonic** A method used to remember something. In the case of assembly language, each machine instruction is given a name so the programmer need not memorize the binary values of each machine instruction. [1](#), [82](#)
- MSB** Most Significant Bit. [10](#), [12](#), [13](#), [16](#), [17](#), [33](#), [34](#), [38](#), [82](#), [83](#)
- nybble** Half of a **byte** is a *nybble* (sometimes spelled nibble.) Another word for *hit*. [10](#), [82](#)
- overflow** The situation where the result of an addition or subtraction operation is approaching positive or negative infinity and exceeds the number of bits allotted to contain the result. This is typically caused by high-order truncation. [70](#), [82](#)
- program** A ordered list of one or more instructions. [1](#), [82](#)
- quadword** A **binary** value represented by 128 **bits**. [82](#)
- RAM** Random Access Memory. [2](#), [82](#)
- register** A unit of storage inside a CPU with the capacity of **XLEN bits**. [2](#), [82](#), [84](#)
- ROM** Read Only Memory. [2](#), [82](#)
- RV32** Short for RISC-V 32. The number 32 refers to the **XLEN**. [41](#), [82](#)
- RV64** Short for RISC-V 64. The number 64 refers to the **XLEN**. [82](#)
- rvddt** A RV32I simulator and debugging tool inspired by the simplicity of the Dynamic Debugging Tool (ddt) that was part of the CP/M operating system. [21](#), [82](#)
- thread** An stream of instructions. When plural, it is used to refer to the ability of a CPU to execute multiple instruction streams at the same time. [3](#), [82](#)
- underflow** The situation where the result of an addition or subtraction operation is approaching zero and exceeds the number of bits allotted to contain the result. This is typically caused by low-order truncation. [70](#), [82](#)
- XLEN** The number of bits a RISC-V x integer **register** (such as x0). For RV32 XLEN=32, RV64 XLEN=64 and so on. [38](#), [39](#), [82](#), [84](#)

Index

2759	A	2802	MULH, 62
2760	ALU, 3	2803	MULHS, 63
2761	ASCII, 75	2804	MULHU, 63
		2805	mv, 27
2762	C	2806	nop, 25
2763	CPU, 2	2807	OR, 58
		2808	ORI, 52
2764	H	2809	REM, 63
2765	hart, 3	2810	REMU, 63
		2811	SB, 49
2766	I	2812	SH, 49
2767	Instruction	2813	SLL, 55
2768	ADD, 54	2814	SLLI, 53
2769	ADDI, 50	2815	SLT, 56
2770	addi, 25	2816	SLTI, 50
2771	AND, 58	2817	SLTIU, 51
2772	ANDI, 52	2818	SLTU, 56
2773	AUIPC, 42	2819	SRA, 57
2774	BEQ, 45	2820	SRAI, 54
2775	BGE, 46	2821	SRL, 57
2776	BGEU, 47	2822	SRLI, 53
2777	BLT, 46	2823	SUB, 55
2778	BLTU, 47	2824	SW, 50
2779	BNE, 45	2825	XOR, 56
2780	CSRRC, 61	2826	XORI, 51
2781	CSRRCI, 62	2827	instruction cycle, 4
2782	CSRRS, 60	2828	instruction decode, 5
2783	CSRRSI, 61	2829	instruction execute, 5
2784	CSRRAW, 60	2830	instruction fetch, 5
2785	CSRRAWI, 61	2831	ISA, 4
2786	DIV, 63		
2787	DIVU, 63	2832	L
2788	EBREAK, 60	2833	Least significant bit, 10
2789	ebreak, 24	2834	LSB, <i>see</i> Least significant bit
2790	ECALL, 60		
2791	FENCE, 59	2835	M
2792	FENCE.I, 59	2836	Most significant bit, 10
2793	JAL, 43	2837	MSB, <i>see</i> Most significant bit
2794	JALR, 44		
2795	LB, 47	2838	O
2796	LBU, 48	2839	objdump, 26
2797	LH, 48		
2798	LHU, 49	2840	R
2799	LUI, 42	2841	register, 2, 3
2800	LW, 48	2842	RV32, 33
2801	MUL, 62	2843	RV32A, 4, 64

2844 RV32C, 4
2845 RV32D, 4, 64
2846 RV32F, 4, 64
2847 RV32G, 4
2848 RV32I, 4, 41
2849 RV32M, 4, 62
2850 RV32Q, 4
2851 rvddt, 21

2852 **S**
2853 sign extension, 16